

# Working with RDDs – Resilient Distributed Datasets

---



**Mohit Batra**

Founder, Crystal Talks

[linkedin.com/in/mohitbatra](https://www.linkedin.com/in/mohitbatra)

# Overview



**Understand RDDs**

**Create RDDs**

**Work with Pair RDDs**

**Apply operations on RDDs**

**Use narrow transformations**

**Use wide transformations and data shuffling**

**Spark application concepts**

# Understanding RDDs

---

**Resilient Distributed Dataset (RDD) is  
the native Data Structure of Spark**

# Resilient Distributed Datasets

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

CSV File in Storage

Read File with  
.....→  
RDD API

1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

RDD in Spark

**In-memory objects**

**Do not have schema**

**Distributed collection of elements**

**All processing in Spark happens on RDDs**



1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



Id, City, Amount		
1, Seattle, 600		
2, London, 300		
3, Delhi, 700		
4, Seattle, 400		
5, Paris, 900		
6, Delhi, 200		
7, Seattle, 900		

Storage

App 1

Job

Driver

Spark Session

Executor 1

Core

Core


Executor 2

Core

Core



1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



Id, City, Amount		
1, Seattle, 600		
2, London, 300		
3, Delhi, 700		
4, Seattle, 400		
5, Paris, 900		
6, Delhi, 200		
7, Seattle, 900		

Storage

App 1

Job

Driver

Spark Session

Logically split  
file1 in 3 parts



Executor 1

Core

Core

Executor 2

Core

Core



1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



1, Seattle, 600

2, London, 300

3, Delhi, 700

4, Seattle, 400

5, Paris, 900

6, Delhi, 200

7, Seattle, 900

Storage



App 1

Job

Driver

Spark Session

Logically split  
file1 in 3 parts

Executor 1

Core

Core

Executor 2

Core

Core





1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



1, Seattle, 600

2, London, 300

3, Delhi, 700

4, Seattle, 400

5, Paris, 900

6, Delhi, 200

7, Seattle, 900

Storage

App 1

Job

Driver

Spark Session

Split Job into Tasks

No. of Tasks =  
No. of Partitions



Executor 1

Core

Core

Executor 2

Core

Core



1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



1, Seattle, 600

2, London, 300

3, Delhi, 700

4, Seattle, 400

5, Paris, 900

6, Delhi, 200

7, Seattle, 900

Storage

App 1

Job

Driver

Spark Session

Split Job into Tasks

No. of Tasks =  
No. of Partitions

Task 1

Task 2

Task 3

Executor 1

Core

Core

Executor 2

Core

Core



1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



1, Seattle, 600

2, London, 300

3, Delhi, 700

4, Seattle, 400

5, Paris, 900

6, Delhi, 200

7, Seattle, 900

Storage

App 1

Job

Driver

Spark Session

Executor 1

Core

Task 1

Core

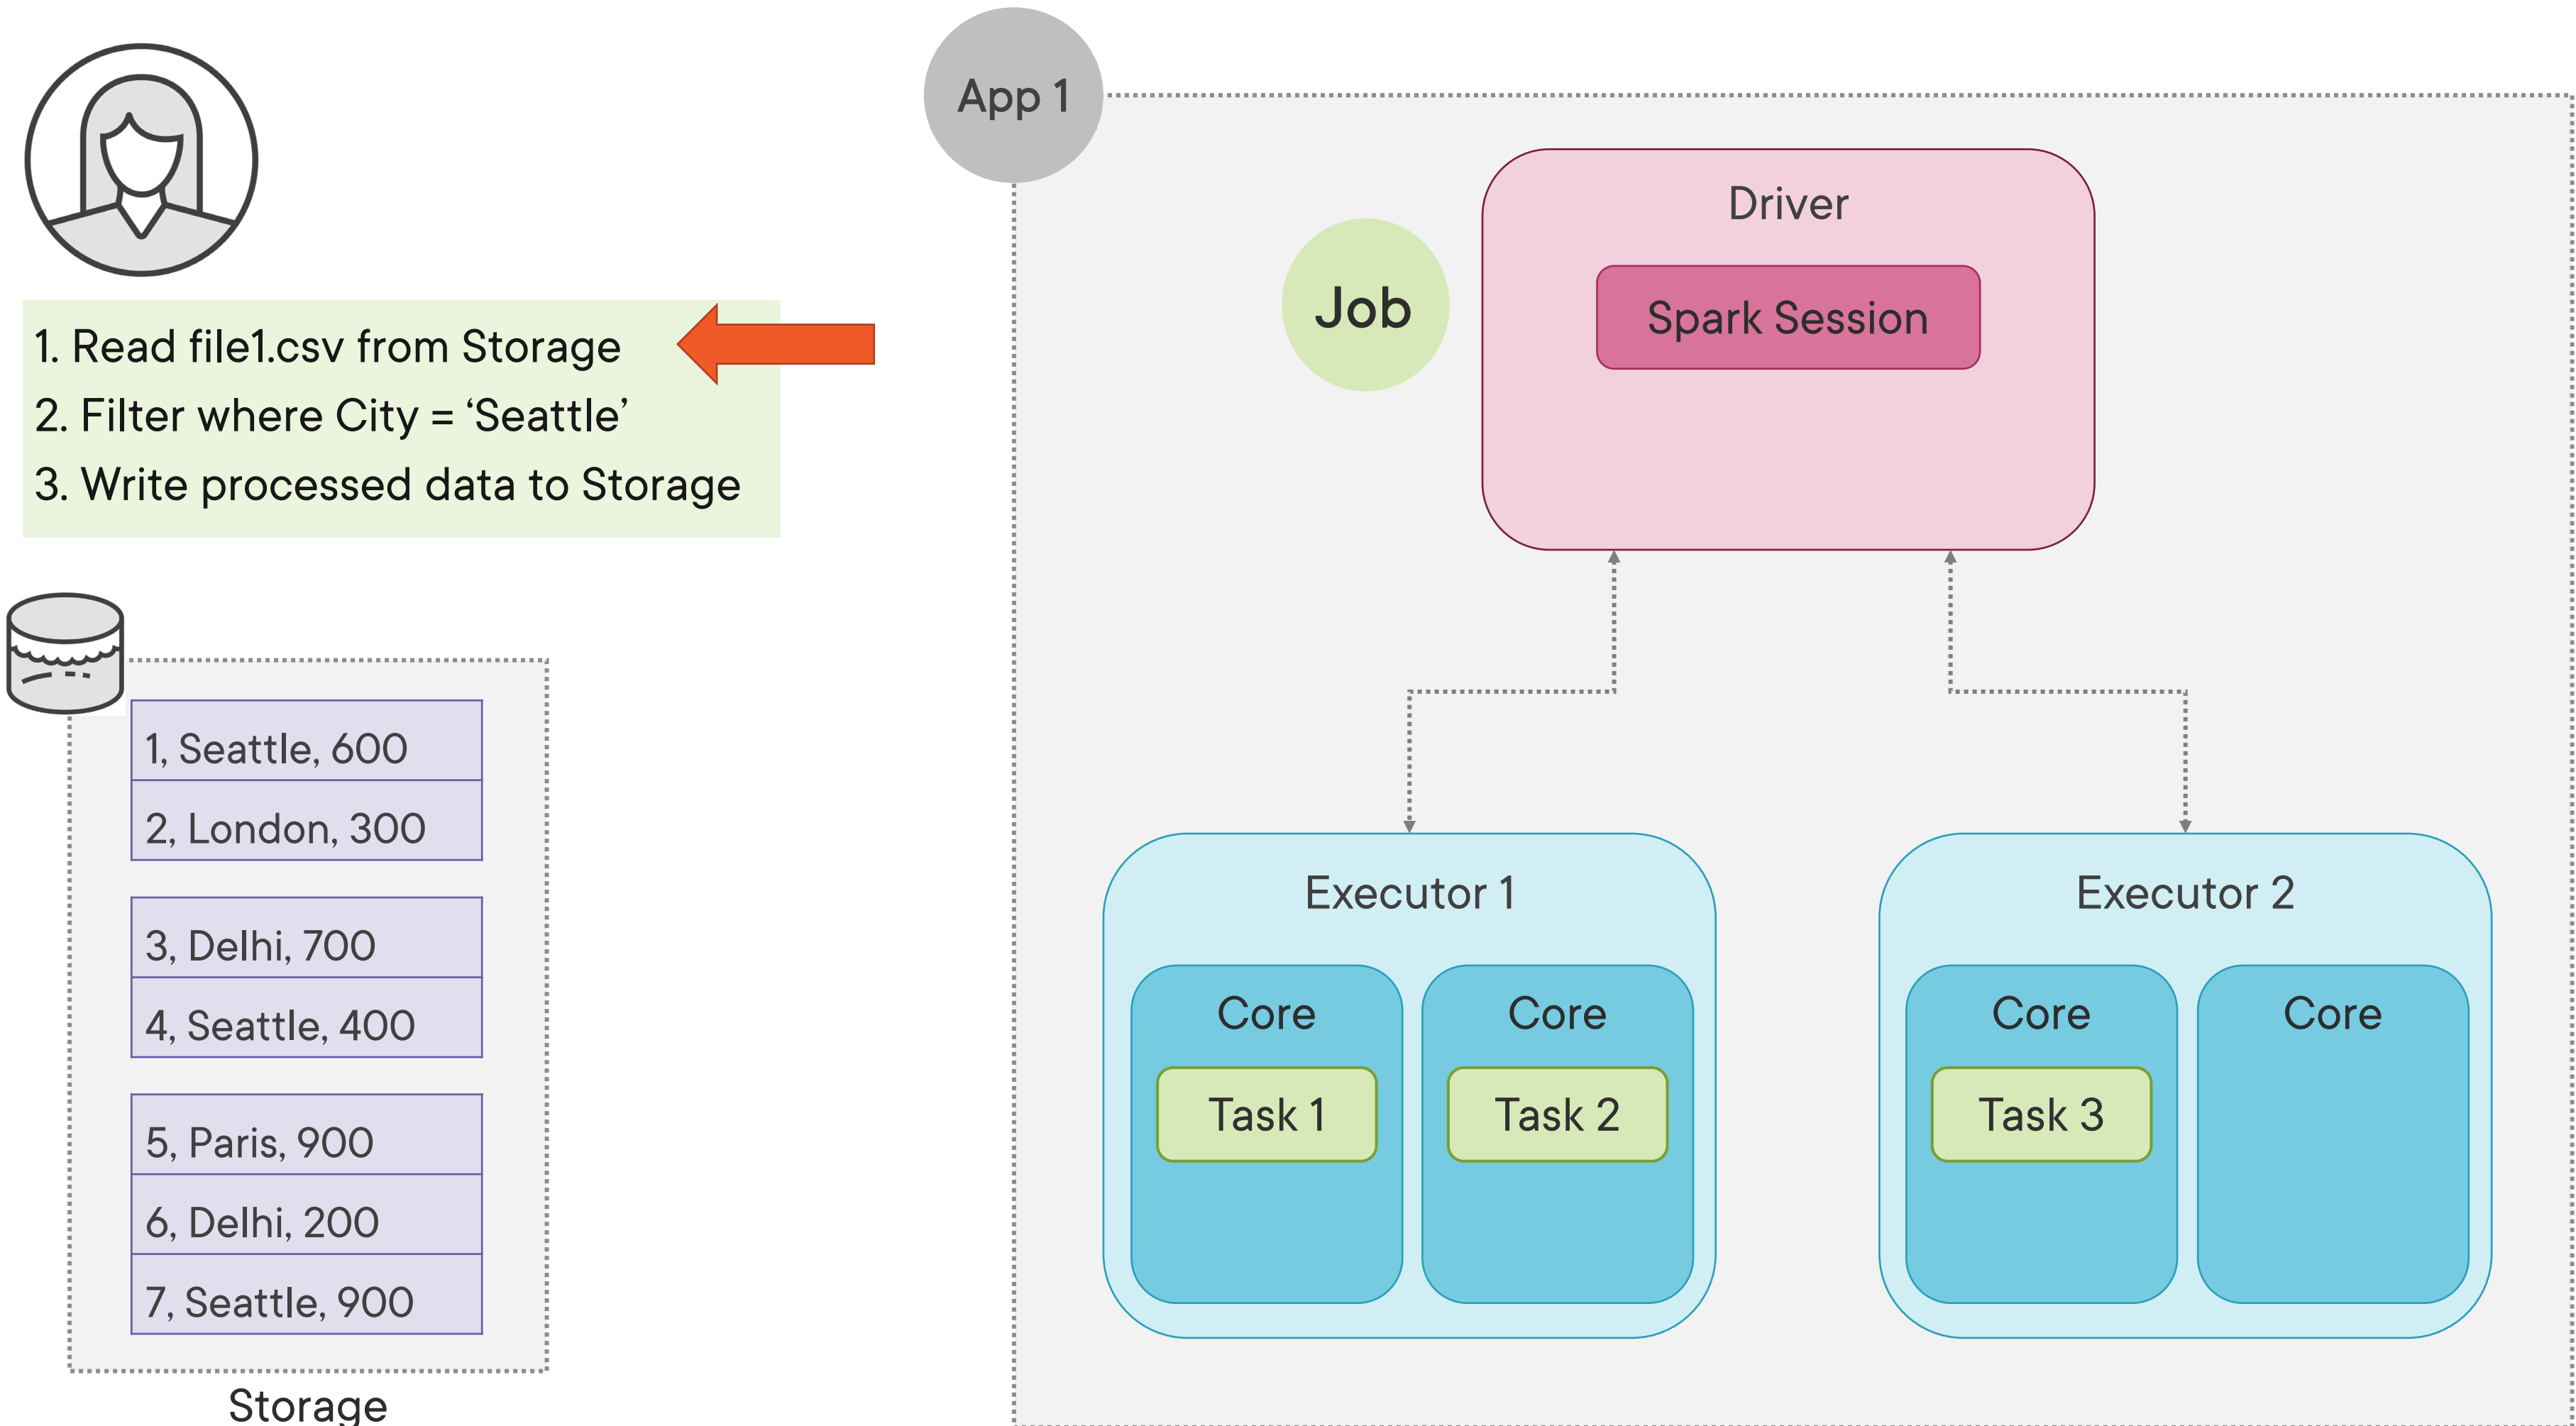
Task 2

Executor 2

Core



Task 3

Core





1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



Storage

App 1

Job

Driver

Spark Session

Executor 1

Core

Task 1

1, Seattle, 600

2, London, 300

Core

Task 2

3, Delhi, 700

4, Seattle, 400

Executor 2

Core

Task 3

5, Paris, 900


6, Delhi, 200

7, Seattle, 900

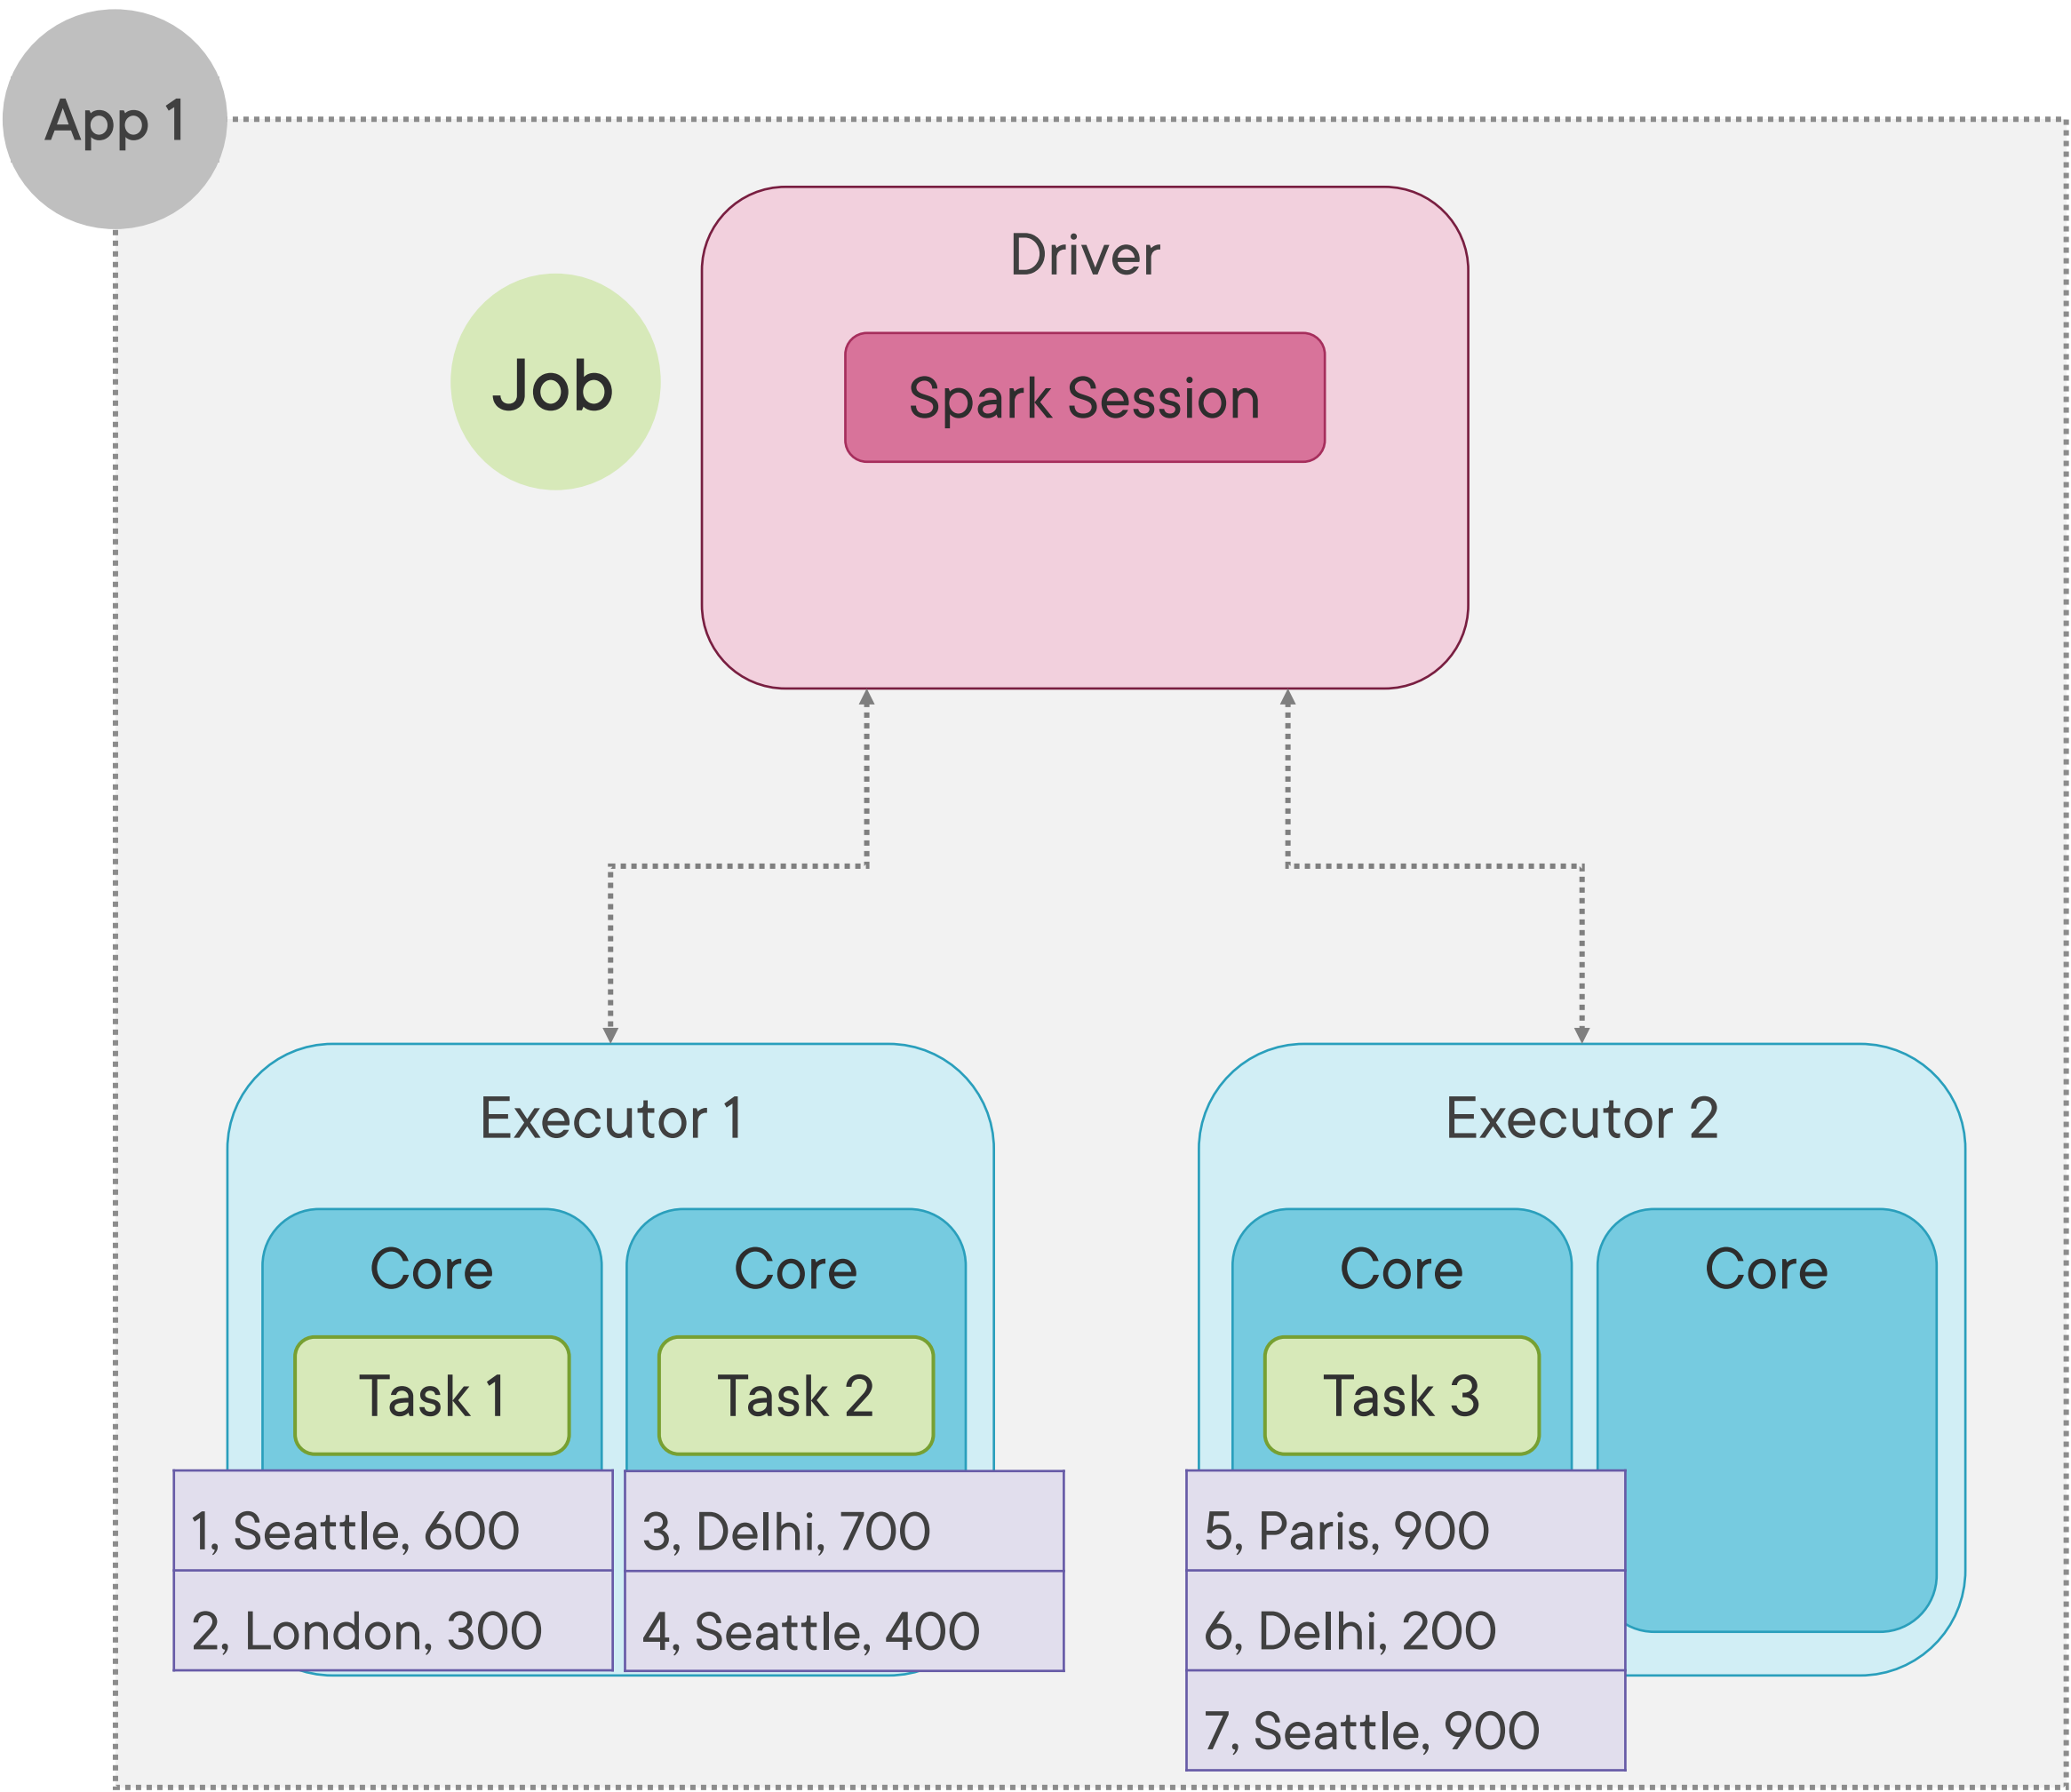
Core



- 1. Read file1.csv from Storage
- 2. Filter where City = 'Seattle'
- 3. Write processed data to Storage



Storage





1. Read file1.csv from Storage
2. Filter where City = 'Seattle' ←
3. Write processed data to Storage



Storage

App 1

Job

Driver

Spark Session

Executor 1

Core

Core

1, Seattle, 600

4, Seattle, 400

Executor 2

Core

Core

7, Seattle, 900



1. Read file1.csv from Storage
2. Filter where City = 'Seattle'
3. Write processed data to Storage



1, Seattle, 600

4, Seattle, 400

7, Seattle, 900

Storage

App 1

Job

Driver

Spark Session

Executor 1

Core

Core

Executor 2

Core

Core

# Features of RDDs

## **In-memory**

Resides in the memory of cluster

## **Partitioned**

Split into partitions, processed by tasks



# Read File

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Task 1


Task 2

Task 3

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

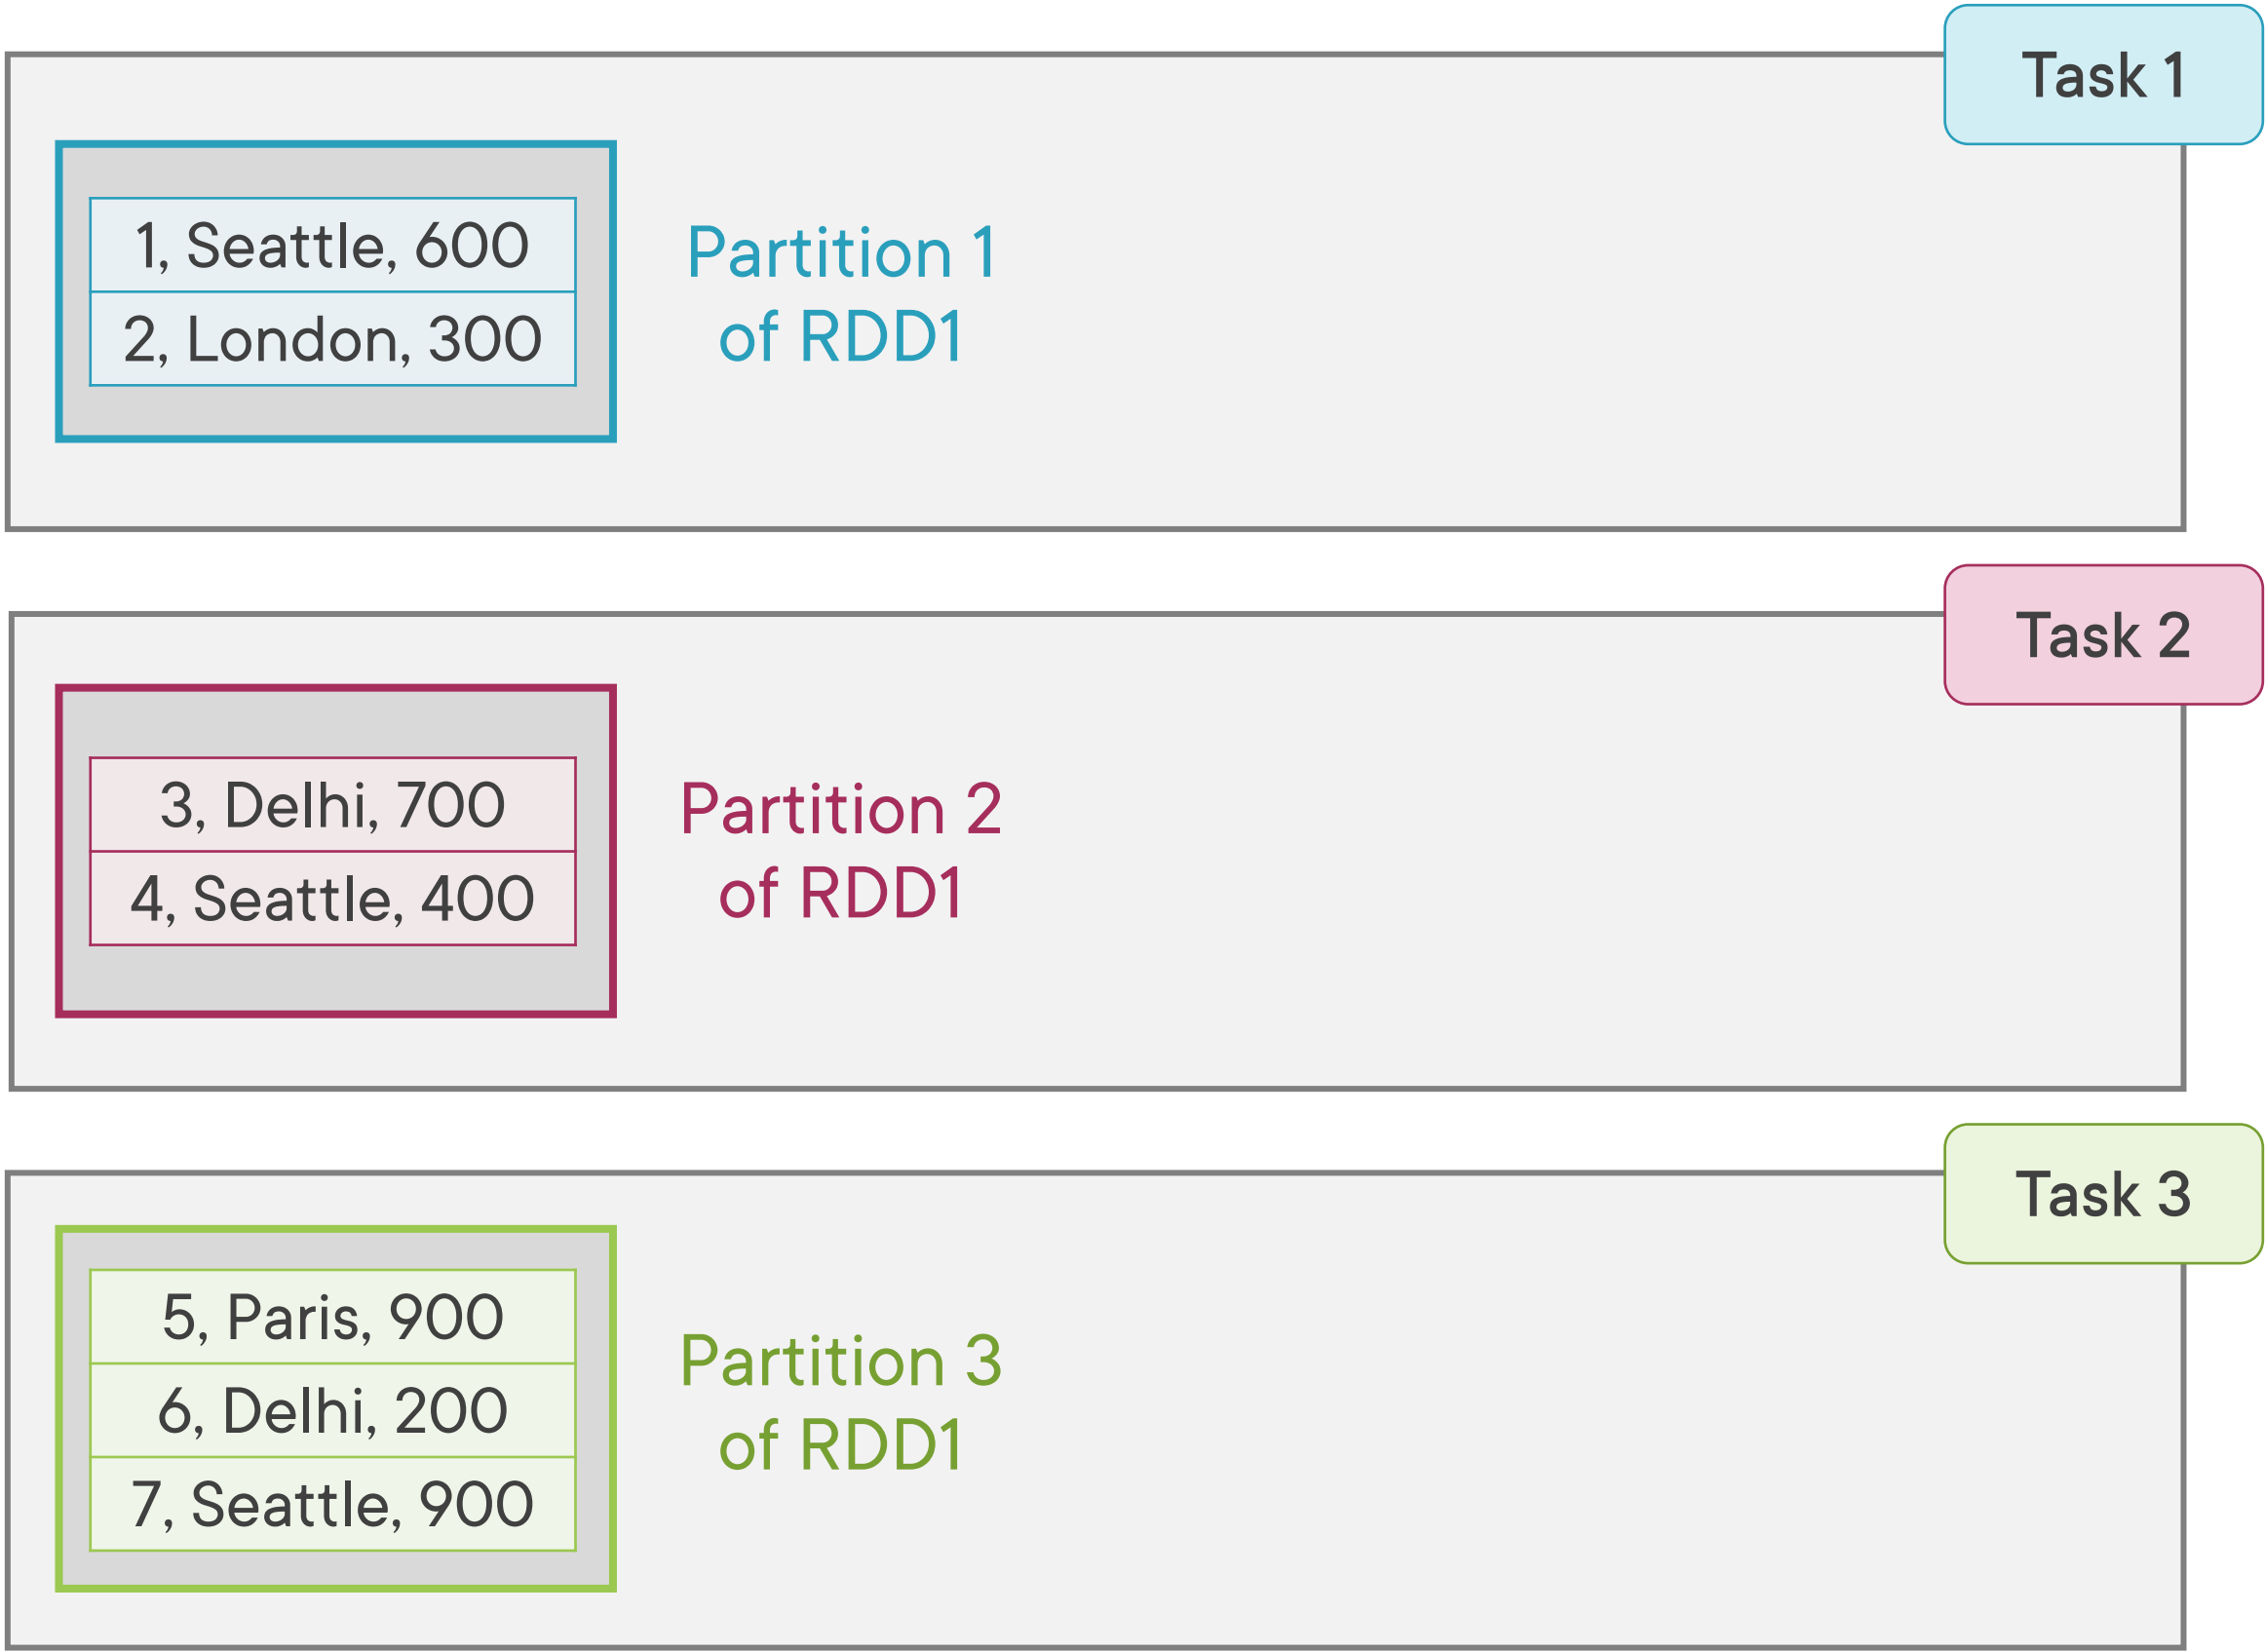
Read File

RDD 1



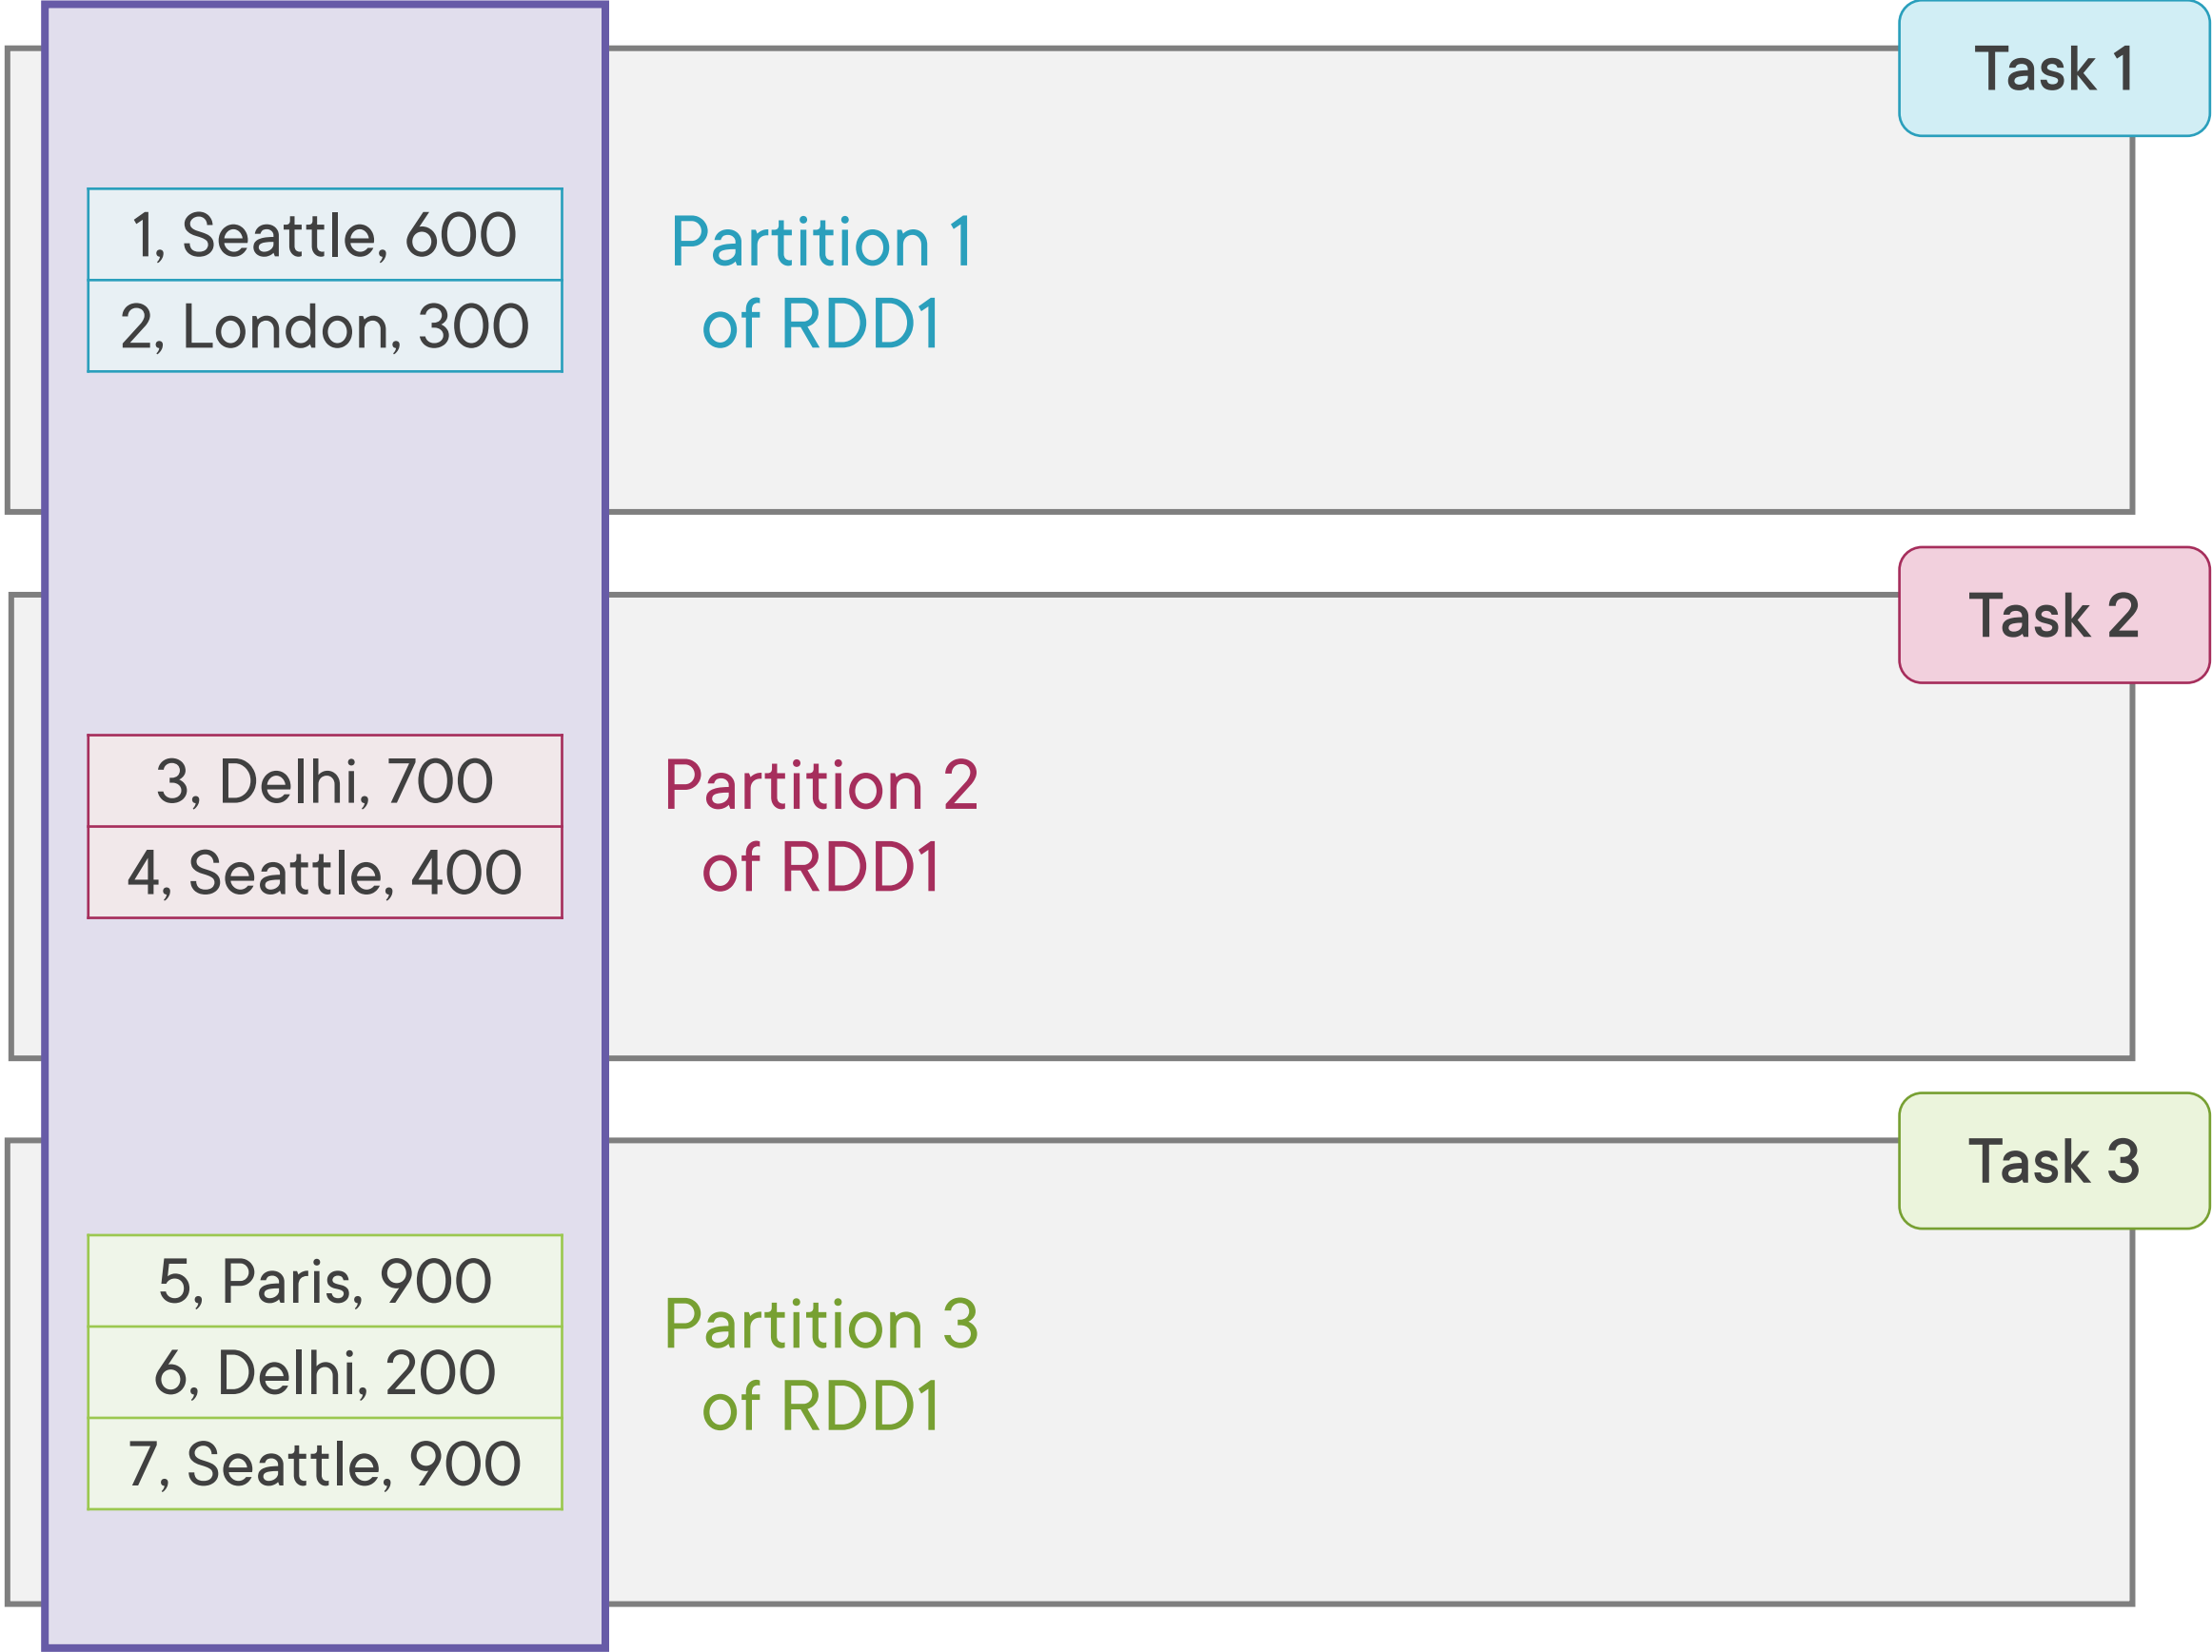
<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read File



<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

**Read File**  
**RDD 1**



# Features of RDDs

## **In-memory**

Resides in the memory of cluster

## **Partitioned**

Split into partitions, processed by tasks

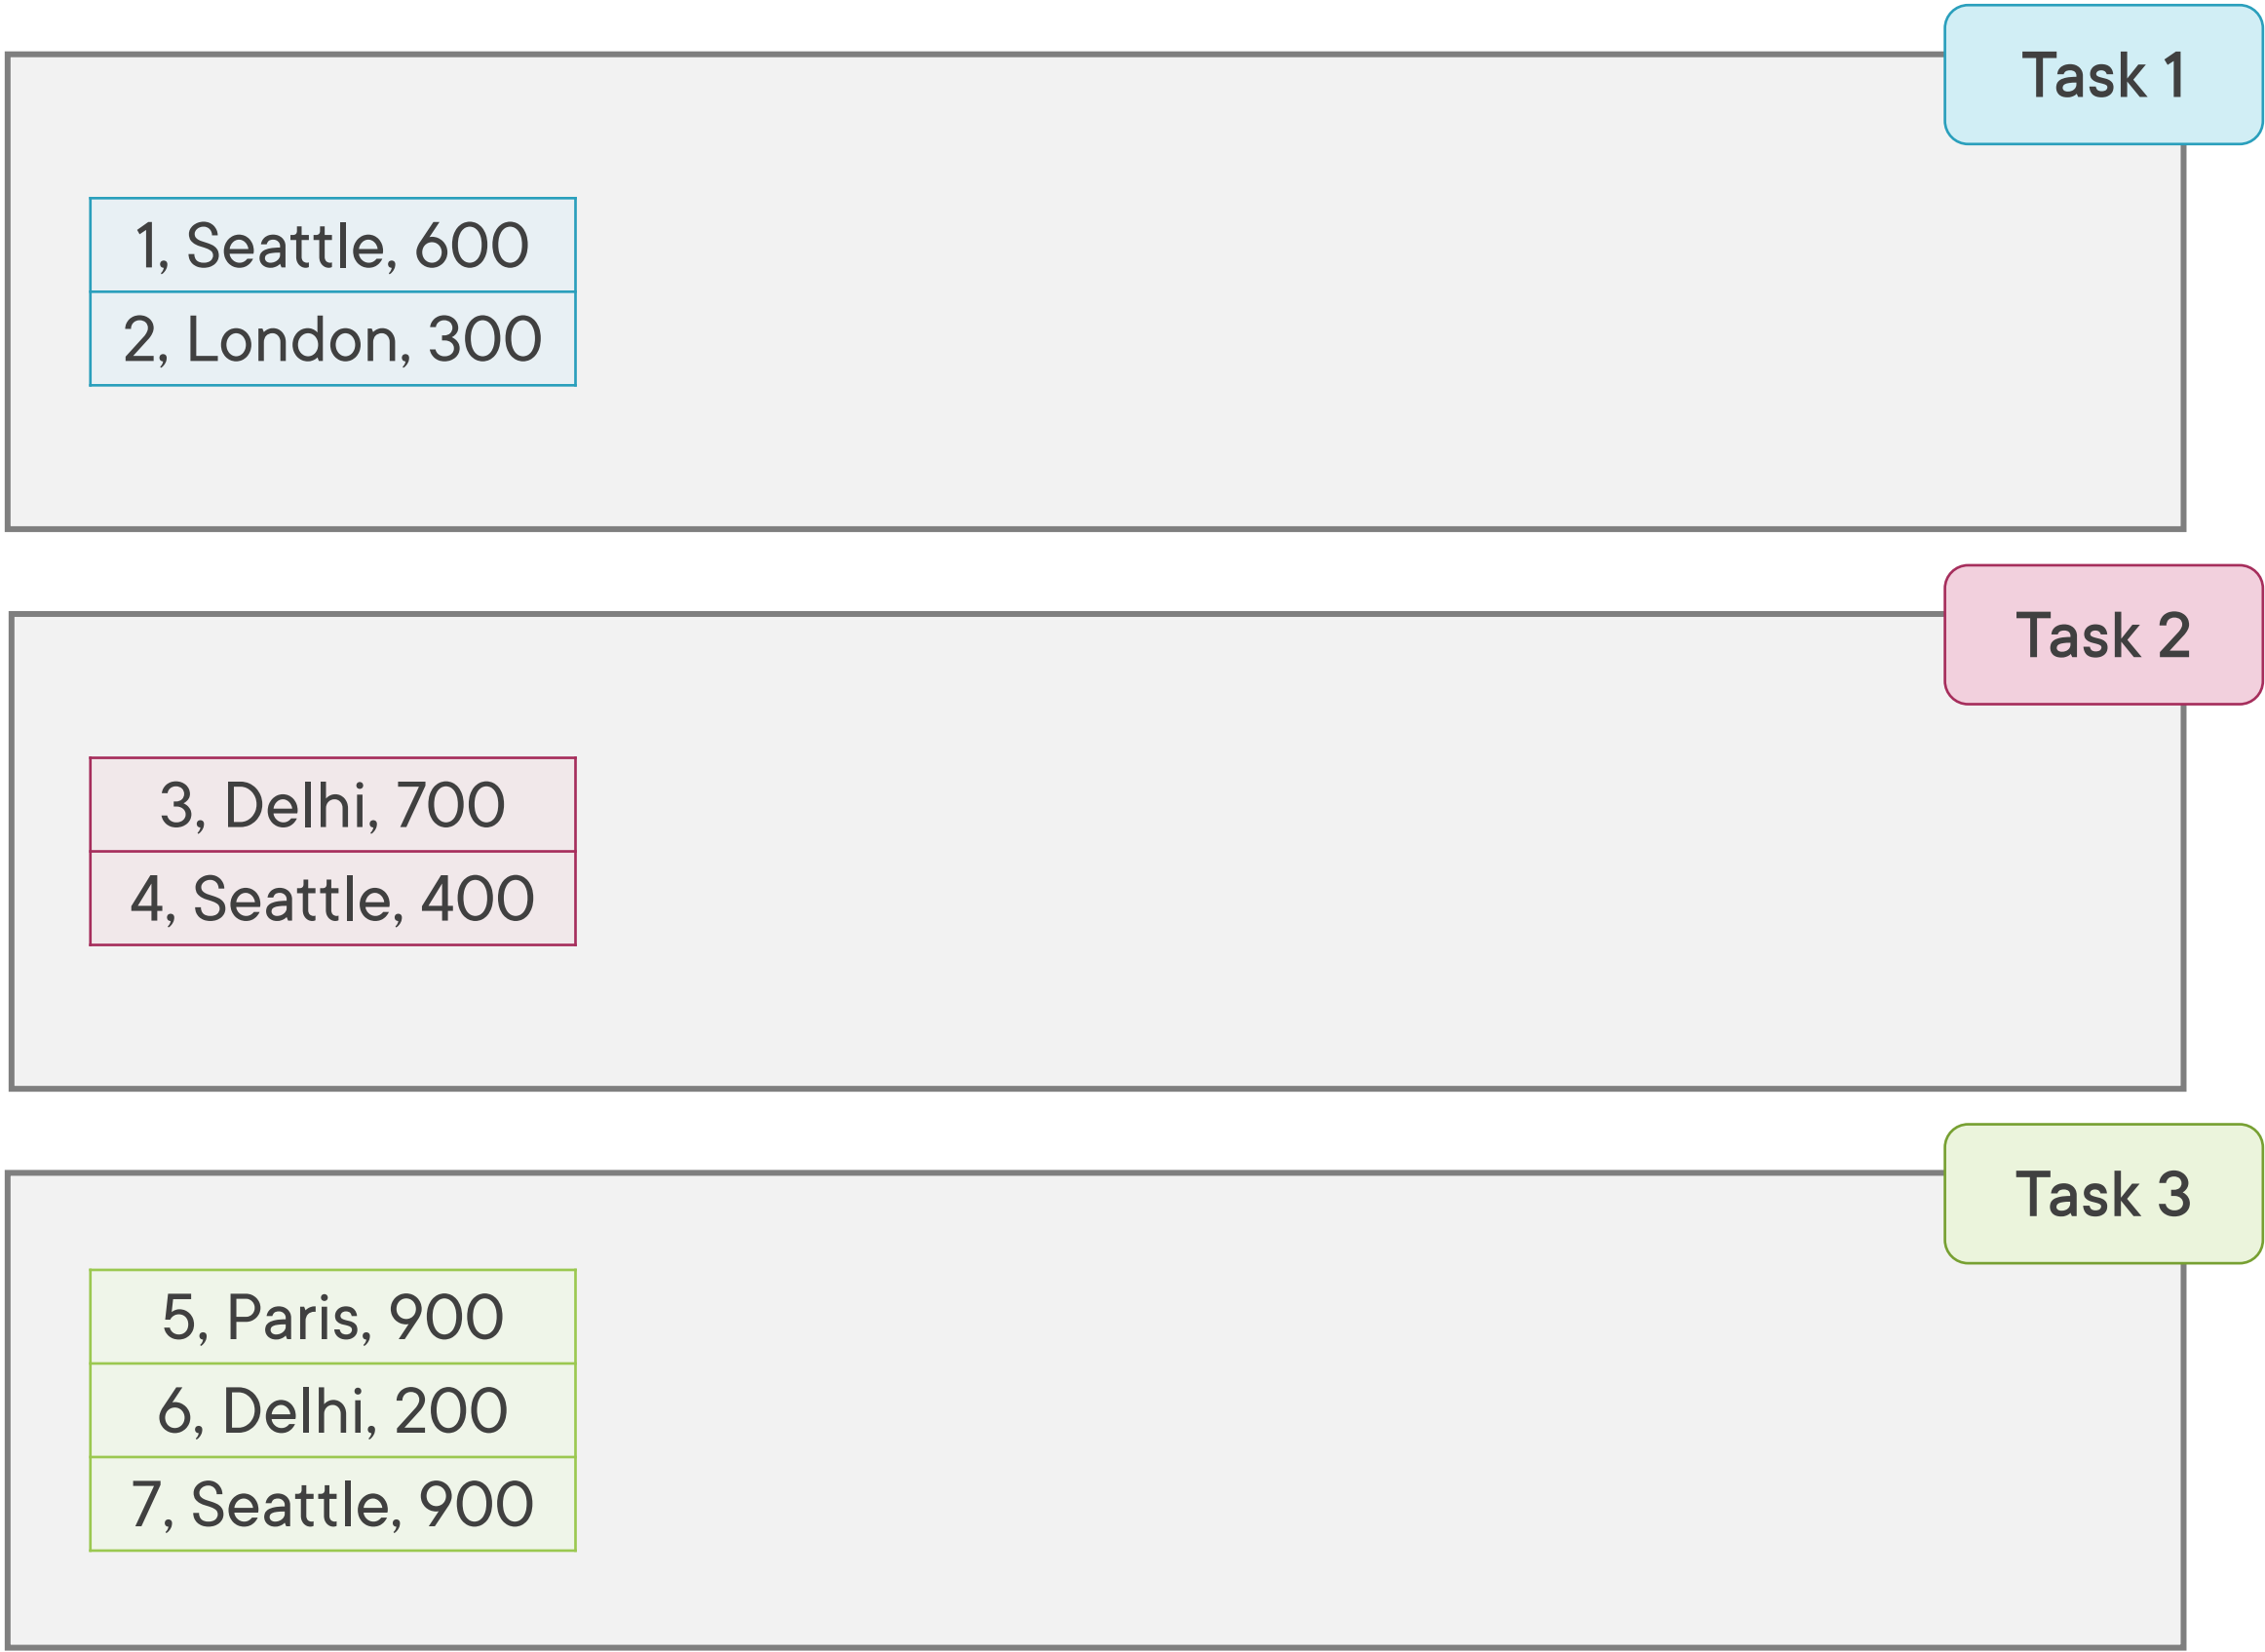
## **Read-only**

Transformed into another RDD or result

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read File      .....→      Split by Comma

**RDD 1**



<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

**Read File**

**RDD 1**

**Split by Comma**

**RDD 2**

**Task 1**

1, Seattle, 600
2, London, 300



1	Seattle	600
2	London	300

**Task 2**

3, Delhi, 700
4, Seattle, 400



3	Delhi	700
4	Seattle	400

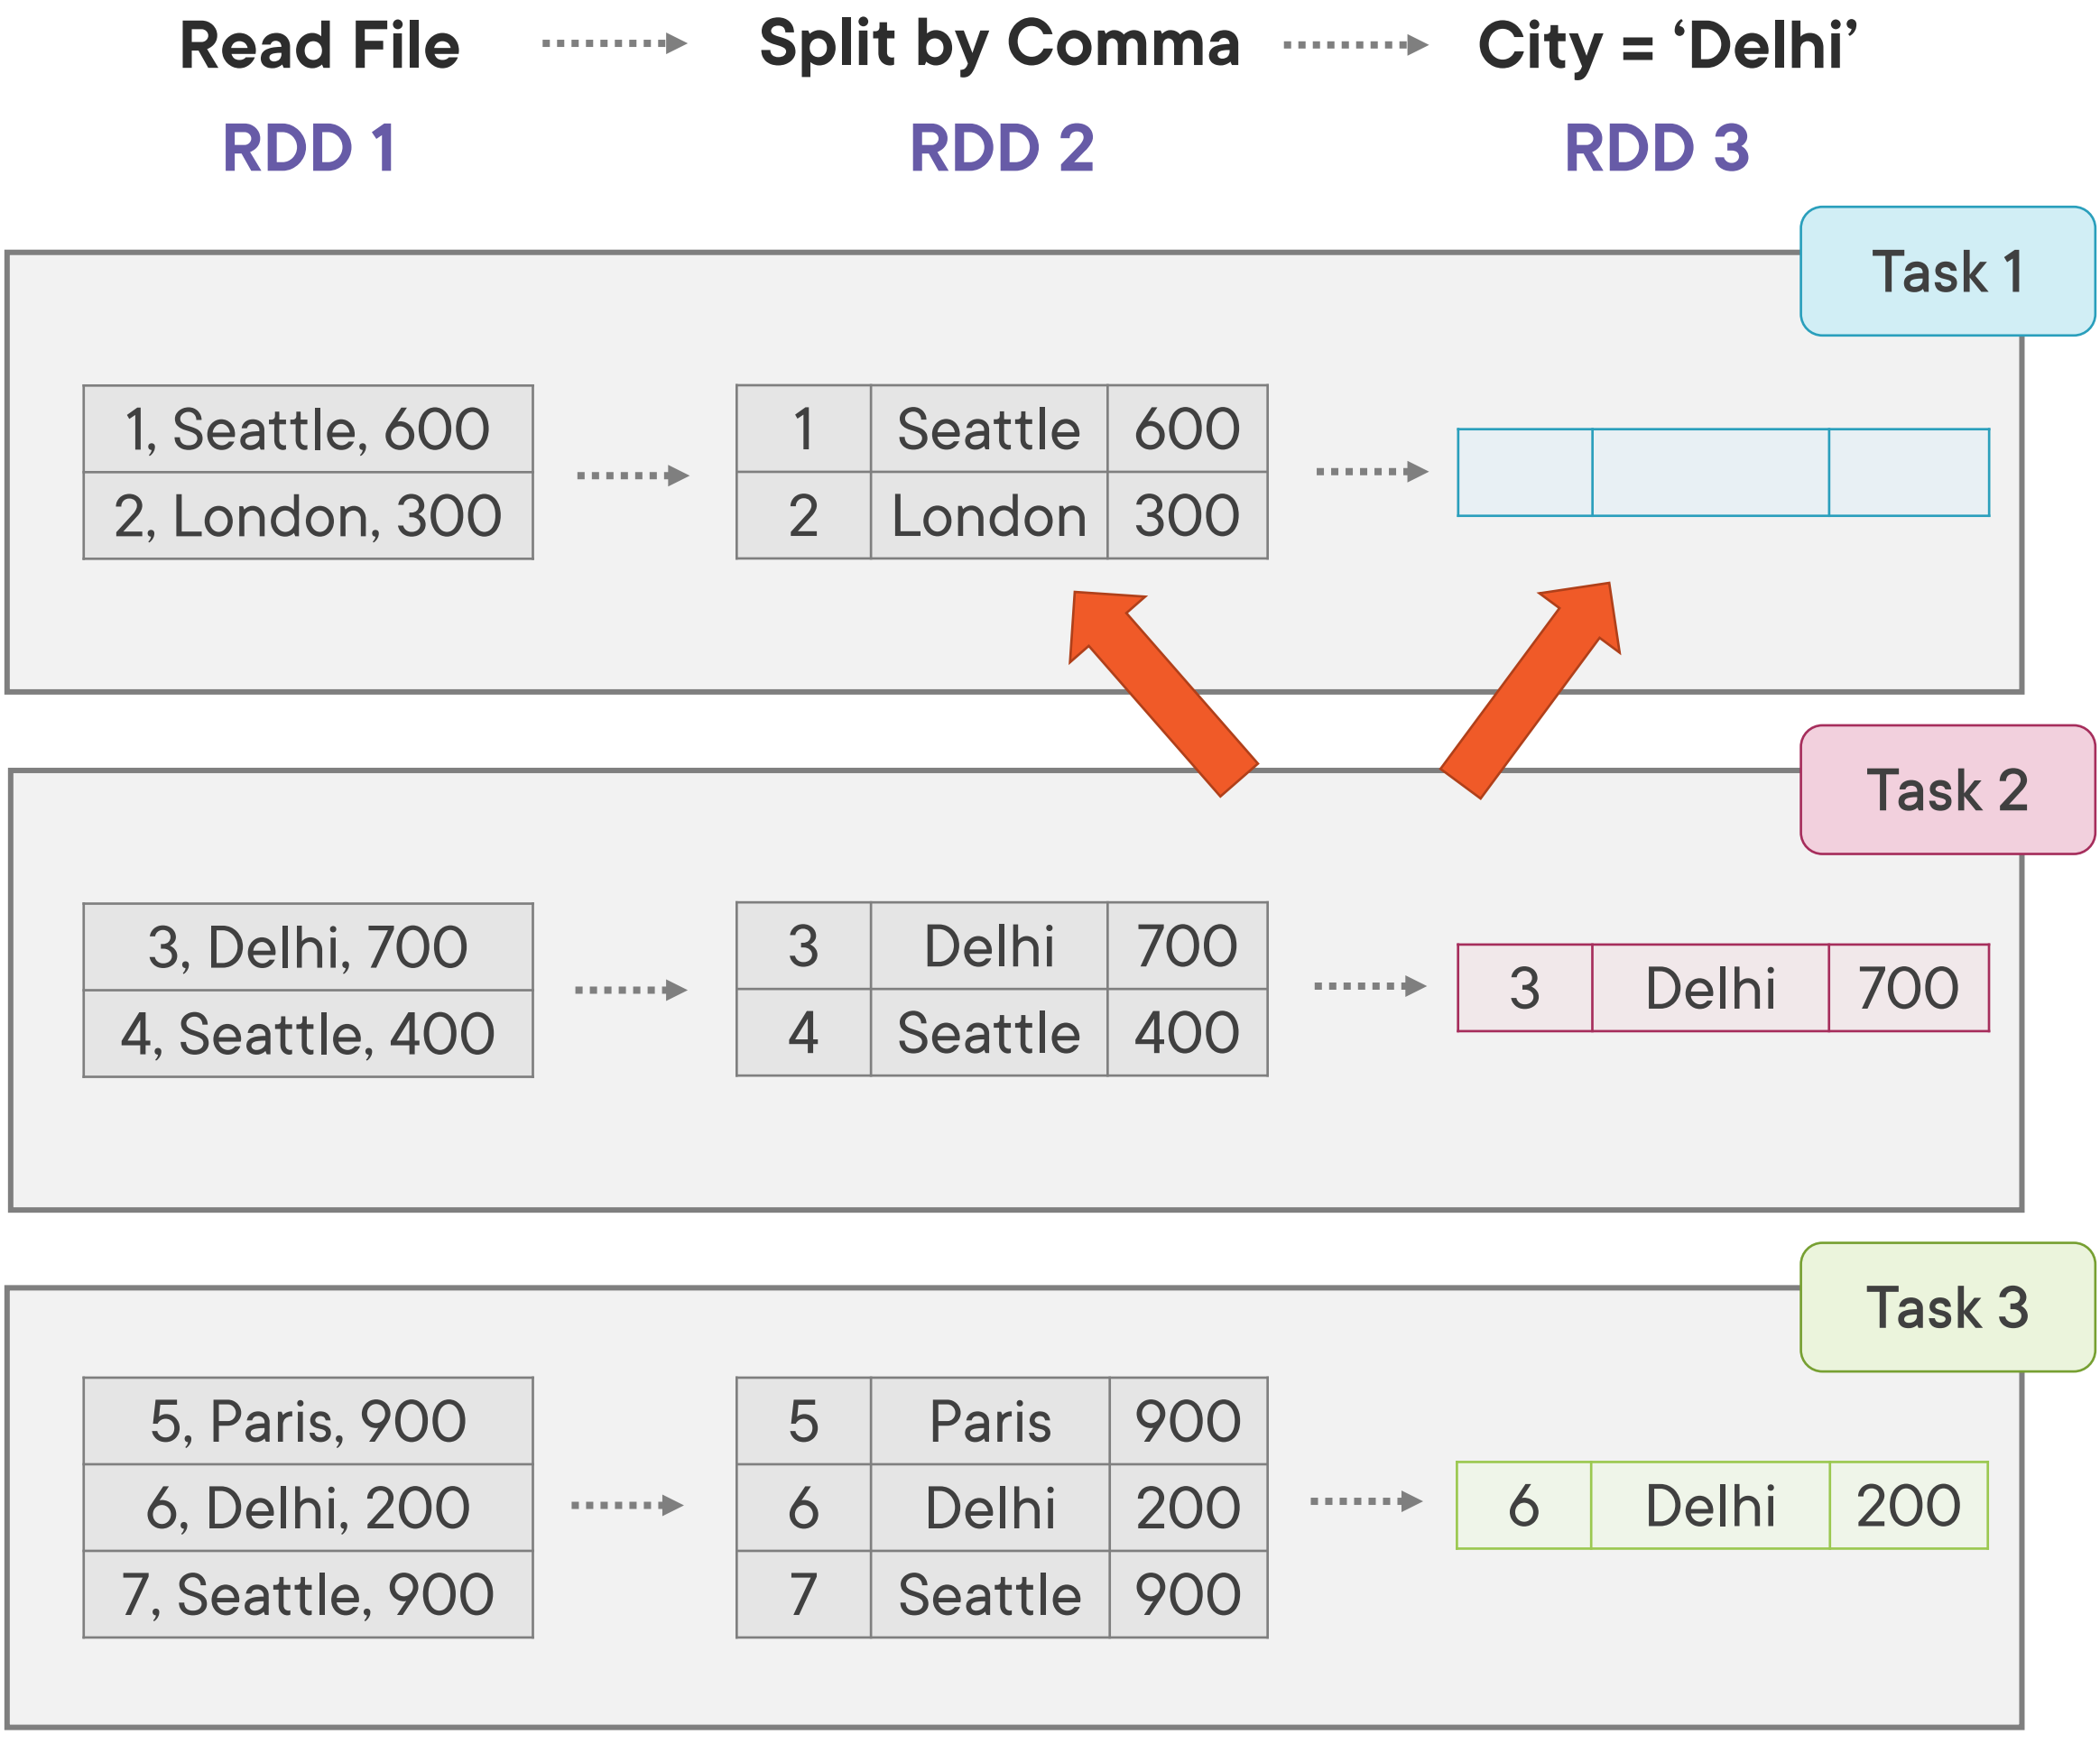
**Task 3**

5, Paris, 900
6, Delhi, 200
7, Seattle, 900



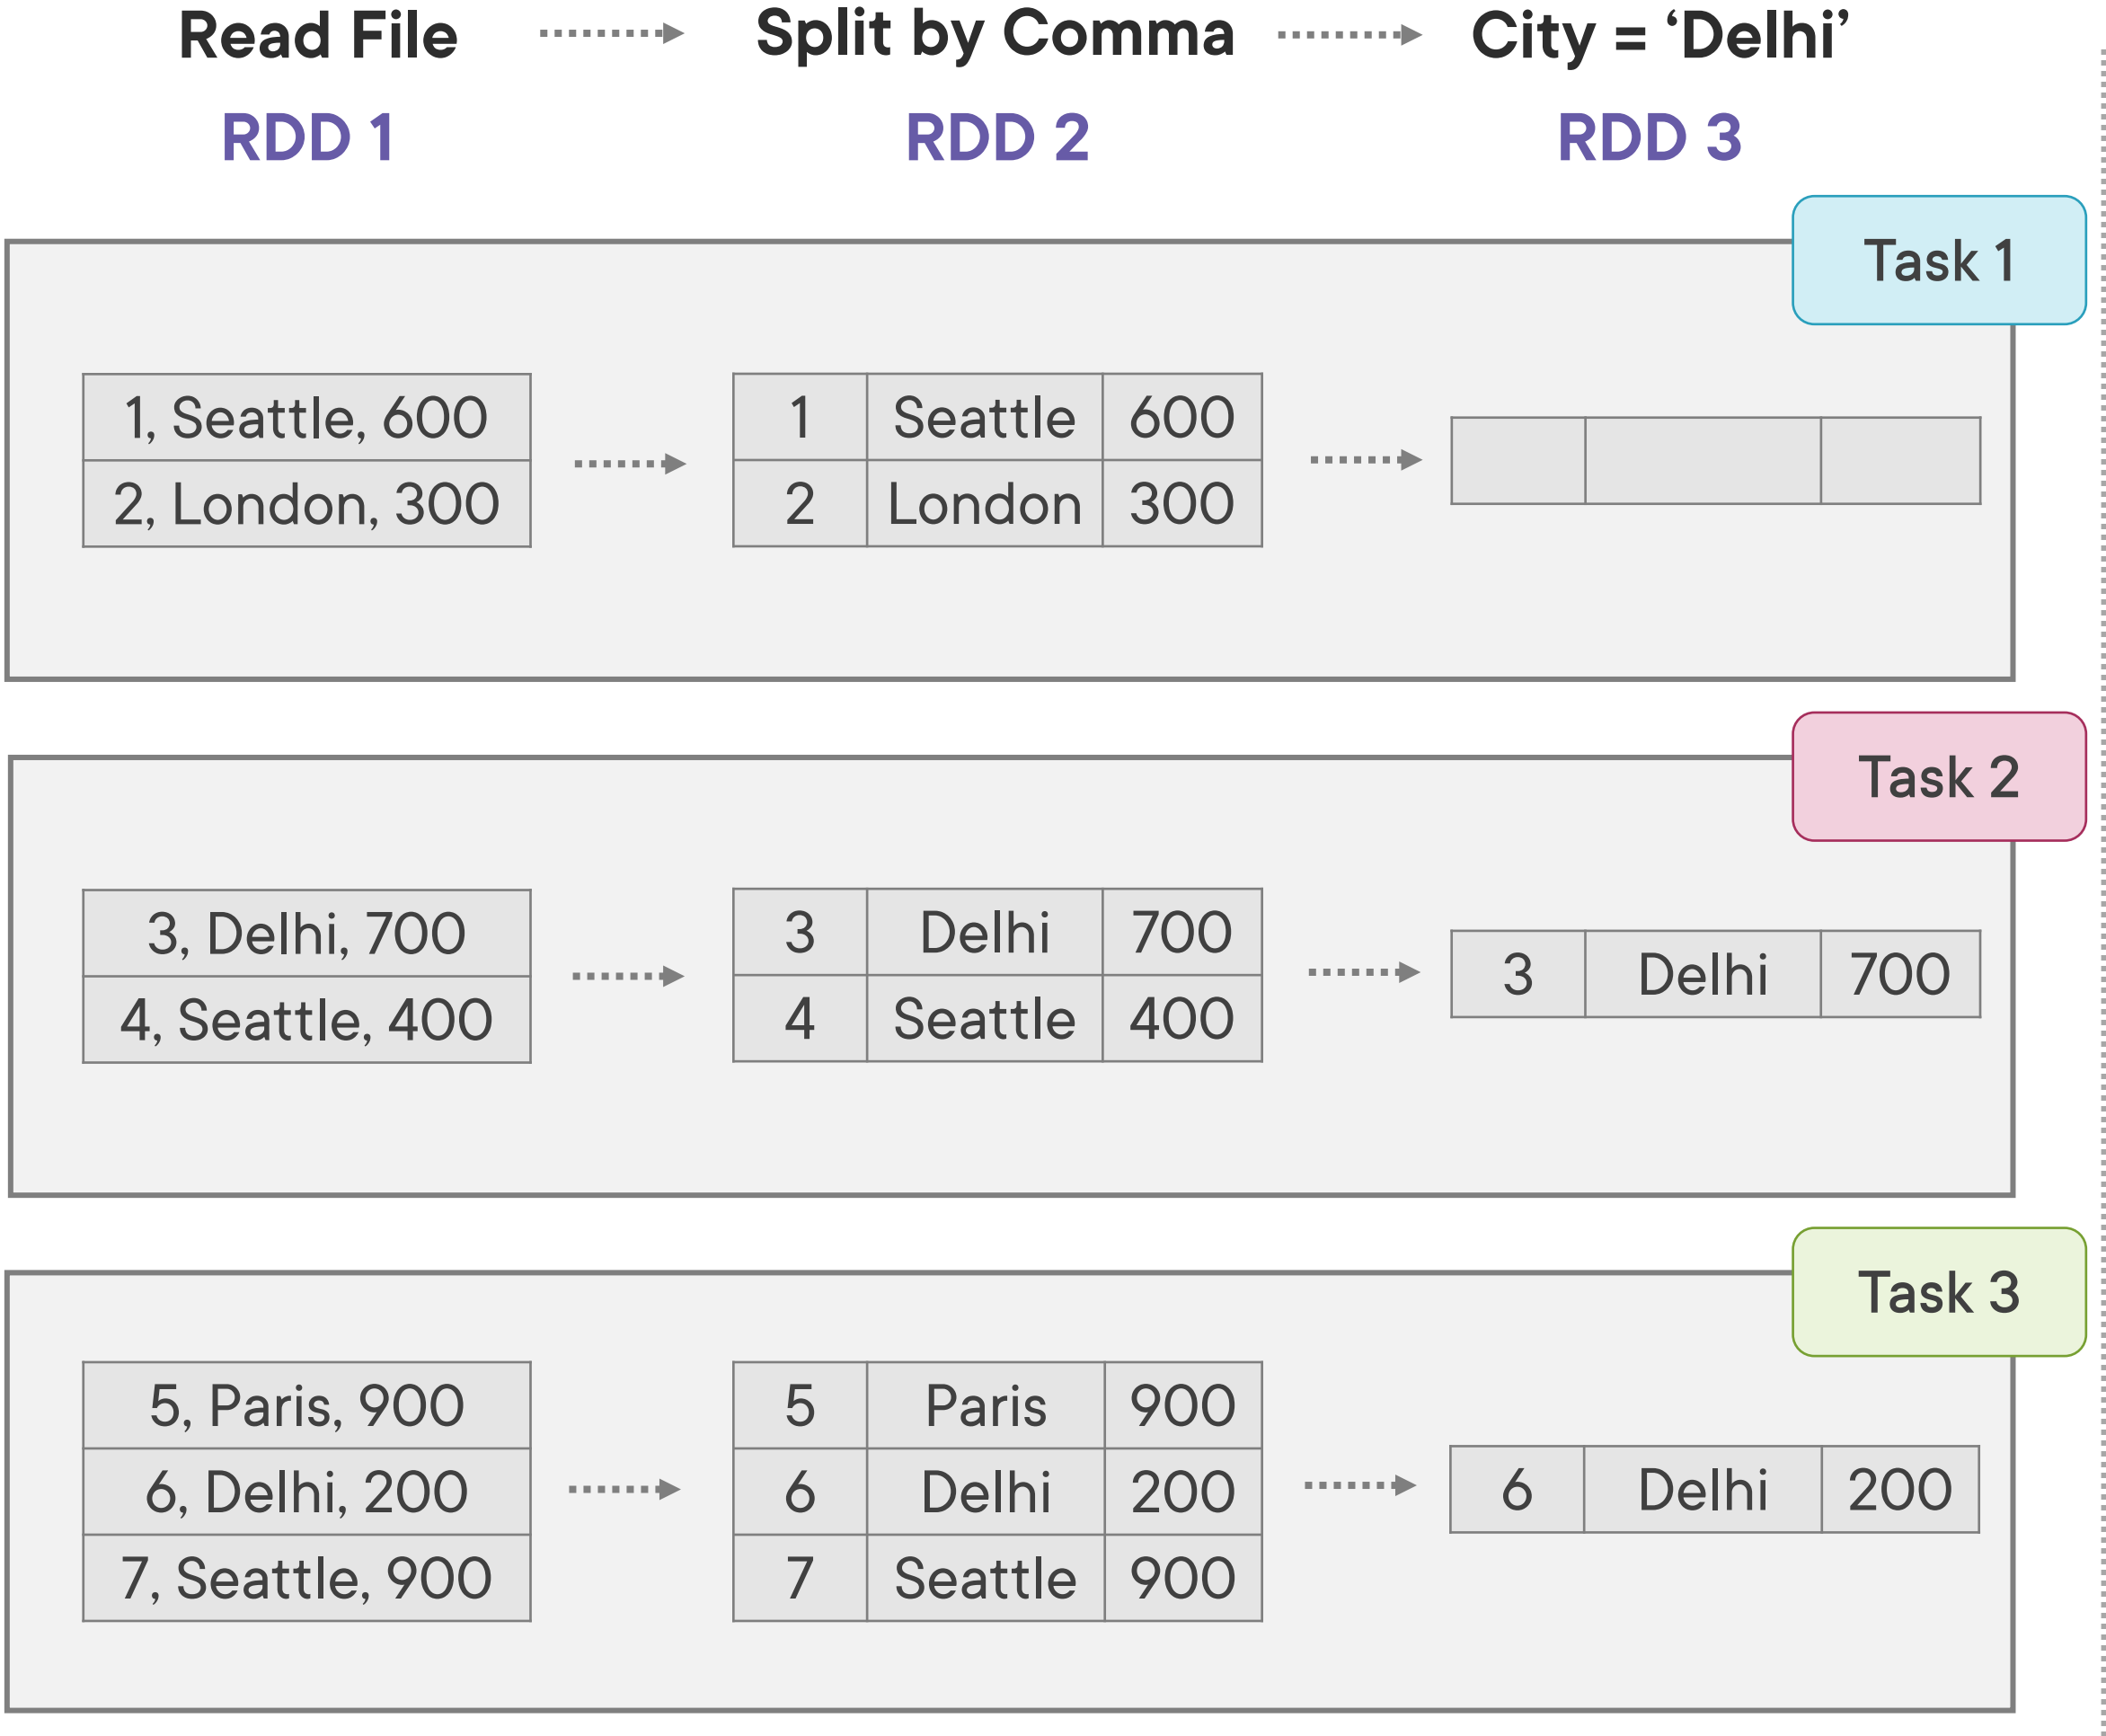
5	Paris	900
6	Delhi	200
7	Seattle	900

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900





<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900



3	Delhi	700
6	Delhi	200

# Features of RDDs

## **In-memory**

Resides in the memory of cluster

## **Partitioned**

Split into partitions, processed by tasks

## **Read-only**


Transformed into another RDD or result

## **Resilient**

Auto recover in case of a failure

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

1. Read file1.csv from Storage
2. Split data by comma
3. Filter where City='Delhi'
4. Write processed data to storage

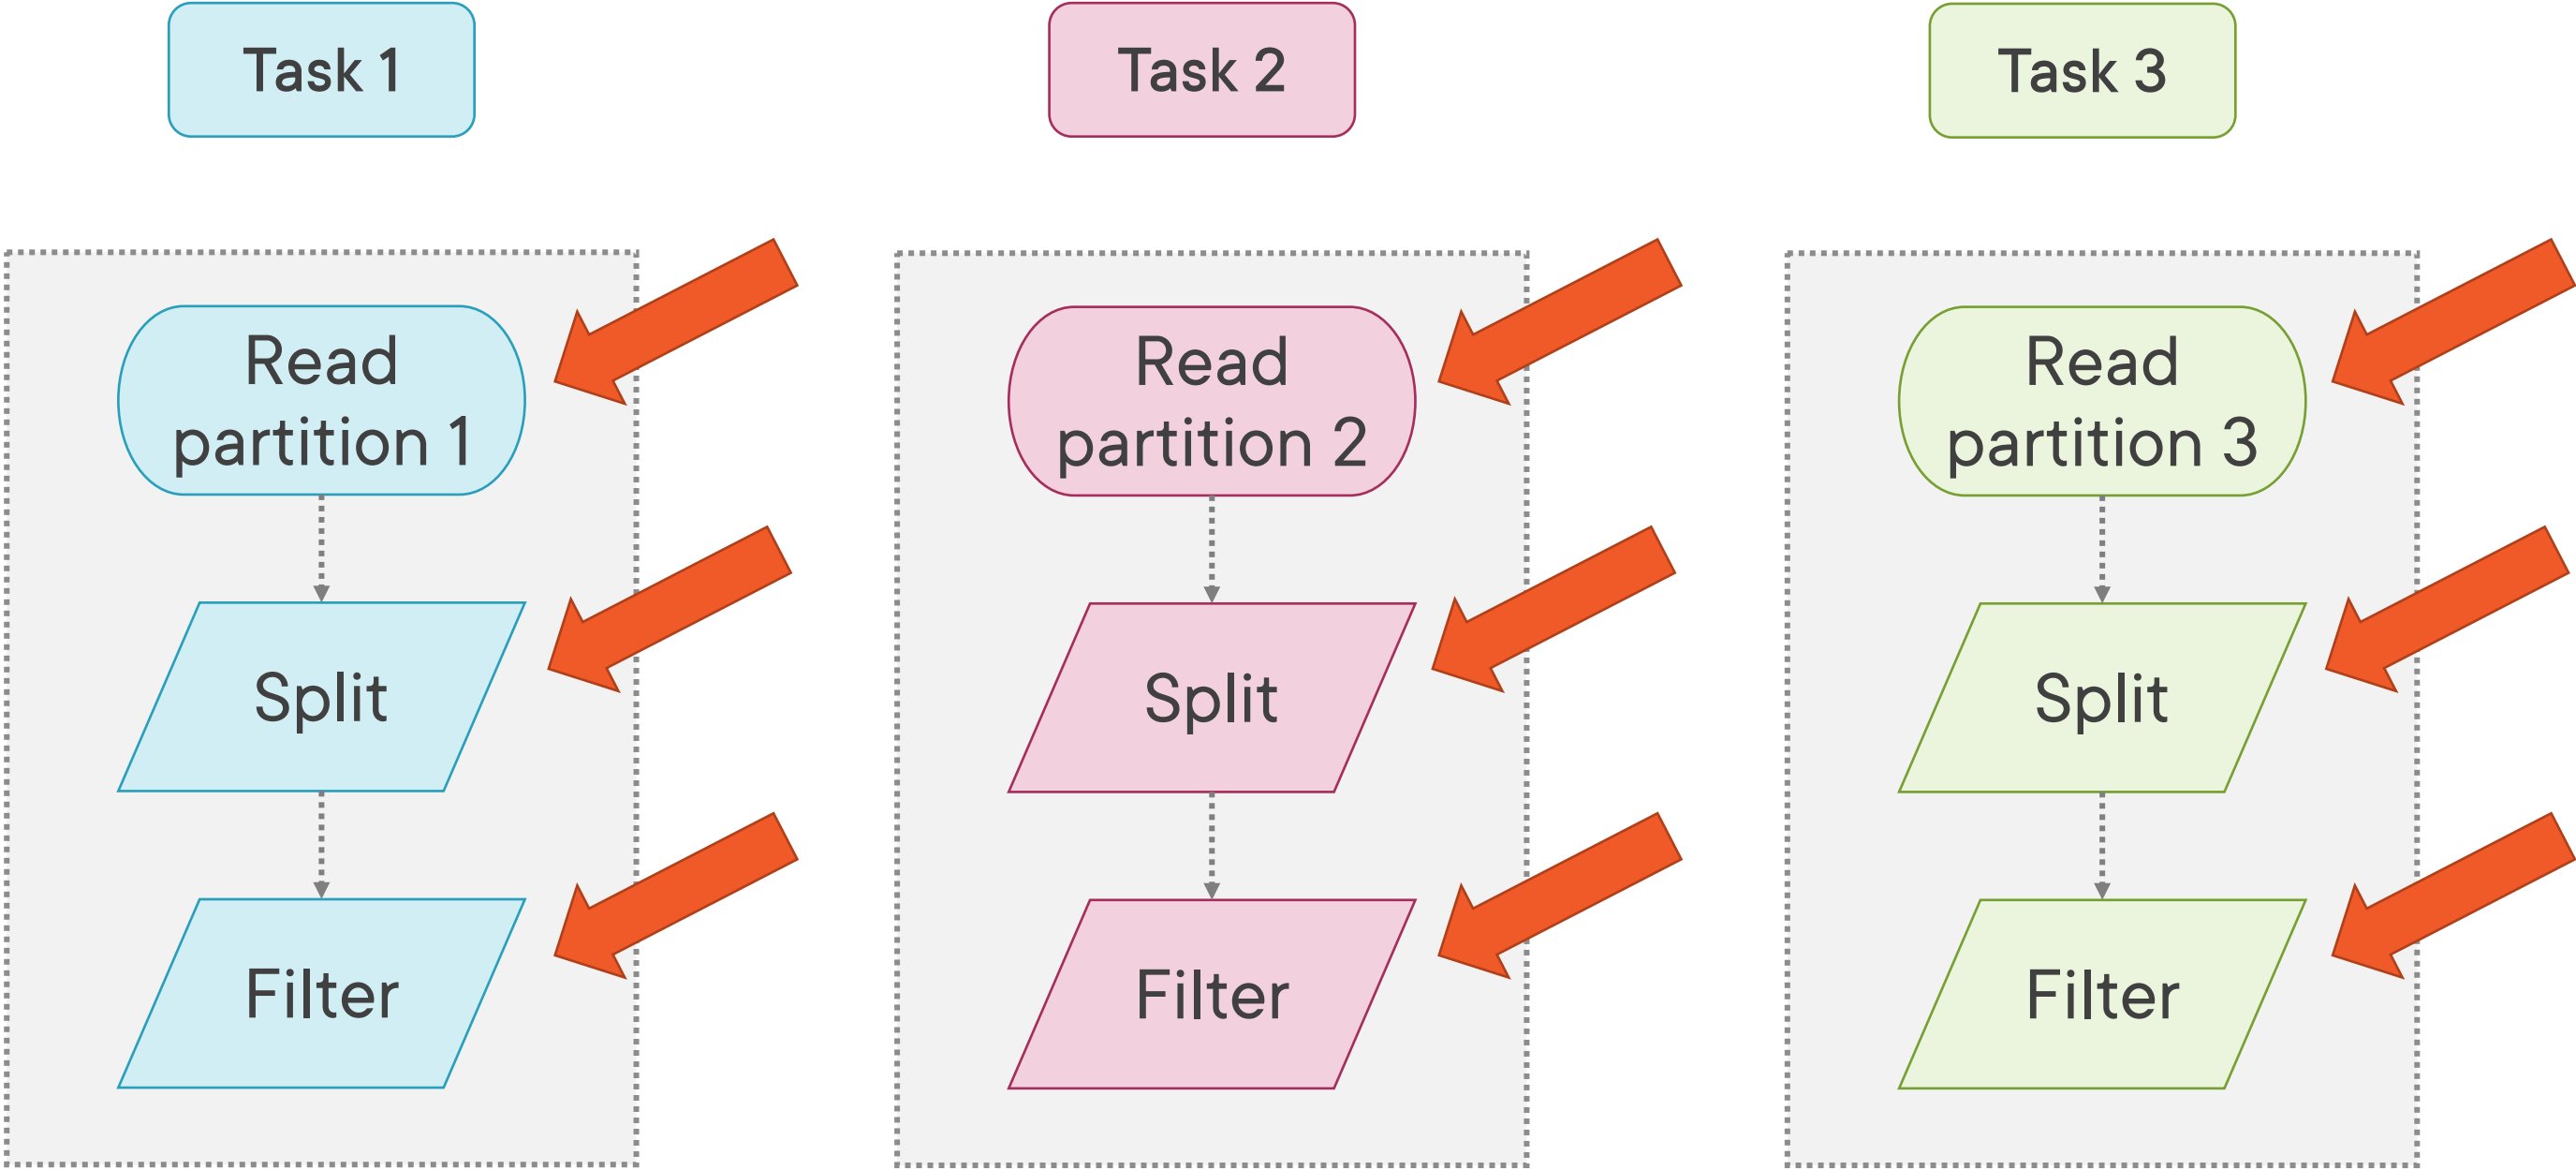


## Lineage Graph

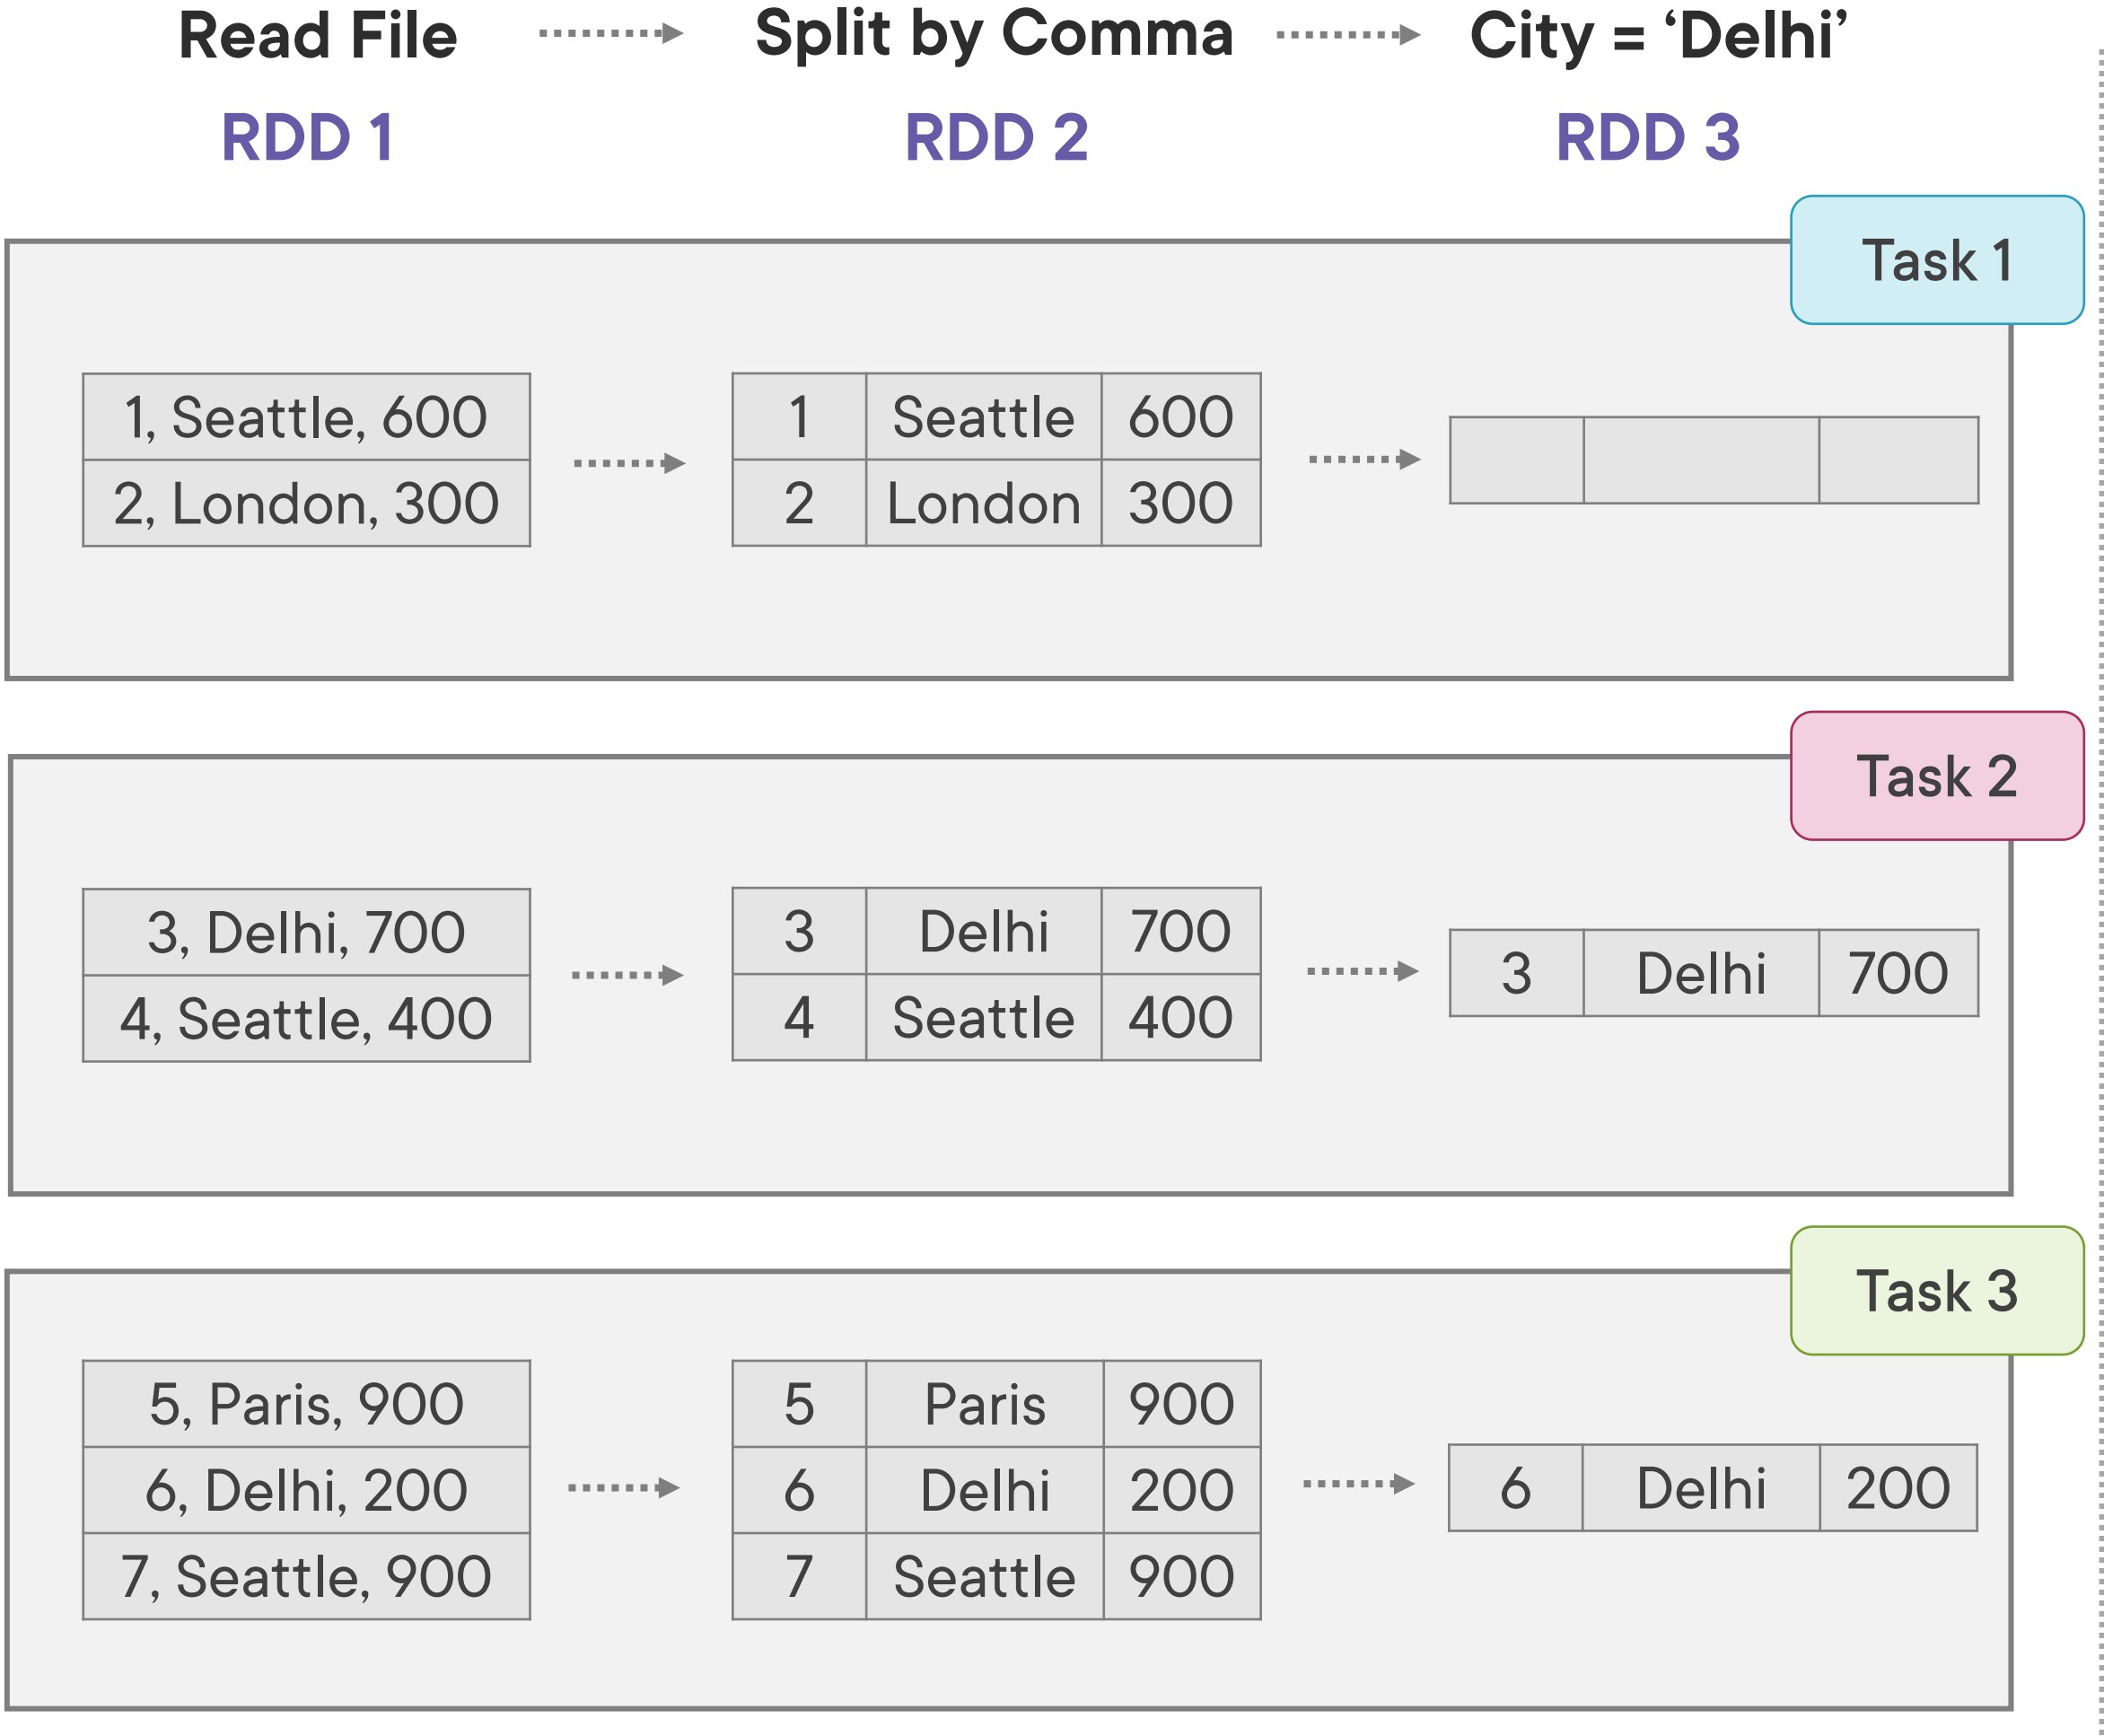
Defines steps that are used for creation of an RDD

Id, City, Amount		
1	Seattle	600
2	London	300
3	Delhi	700
4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900

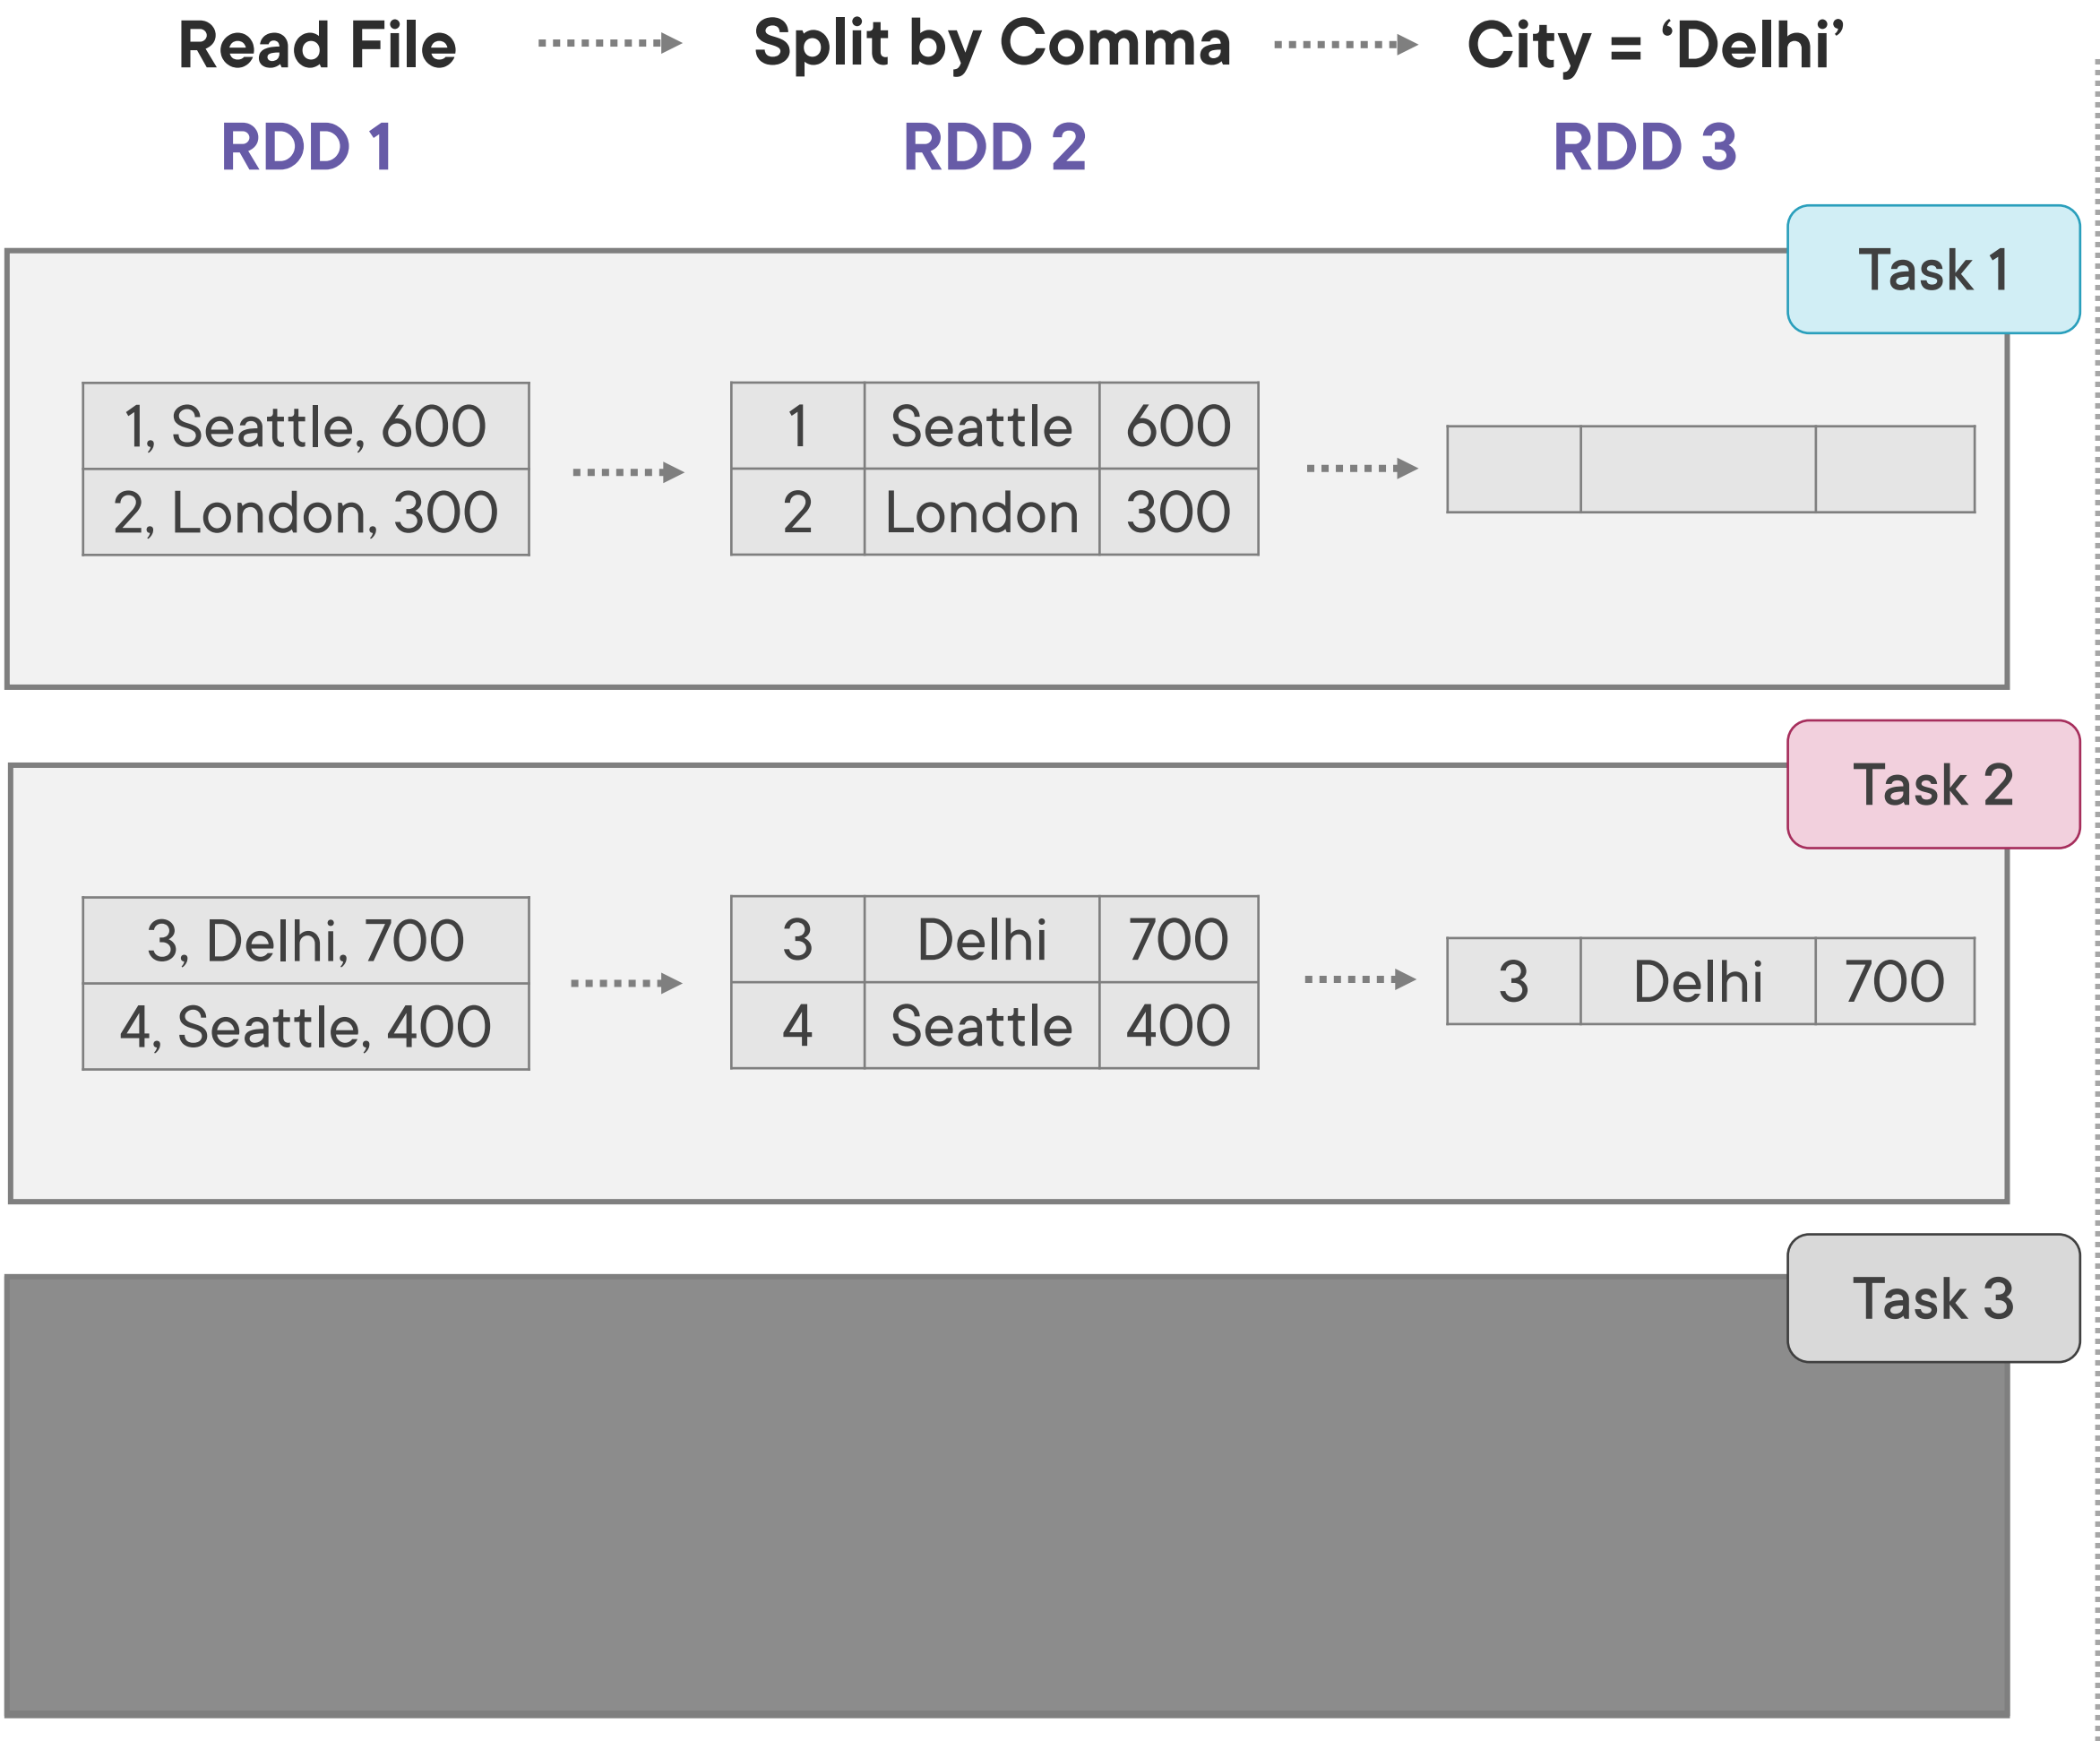
- 1. Read file1.csv from Storage
- 2. Split data by comma
- 3. Filter where City='Delhi'
- 4. Write processed data to storage



<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900



<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900




3	Delhi	700

Spark will **re-execute** failed Tasks  
using their Lineage Graph

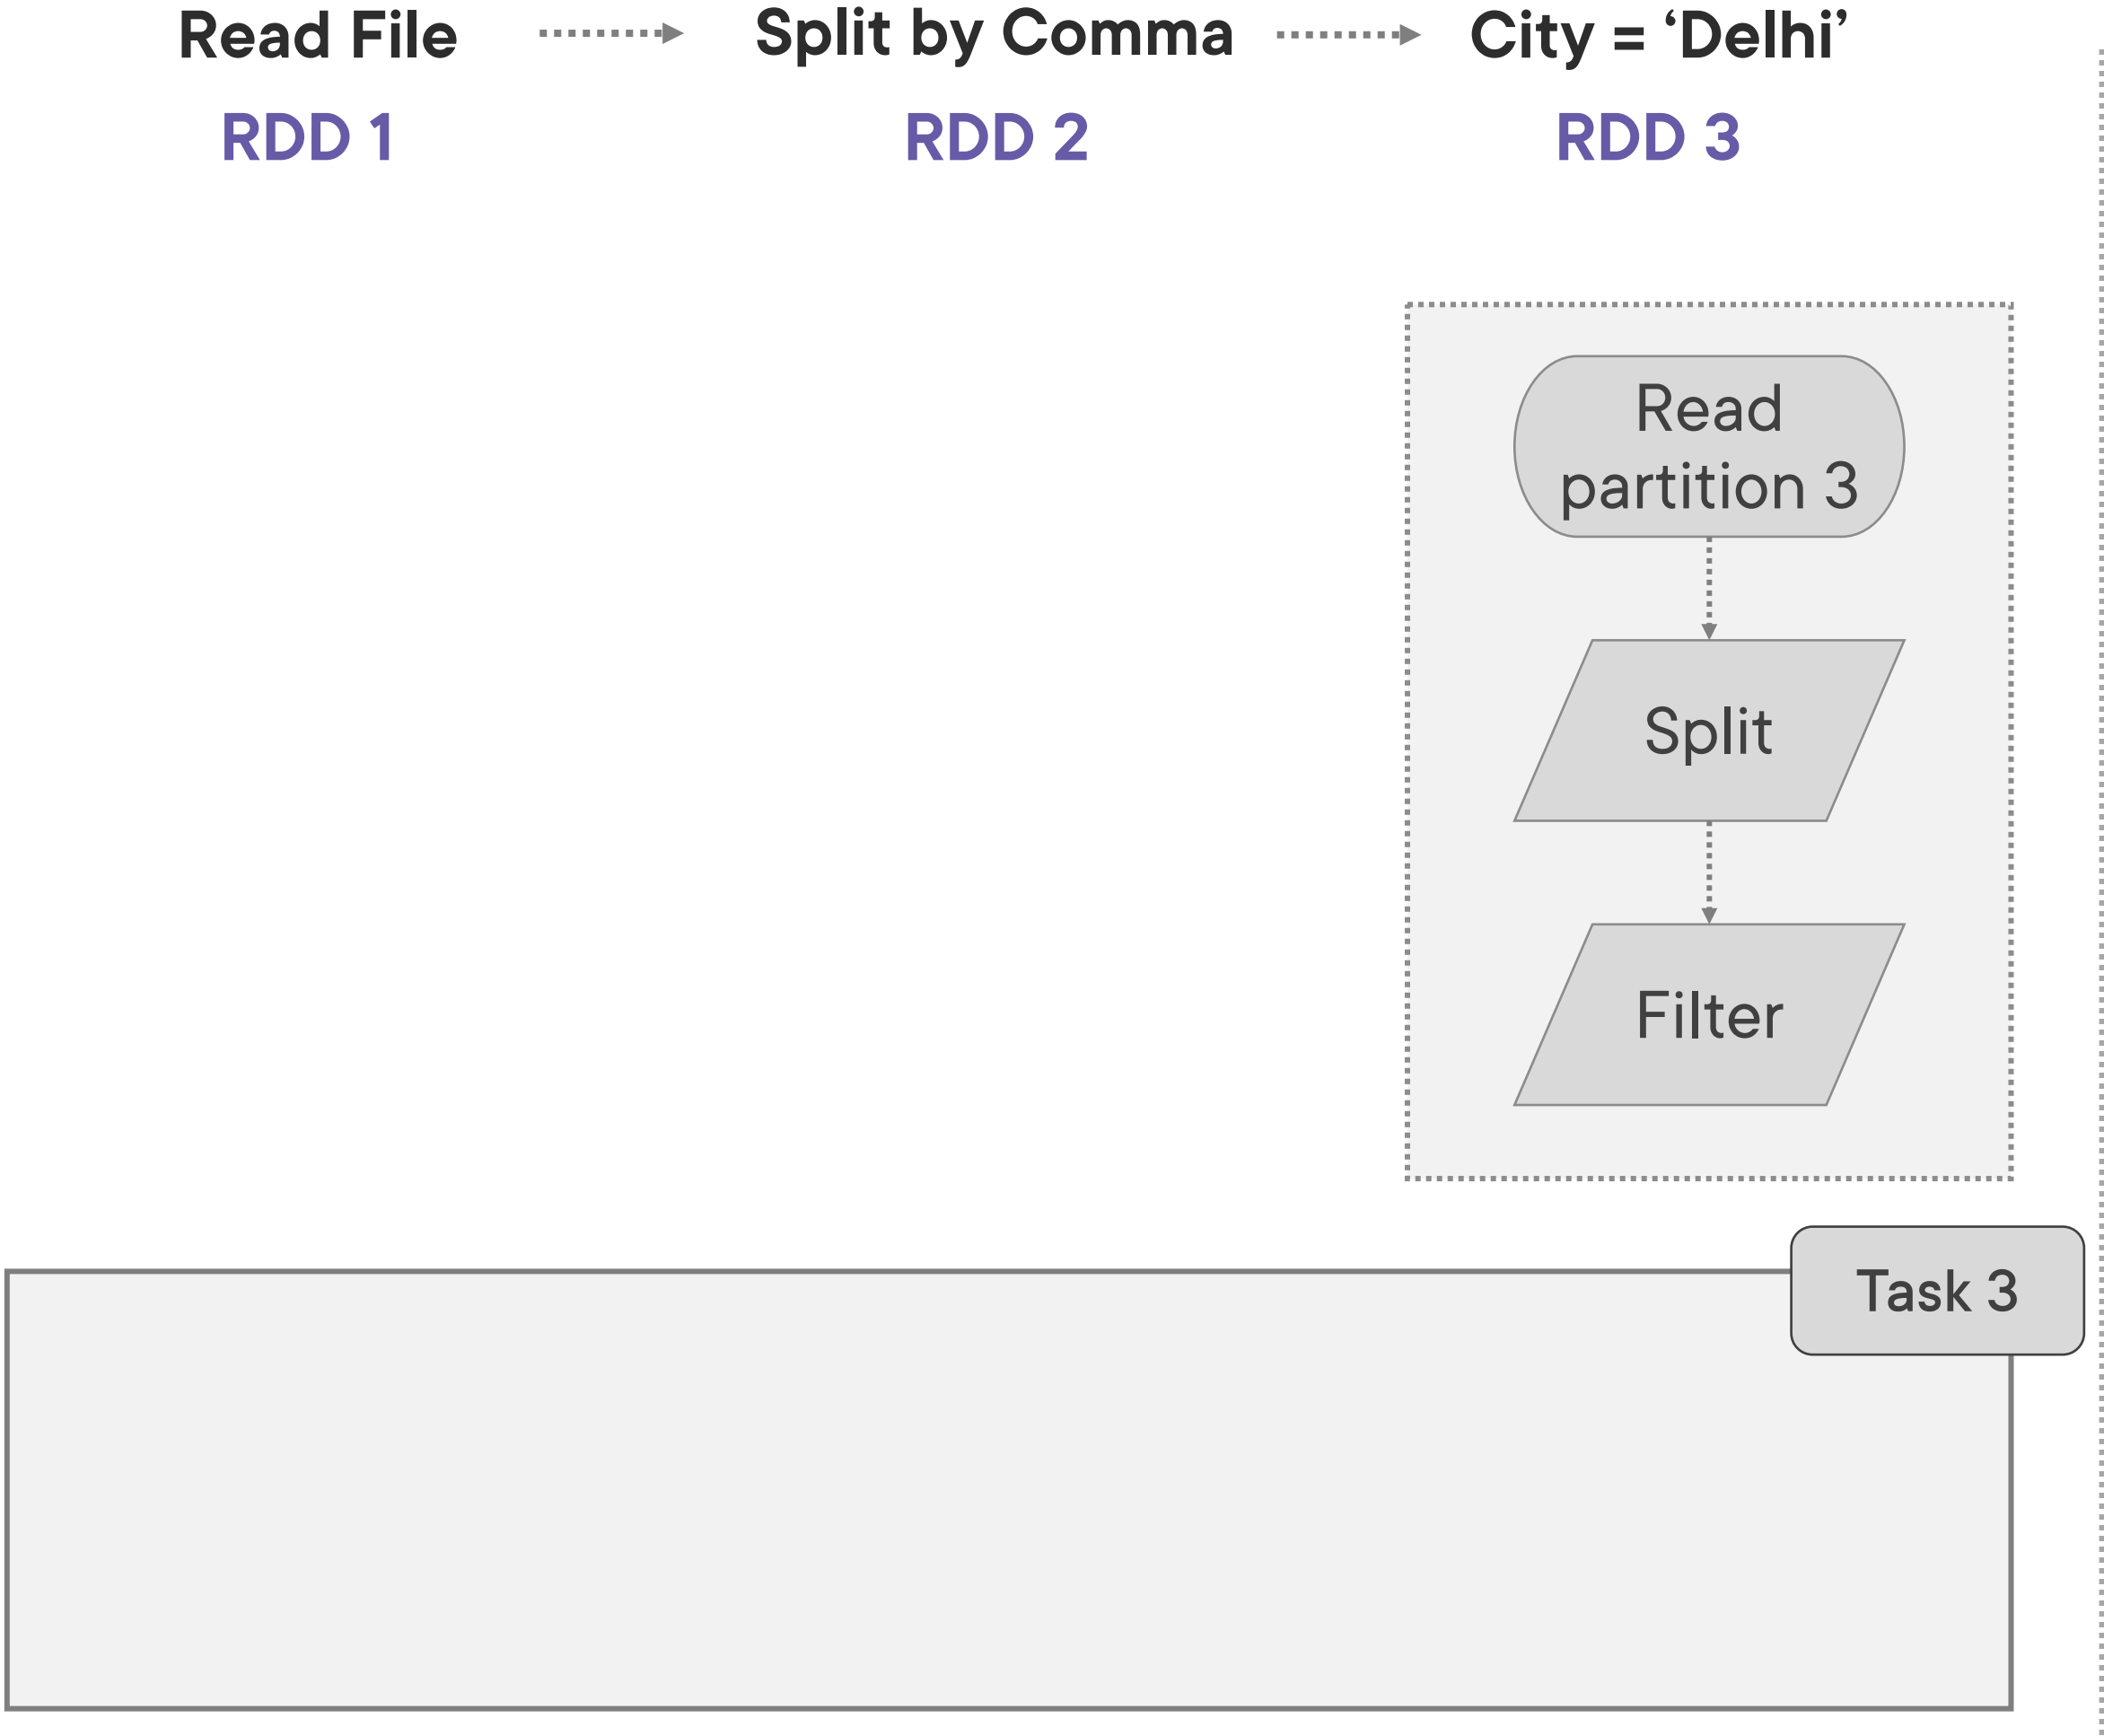
Only if the failure is transient,  
& not permanent (like divide by zero)

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900





<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900



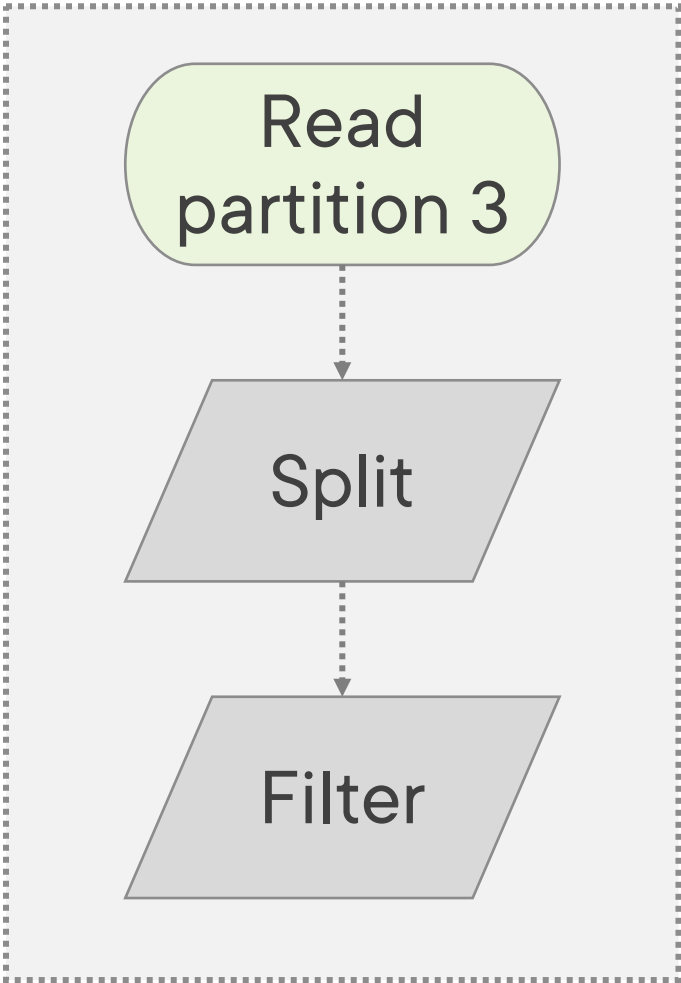
3	Delhi	700

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

**Read File**  
**RDD 1**

**Split by Comma**  
**RDD 2**

**City = 'Delhi'**  
**RDD 3**



3	Delhi	700



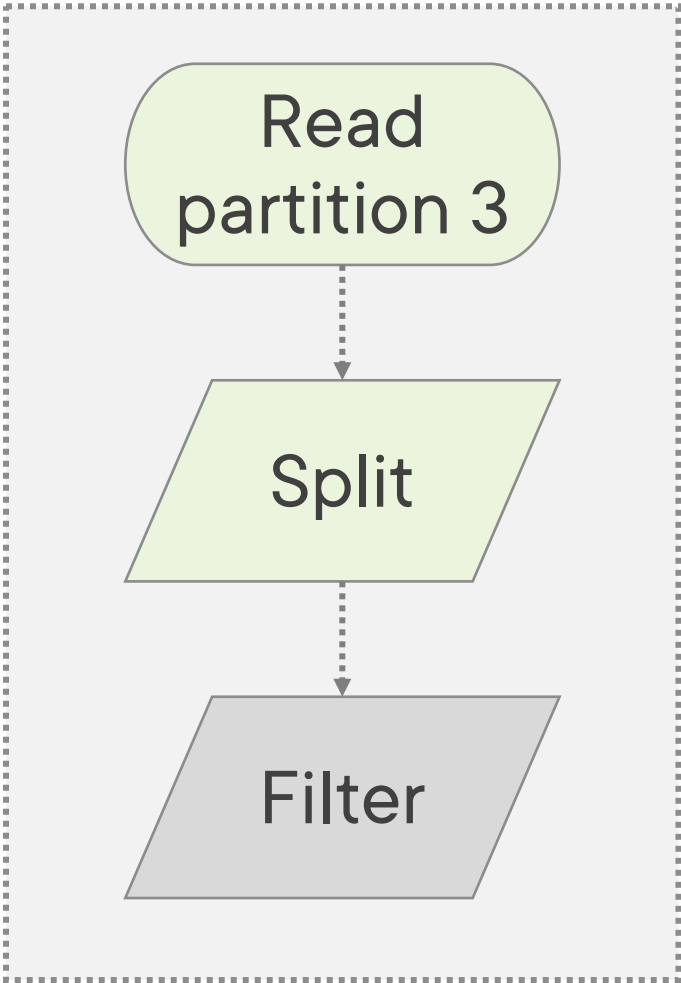
**Task 3**

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

**Read File**  
**RDD 1**

**Split by Comma**  
**RDD 2**

**City = 'Delhi'**  
**RDD 3**



3	Delhi	700

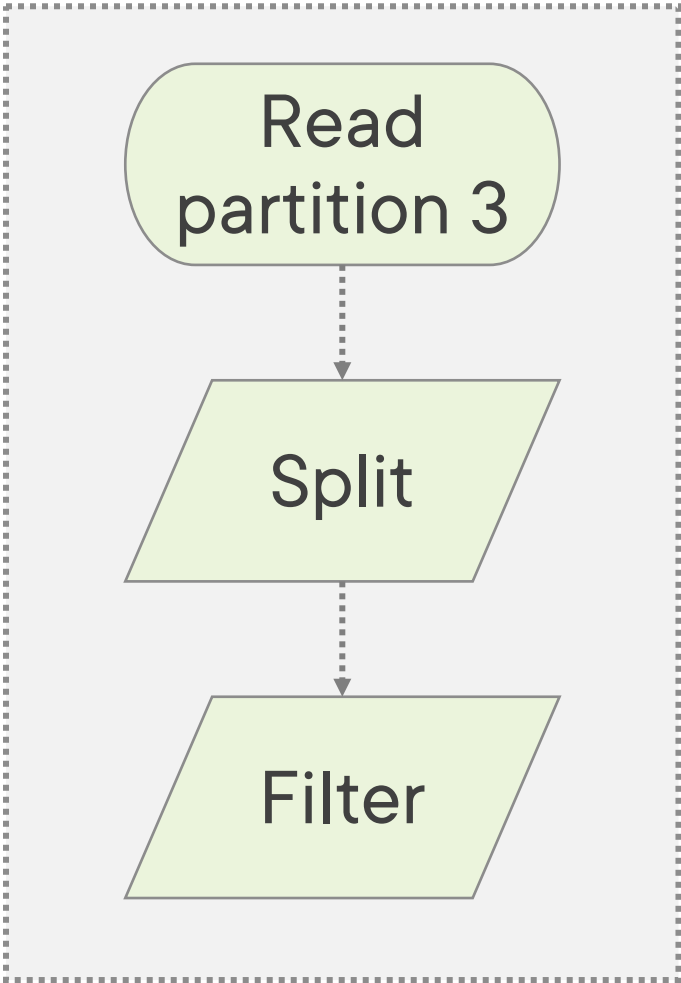


<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

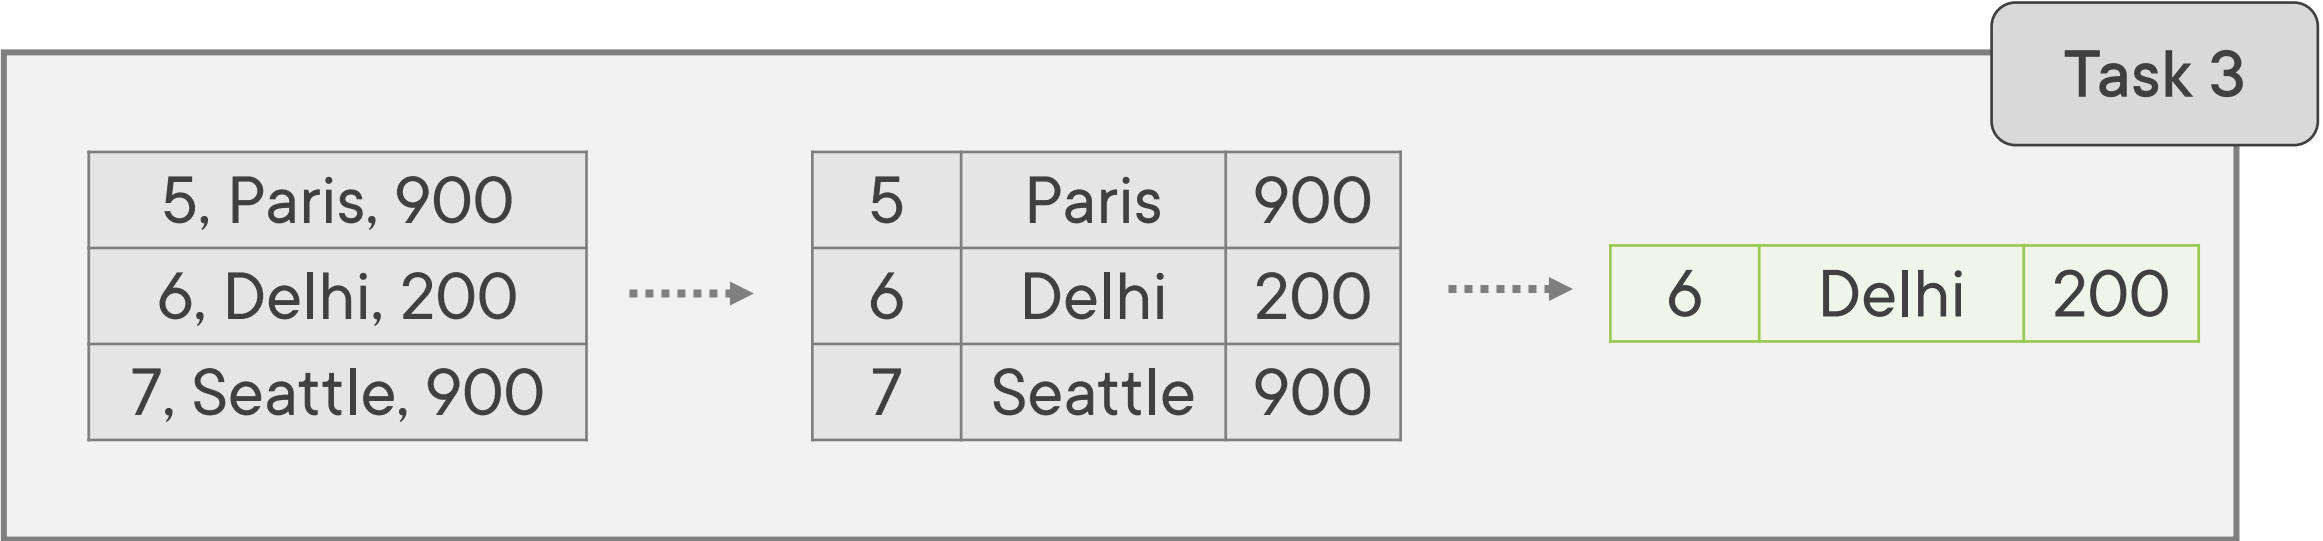
**Read File**  
**RDD 1**

**Split by Comma**  
**RDD 2**

**City = 'Delhi'**  
**RDD 3**



3	Delhi	700

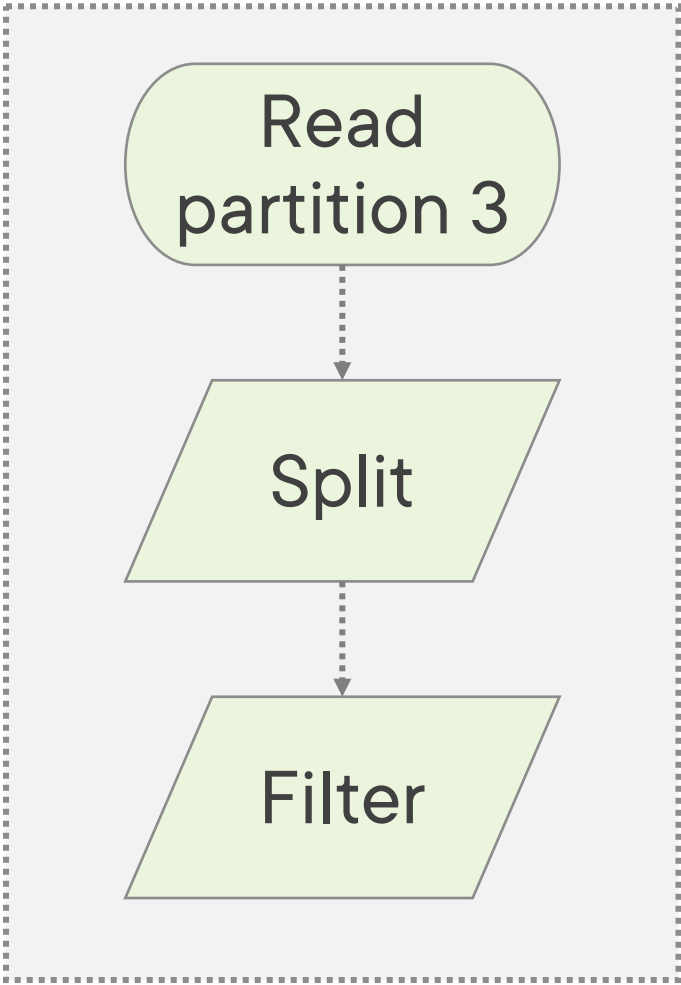


Id, City, Amount		
1	Seattle	600
2	London	300
3	Delhi	700
4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900

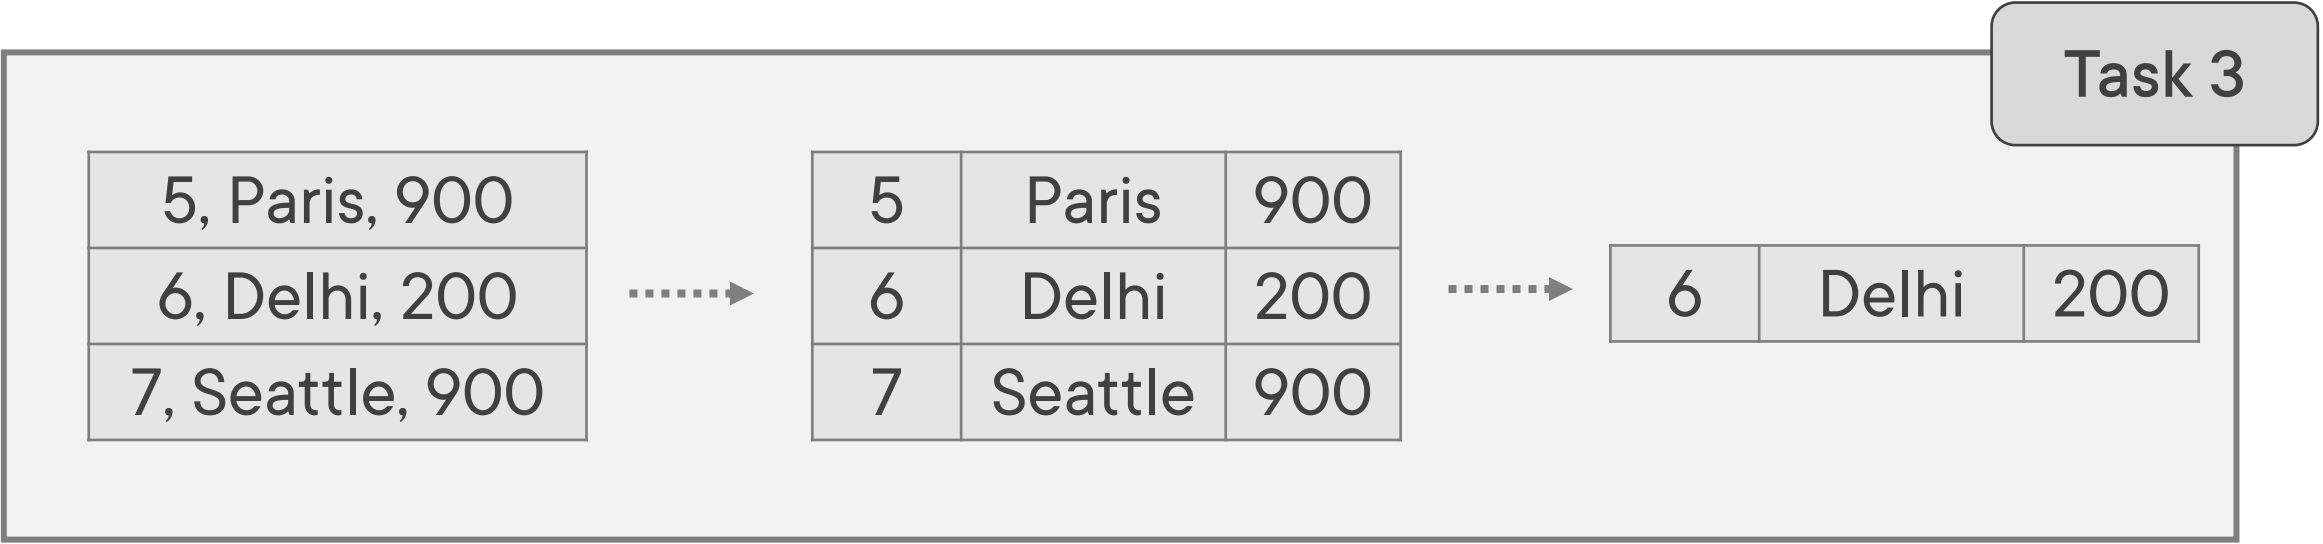
Read File  
RDD 1

Split by Comma  
RDD 2

City = 'Delhi'  
RDD 3



3	Delhi	700
6	Delhi	200



# Features of RDDs

## **In-memory**

Resides in the memory of cluster

## **Partitioned**

Split into partitions, processed by tasks

## **Read-only**

Transformed into another RDD or result

## **Resilient**

Auto recover in case of a failure

# Creating RDDs

---

# Options to Create RDD

**Parallelize a  
Collection**

**Read a File**

**From another RDD**



Setup Instructions, and Code & Data Files  
are available in Exercise Files section  
of the course

# Working with Pair RDDs

---

# Pair RDDs

2
3
4
5
6

**RDD1**  
type: integer

math.sqrt



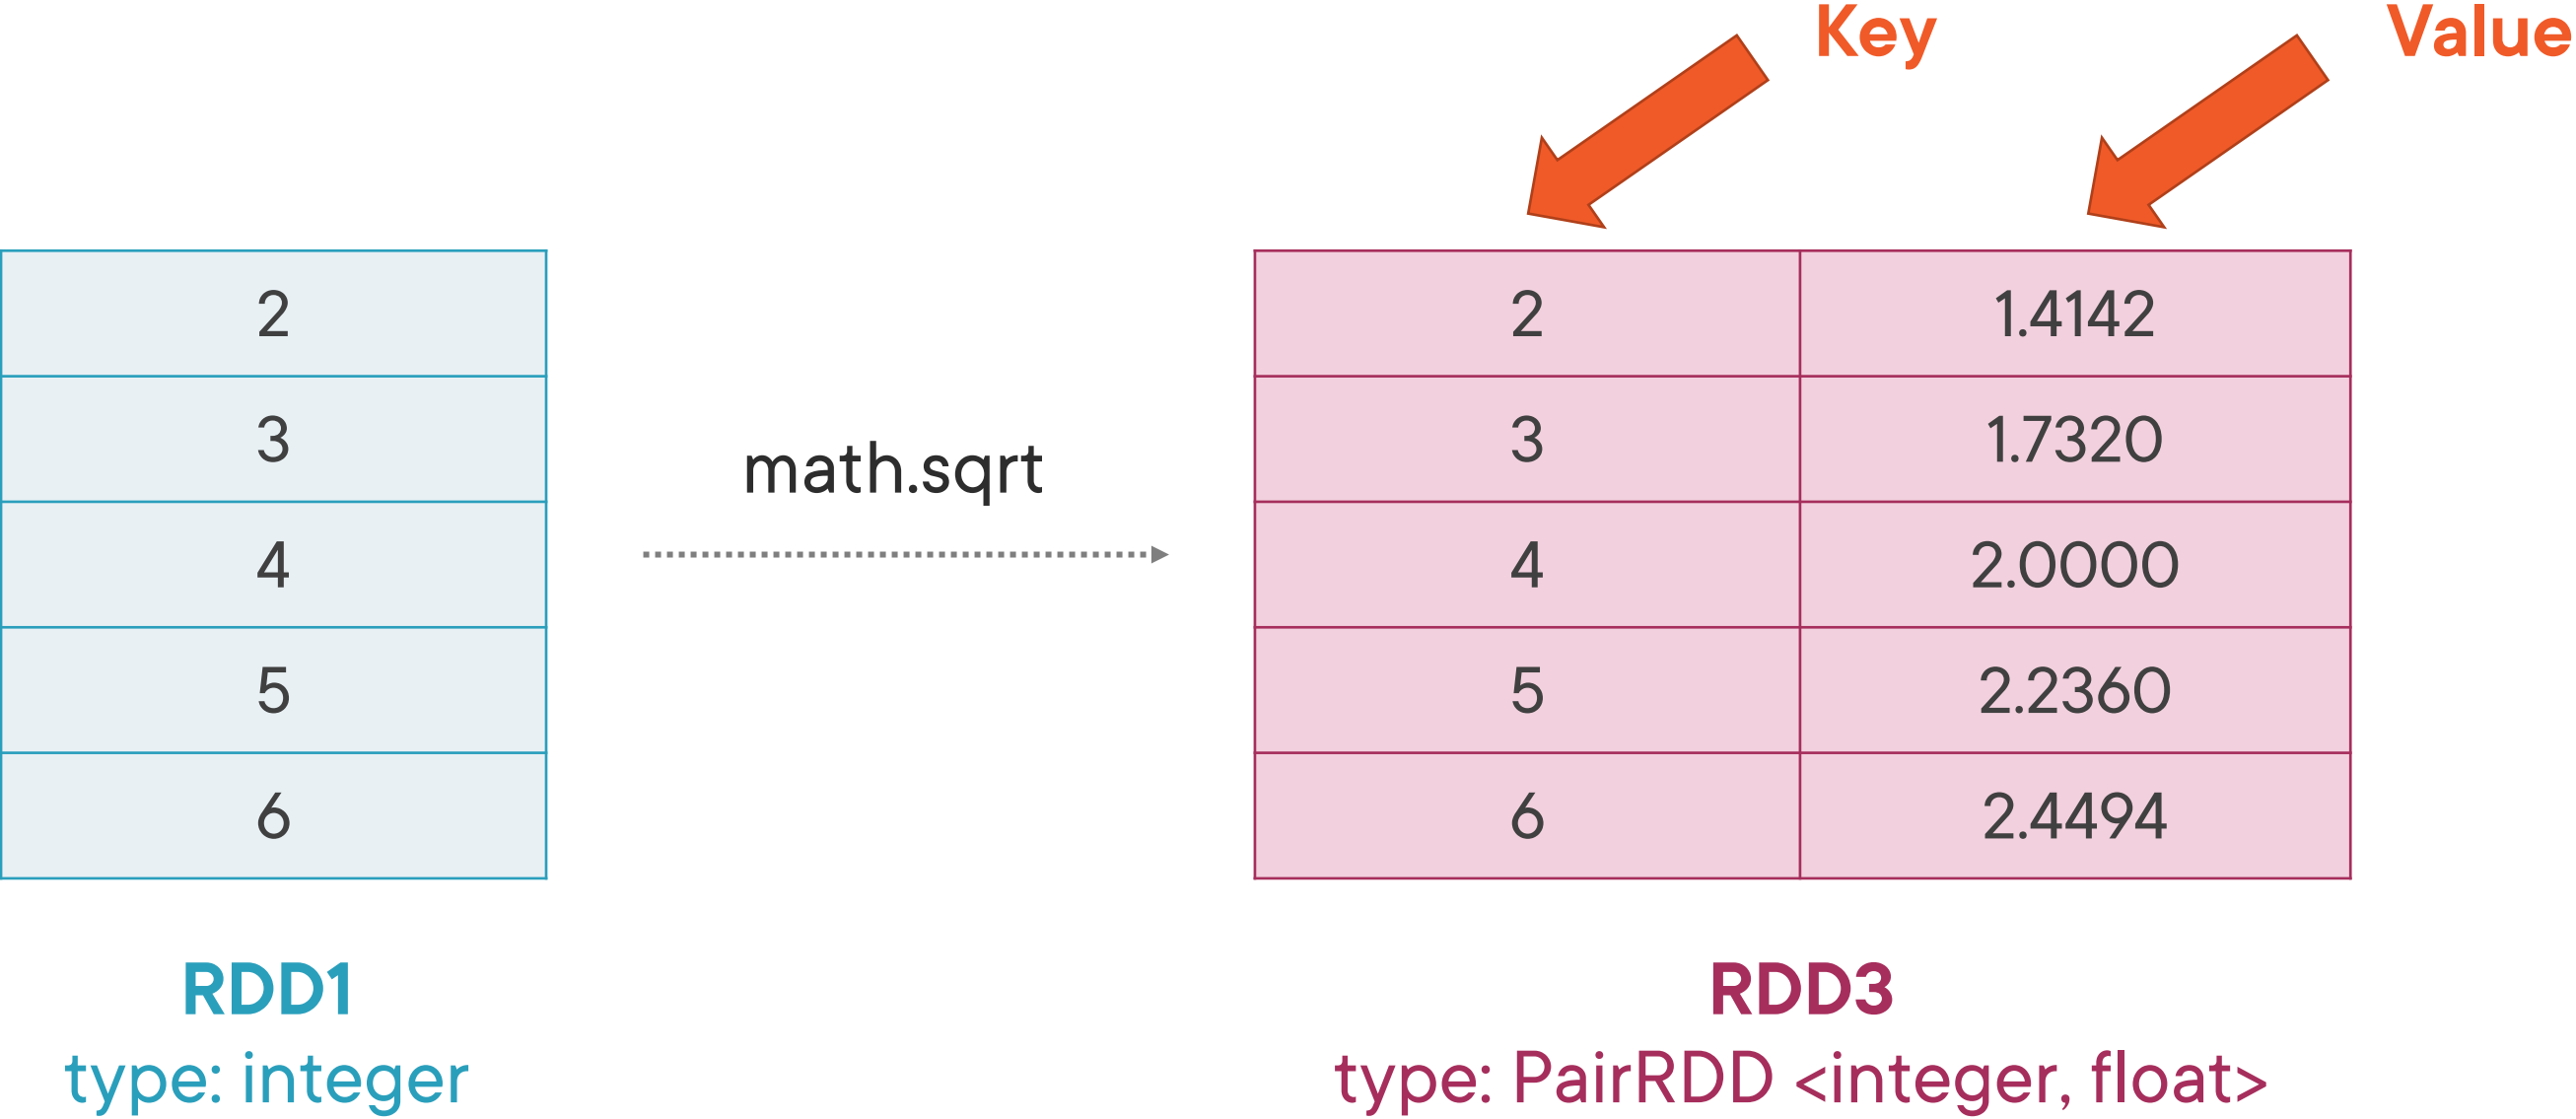
1.4142
1.7320
2.0000
2.2360
2.4494

**RDD2**  
type: float




Whose square  
root is it?

# Pair RDDs



# Pair RDDs



Key (CustomerID)	Value (Customer Info)
1	['Mohit', '20000', 'Hyderabad']
2	['Ivan', '23000', 'Seattle']
3	['Kerri', '27000', 'London']
4	['Andrew', '18000', 'Paris']
5	['Neha', '32000', 'Delhi']

**Customer Data – PairRDD <integer, array>**

# Pair RDDs

## RDDs with Key-Value pairs

- Two items are linked together
- Key is identifier. Value is corresponding data

## Key need not be unique

## Spark has special operations for Pair RDDs

- reduceByKey, sortByKey, countByKey etc.
- Some operations are costly

1	['Mohit', '20000', 'Hyderabad']
2	['Ivan', '23000', 'Seattle']
3	['Kerri', '27000', 'London']
4	['Andrew', '18000', 'Paris']
5	['Neha', '32000', 'Delhi']

# Calculate sum of Amount by City

<b>Id, City, Amount</b>
1, Seattle, 600
2, Seattle, 300
3, Delhi, 700
4, Seattle, 400
5, Delhi, 900
6, Delhi, 200
7, Seattle, 900

**CSV File**

Create  
Pair RDD

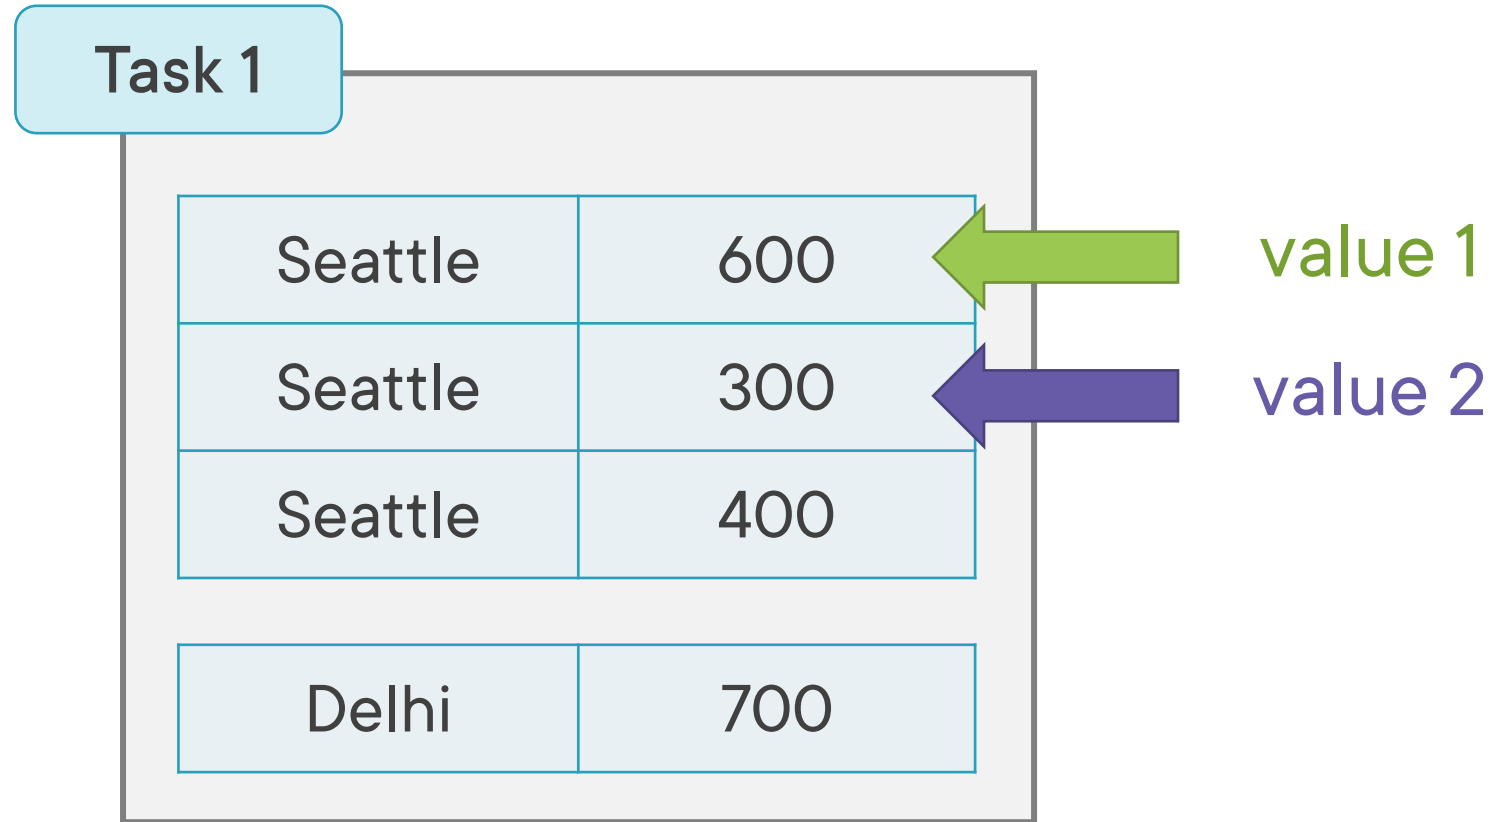
Seattle	600
Seattle	300
Delhi	700
Seattle	400

Delhi	900
Delhi	200
Seattle	900

**PairRDD <string, integer>**

```
cityPairRDD
  .reduceByKey (lambda value1, value2: value1 + value2)
```

## Calculate sum of Amount by City



Seattle	600
Seattle	300
Delhi	700
Seattle	400

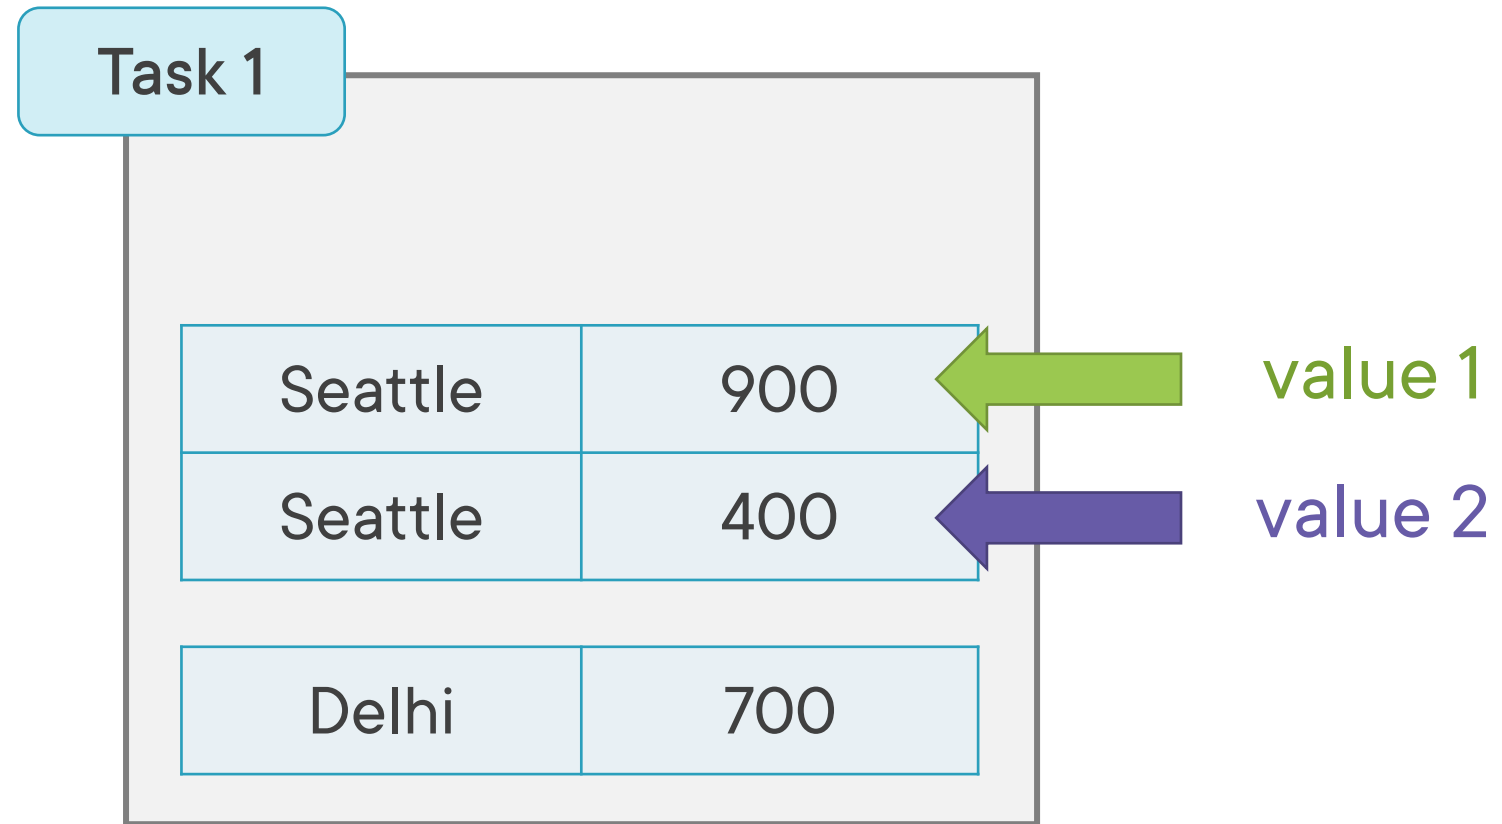
Delhi	900
Delhi	200
Seattle	900

### Lambda Function

```
value1, value2: value1 + value2
```



# Calculate sum of Amount by City




Delhi	900
Delhi	200
Seattle	900

## Lambda Function

```
value1, value2: value1 + value2
```

# Calculate sum of Amount by City



Delhi	900
Delhi	200
Seattle	900

## Lambda Function

```
value1, value2: value1 + value2
```

# Calculate sum of Amount by City

Task 1	
Seattle	1300
Delhi	700

Task 2	
Seattle	900
Delhi	900
Delhi	200

value 1

value 2

## Lambda Function

```
value1, value2: value1 + value2
```

# Calculate sum of Amount by City

Task 1

Seattle

1300

Delhi

700

Task 2

Seattle

900

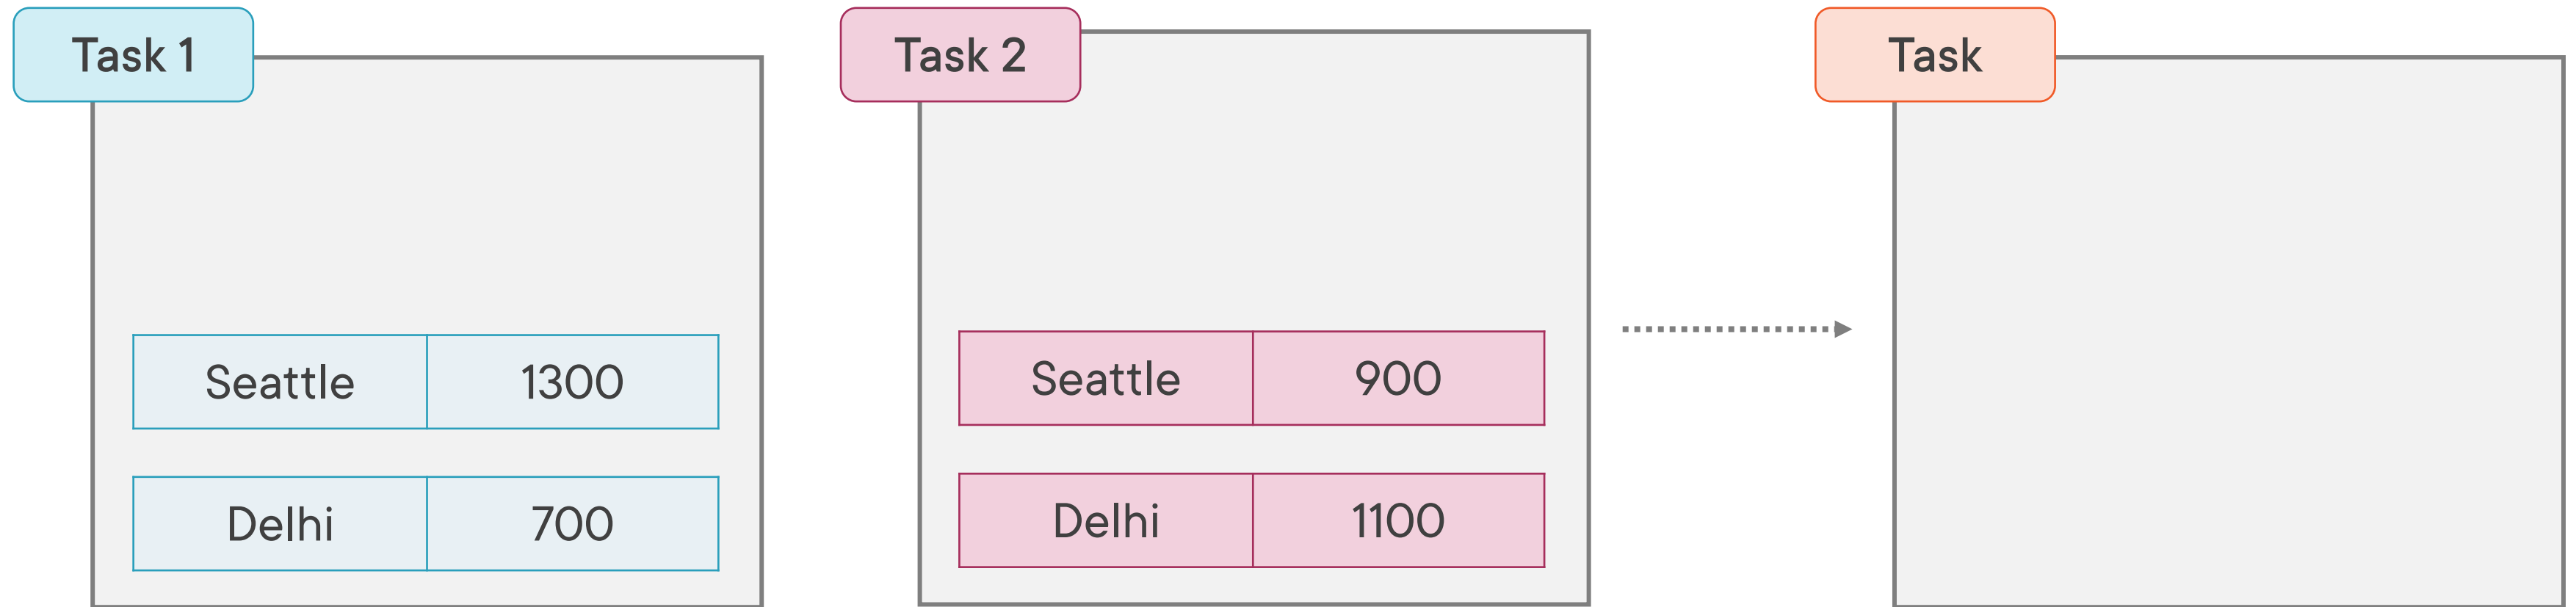
Delhi

1100

**Lambda Function**

```
value1, value2: value1 + value2
```

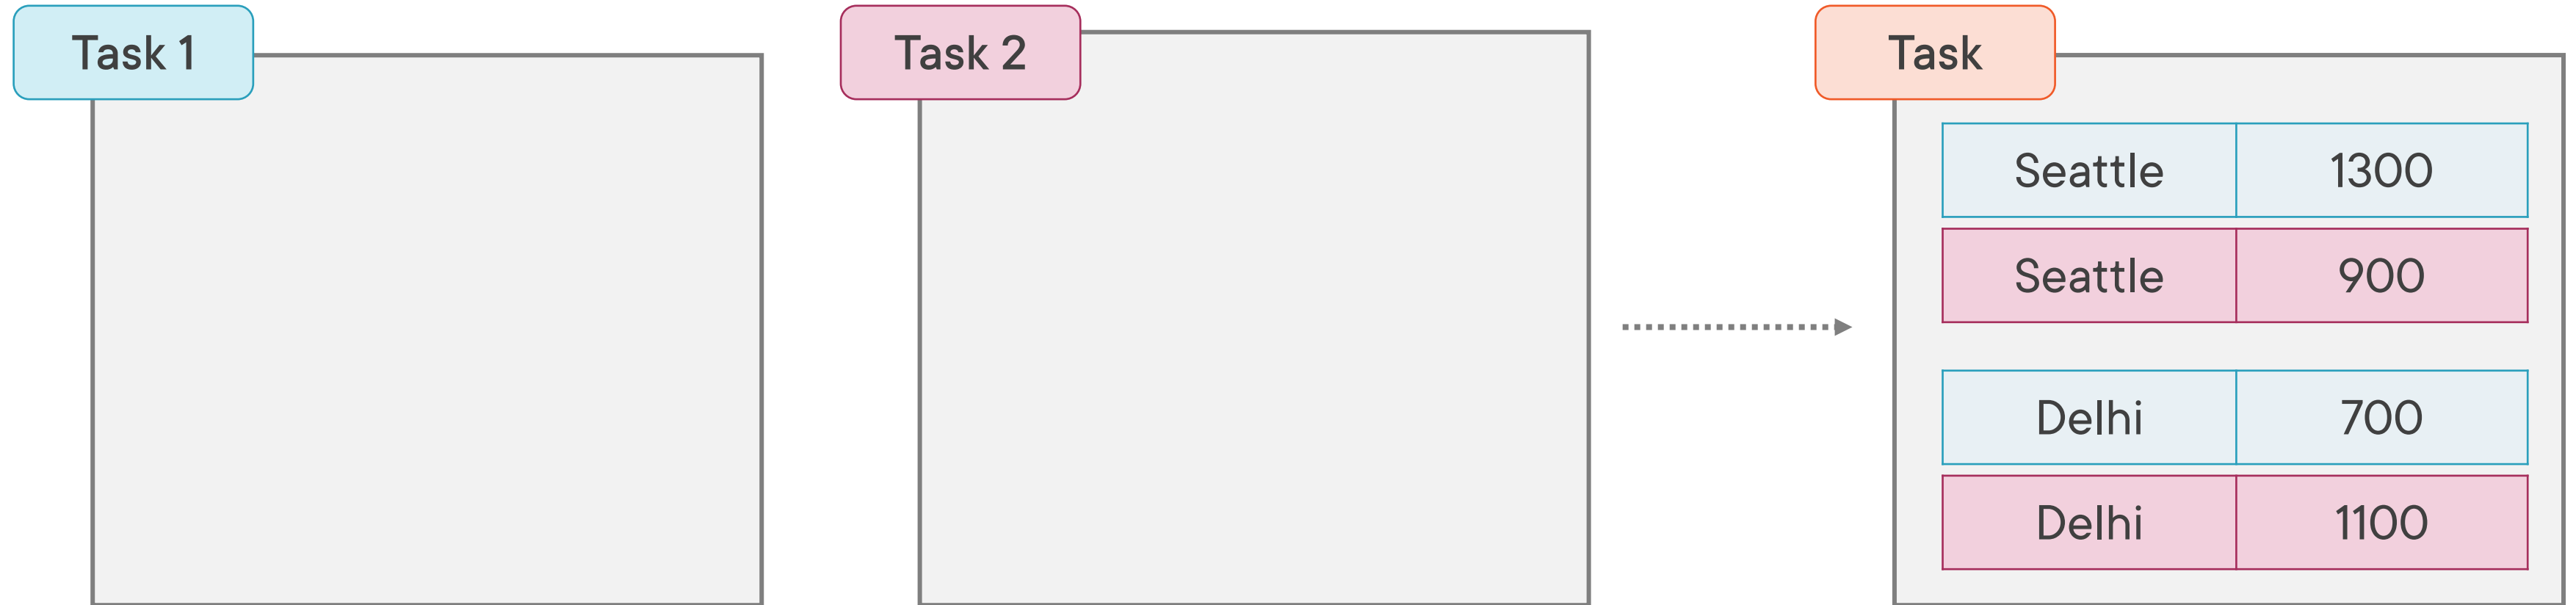
# Calculate sum of Amount by City



## Lambda Function

```
value1, value2: value1 + value2
```

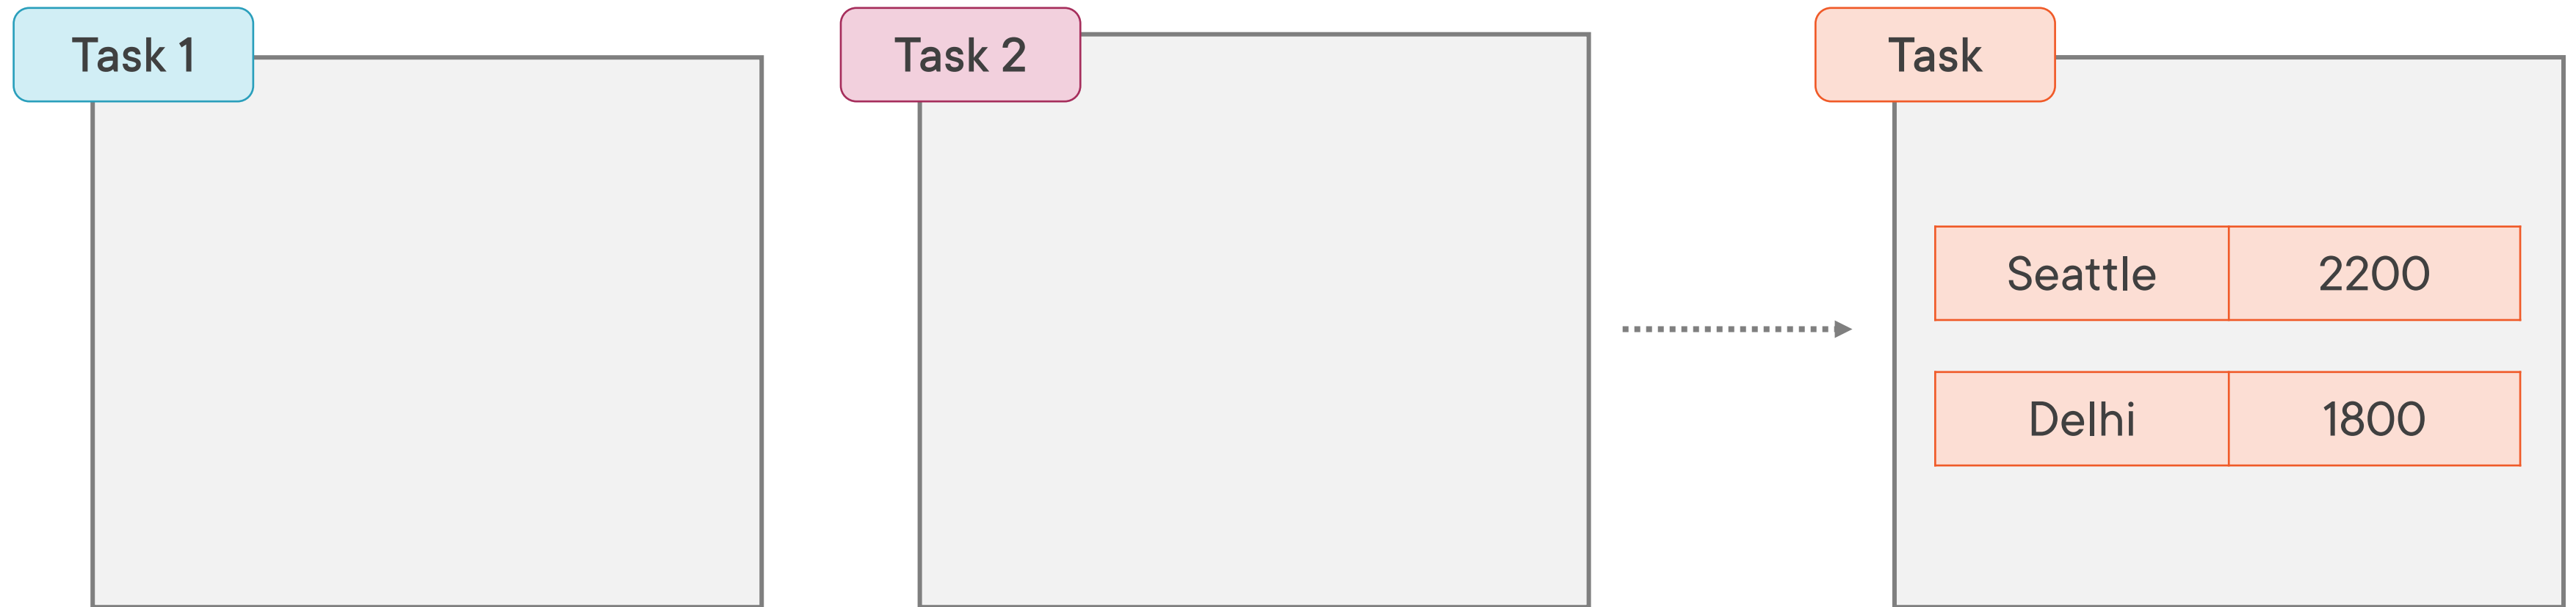
# Calculate sum of Amount by City



## Lambda Function

```
value1, value2: value1 + value2
```

# Calculate sum of Amount by City



## Lambda Function

```
value1, value2: value1 + value2
```

# Applying Operations on RDDs

---




# RDD Operations

**Transformation**

**Action**

# Transformation Operation



**Function that produces new RDD from existing RDDs**

**Transformation operations help in building a Lineage Graph**

**Examples:**

- Read a file
- Convert sales amount from INR to USD
- Filter records with sales amount greater than 1000

Transformations are Lazy operations,  
which are only executed when an  
Action operation is applied


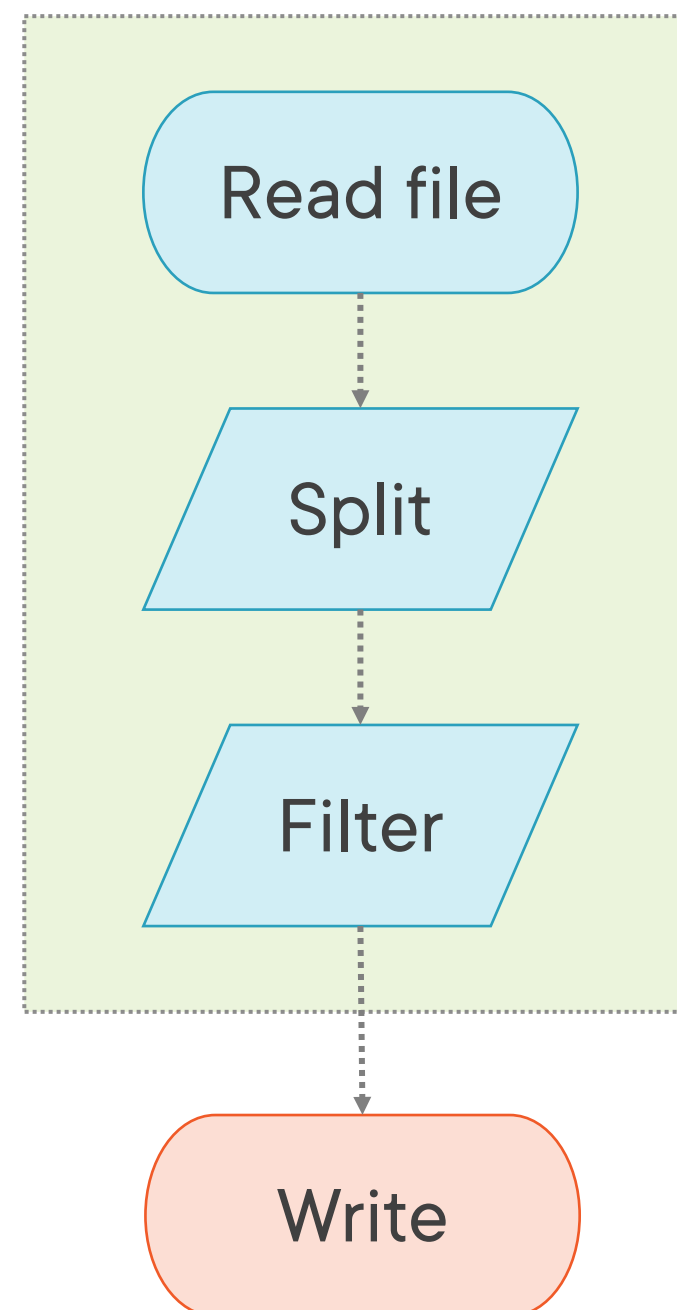
<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

1. Read file1.csv from Storage
2. Split data by comma
3. Filter where City='Delhi'
4. Write processed data to storage


## Lineage Graph

Built using transformation operations

Action operation



# Action Operation



**Returns result of RDD computations**

**Triggers execution of transformations using Lineage Graph**

**Examples:**

- Write the output to storage
- Print the output of operations
- Display the count

# Using Narrow Transformations

---

# Transformation Operations

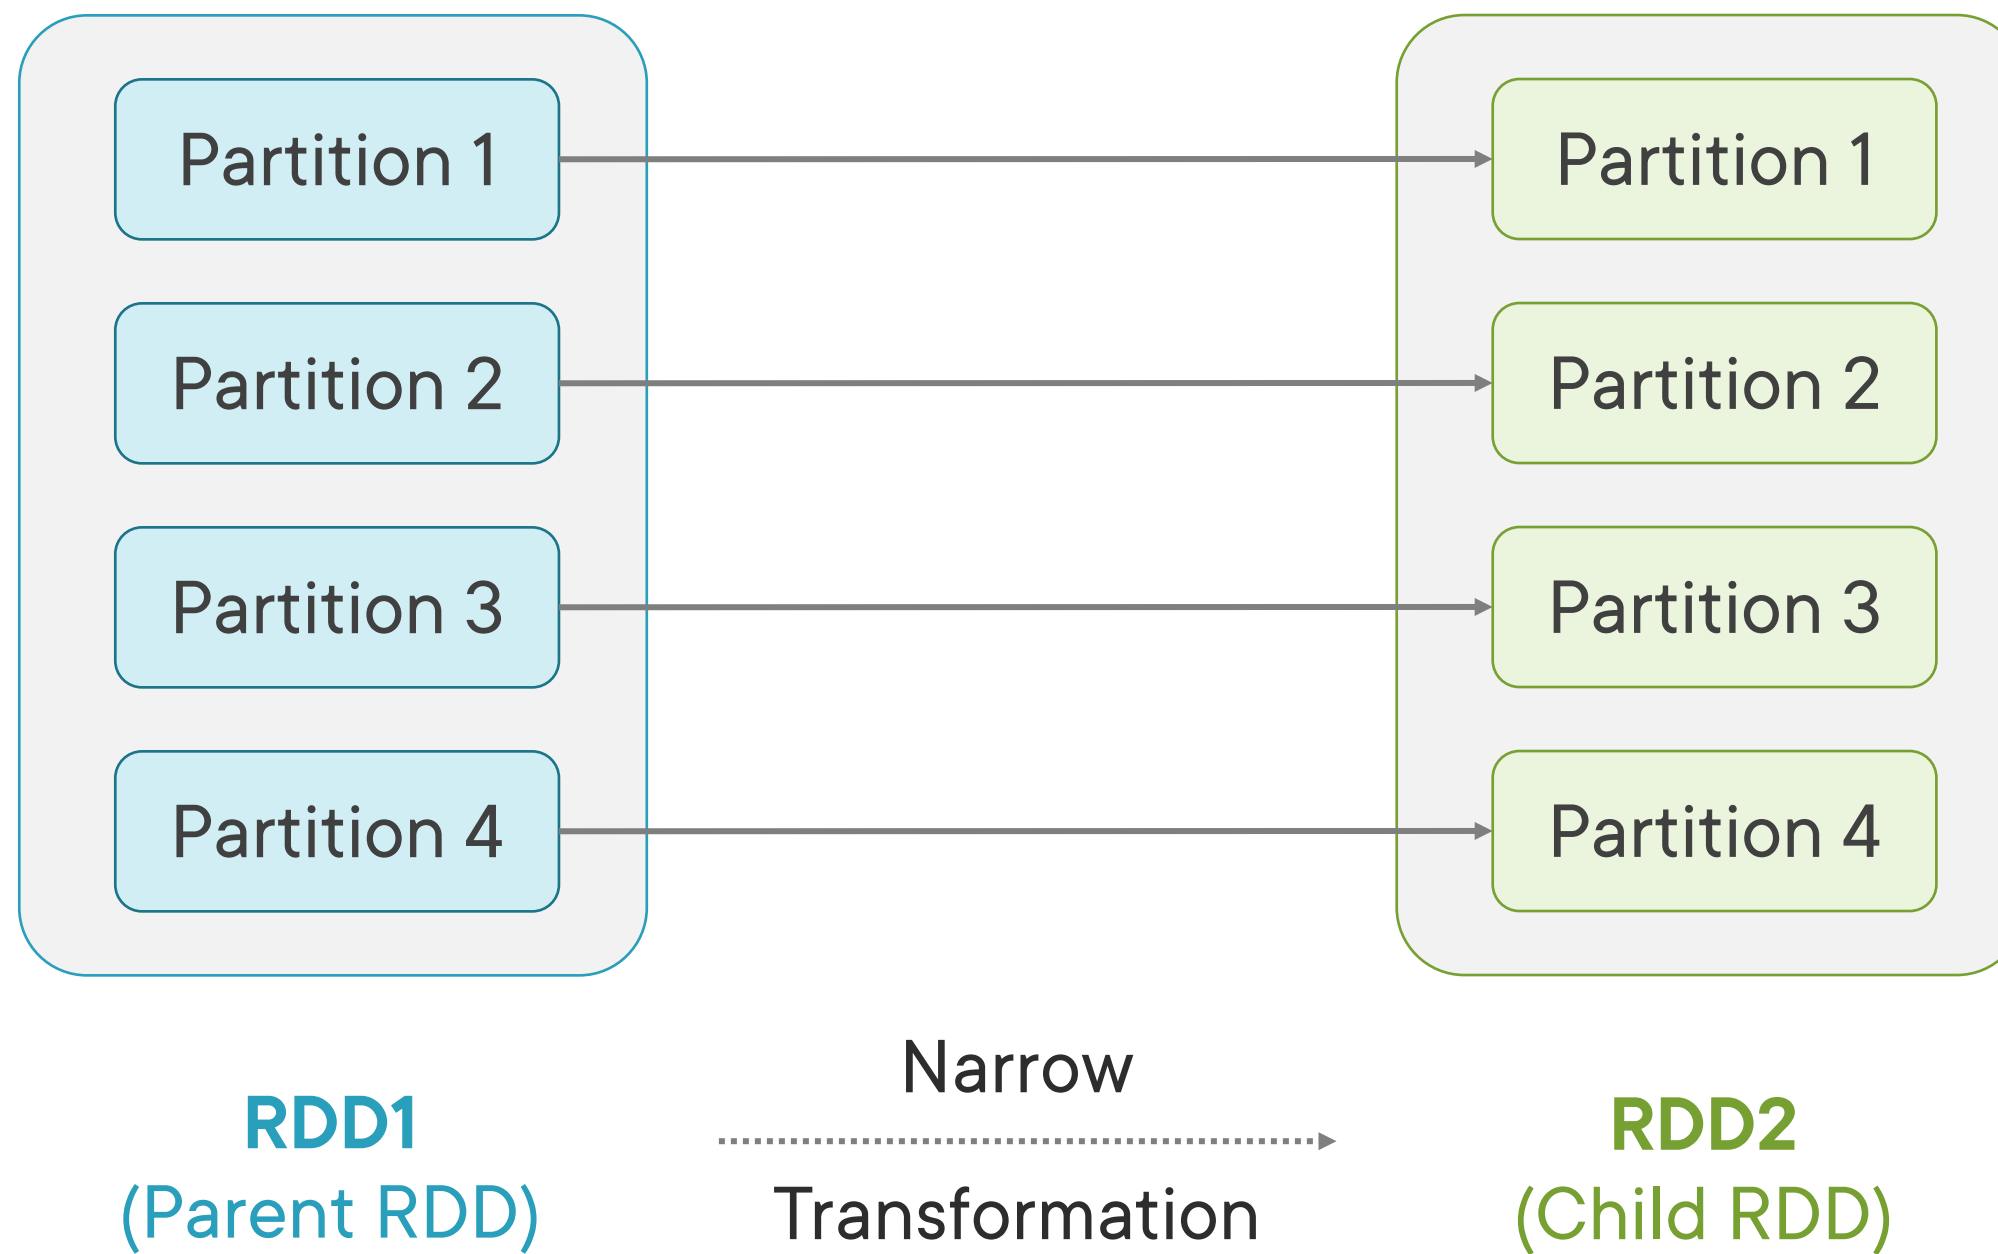
**Narrow Dependency  
Transformation**

**Wide Dependency  
Transformation**

A transformation where **each input partition** is used **at-most once** to produce output partitions



Each input partition is used **at-most once** to produce output partitions



**Filter & Map  
operations are  
examples of this**

Read file,  
split by comma  
RDD 1

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Task 1

Task 2

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Filter  
City = 'Seattle'  
**RDD 2**

Task 1

1	Seattle	600
2	London	300
3	Delhi	700

Partition 1

Task 2

4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900

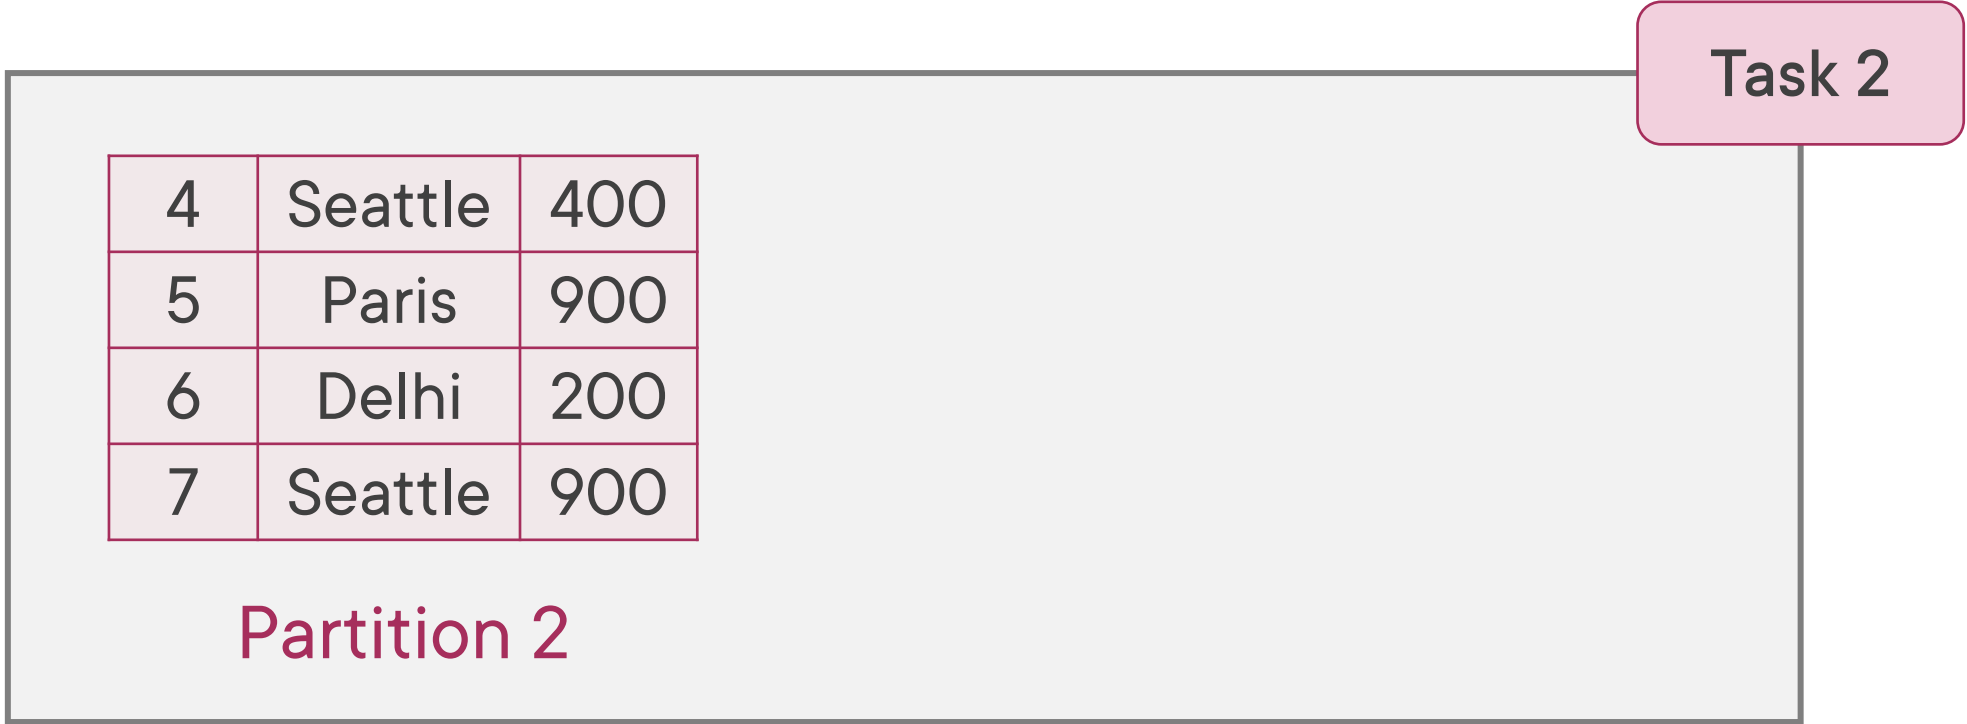
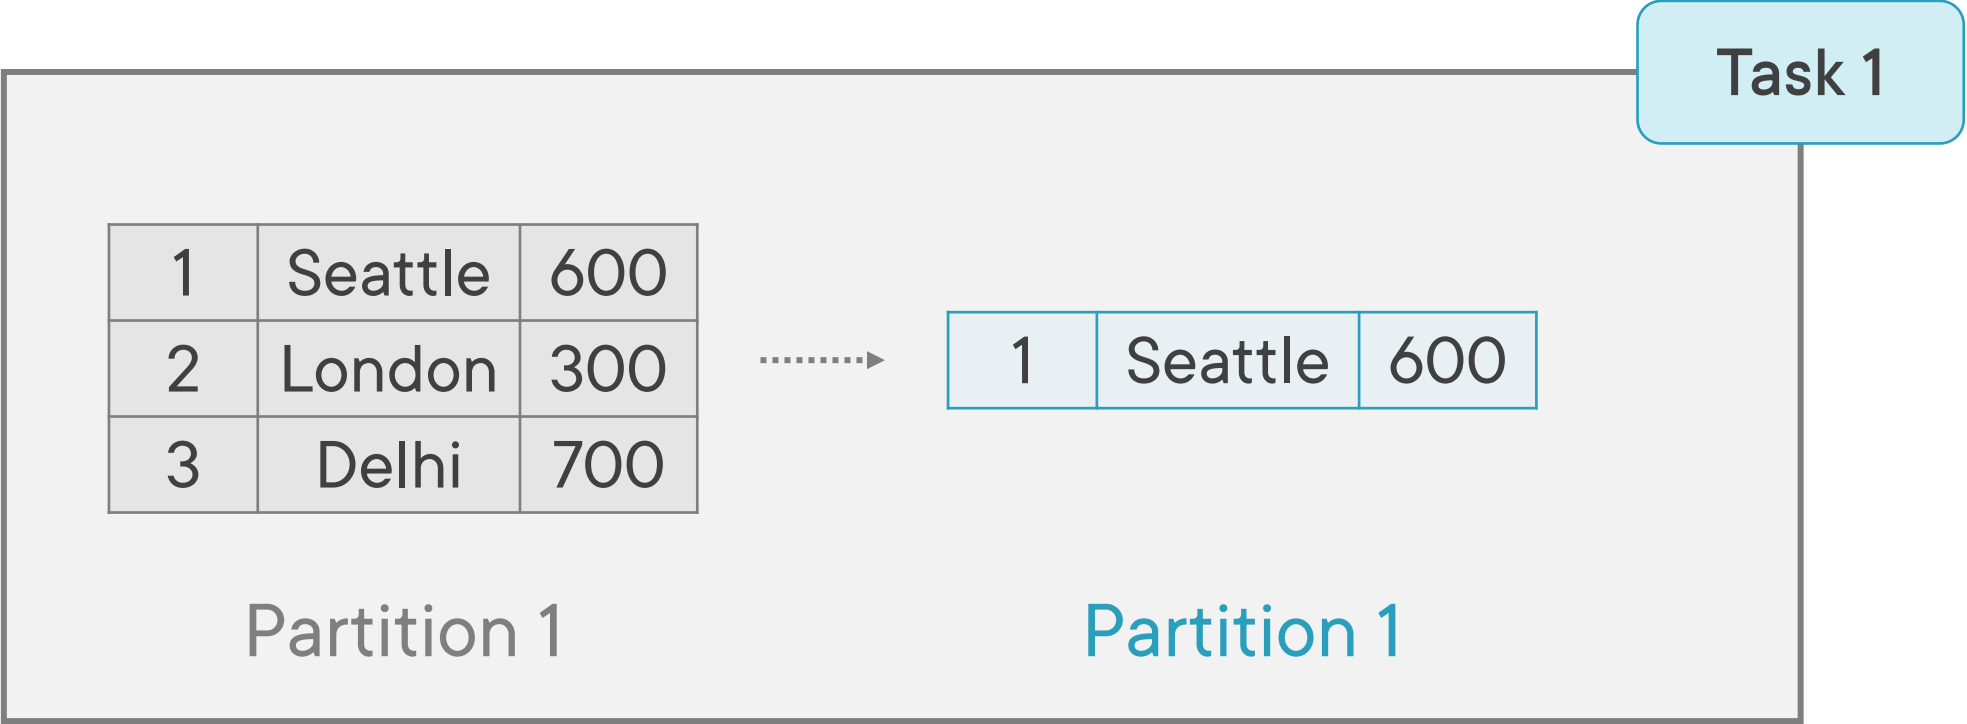
Partition 2

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Filter  
City = 'Seattle'  
**RDD 2**

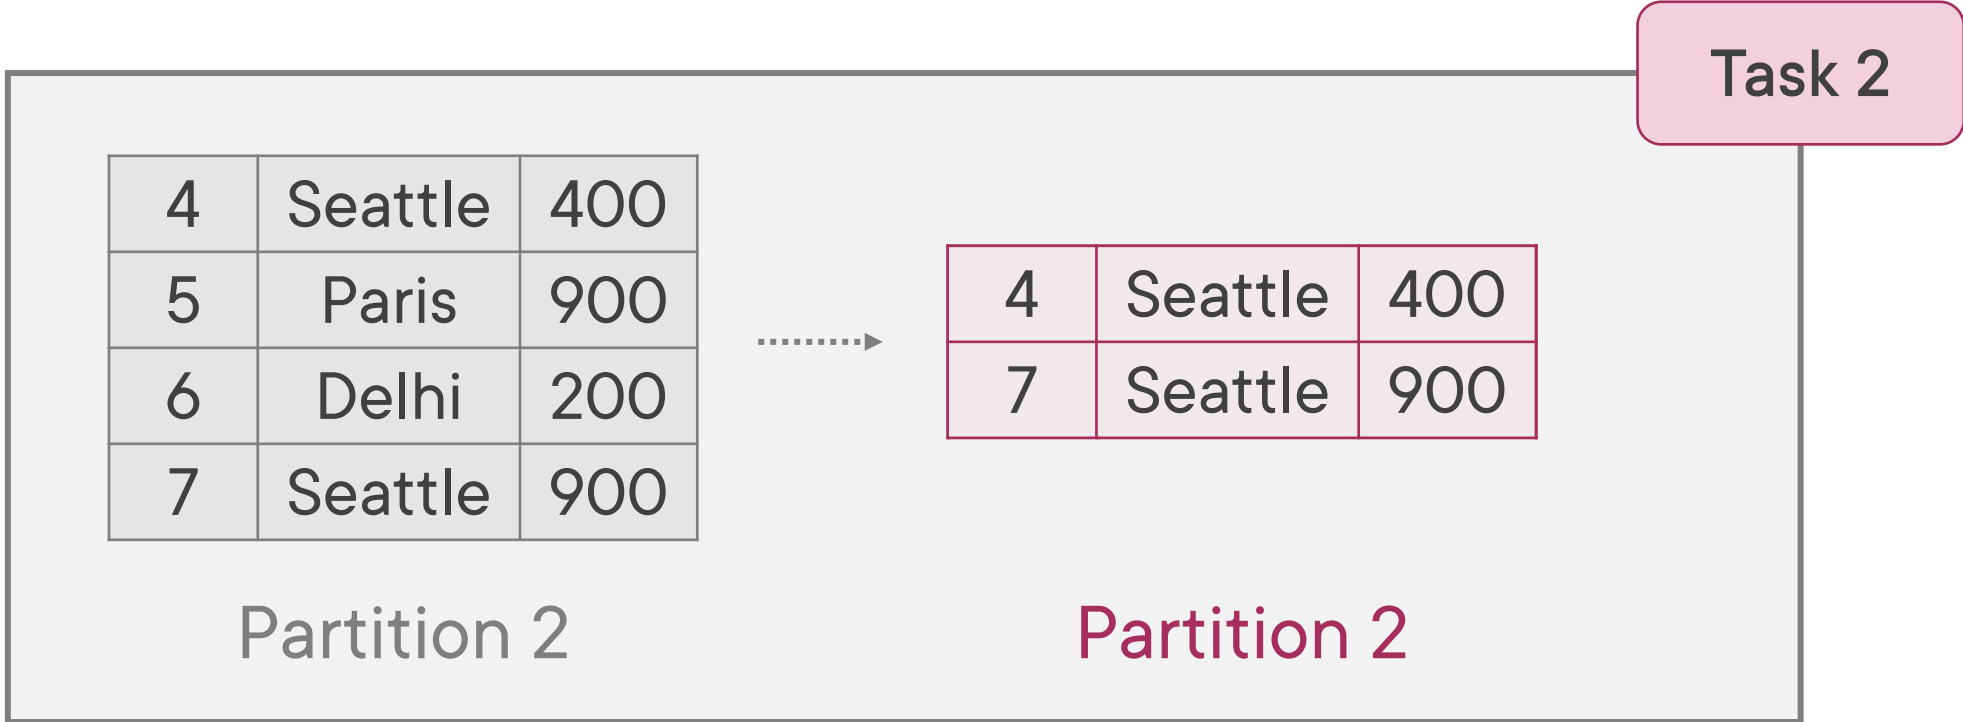
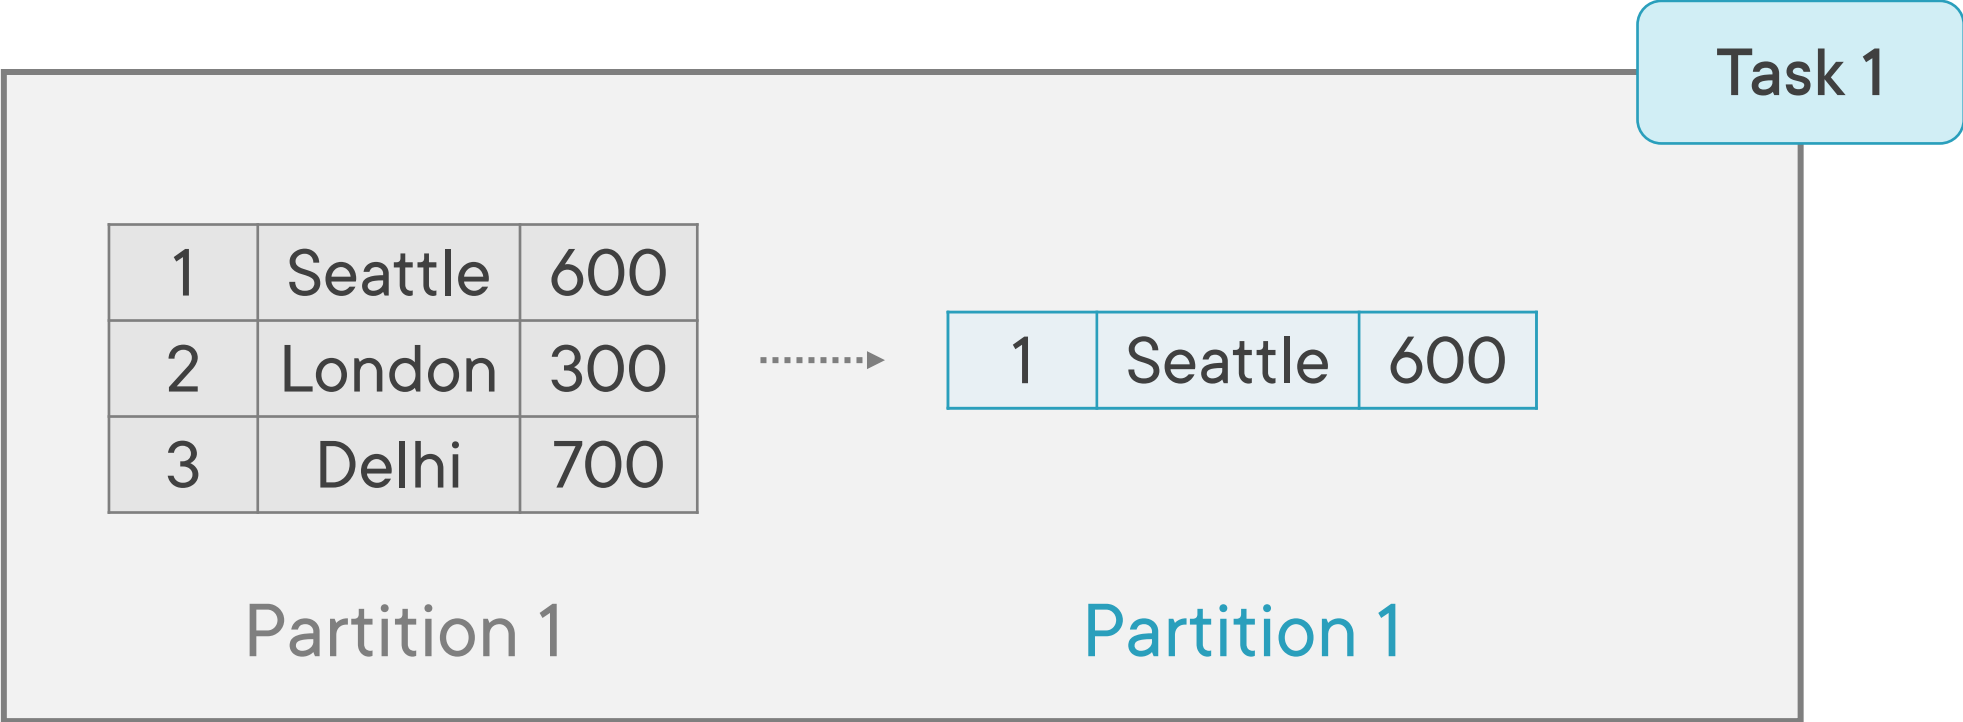


<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Filter  
City = 'Seattle'  
**RDD 2**



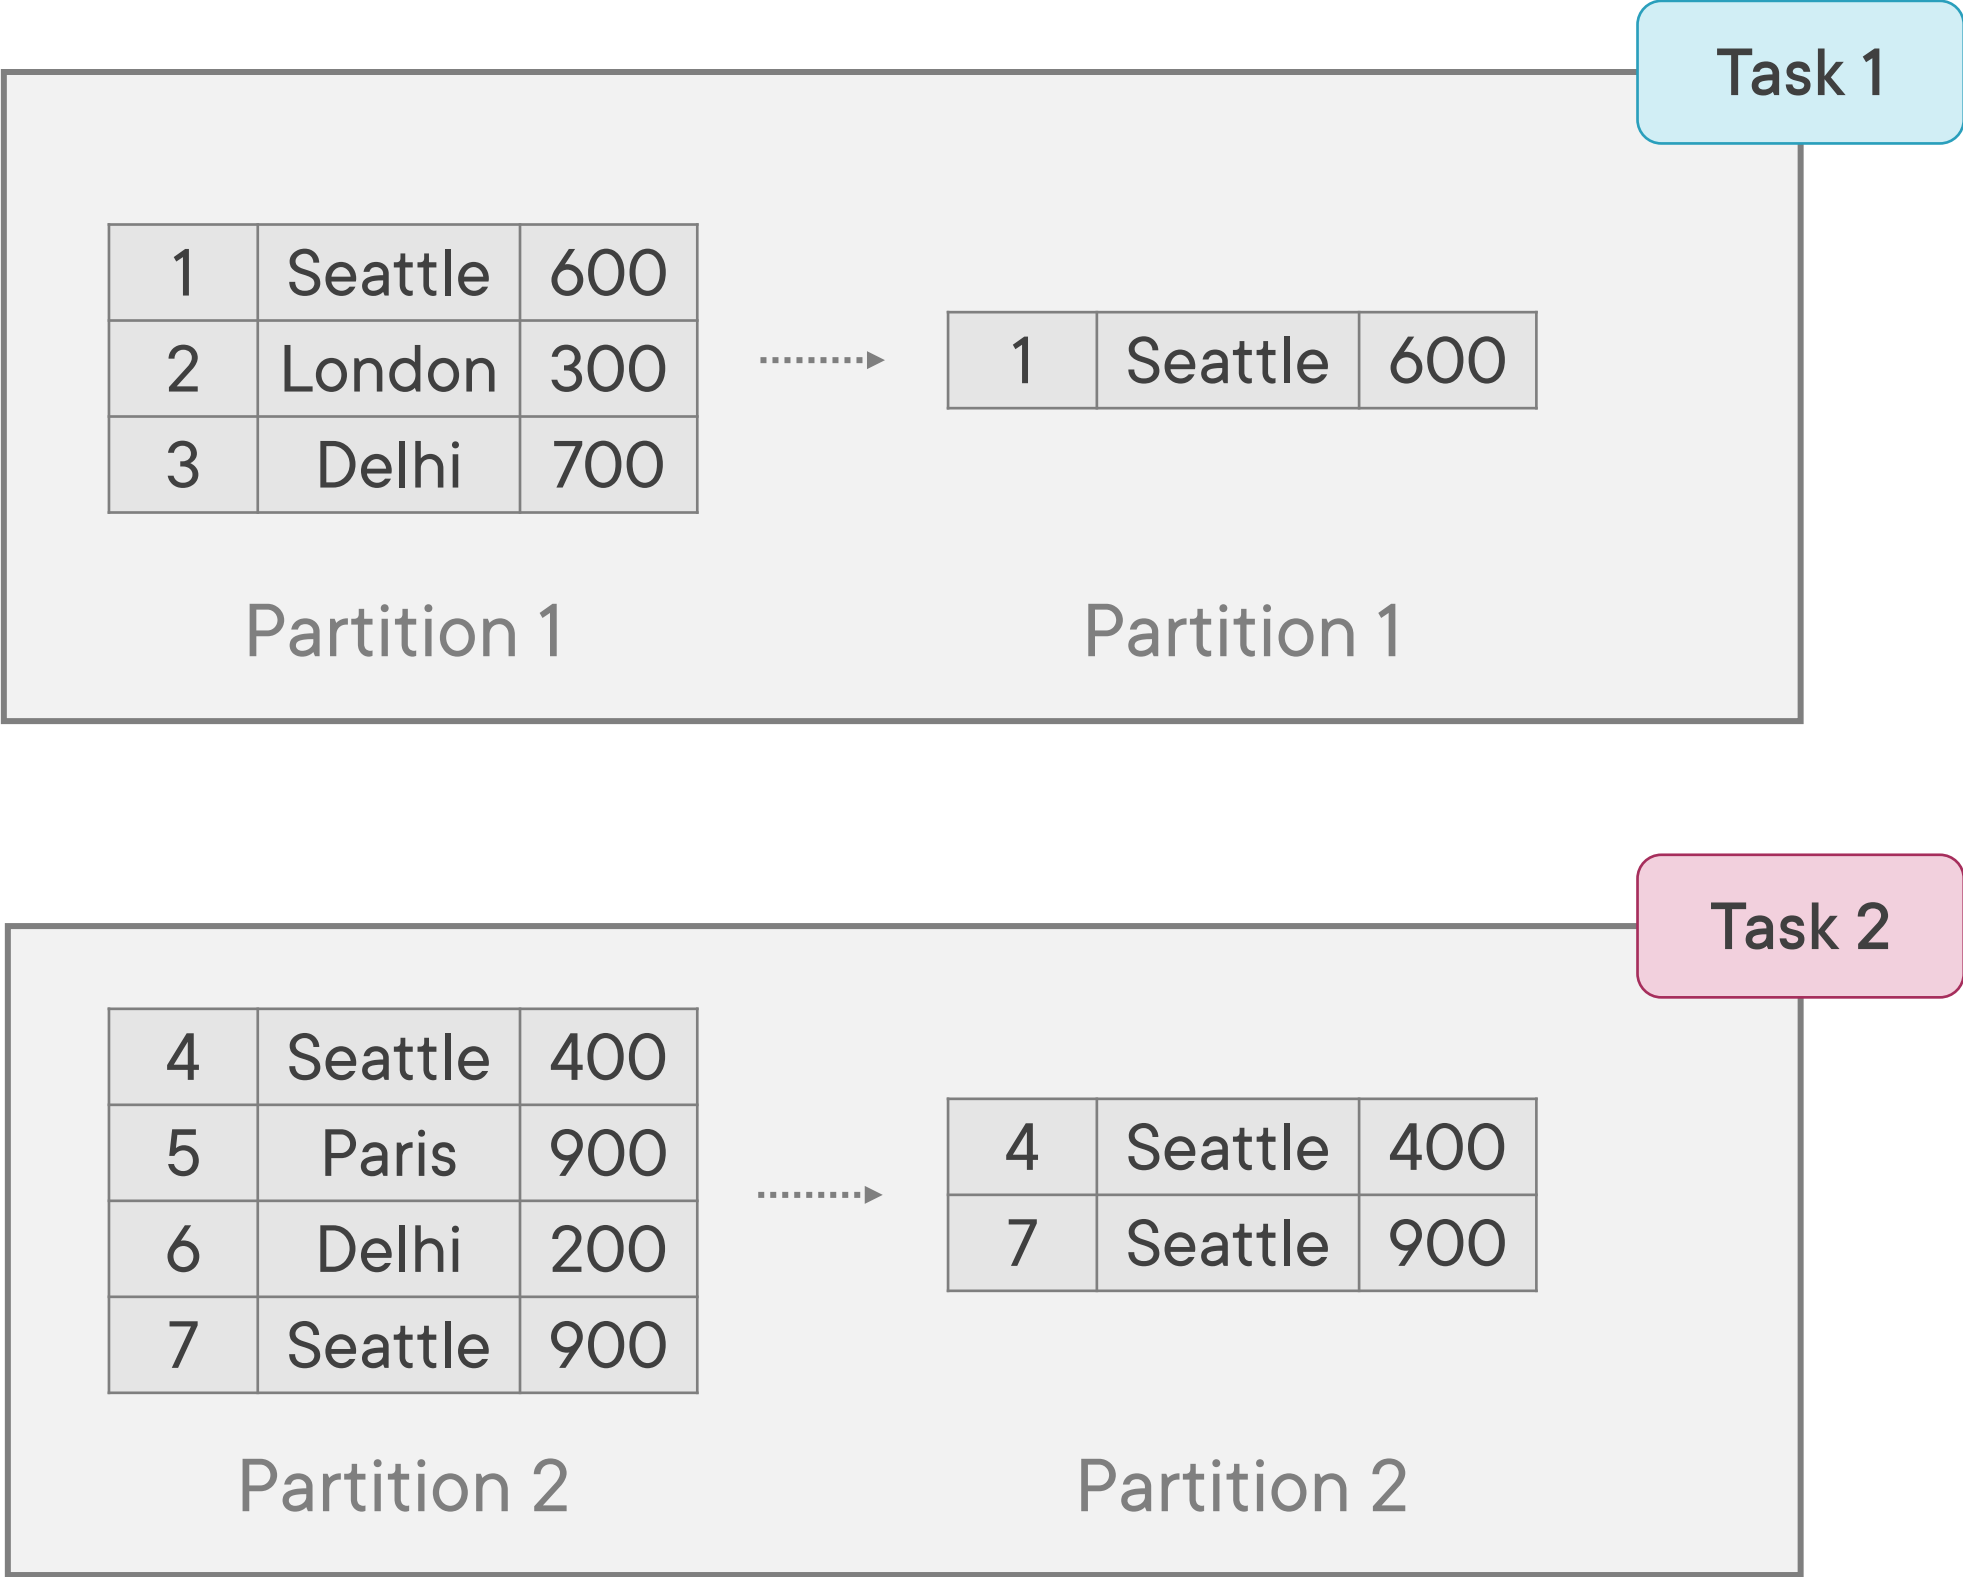
Filter is a Narrow Transformation

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Filter  
City = 'Seattle'  
**RDD 2**



**Filter is a Narrow Transformation**

1	Seattle	600
4	Seattle	400
7	Seattle	900

Read file,  
split by comma  
RDD 1

Id, City, Amount
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Task 1

Task 2

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Use map() to derive  
annual amount  
**RDD 2**

Task 1

1	Seattle	600
2	London	300
3	Delhi	700

Partition 1

Task 2

4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900

Partition 2



<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Use map() to derive  
annual amount  
**RDD 2**

Task 1

1	Seattle	600
2	London	300
3	Delhi	700



1	Seattle	600	7200
2	London	300	3600
3	Delhi	700	8400

Partition 1

Partition 1

Task 2

4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900



4	Seattle	400	4800
5	Paris	900	10800
6	Delhi	200	2400
7	Seattle	900	10800

Partition 2

Partition 2

Map is a Narrow Transformation

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Use map() to derive  
annual amount  
**RDD 2**

Task 1

1	Seattle	600
2	London	300
3	Delhi	700



1	Seattle	600	7200
2	London	300	3600
3	Delhi	700	8400

Partition 1

Partition 1

Task 2

4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900



4	Seattle	400	4800
5	Paris	900	10800
6	Delhi	200	2400
7	Seattle	900	10800

Partition 2

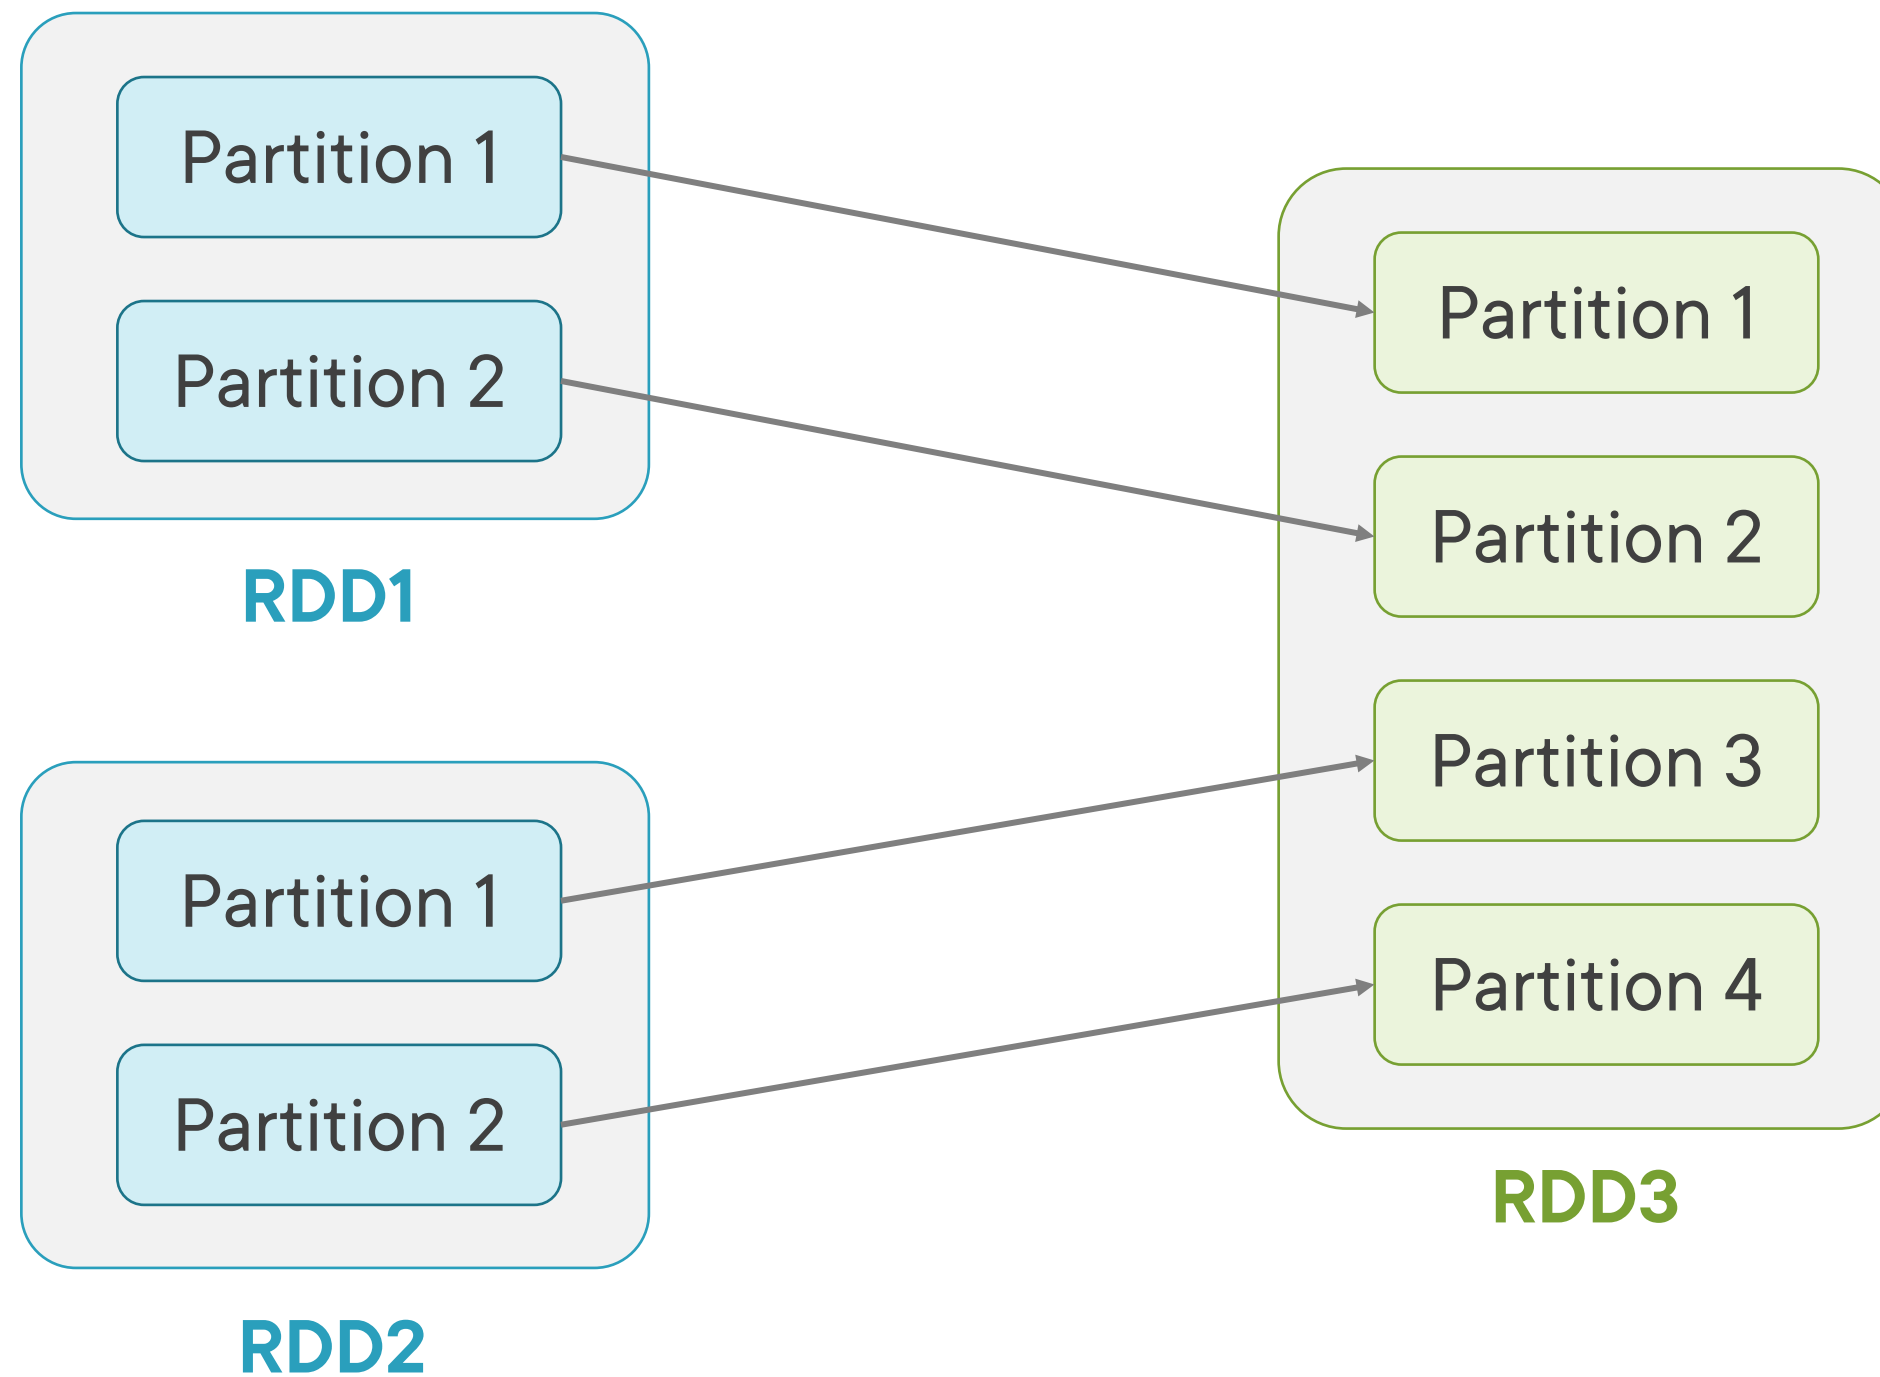
Partition 2

Map is a Narrow Transformation

1	Seattle	600	7200
2	London	300	3600
3	Delhi	700	8400


4	Seattle	400	4800
5	Paris	900	10800
6	Delhi	200	2400
7	Seattle	900	10800

Each input partition is used **at-most once** to produce output partitions



**Union operation is an example of this**

# Narrow Transformation



**Extremely fast**

**No data movement between partitions / No shuffling**

**Examples**

- Filter, Map, FlatMap, MapPartition, Sample, Union etc.

# Wide Transformations and Data Shuffling

---

Read file,  
split by comma  
RDD 1

Id, City, Amount
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Task 1

Task 2

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Group by City,  
calculate count  
**RDD 2**

Task 1

1	Seattle	600
2	London	300
3	Delhi	700

Partition 1

Task 2

4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900

Partition 2

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
split by comma  
**RDD 1**



Group by City,  
calculate count  
**RDD 2**

Task 1

1	Seattle	600
2	London	300
3	Delhi	700

Partition 1

.....>

Seattle	1
London	1
Delhi	1

Partition 1

Task 2

4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900

Partition 2

.....>

Seattle	2
Paris	1
Delhi	1

Partition 2



Id, City, Amount		
1	Seattle	600
2	London	300
3	Delhi	700
4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900

Read file,  
split by comma  
**RDD 1**



Group by City,  
calculate count  
**RDD 2**

Task 1

1	Seattle	600
2	London	300
3	Delhi	700



Seattle	1
London	1
Delhi	1

Partition 1

Partition 1

Task 2

4	Seattle	400
5	Paris	900
6	Delhi	200
7	Seattle	900



Seattle	2
Paris	1
Delhi	1

Partition 2

Partition 2

**Incorrect Output!**

Data is not  
completely grouped

Seattle	1
---------	---

London	1
--------	---

Delhi	1
-------	---

Seattle	2
---------	---

Paris	1
-------	---

Delhi	1
-------	---

# Transformation Operations

**Narrow Dependency  
Transformation**

**Wide Dependency  
Transformation**

Wide Transformation is a **two-step** process and requires **shuffling** of data

<b>Id, City, Amount</b>
1, Seattle, 600
2, London, 300
3, Delhi, 700
4, Seattle, 400
5, Paris, 900
6, Delhi, 200
7, Seattle, 900

Read file,  
get City & Amount  
RDD 1

Seattle	600
London	300
Delhi	700

Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700


Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

.....→  
Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange



Shuffle Write

Read file,  
get City & Amount  
RDD 1

Seattle	600
London	300
Delhi	700

Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
RDD 2

Step 1 – Group: Shuffle / Exchange



Seattle	600
London	300
Delhi	700

Shuffle Block 1

Shuffle Write

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

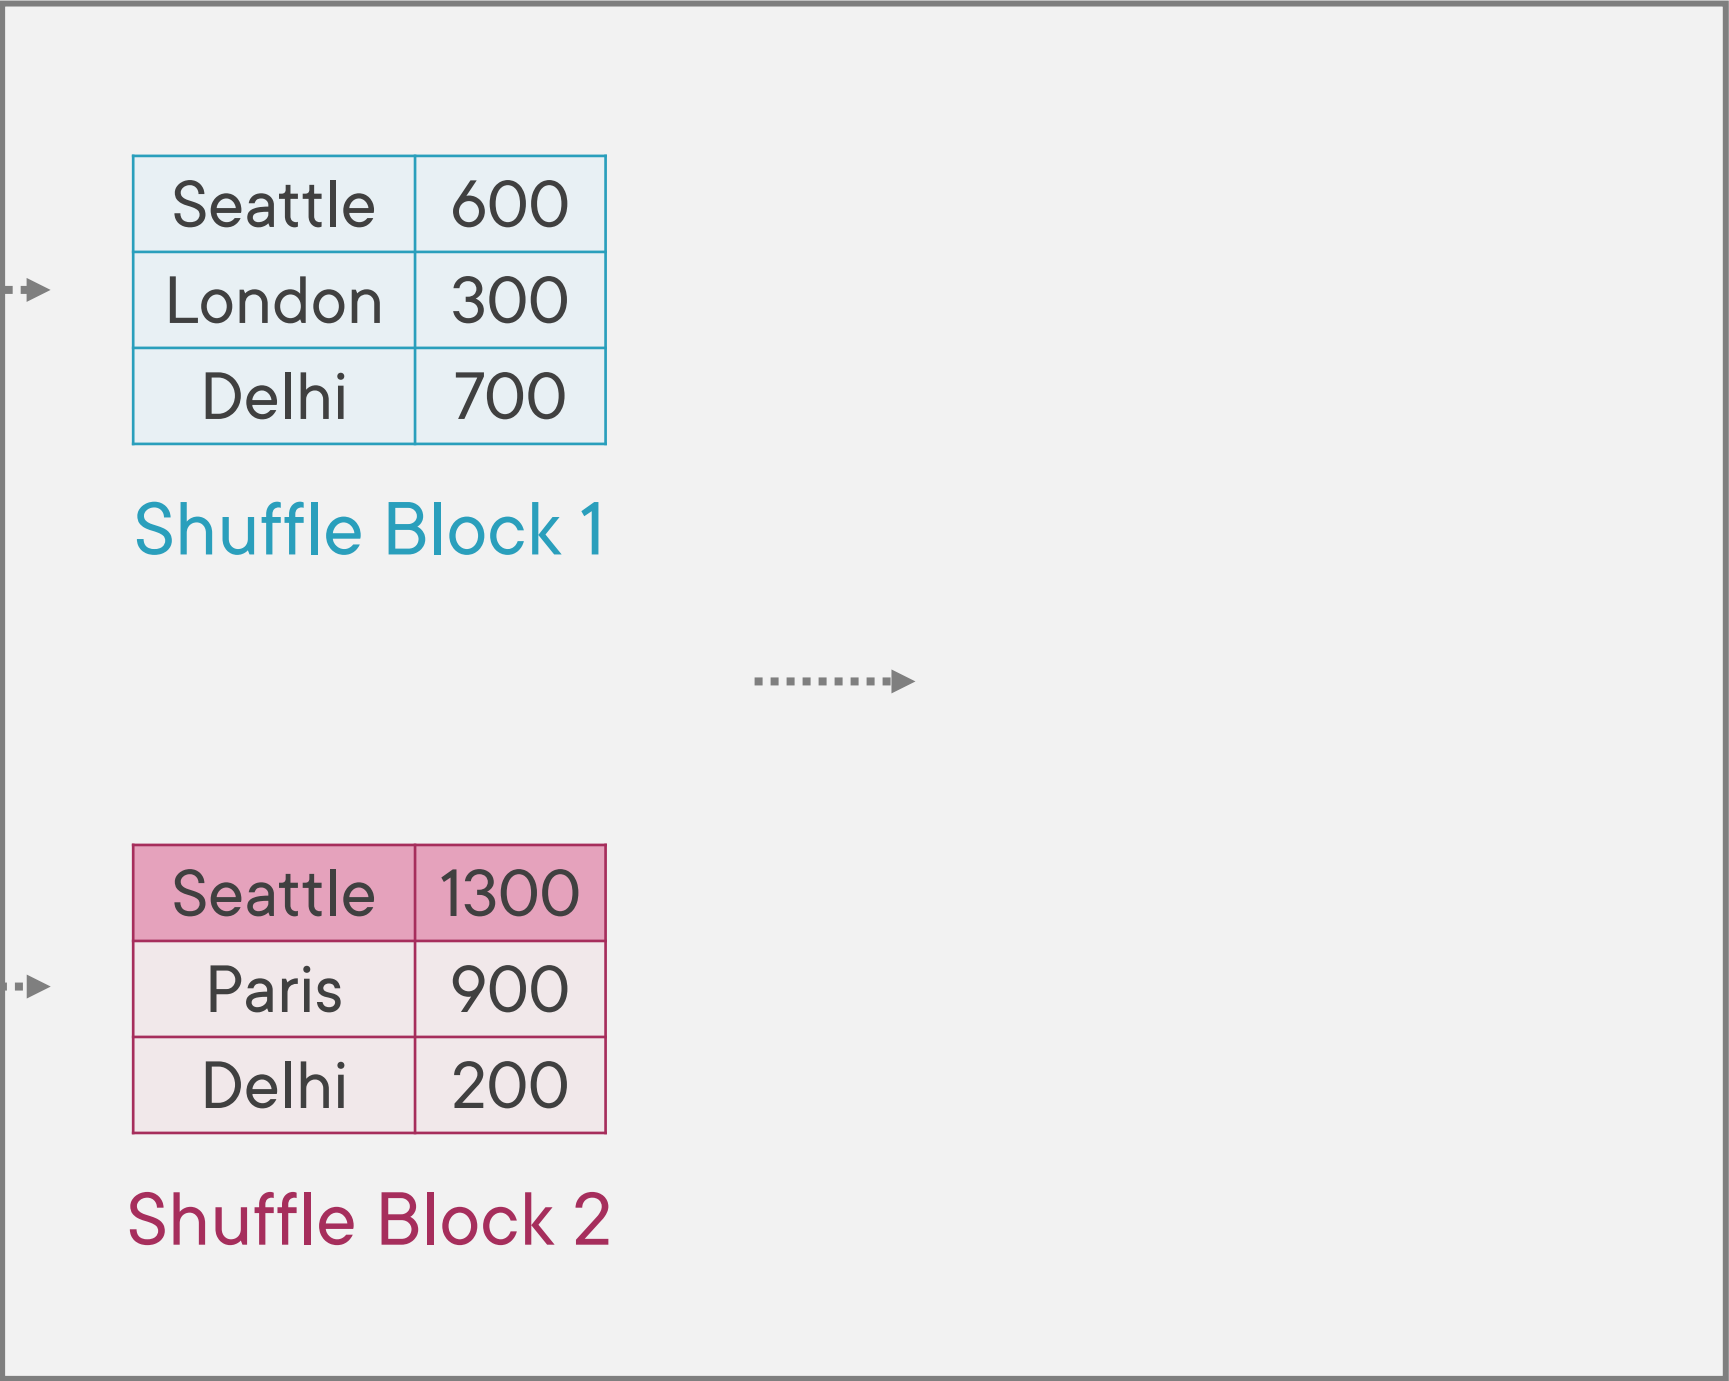
Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange



Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

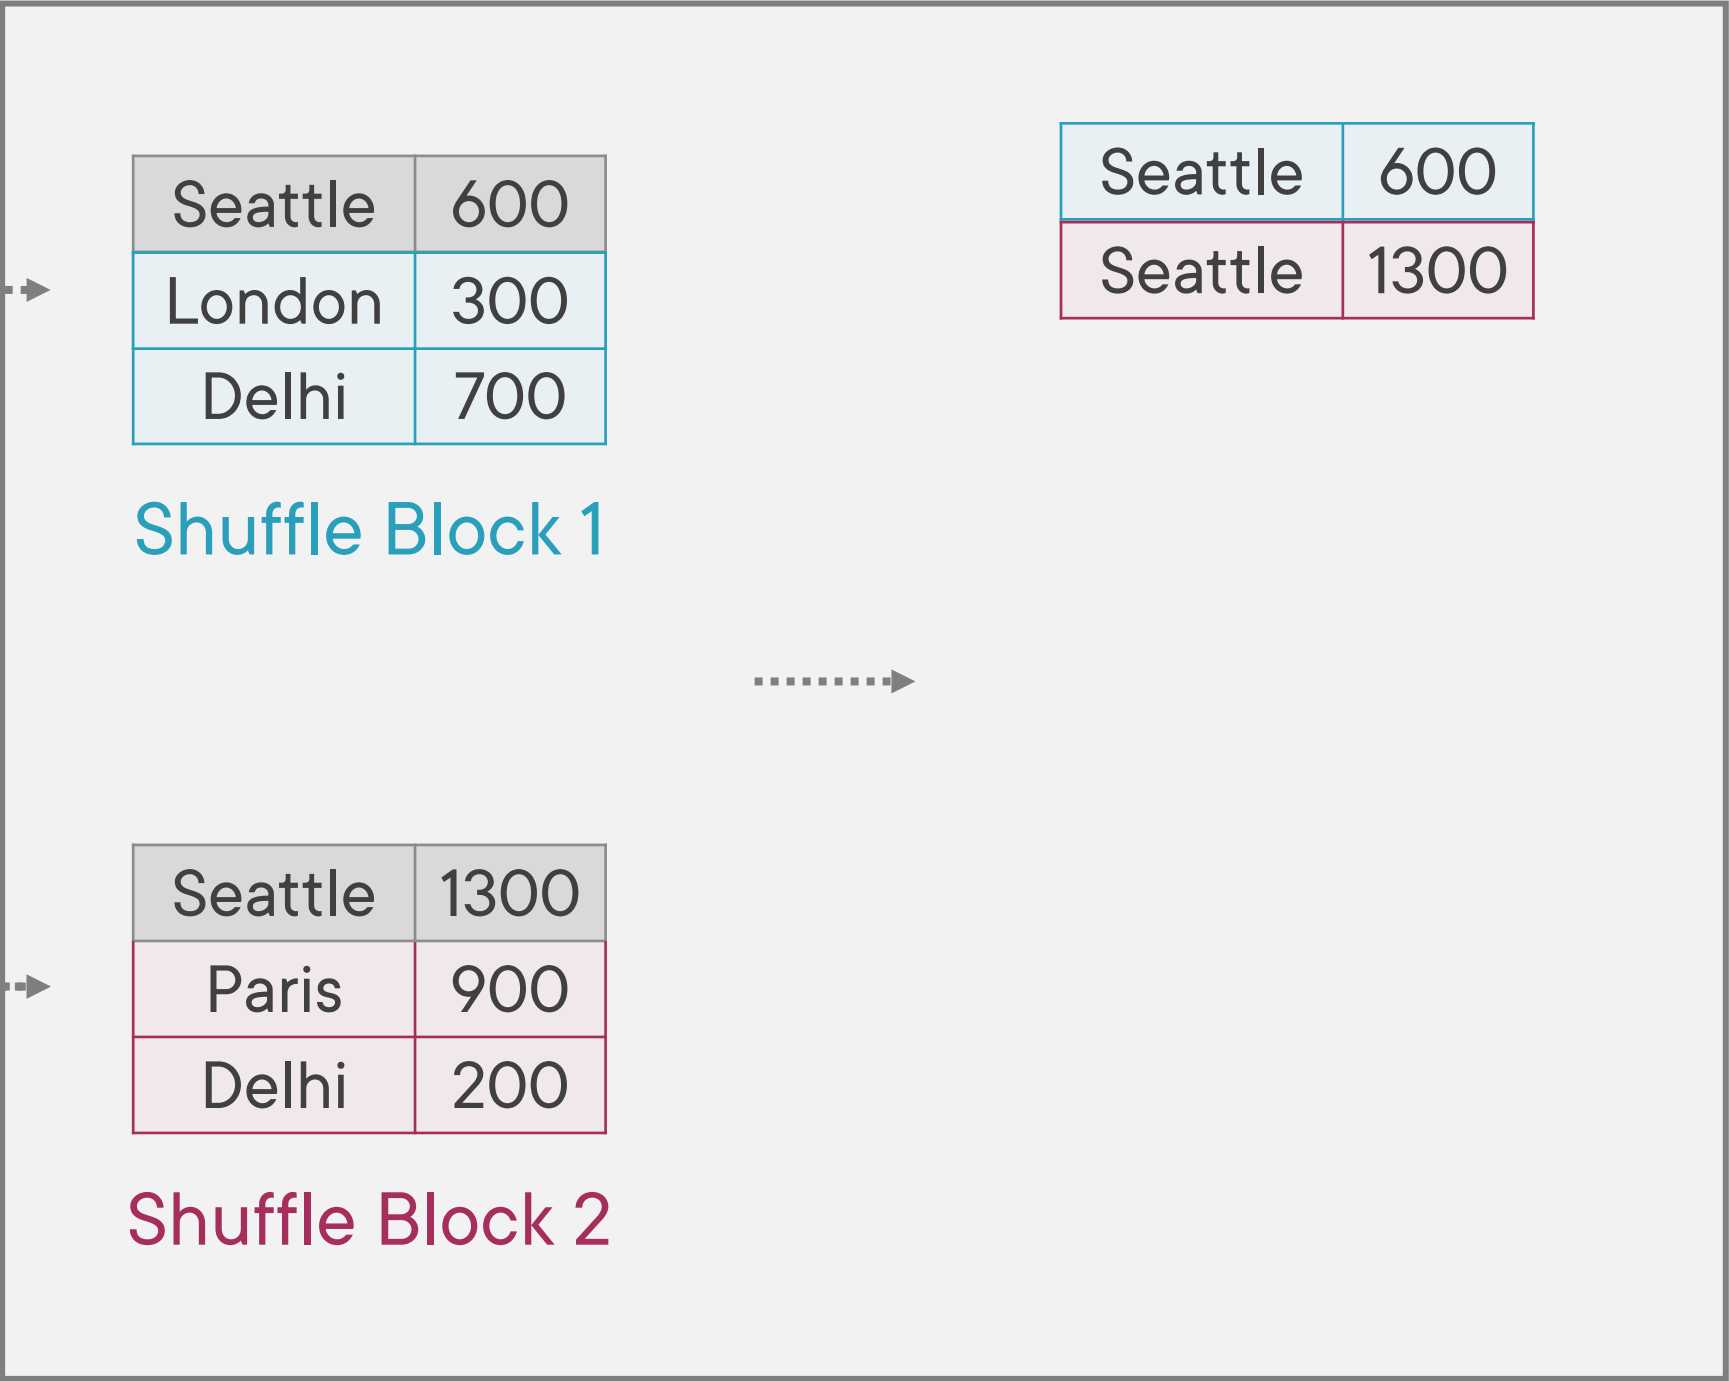
Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange



Seattle	600
Seattle	1300

Shuffle Block 1

Seattle	1300
Paris	900
Delhi	200

Shuffle Block 2

Shuffle Write

Shuffle Read



Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

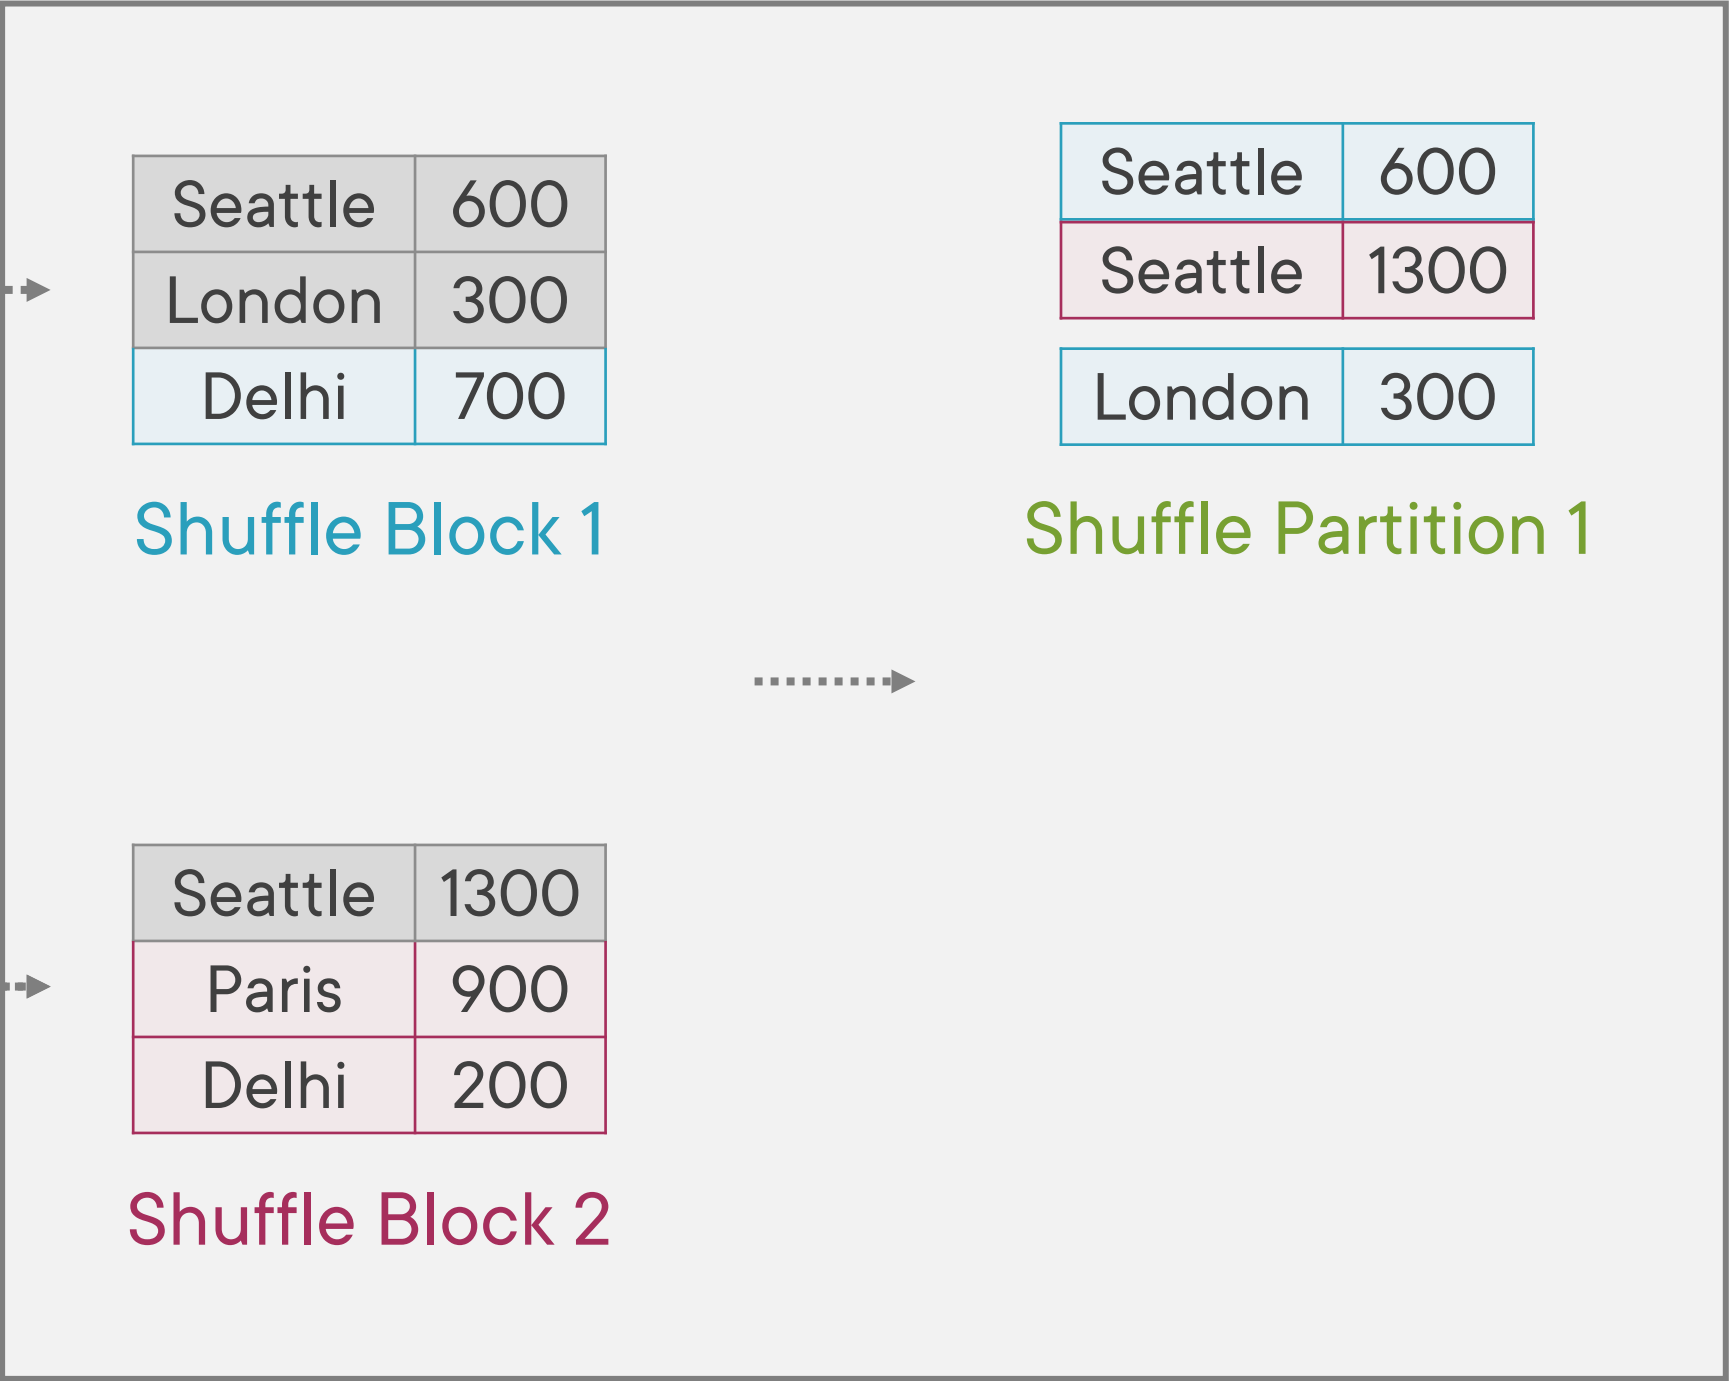
Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange



Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

Partition 1

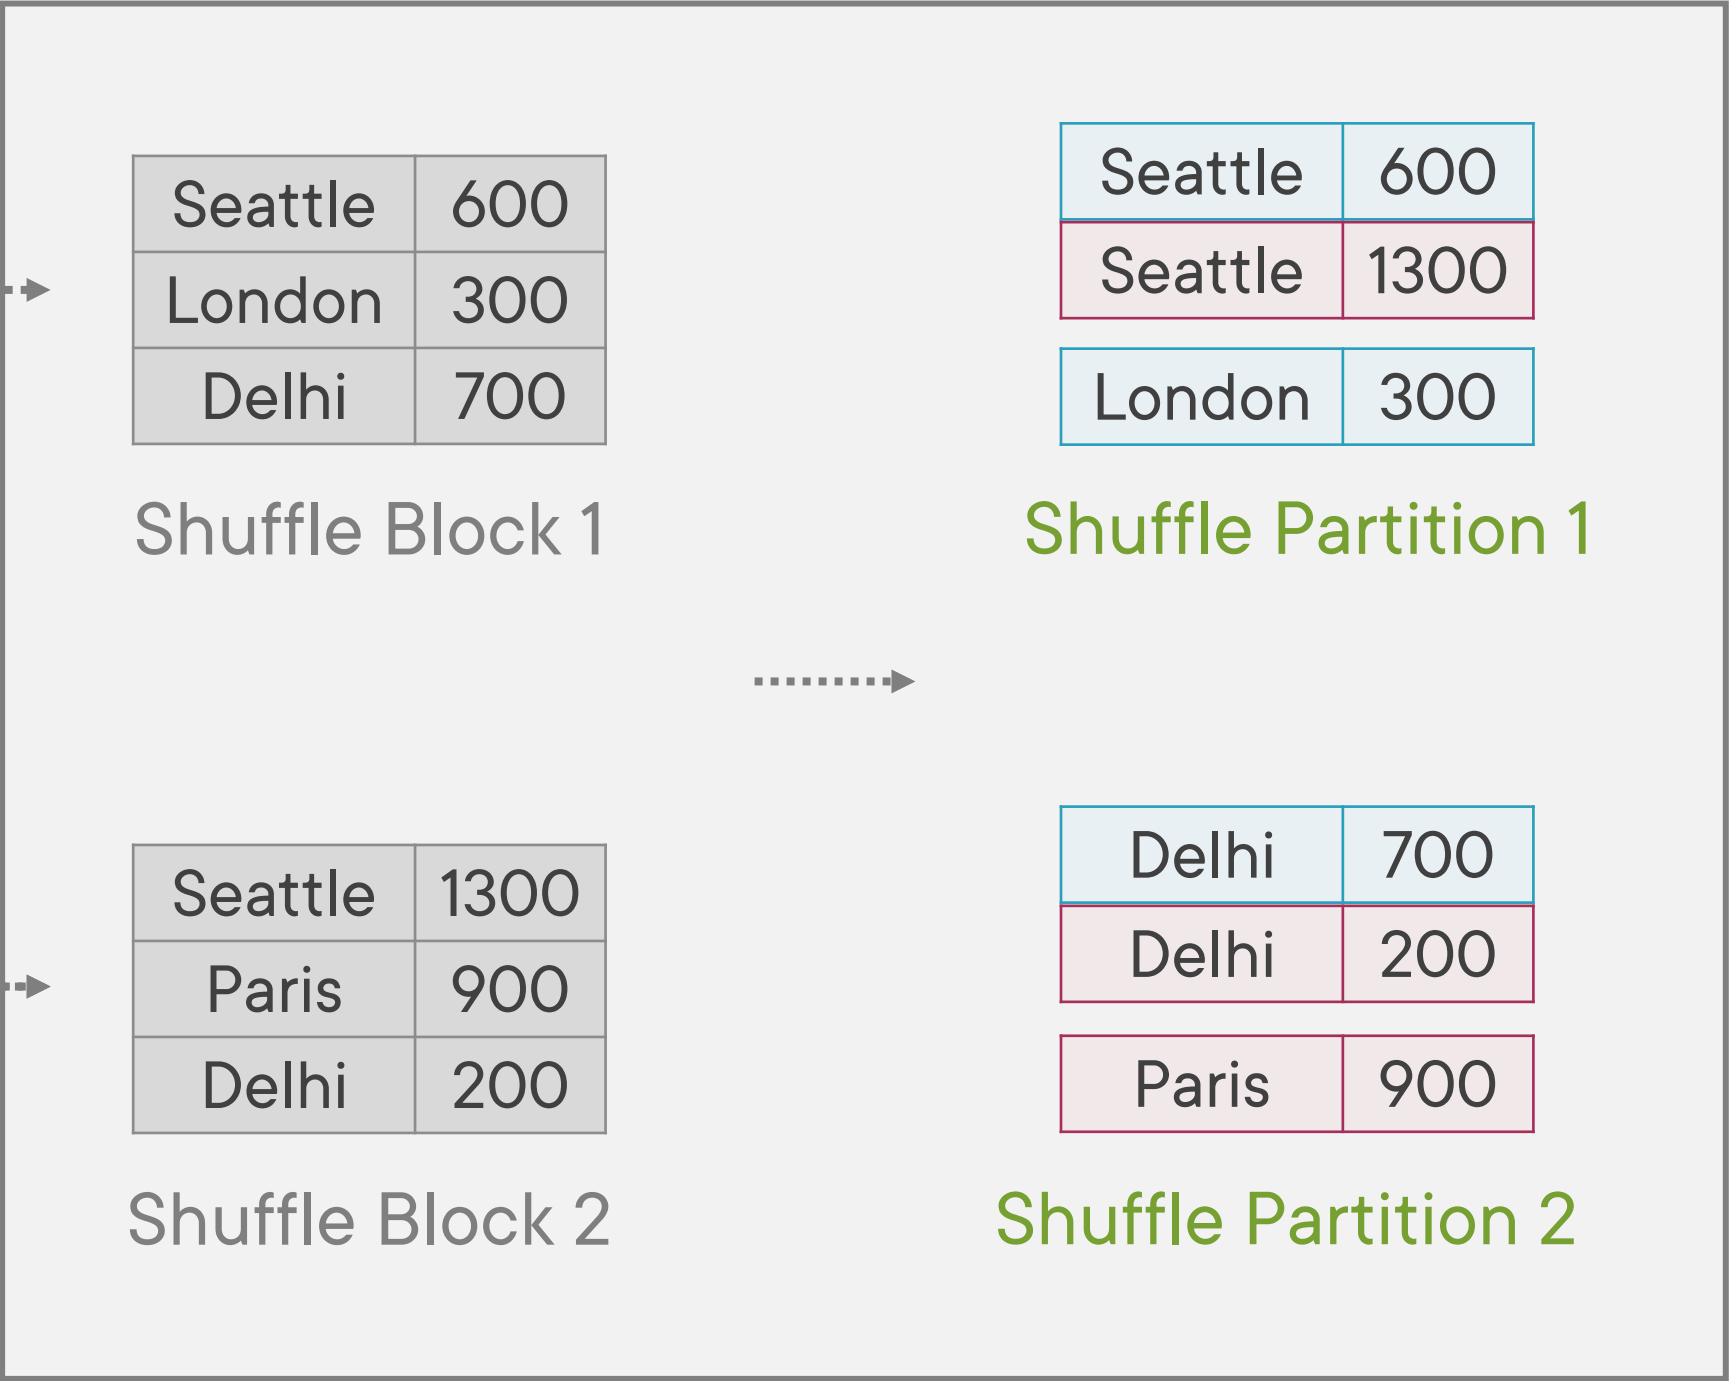
Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange

Step 2 – Sum



Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700


Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange



Step 2 – Sum

Seattle	1900
London	300

Partition 1

Delhi	900
Paris	900

Partition 2

Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

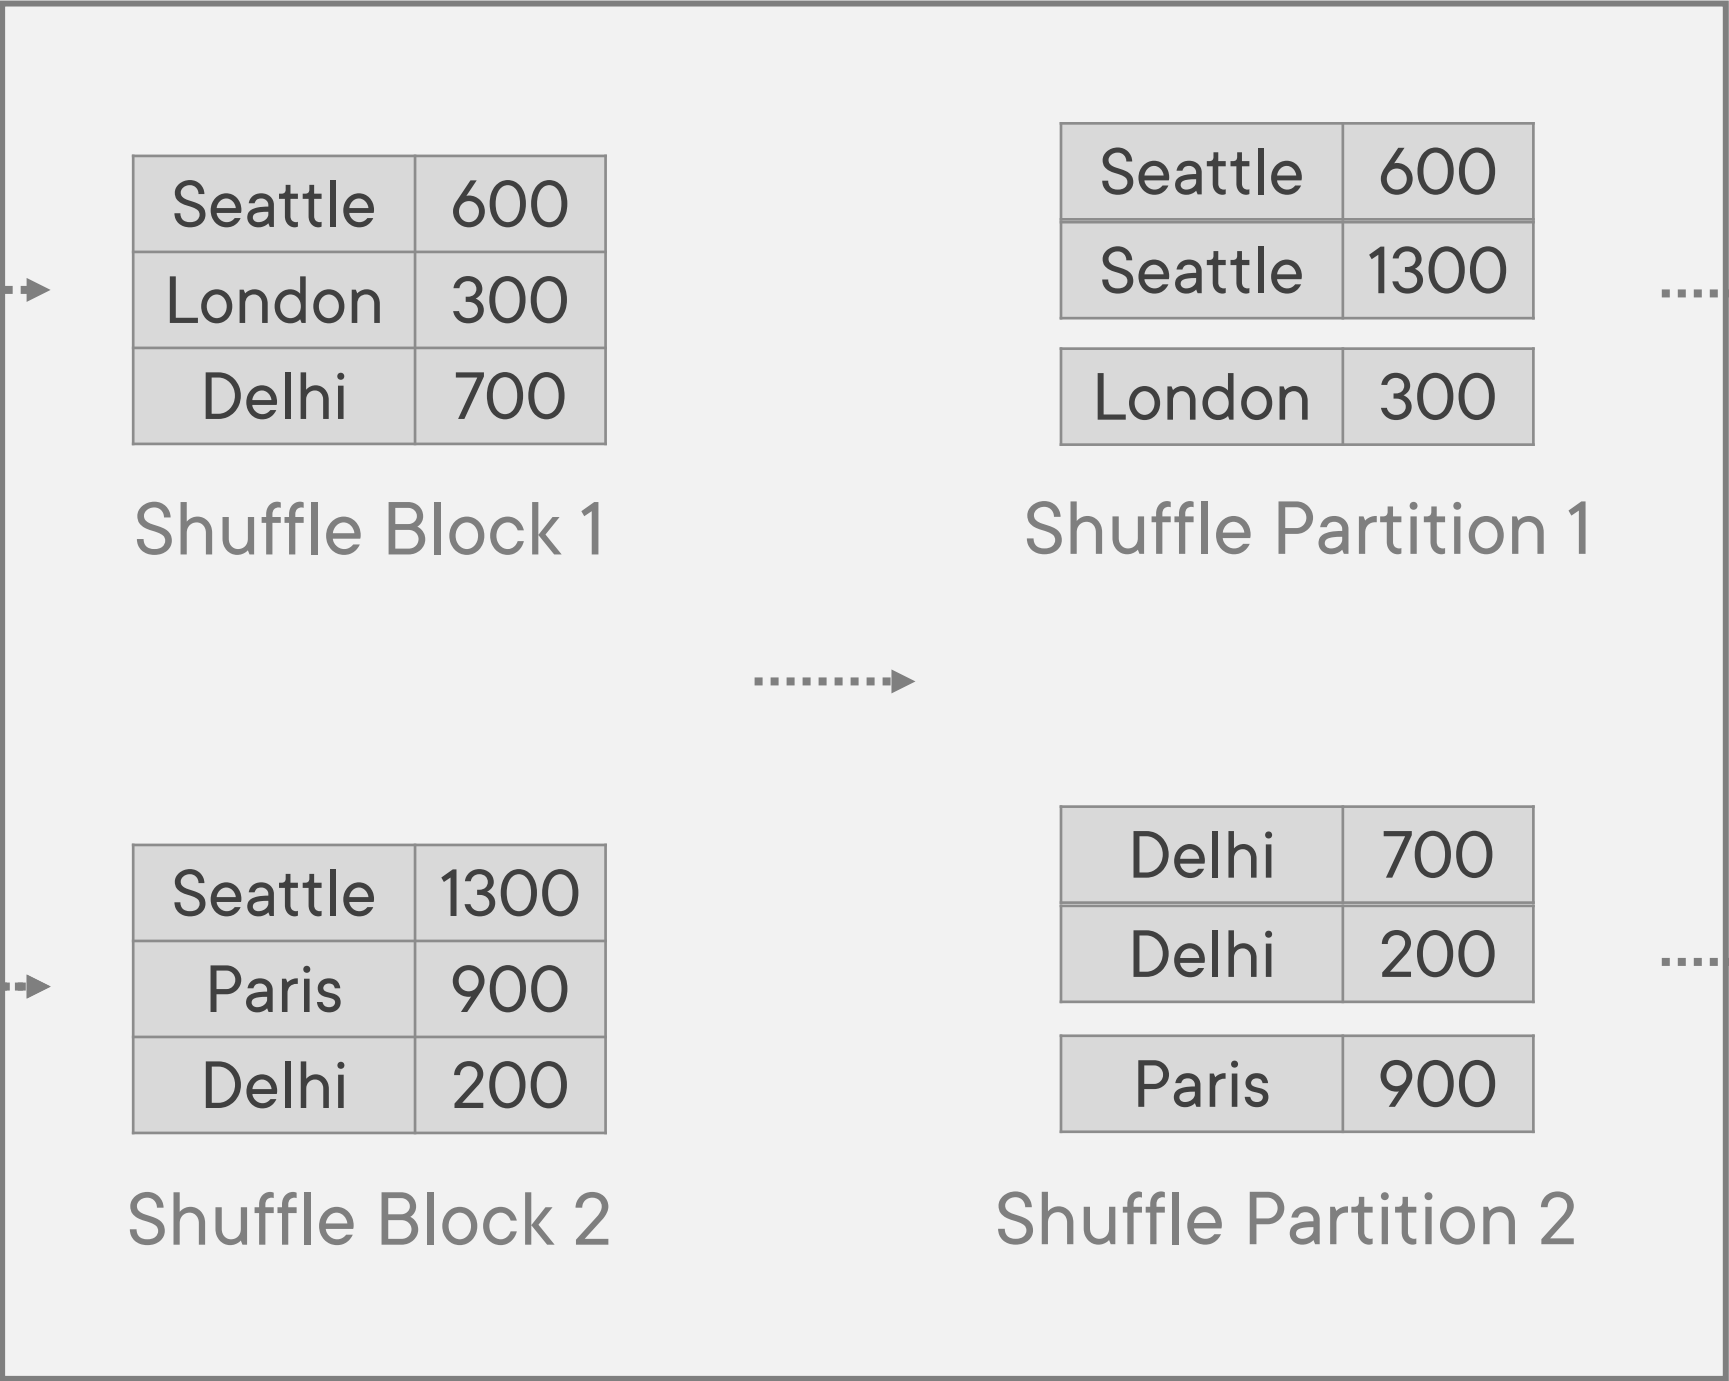
Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange



Step 2 – Sum

Seattle	1900
London	300

Partition 1

Delhi	900
Paris	900

Partition 2

**Correct  
Output!**

Seattle	1900
London	300
Delhi	900
Paris	900

Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange

Step 2 – Sum

Task 1

Seattle	600
London	300
Delhi	700

Partition 1

.....→

Seattle	600
London	300
Delhi	700

Shuffle Block 1

Task 2

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

.....→

Seattle	1300
Paris	900
Delhi	200

Shuffle Block 2

Seattle	600
Seattle	1300
London	300

Shuffle Partition 1

Seattle	1900
London	300

Partition 1

Delhi	700
Delhi	200
Paris	900

Shuffle Partition 2

Delhi	900
Paris	900

Partition 2

Shuffle Write

Shuffle Read

Aggregate

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Group by City, calculate sum of Amount  
**RDD 2**

Step 1 – Group: Shuffle / Exchange

Step 2 – Sum

Seattle	600
London	300
Delhi	700

Shuffle Block 1

Seattle	1300
Paris	900
Delhi	200

Shuffle Block 2

Task 3

Seattle	600
Seattle	1300
London	300

Seattle	1900
London	300

Shuffle Partition 1

Partition 1

Task 4

Delhi	700
Delhi	200
Paris	900

Delhi	900
Paris	900

Shuffle Partition 2

Partition 2

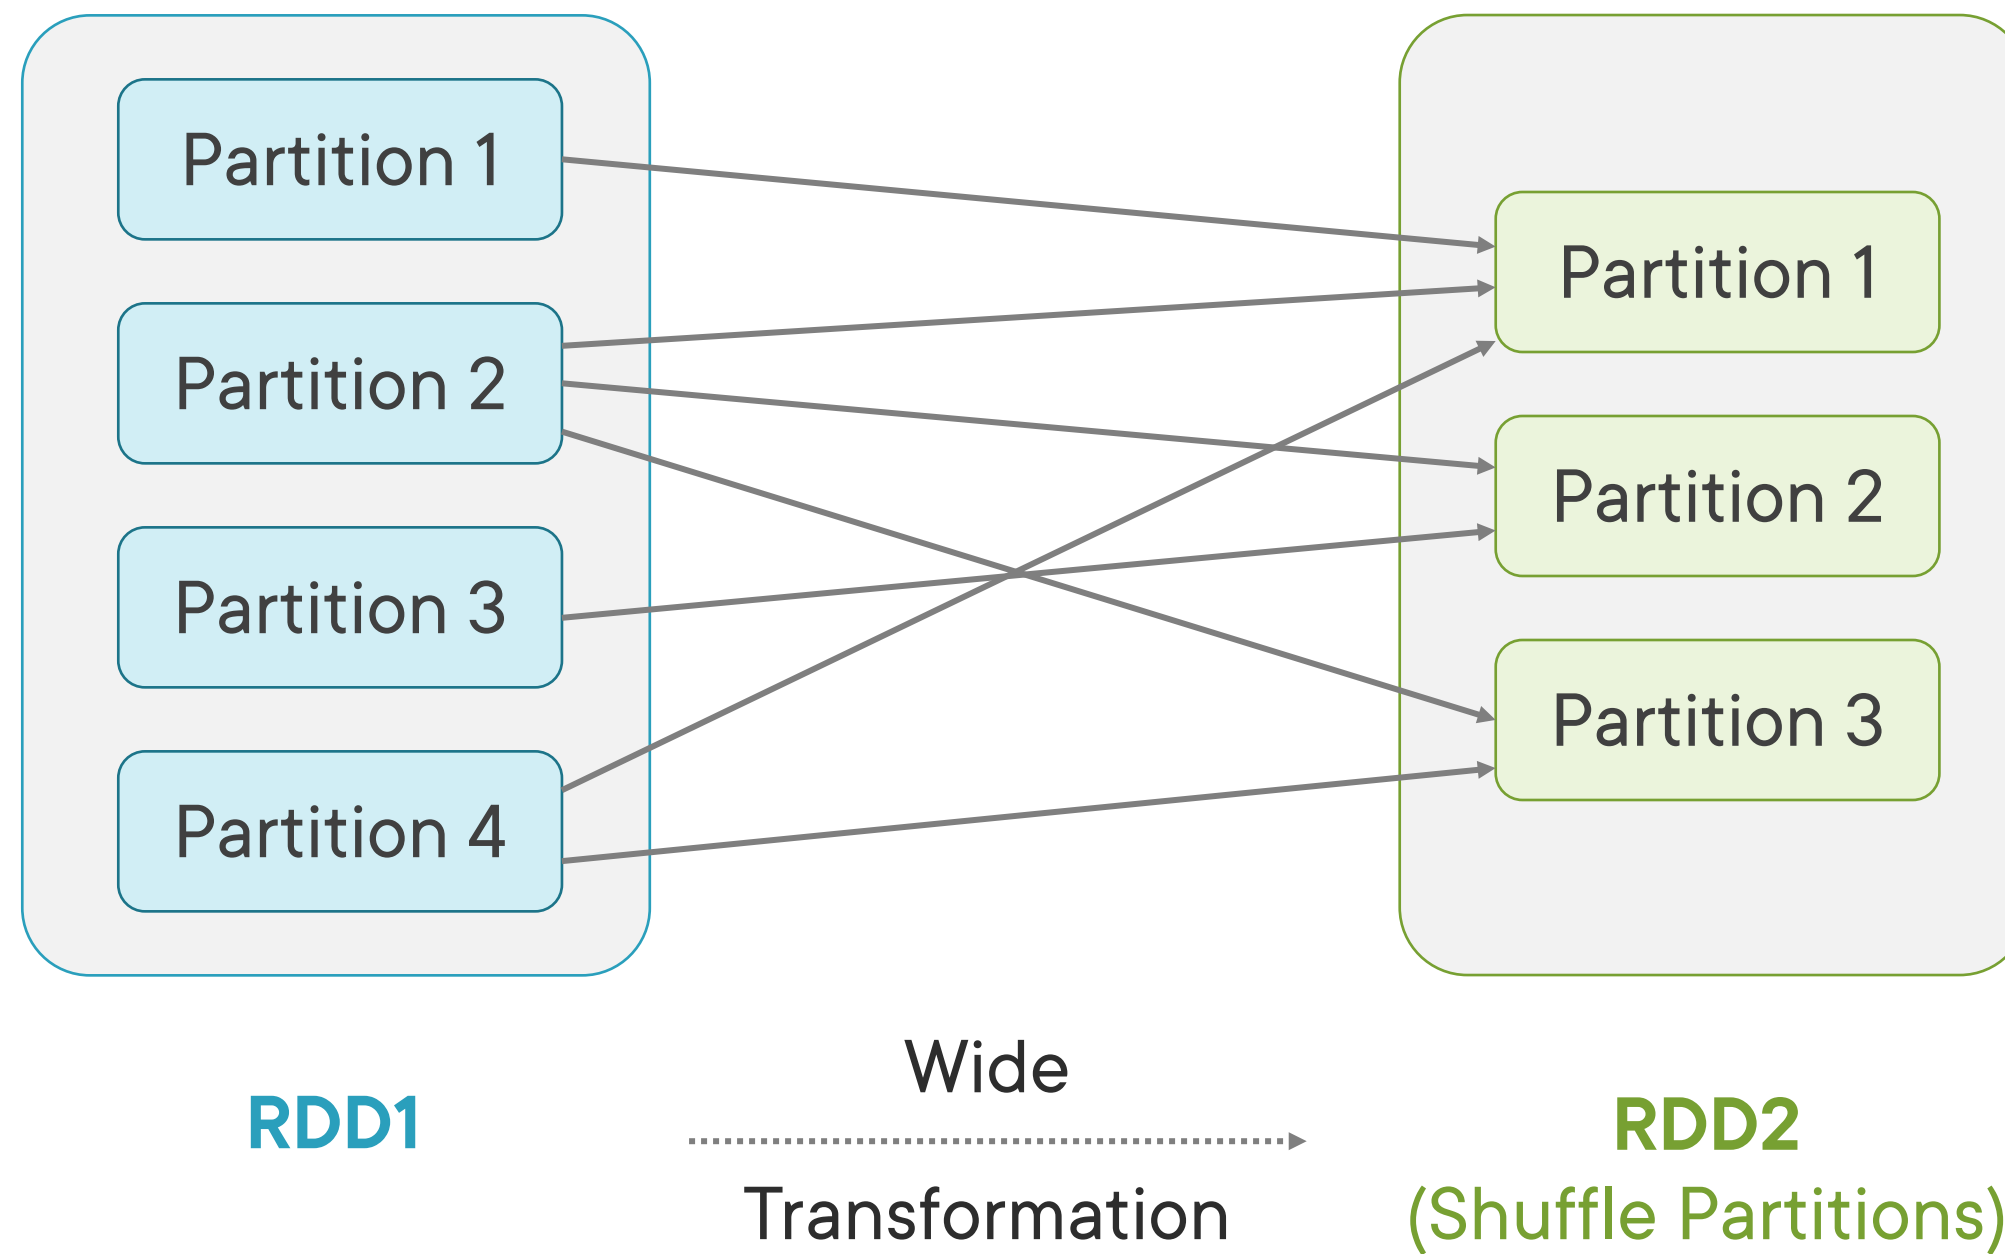
Shuffle Write

Shuffle Read

Aggregate

In Wide Transformation, **one input partition** might be used **multiple times** to produce output partitions

One **input partition** might be used **multiple times** to produce output partitions



**ReduceByKey  
& Distinct operations  
are examples of this**



Read file,  
get City & Amount  
RDD 1

Seattle	600
London	300
Delhi	700

Partition 1


Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2



Get distinct cities  
RDD 2

Step 1 – Group: Shuffle / Exchange



Shuffle Write

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

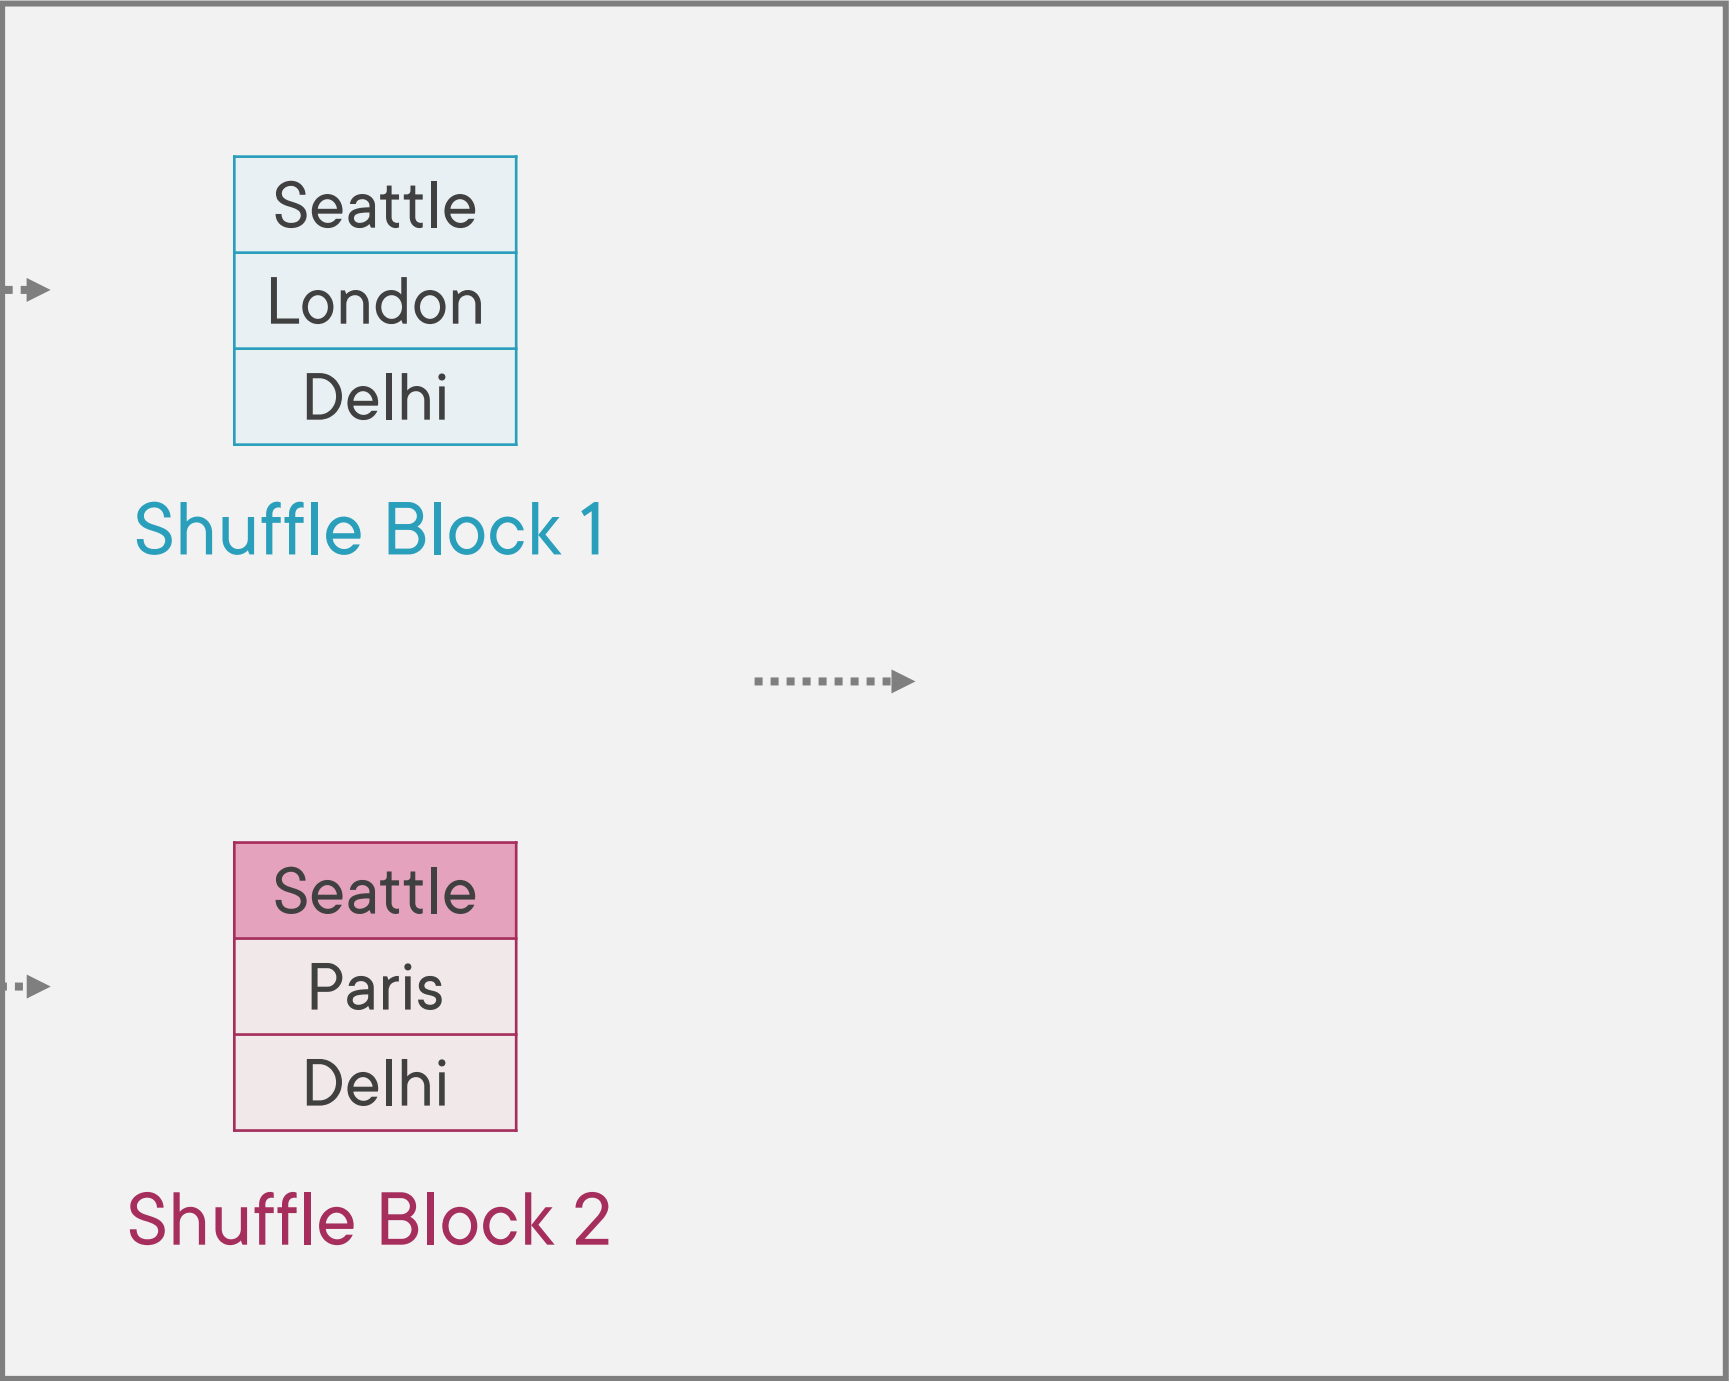
Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Get distinct cities  
**RDD 2**

Step 1 – Group: Shuffle / Exchange



Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Get distinct cities  
**RDD 2**

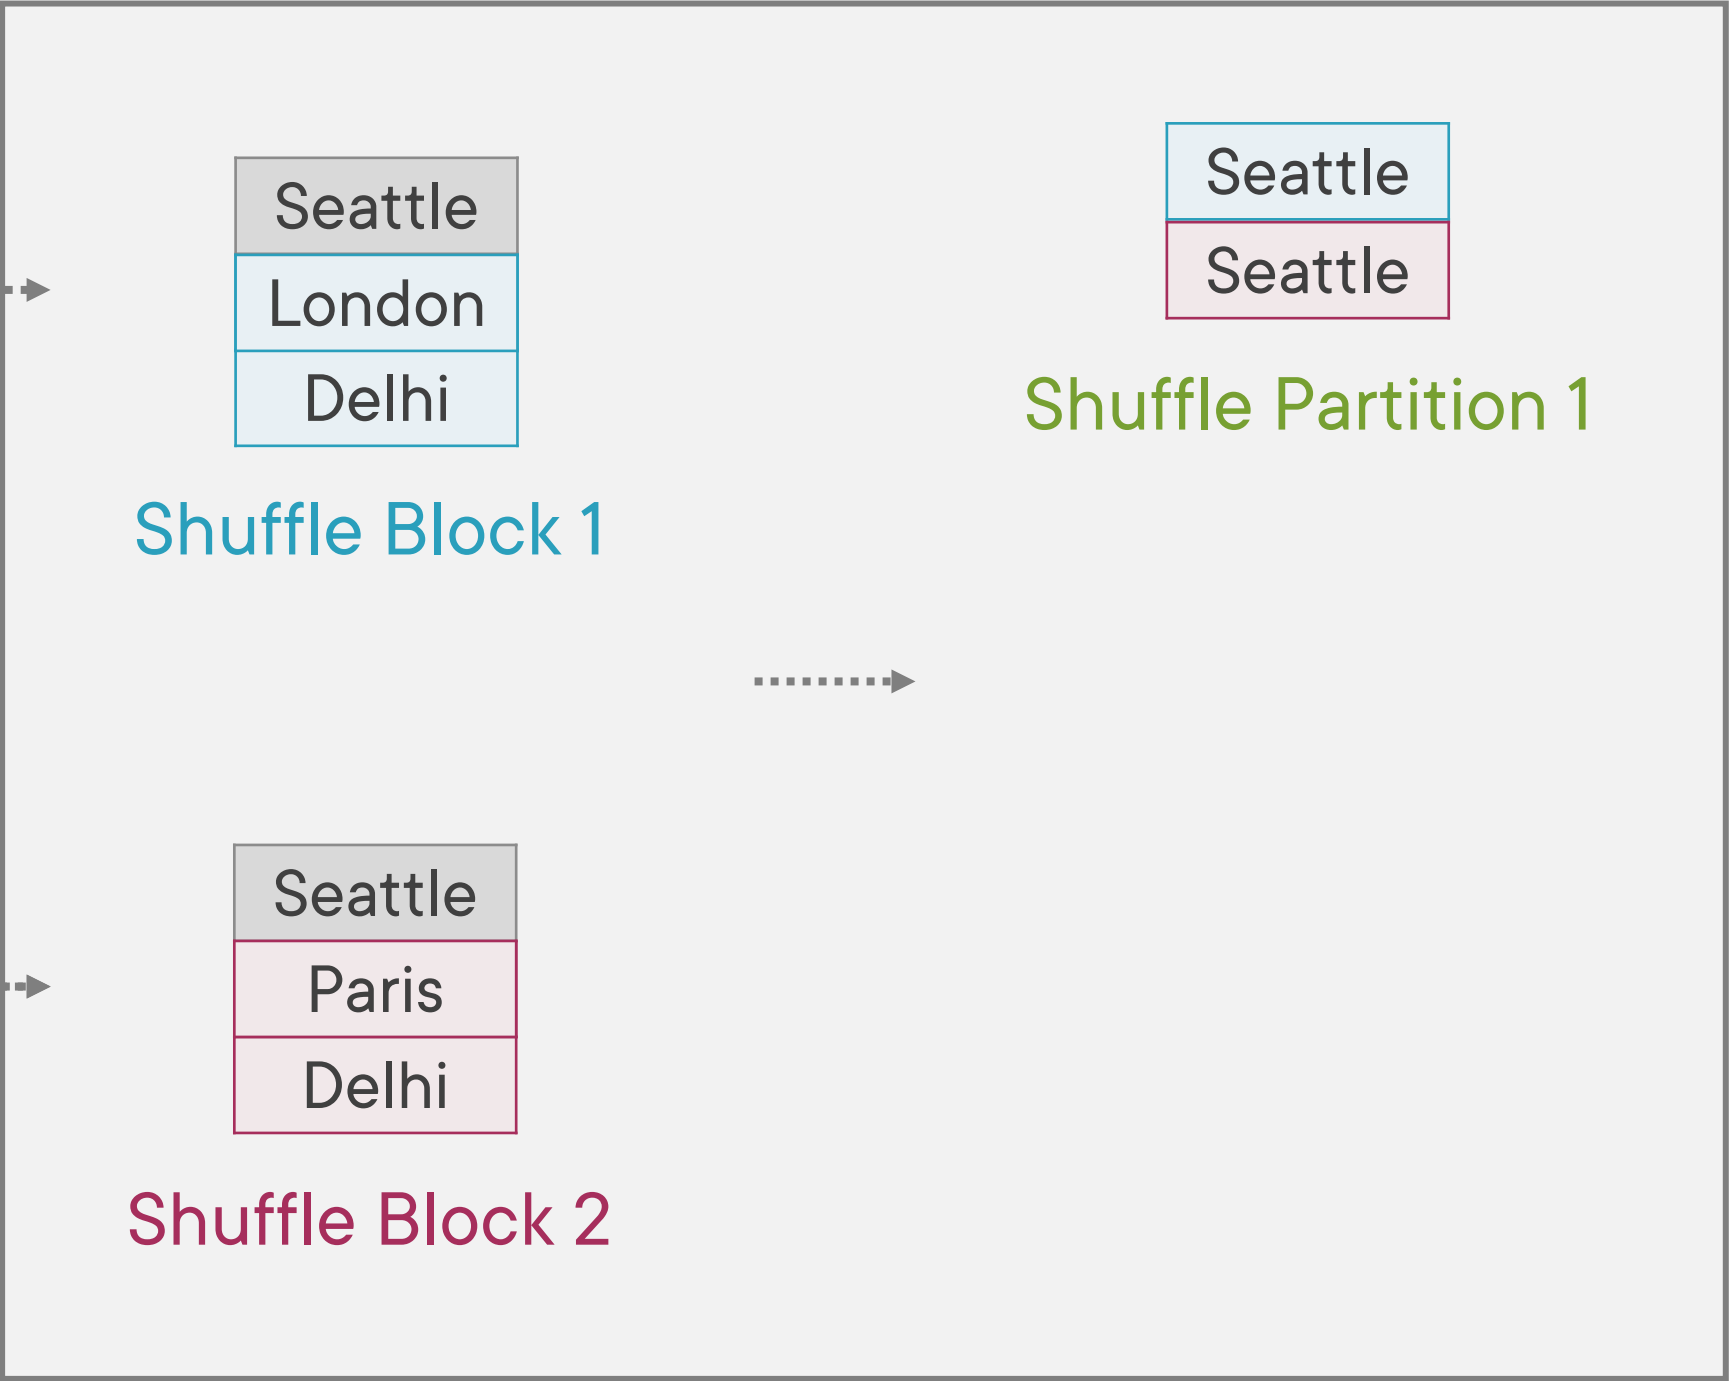
Step 1 – Group: Shuffle / Exchange

Seattle	600
London	300
Delhi	700

Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2



Read file,  
get City & Amount  
**RDD 1**

Get distinct cities  
**RDD 2**

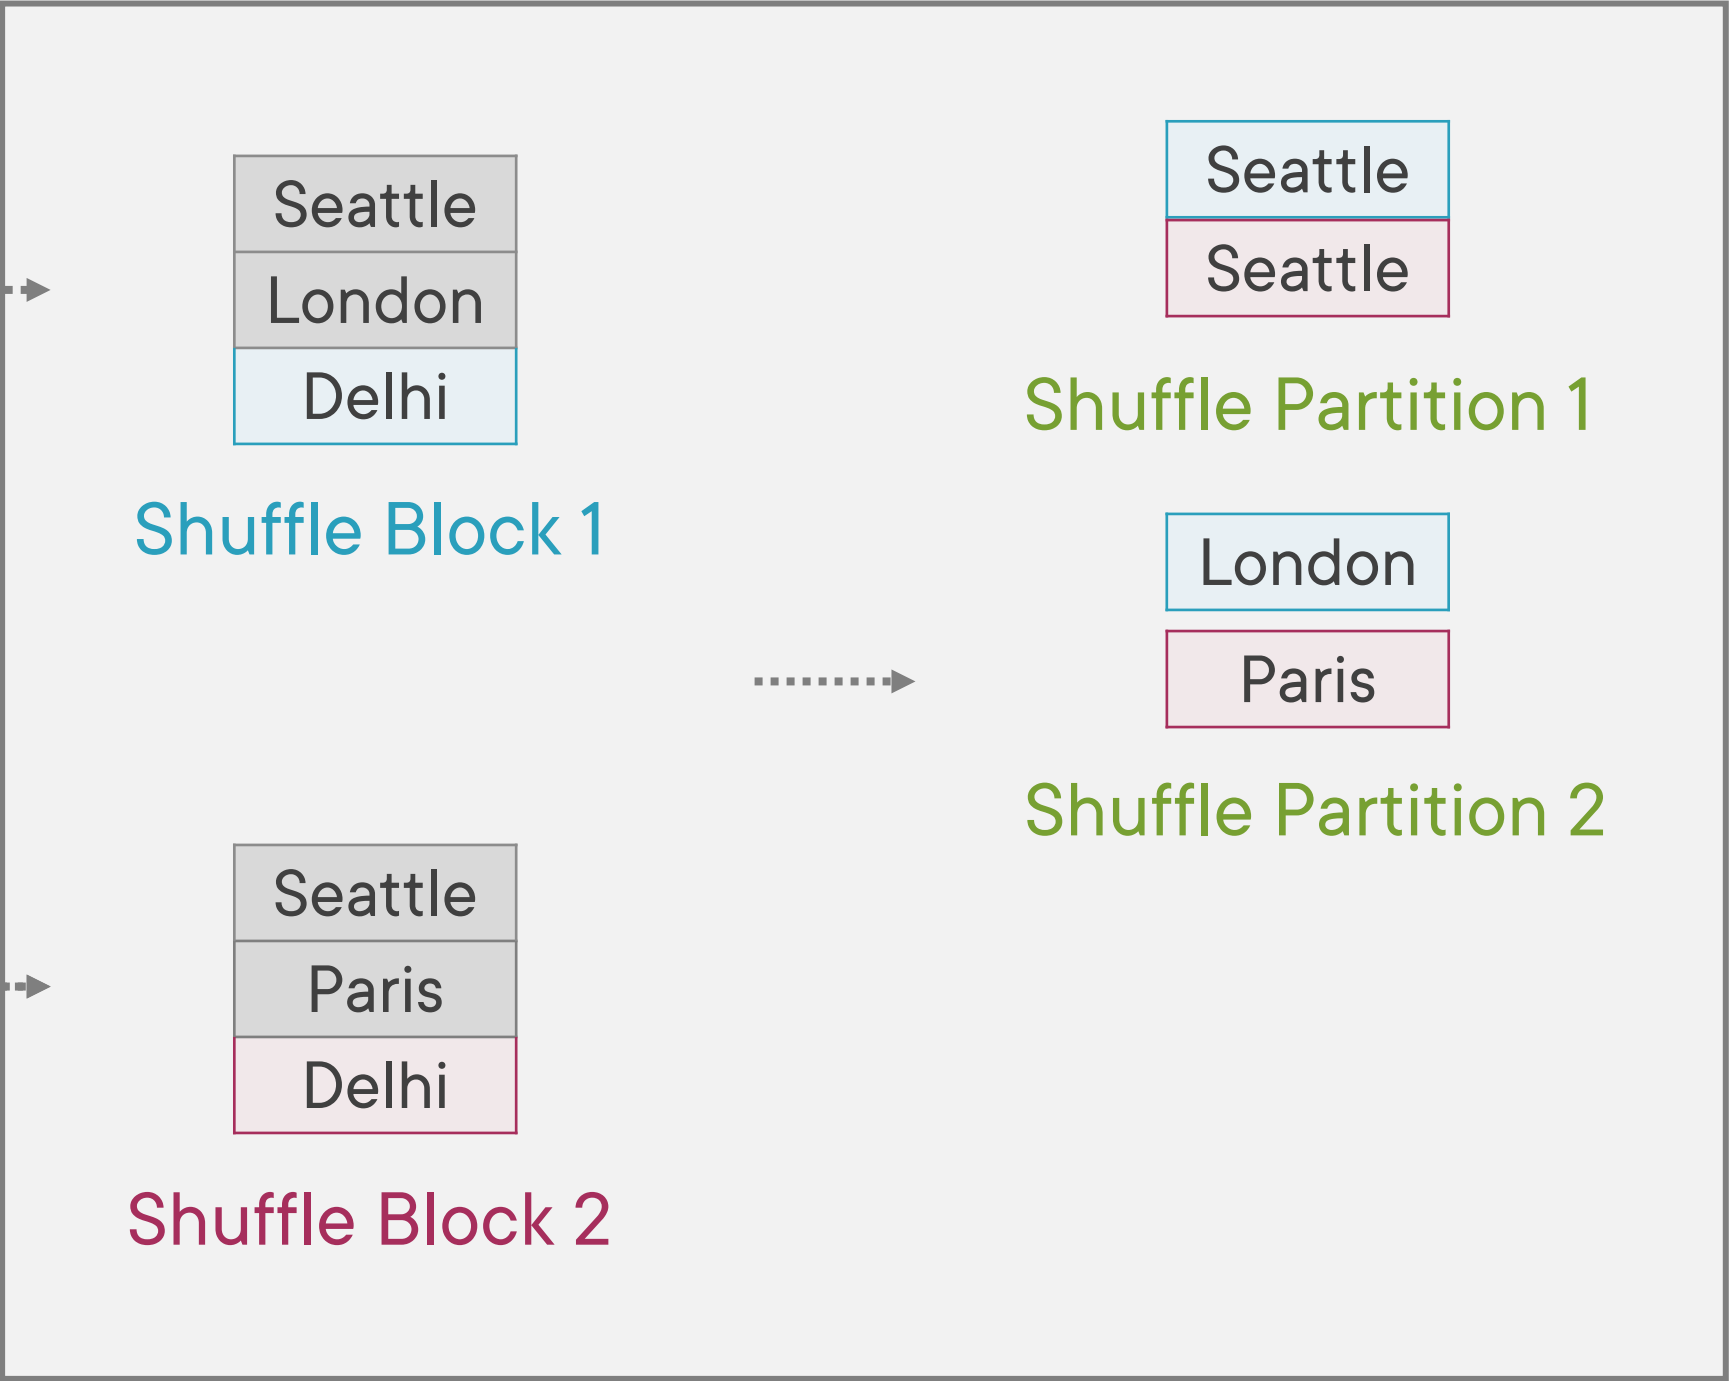
Step 1 – Group: Shuffle / Exchange

Seattle	600
London	300
Delhi	700

Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2



Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Get distinct cities  
**RDD 2**

Step 1 – Group: Shuffle / Exchange

Step 2 – Sum

Seattle	600
London	300
Delhi	700

Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Seattle
London
Delhi

Shuffle Block 1

Seattle
Paris
Delhi

Shuffle Block 2

Seattle
Seattle

Shuffle Partition 1

London
Paris

Shuffle Partition 2

Delhi
Delhi

Shuffle Partition 3

Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Get distinct cities  
**RDD 2**

Step 1 – Group: Shuffle / Exchange

Step 2 – Sum

Seattle	600
London	300
Delhi	700

Partition 1

Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Seattle
London
Delhi

Shuffle Block 1

Seattle
Paris
Delhi

Shuffle Block 2

Seattle
Seattle

Shuffle Partition 1

London
Paris

Shuffle Partition 2

Delhi
Delhi

Shuffle Partition 3

Seattle
---------

Partition 1

London
Paris

Partition 2

Delhi
-------

Partition 3

Shuffle Write

Shuffle Read

Read file,  
get City & Amount  
**RDD 1**

Seattle	600
London	300
Delhi	700

Partition 1

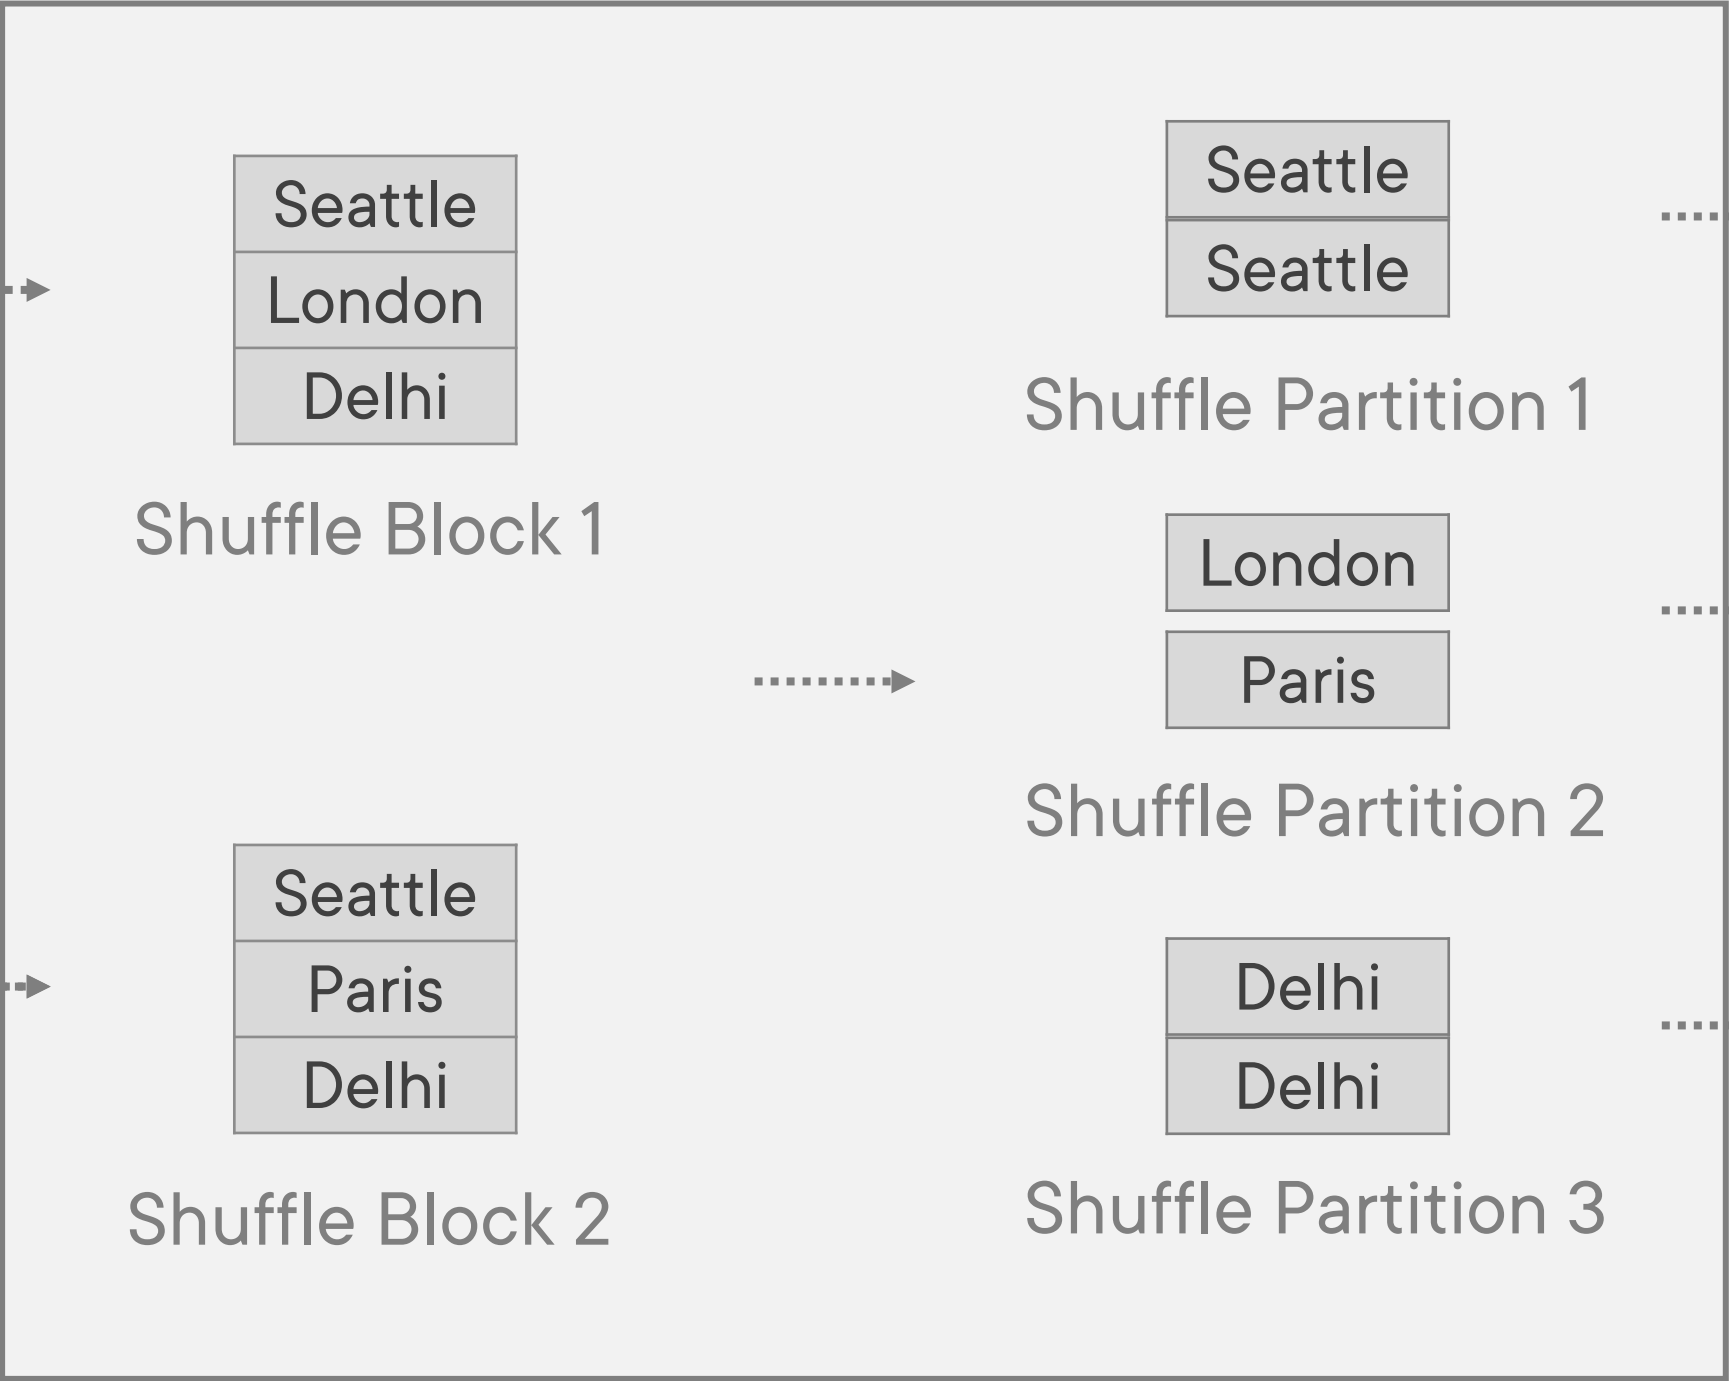
Seattle	400
Paris	900
Delhi	200
Seattle	900

Partition 2

Get distinct cities  
**RDD 2**

Step 1 – Group: Shuffle / Exchange

Step 2 – Sum



Seattle

Partition 1

London

Paris

Partition 2

Delhi

Partition 3

**Correct  
Output!**

Seattle

London

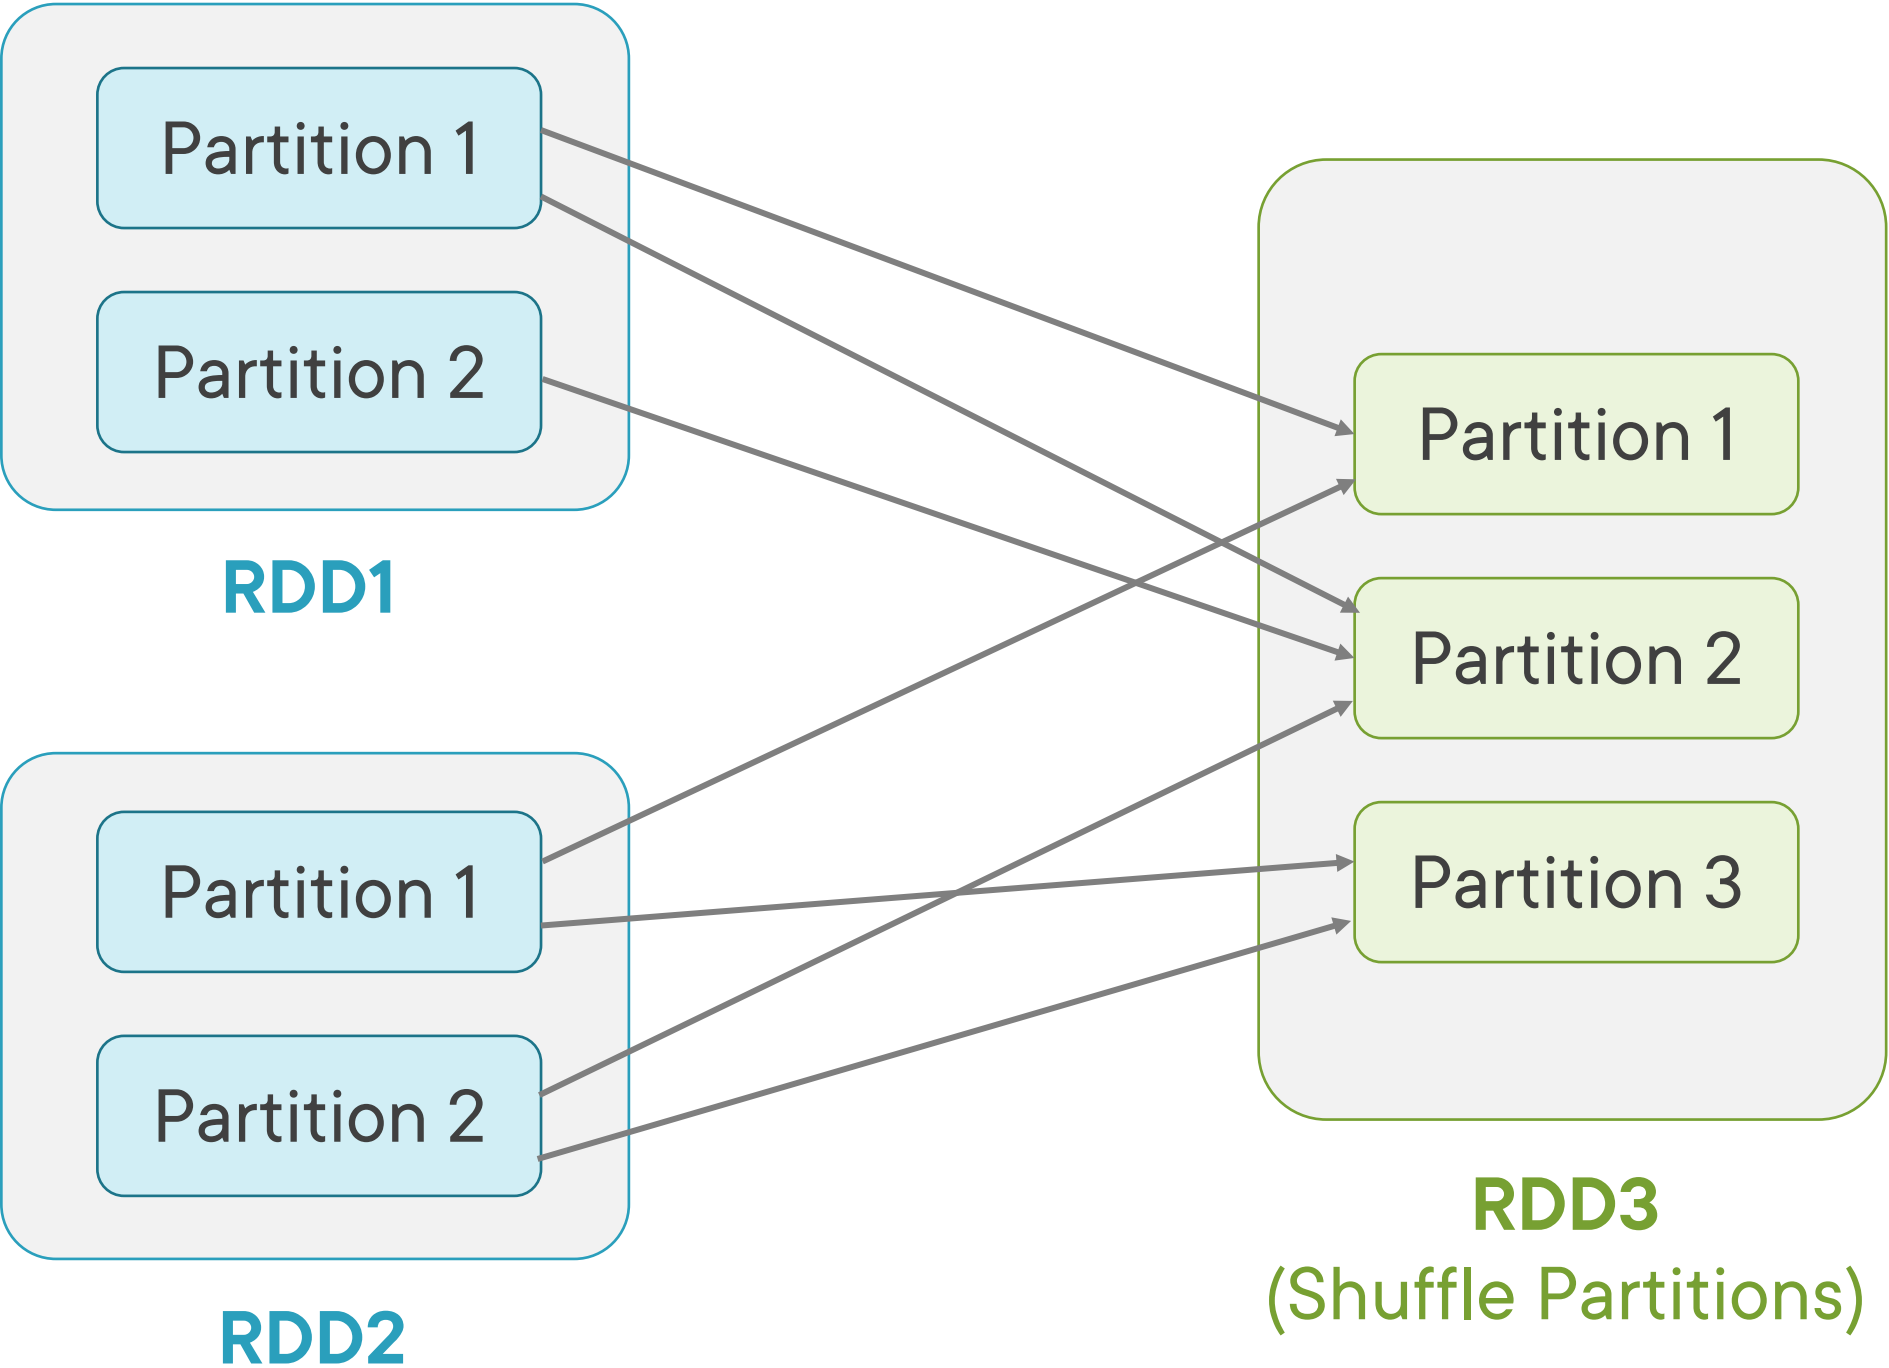
Paris

Delhi

Shuffle Write

Shuffle Read


One **input partition** might be used **multiple times** to produce output partitions



**Join operation is  
an example of this**



# Wide Transformation



**Requires shuffling of data between partitions**

**Expensive operation**

**Shuffle Read can only start when all data is written out as Shuffle Blocks**

**Number of shuffle partitions can be different than parent RDD partitions**

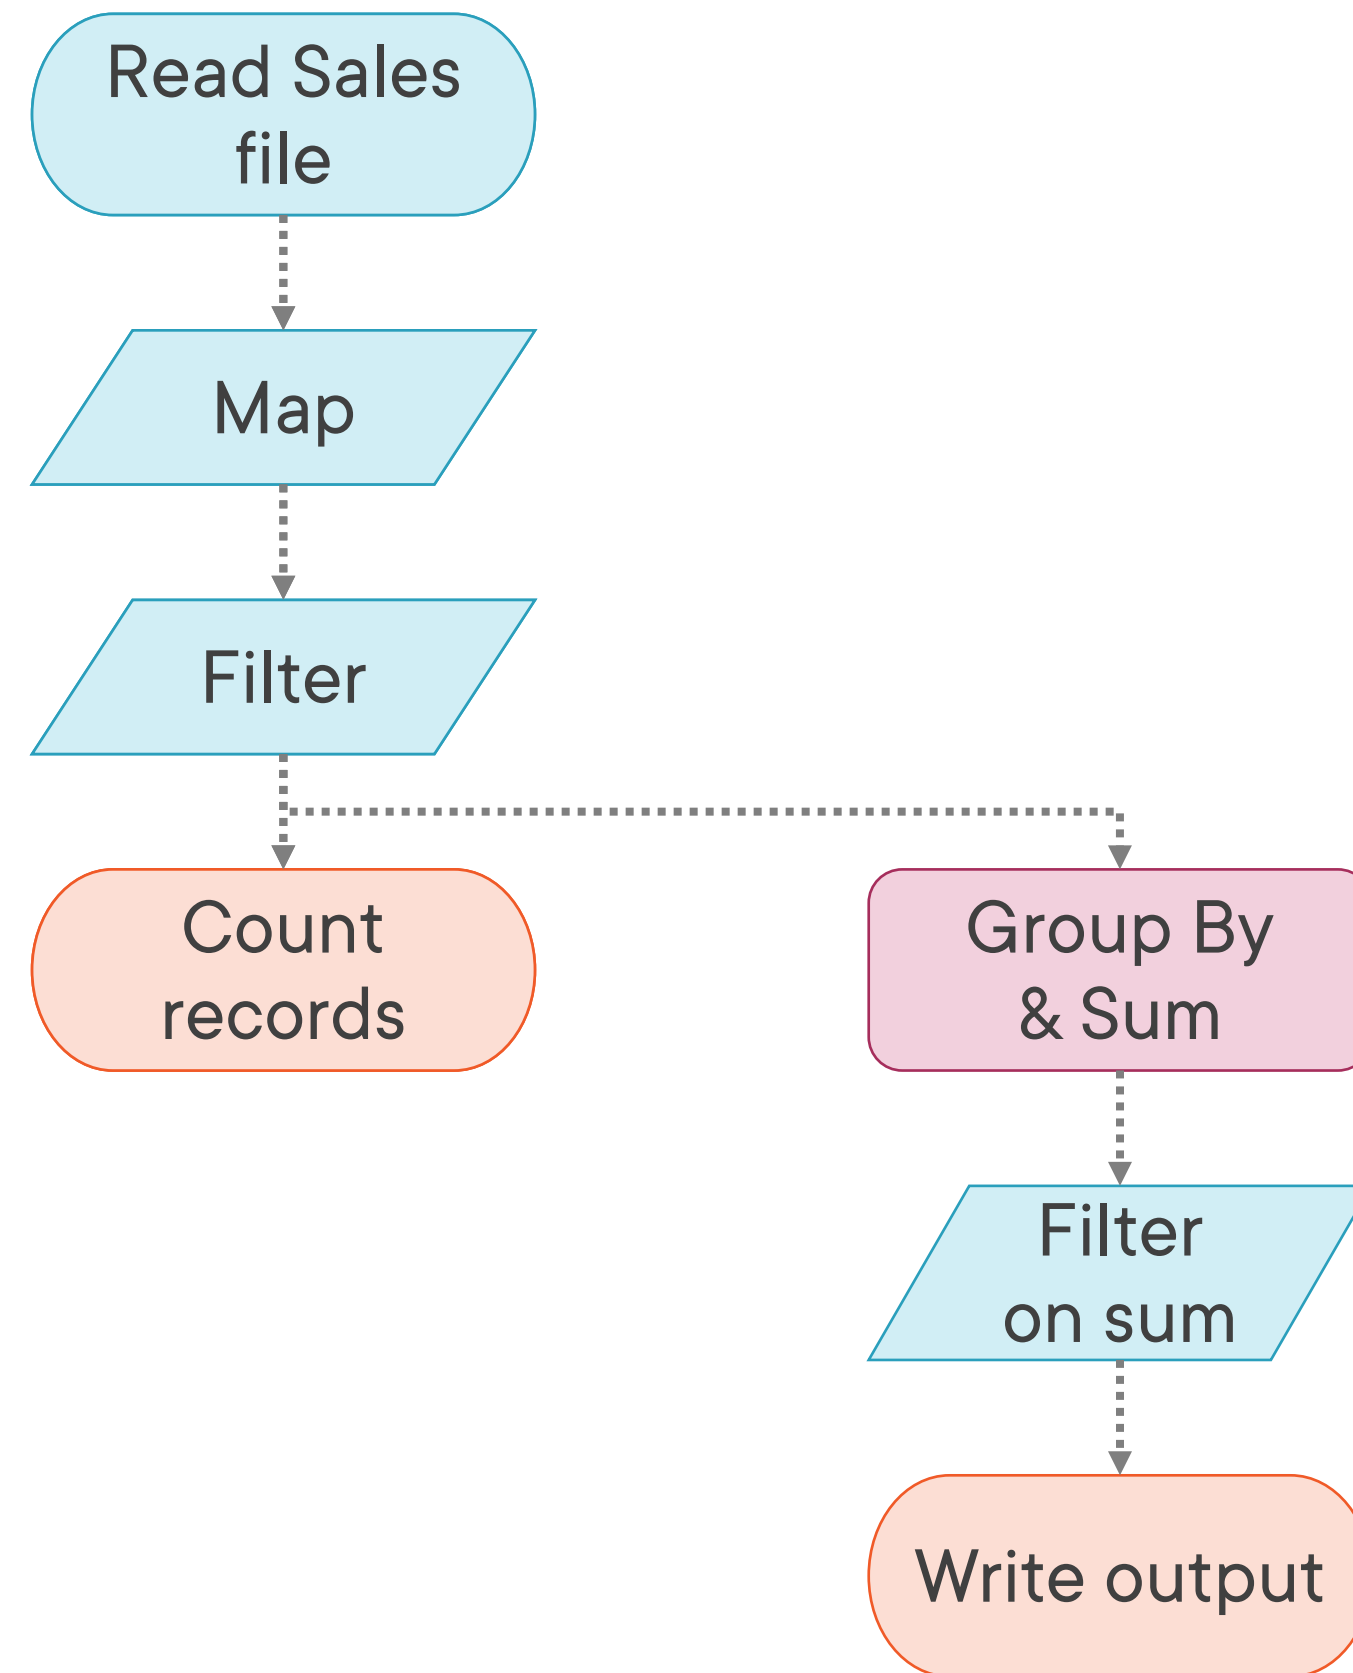
**Examples**

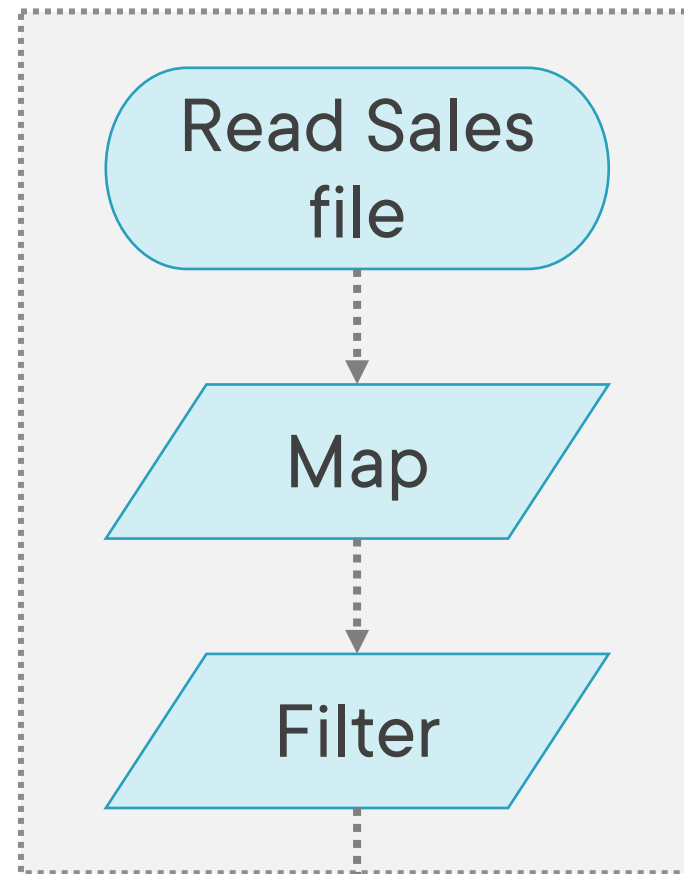
- reduceByKey, aggregateByKey, distinct, join, intersection etc.

# Spark Application Concepts: Jobs, Stages & Tasks

---

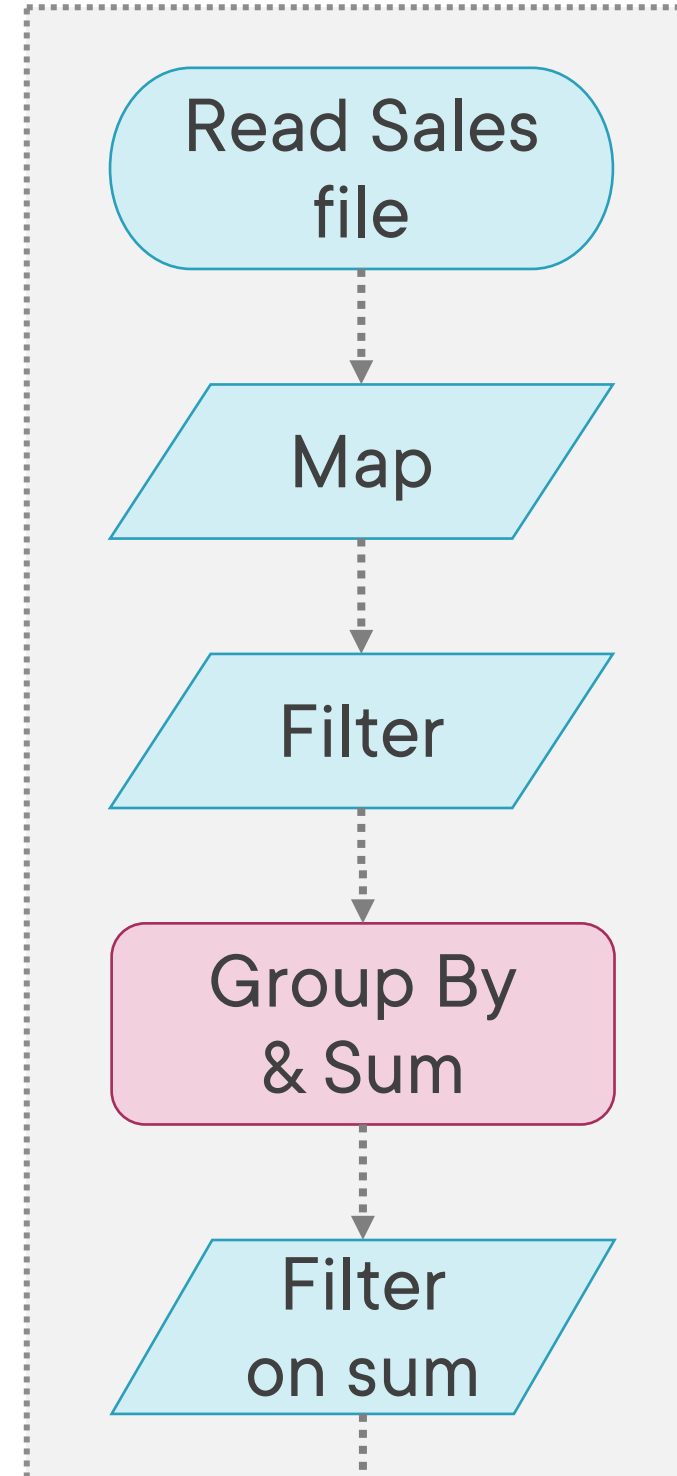
1. Read Sales csv file from Storage
2. Split Sales data by comma & define schema
3. Filter Sales where City = 'Delhi'
4. Count number of Sales
5. Group by Product, calculate sum of Amount
6. Filter grouped data where TotalAmount > '10000'
7. Write processed data to storage





Count records



**Job 1**




Write output


**Job 2**


**Jobs = Action Operations**







**Job 1**






**Job 1**









**Job 1**



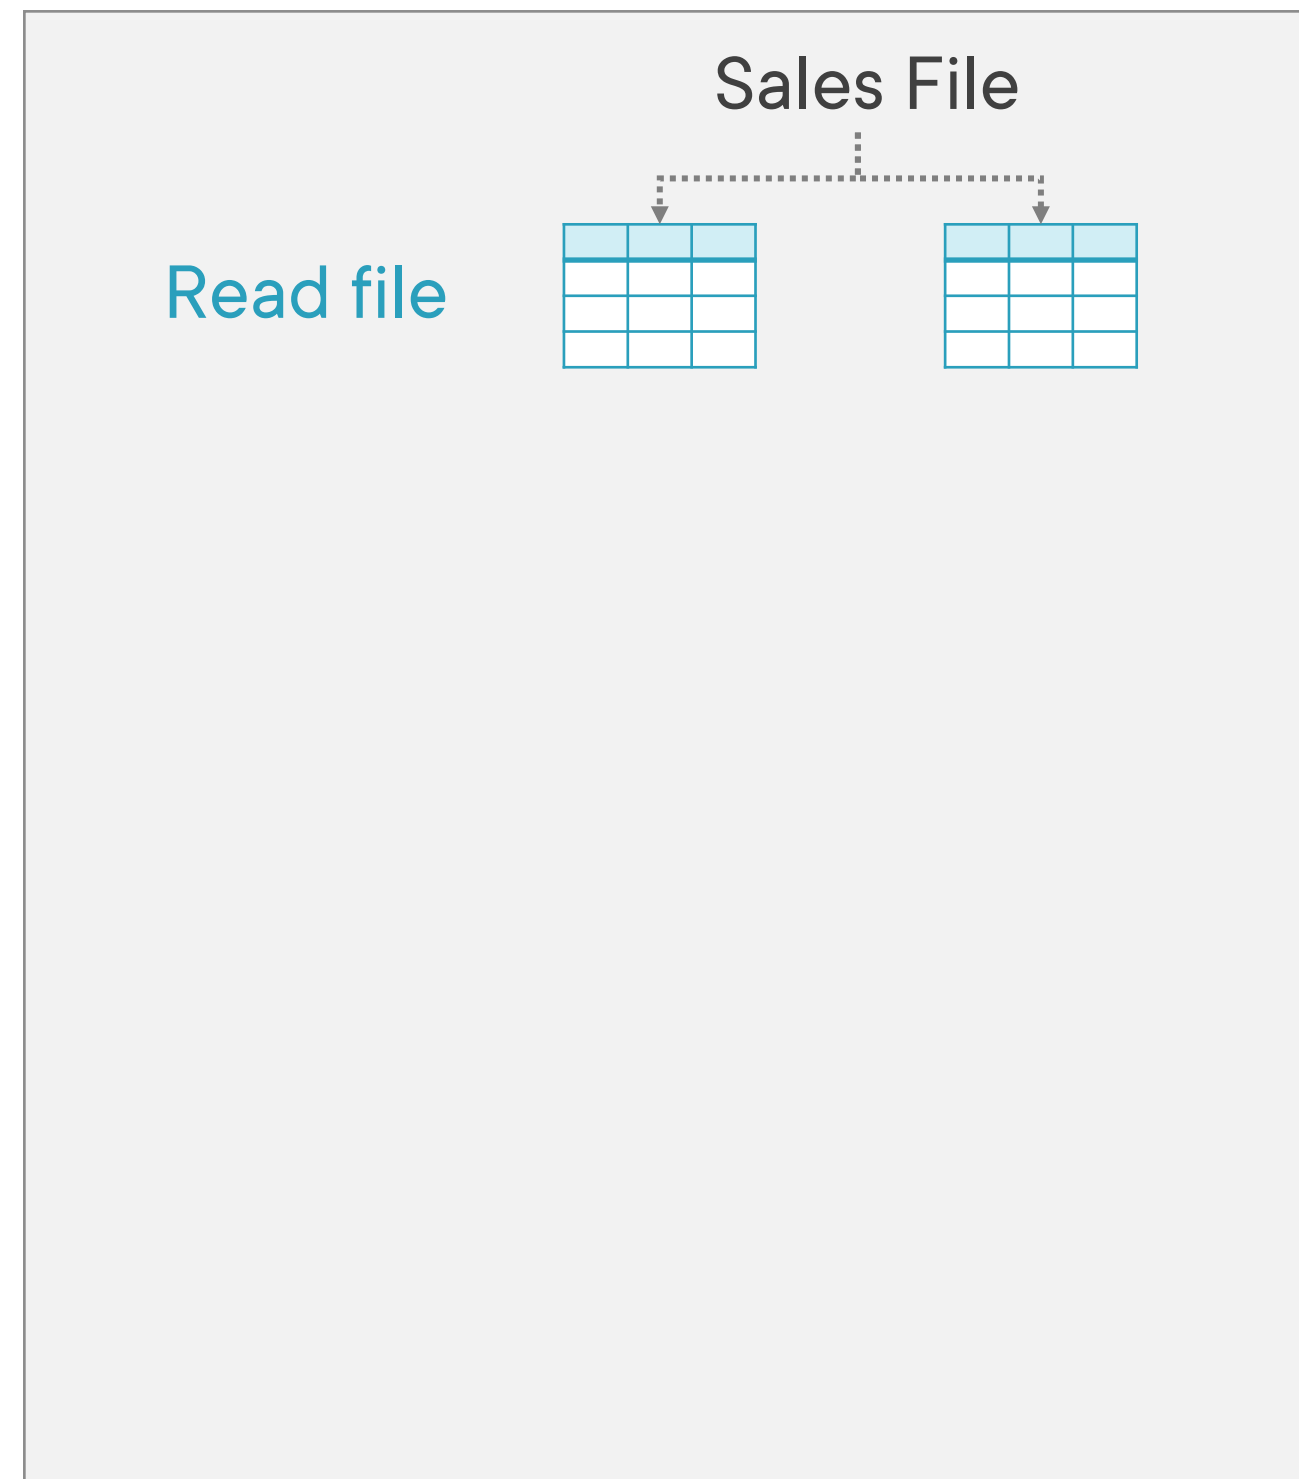
## Narrow Transformation

Number of output partitions are typically the same as input partitions










**Job 2**



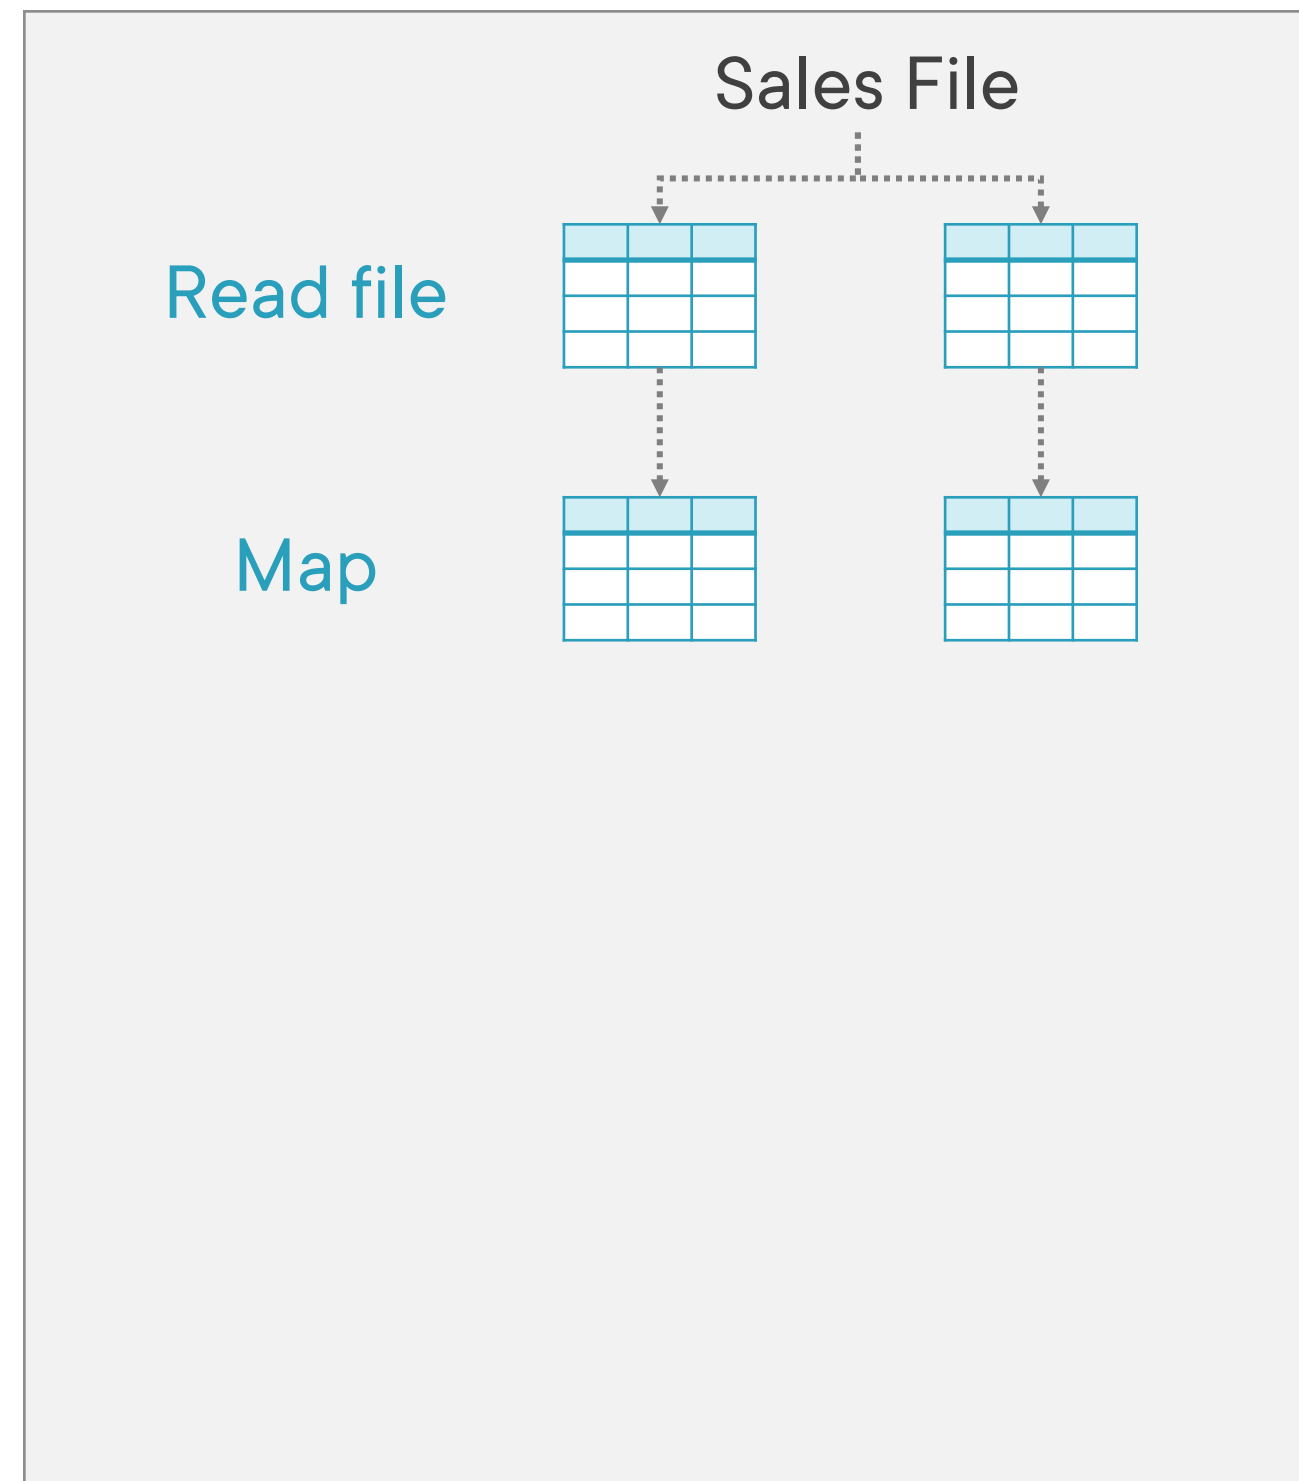
**Stage 1**




**Stage 2**






**Job 2**



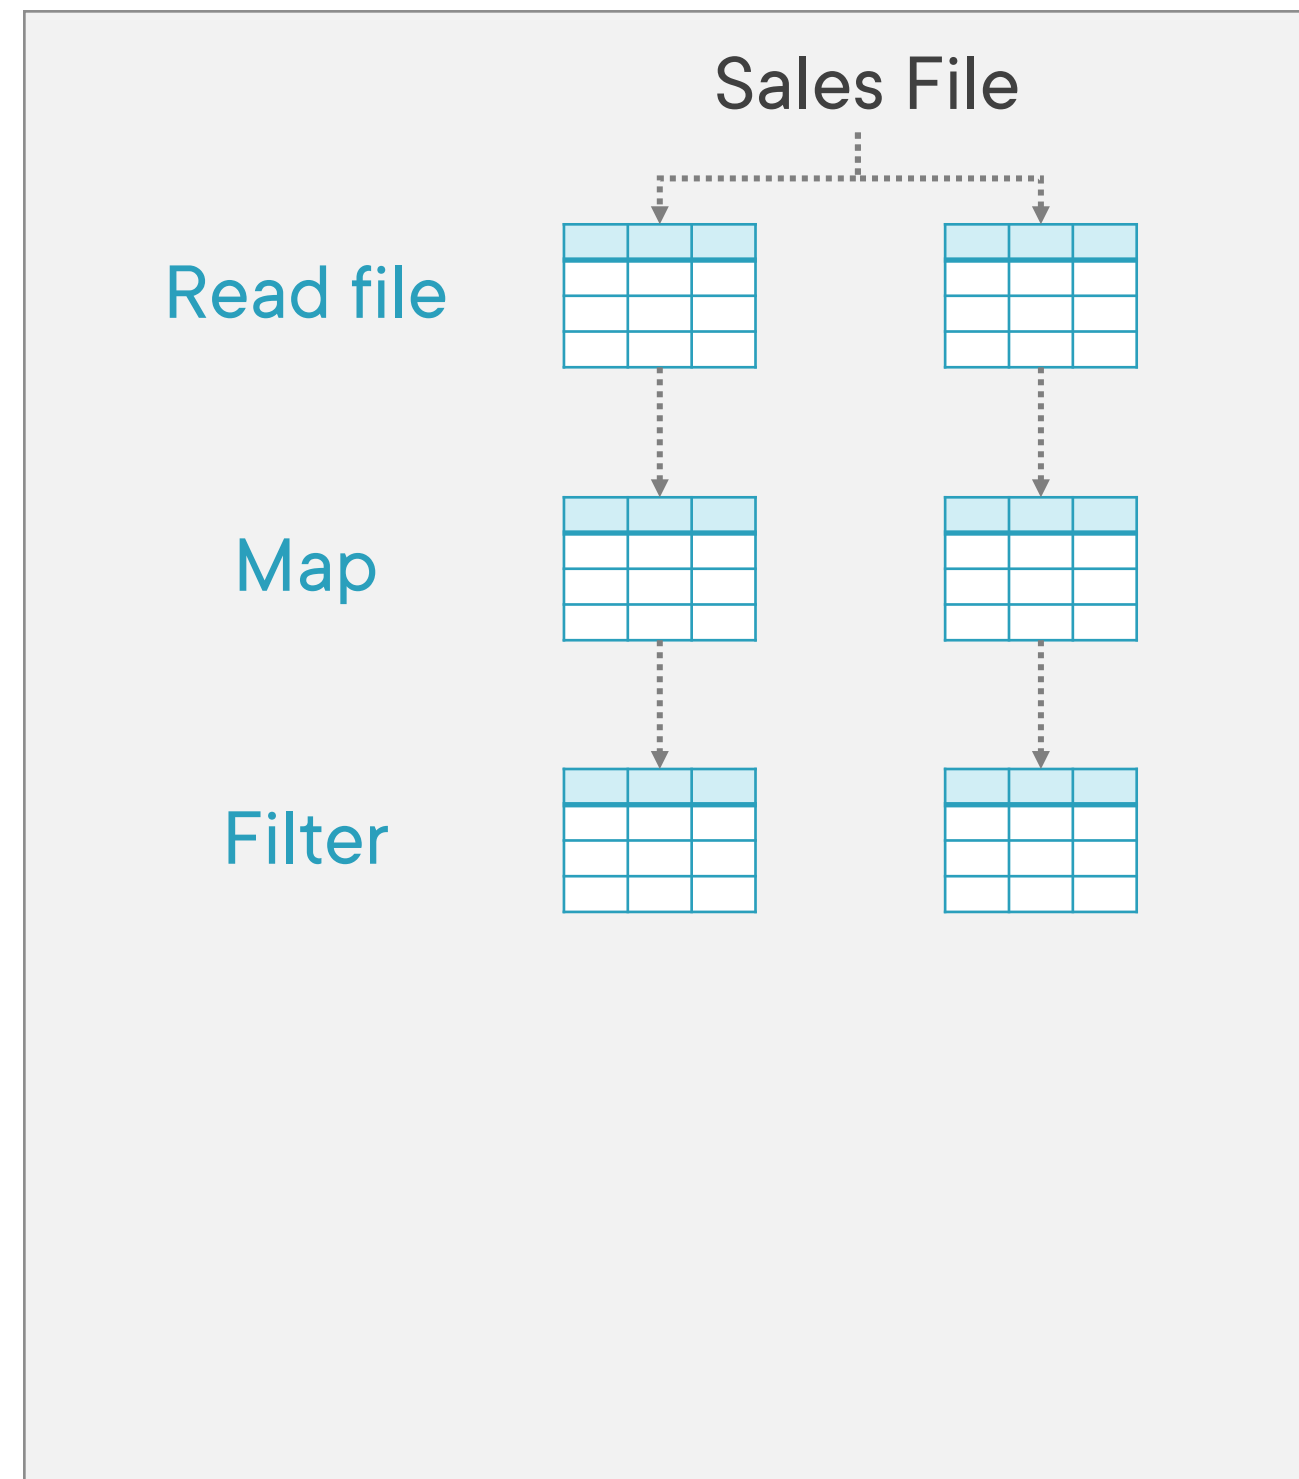
**Stage 1**




**Stage 2**




**Job 2**



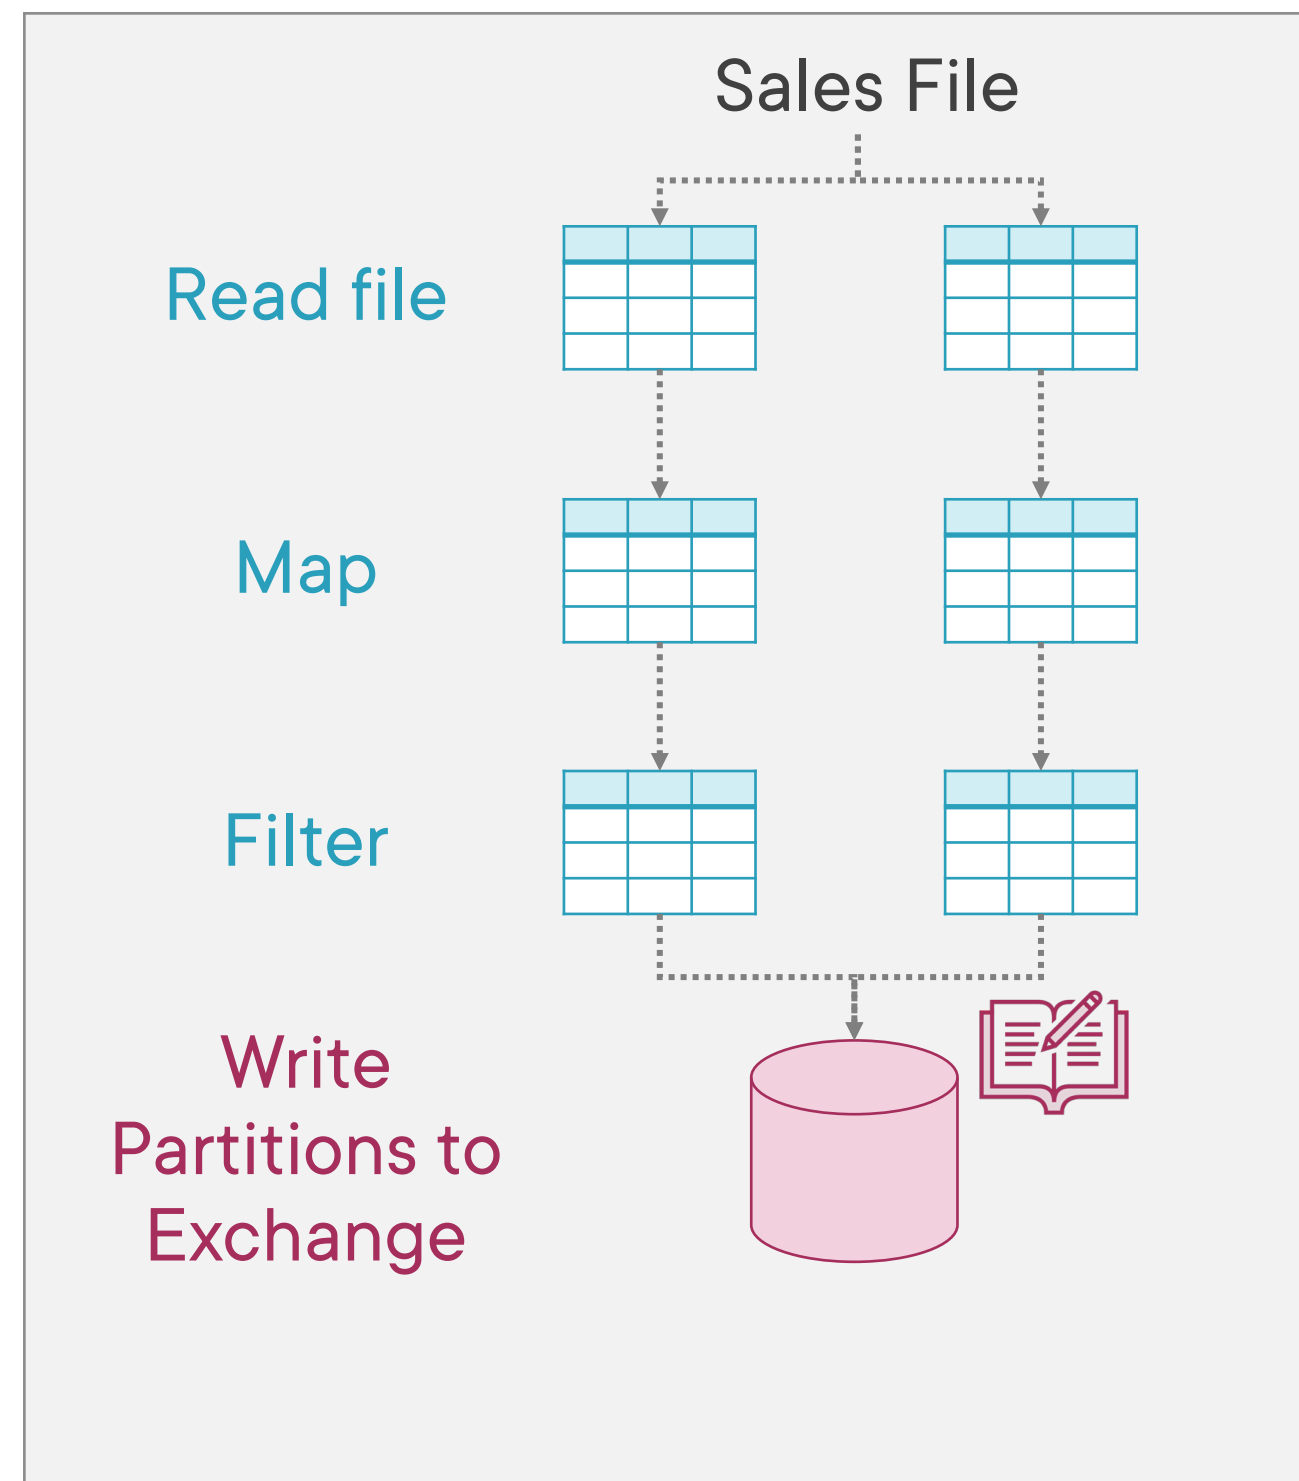
**Stage 1**




**Stage 2**




**Job 2**



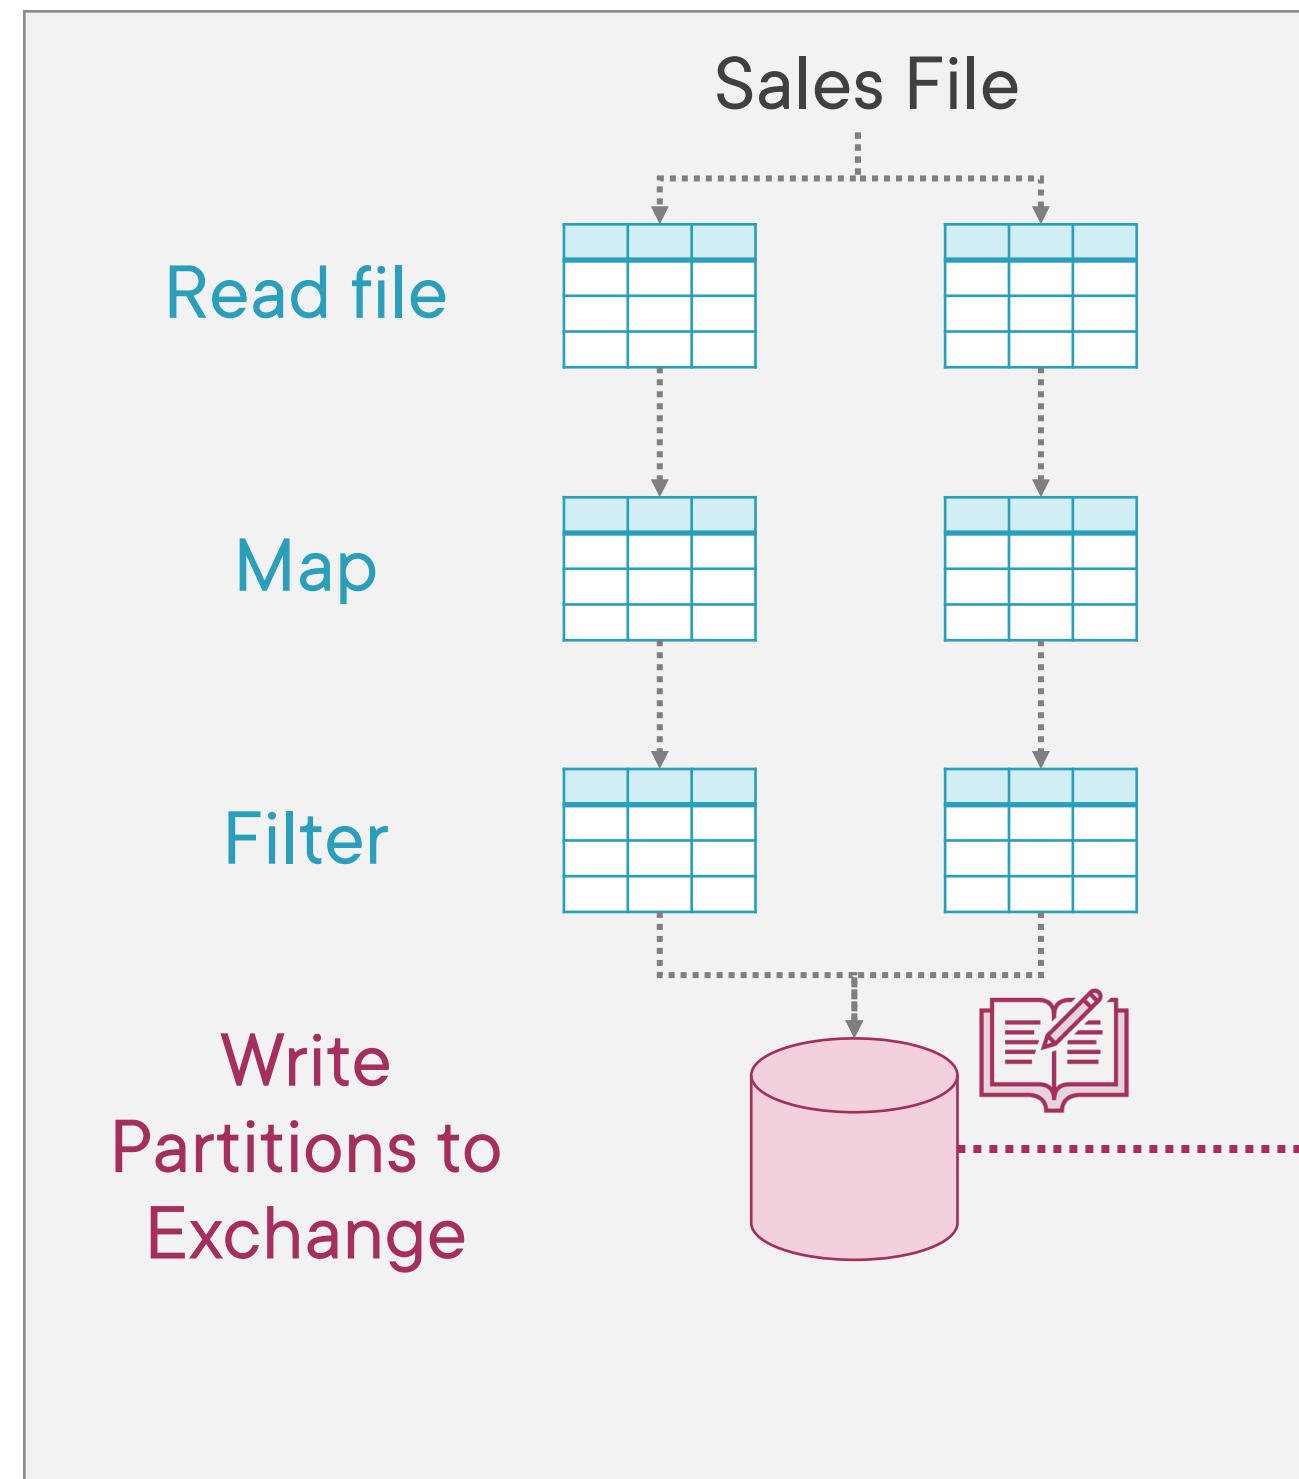
**Stage 1**



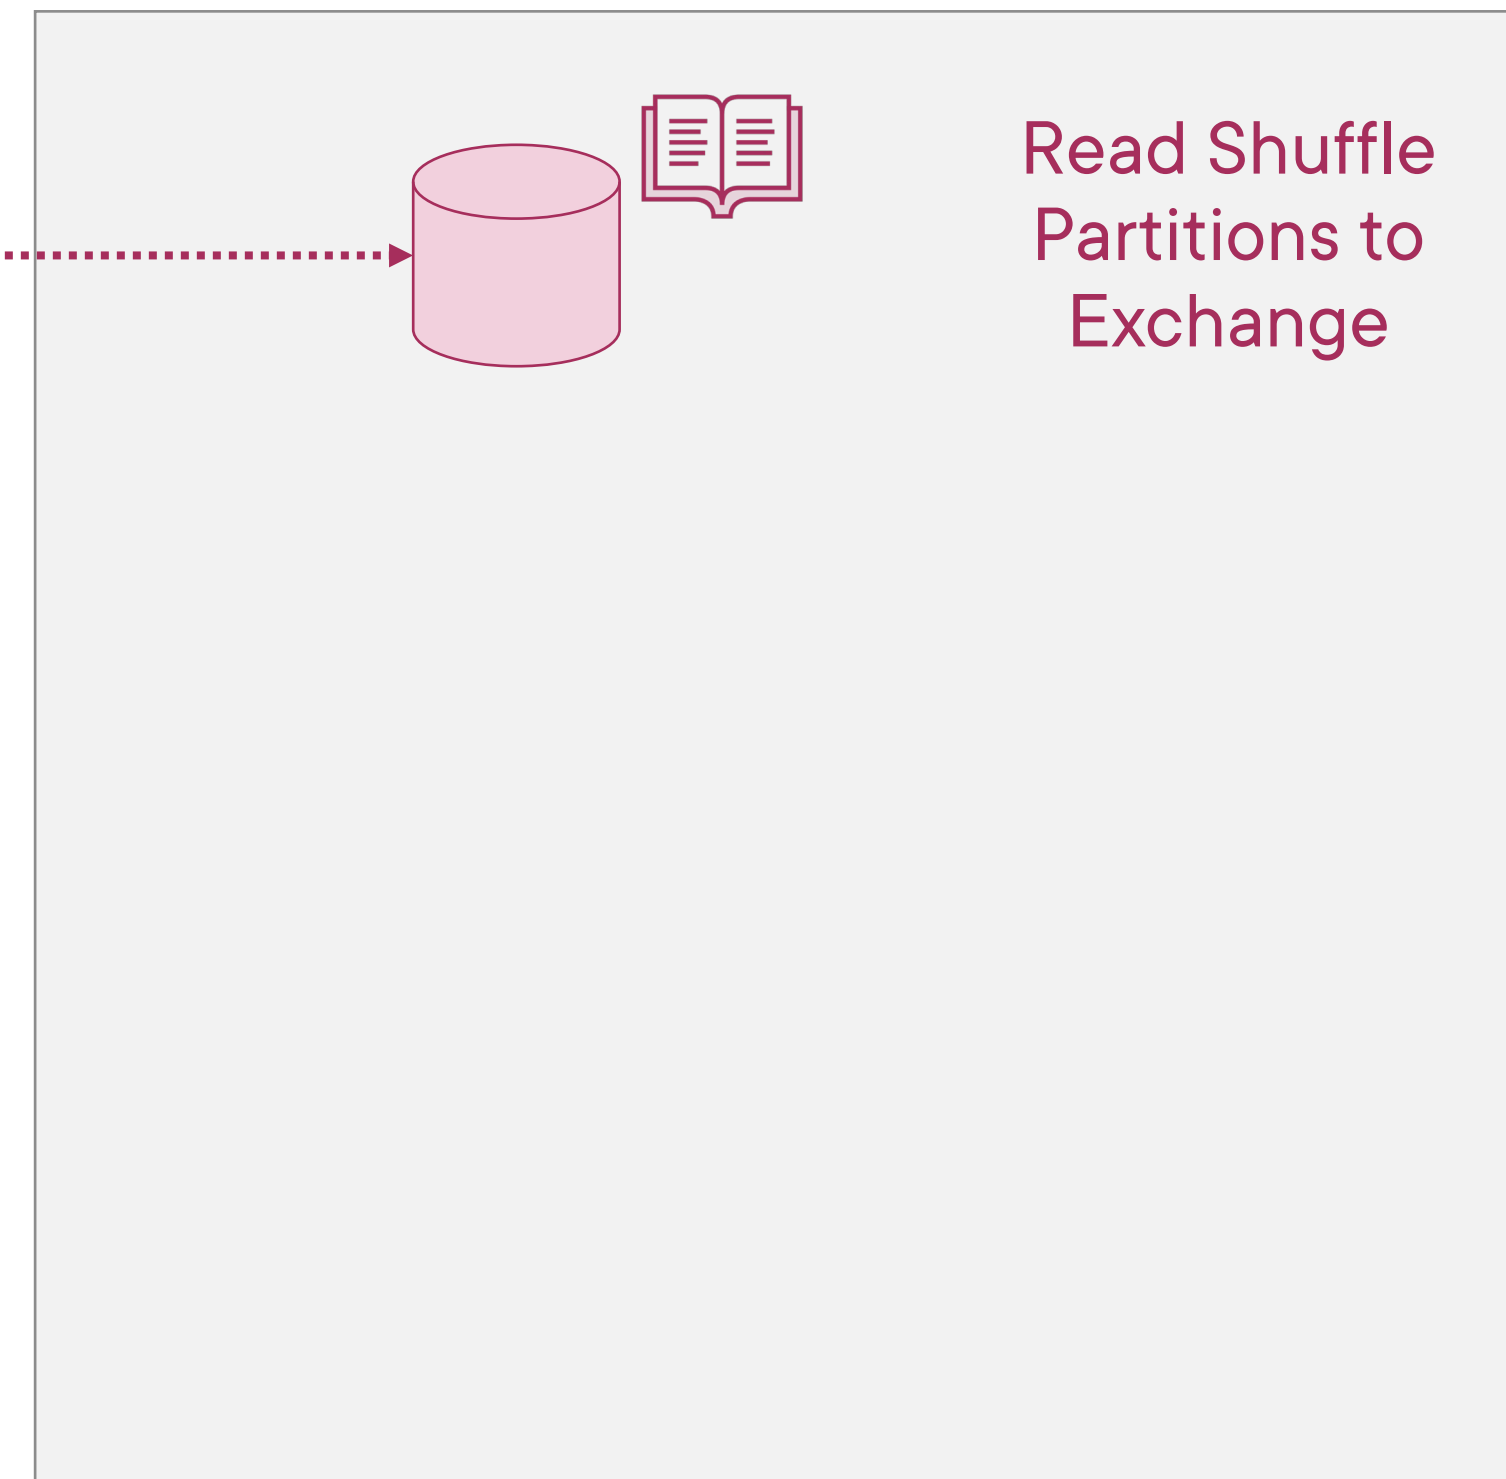
**Stage 2**




**Job 2**



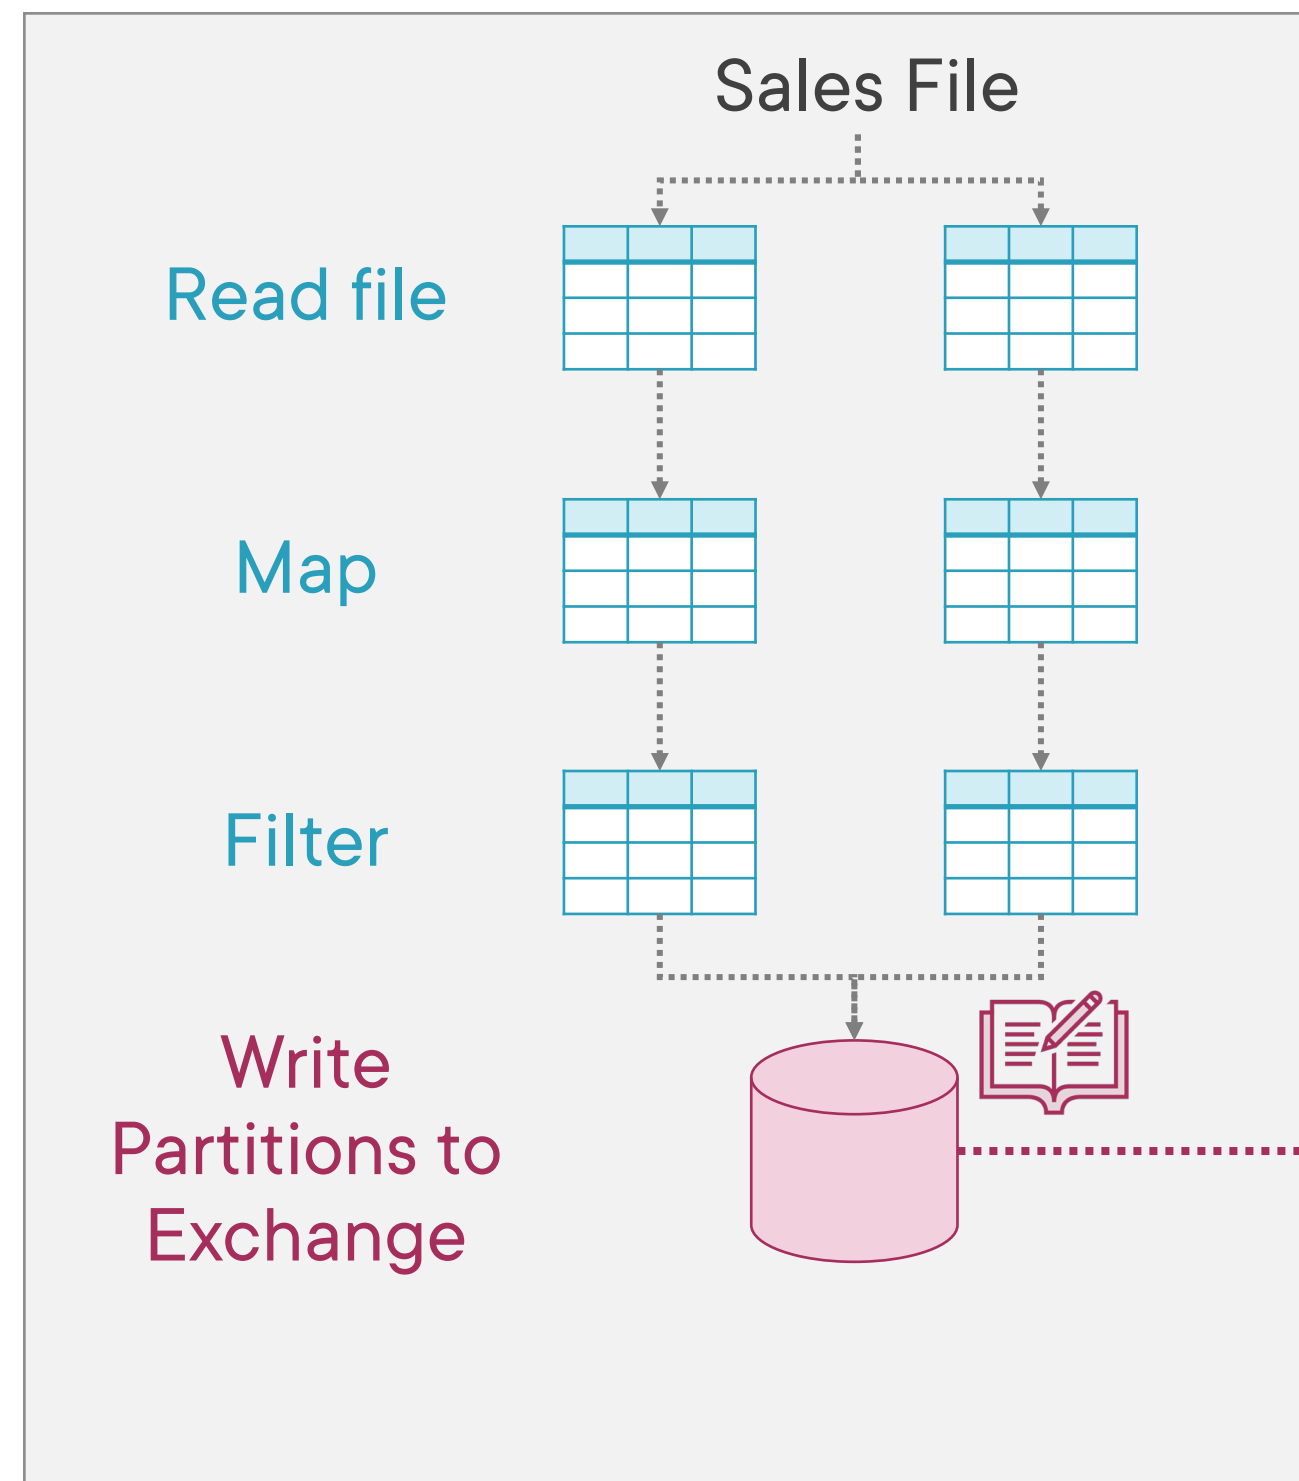
**Stage 1**



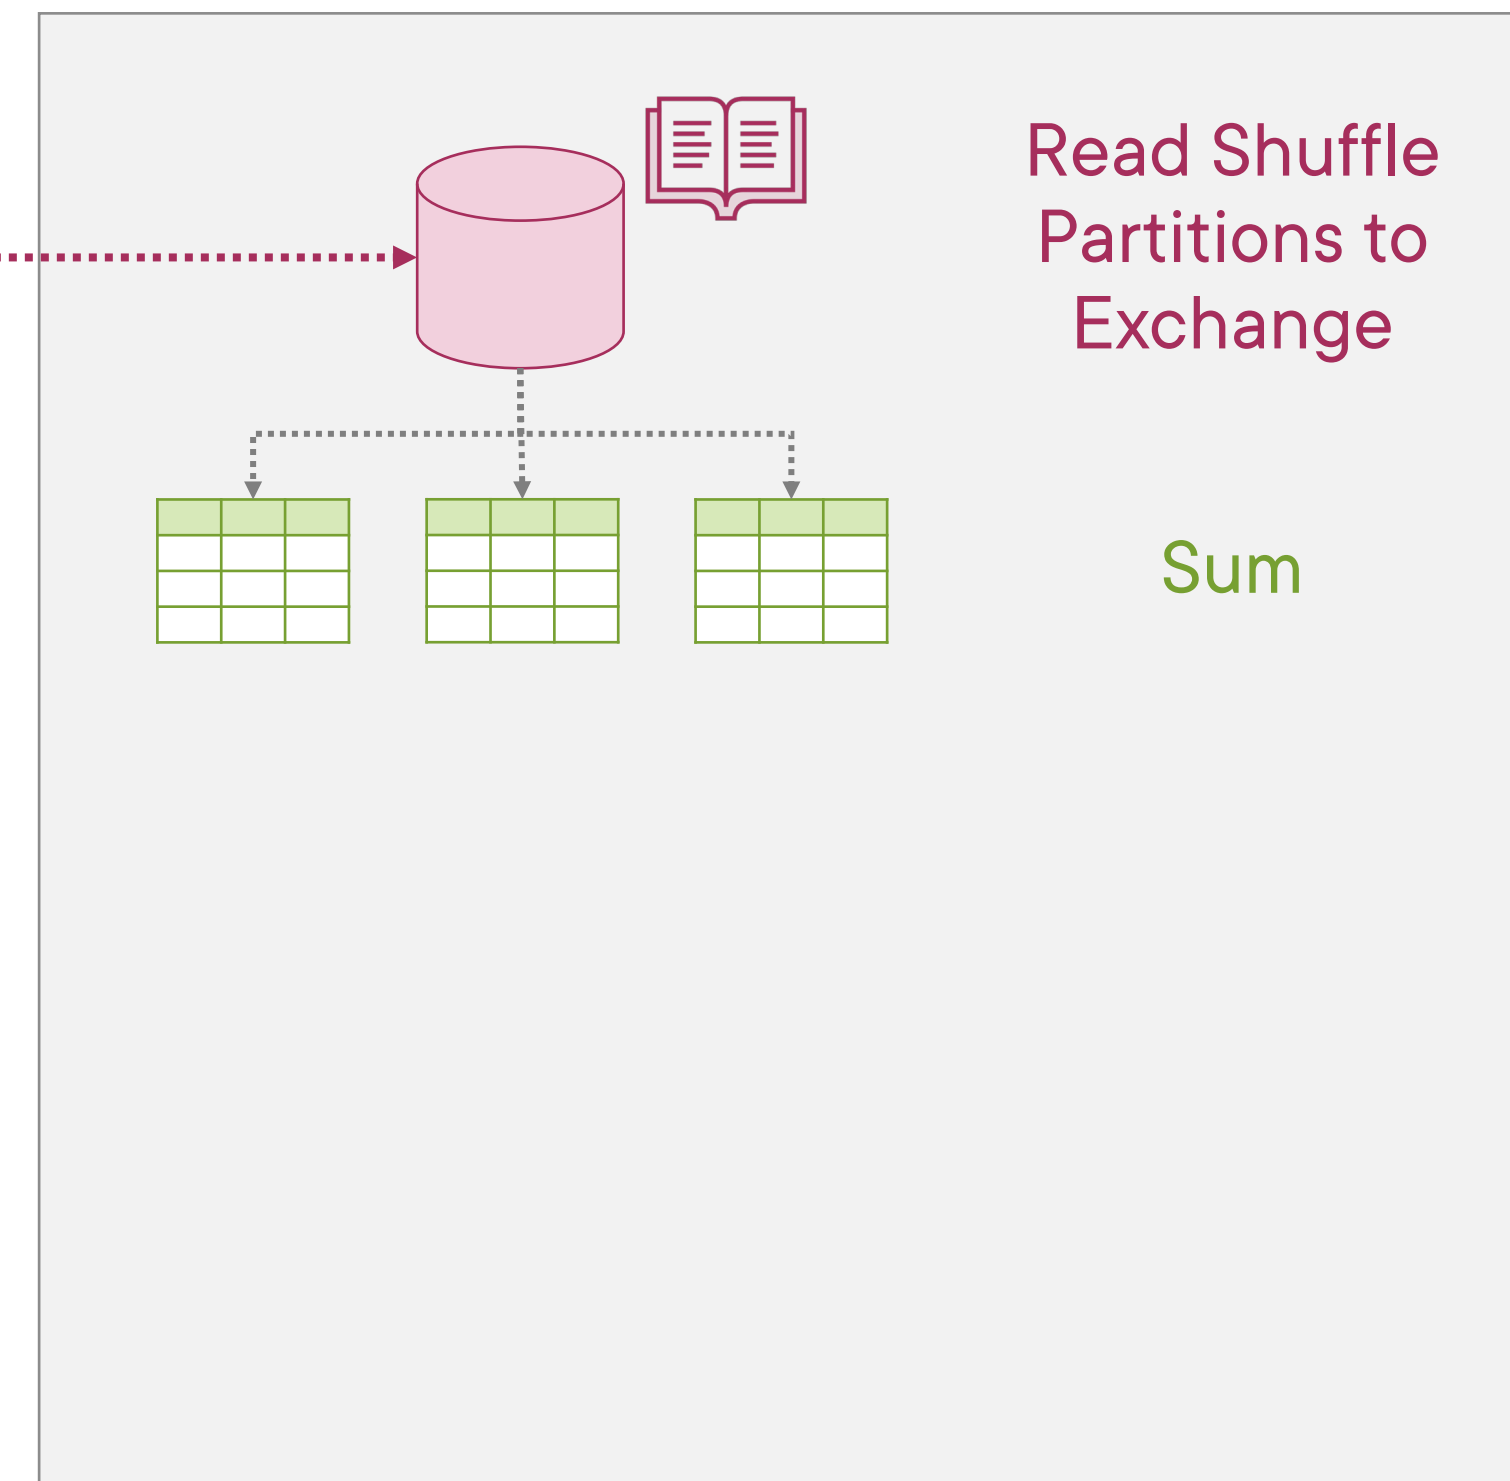
**Stage 2**




**Job 2**



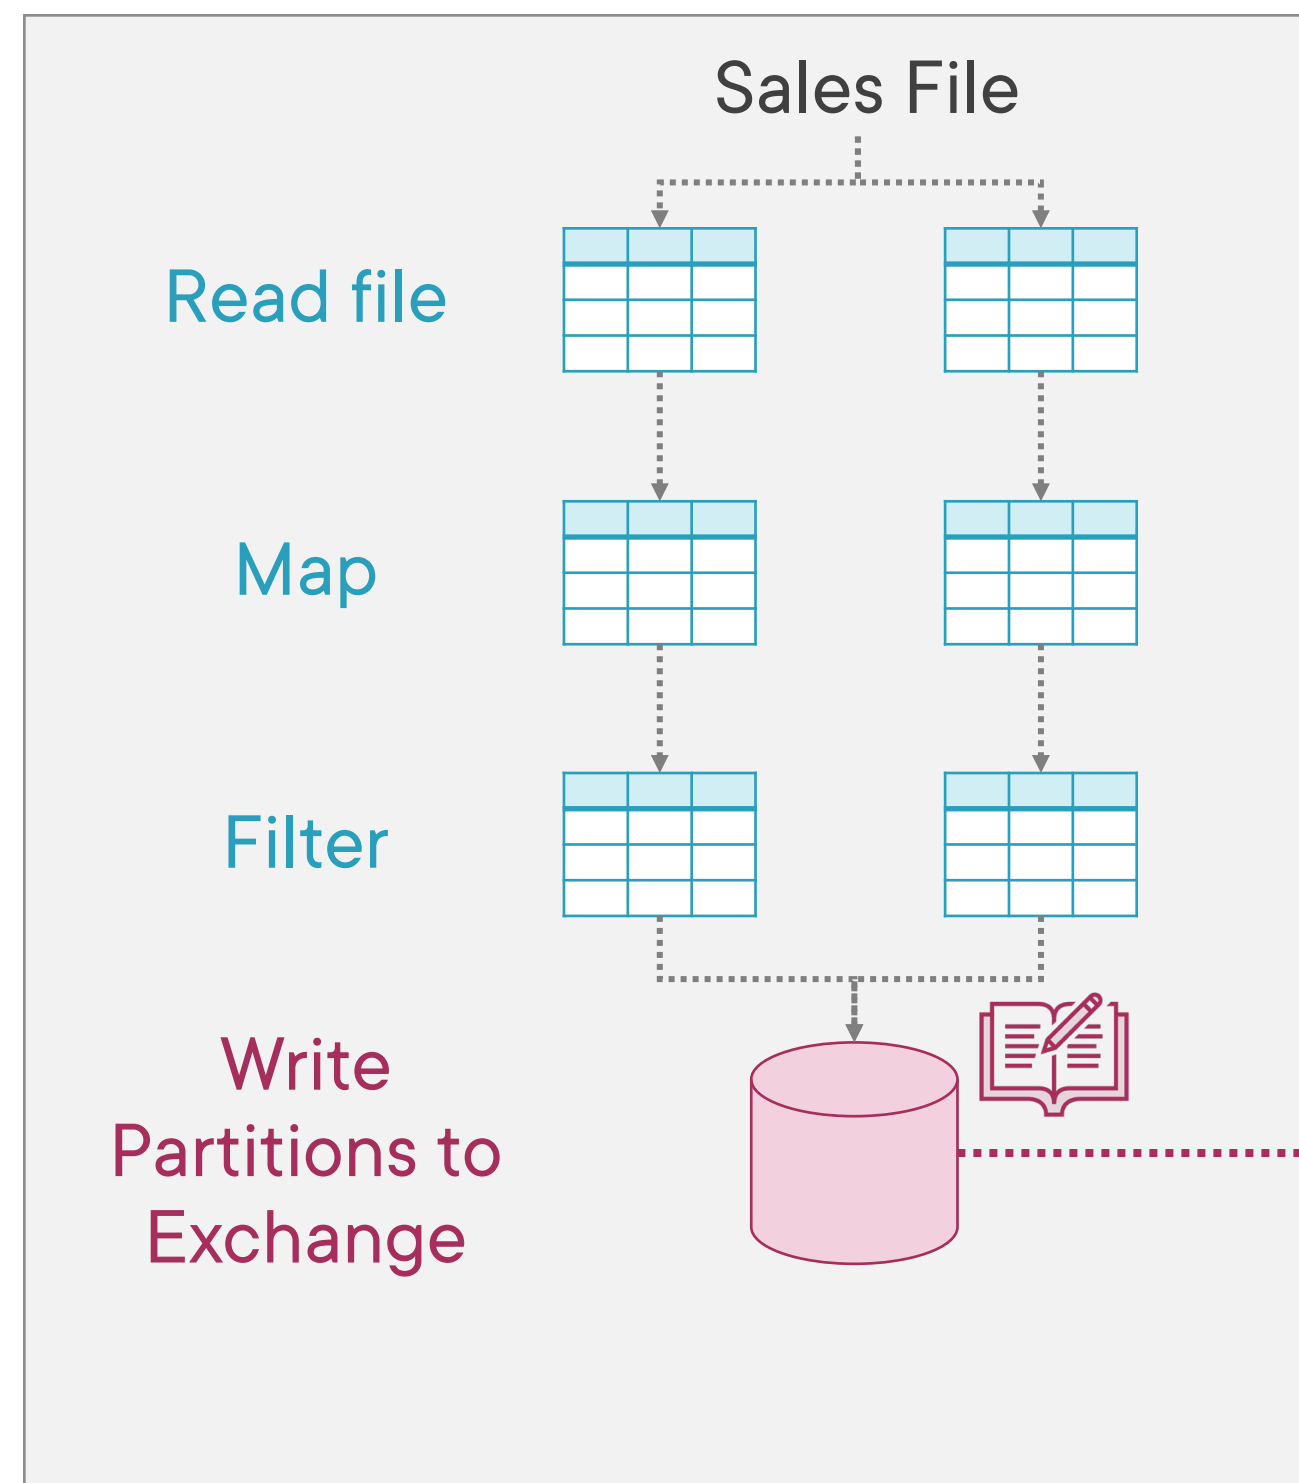
**Stage 1**



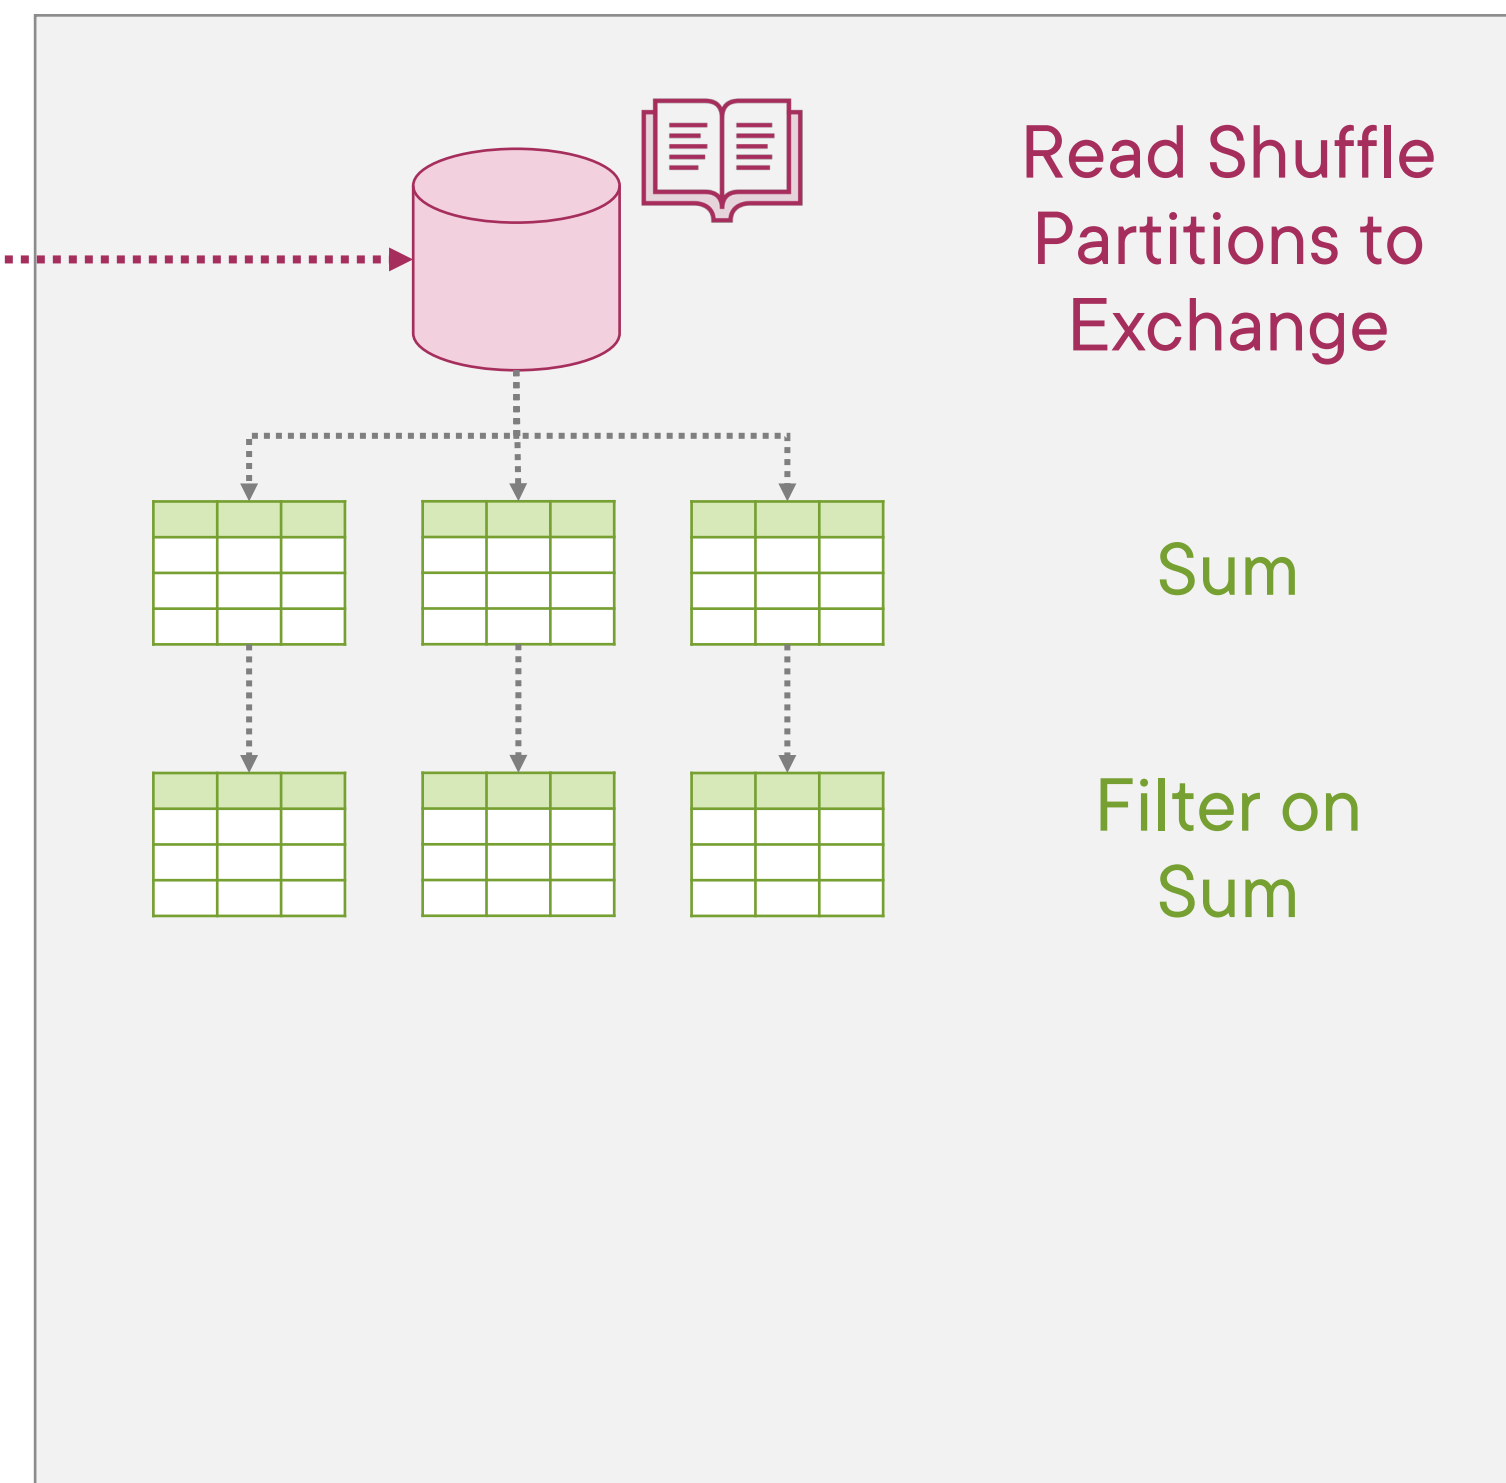
**Stage 2**




**Job 2**




**Stage 1**



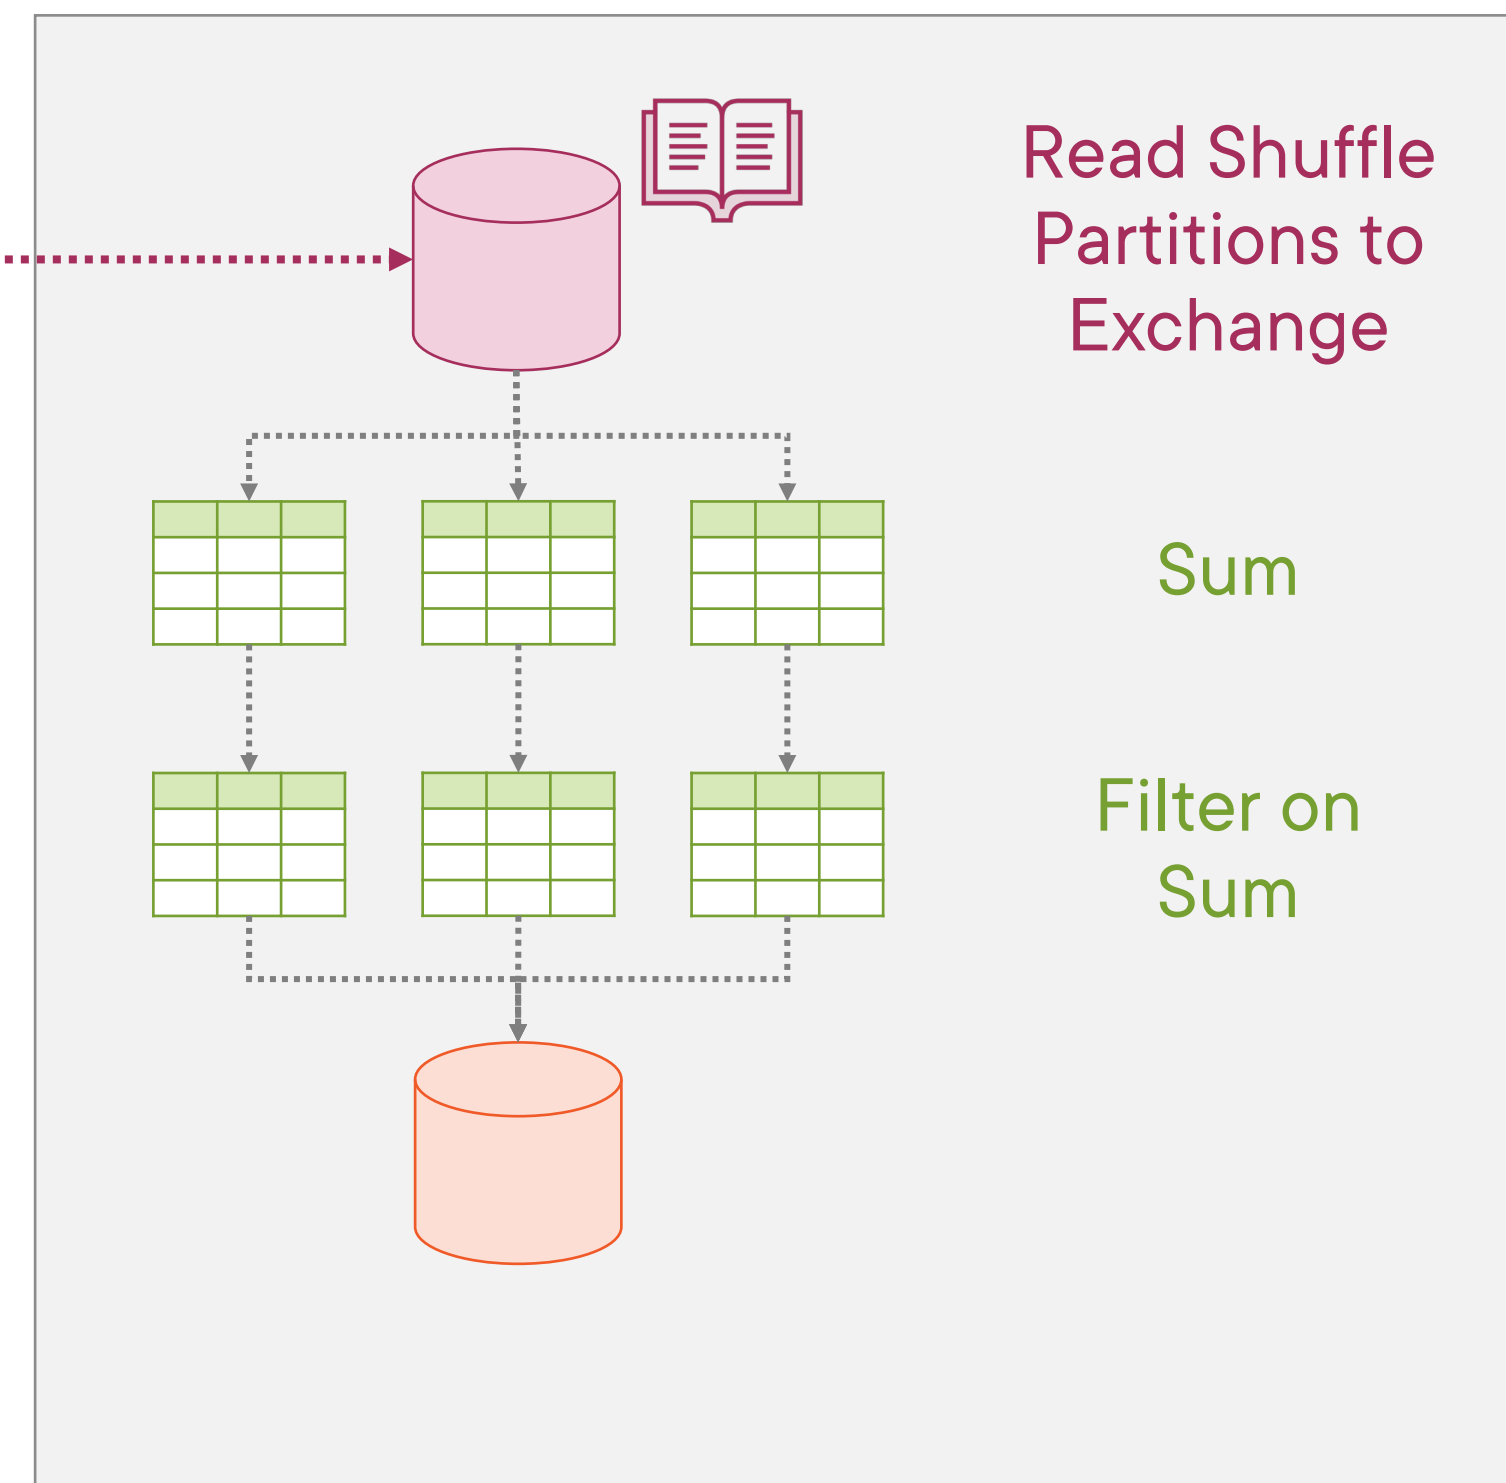
**Stage 2**



**Job 2**



**Stage 1**



**Stage 2**



# Spark Execution Components

## **Spark Application is a set of resources**

- Contains driver and executor processes

## **Multiple jobs can run in an Application**

- Number of jobs = Action operations applied

## **Each job is divided into Stages**

- Number of stages = Wide transformations + 1

## **Stages are typically executed in sequence**

- When one stage finishes, then only next can start
- Exceptions – Join where 2 datasets can be read parallelly as separate stages

## **Each Stage has its own set of Tasks**

- Number of tasks = Number of partitions
- Number of parallel tasks = Number of cores

## Summary



**RDD is the native data structure of Spark**

- In-memory, Partitioned, Read-only & Resilient

**RDD can be created in many ways**

- Parallelize a collection, read a file or convert an RDD

**Pair RDD is an RDD with Key-Value pairs**

**Apply transformation and action operations on RDDs**

- Transformation gives new RDD. Action gives result
- Transformations are lazy operations

**Transformations are of two types**

- Narrow – Each input partition is used at-most once
- Wide – An input partition may be used multiple times

**Data shuffling is a two-step process**

**Application → Jobs → Stages → Tasks**

Up Next:  
Cleaning & Transforming Data with DataFrames

---