

Performing Optimizations in Spark



Mohit Batra

Founder, Crystal Talks

linkedin.com/in/mohitbatra

Overview



Work with Spark partitions

Change DataFrame partitions

Memory management

Persist data

Spark join strategies and broadcast join


Optimize join with bucketing

Dynamic resource allocation

Resource allocation using fair scheduling


Working with Spark Partitions

Types of Partitioning



In-Memory Partitioning

- Chunks of data read in memory
- All partitions together constitute RDD/DataFrame



Disk Partitioning

- Writing output to disk by physically partitioning data based on columns
- Done using partitionBy method

In-Memory Partition Settings

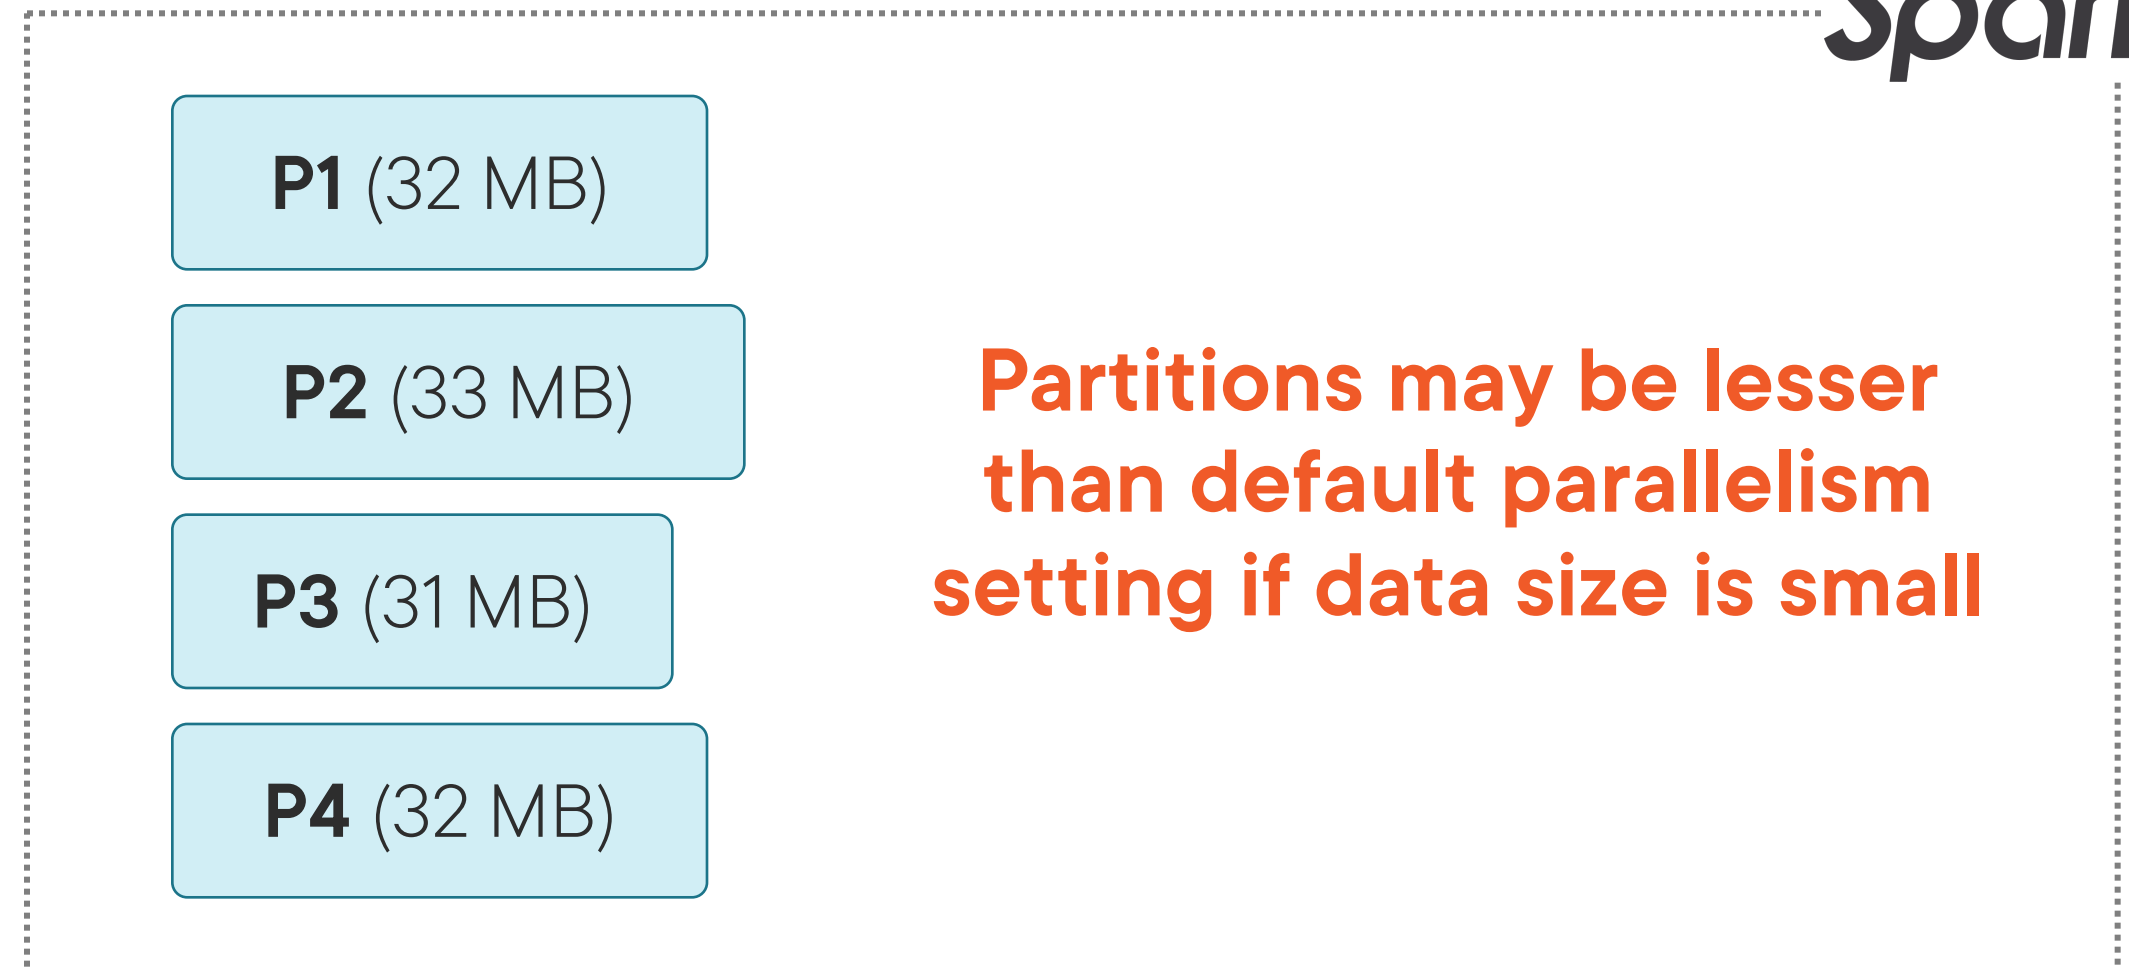

Settings for Reading Data

Settings for Data Shuffling

1. For Reading Data

`spark.default.parallelism = 4`

[default = no. of cores]

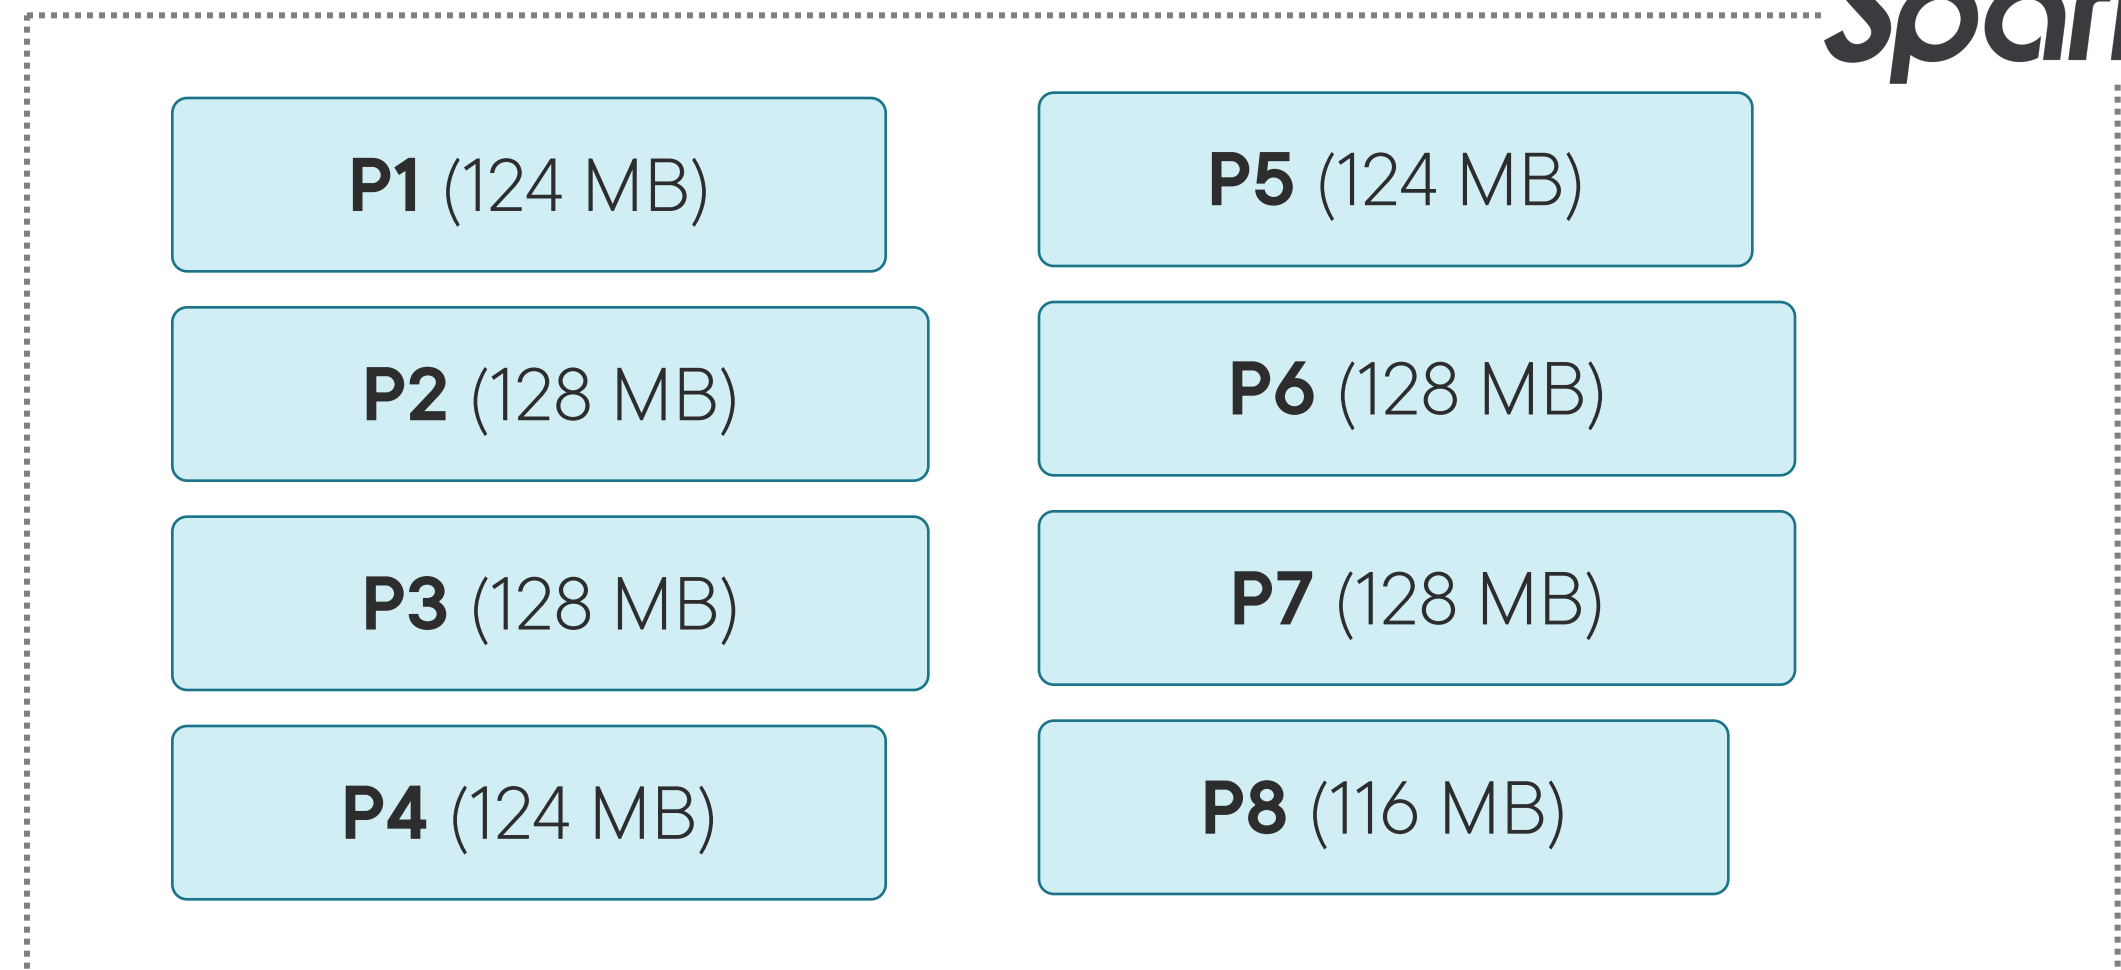



Partitions may be lesser than default parallelism setting if data size is small

1. For Reading Data

`spark.default.parallelism = 4` [default = no. of cores]

`spark.sql.files.maxPartitionBytes = 128 MB` [default = 128 MB]




2. For Shuffling Data

`spark.sql.shuffle.partitions = 3` [default = 200]



P1



P2



P3



P4



Shuffle operation

Ex - Group By

P1




P2



P3



Impact of In-Memory Partitions



Partitions are processed in parallel

Partitions and Cores determine parallelism of a Job

Having very few/big partitions:

- Each task may need to process lot of data
- Cluster resources may be under-utilized

Having lot of/small partitions:

- Too many tasks are created
- Reduces parallelism since tasks go in waiting state

Changing DataFrame Partitions

Methods to Change DataFrame Partitions

Repartition Method

Coalesce Method

Repartition Method

Typically used to increase number of partitions

Wide transformation – performs shuffling

- Avoid using it to decrease partitions

Partitioning Options

- Round Robin – Creates equal sized partitions
- Hash – Co-locates data based on columns
- Range – Sorts & co-locates data based on columns

Use Cases

- Reduce skewness (some partitions have much more data than others)
- Reduce size of partitions (when they are too big)
- Co-locate data based on certain columns

Coalesce Method

Used for decreasing number of partitions of DataFrame

- Cannot increase partitions

Narrow transformation – no shuffling

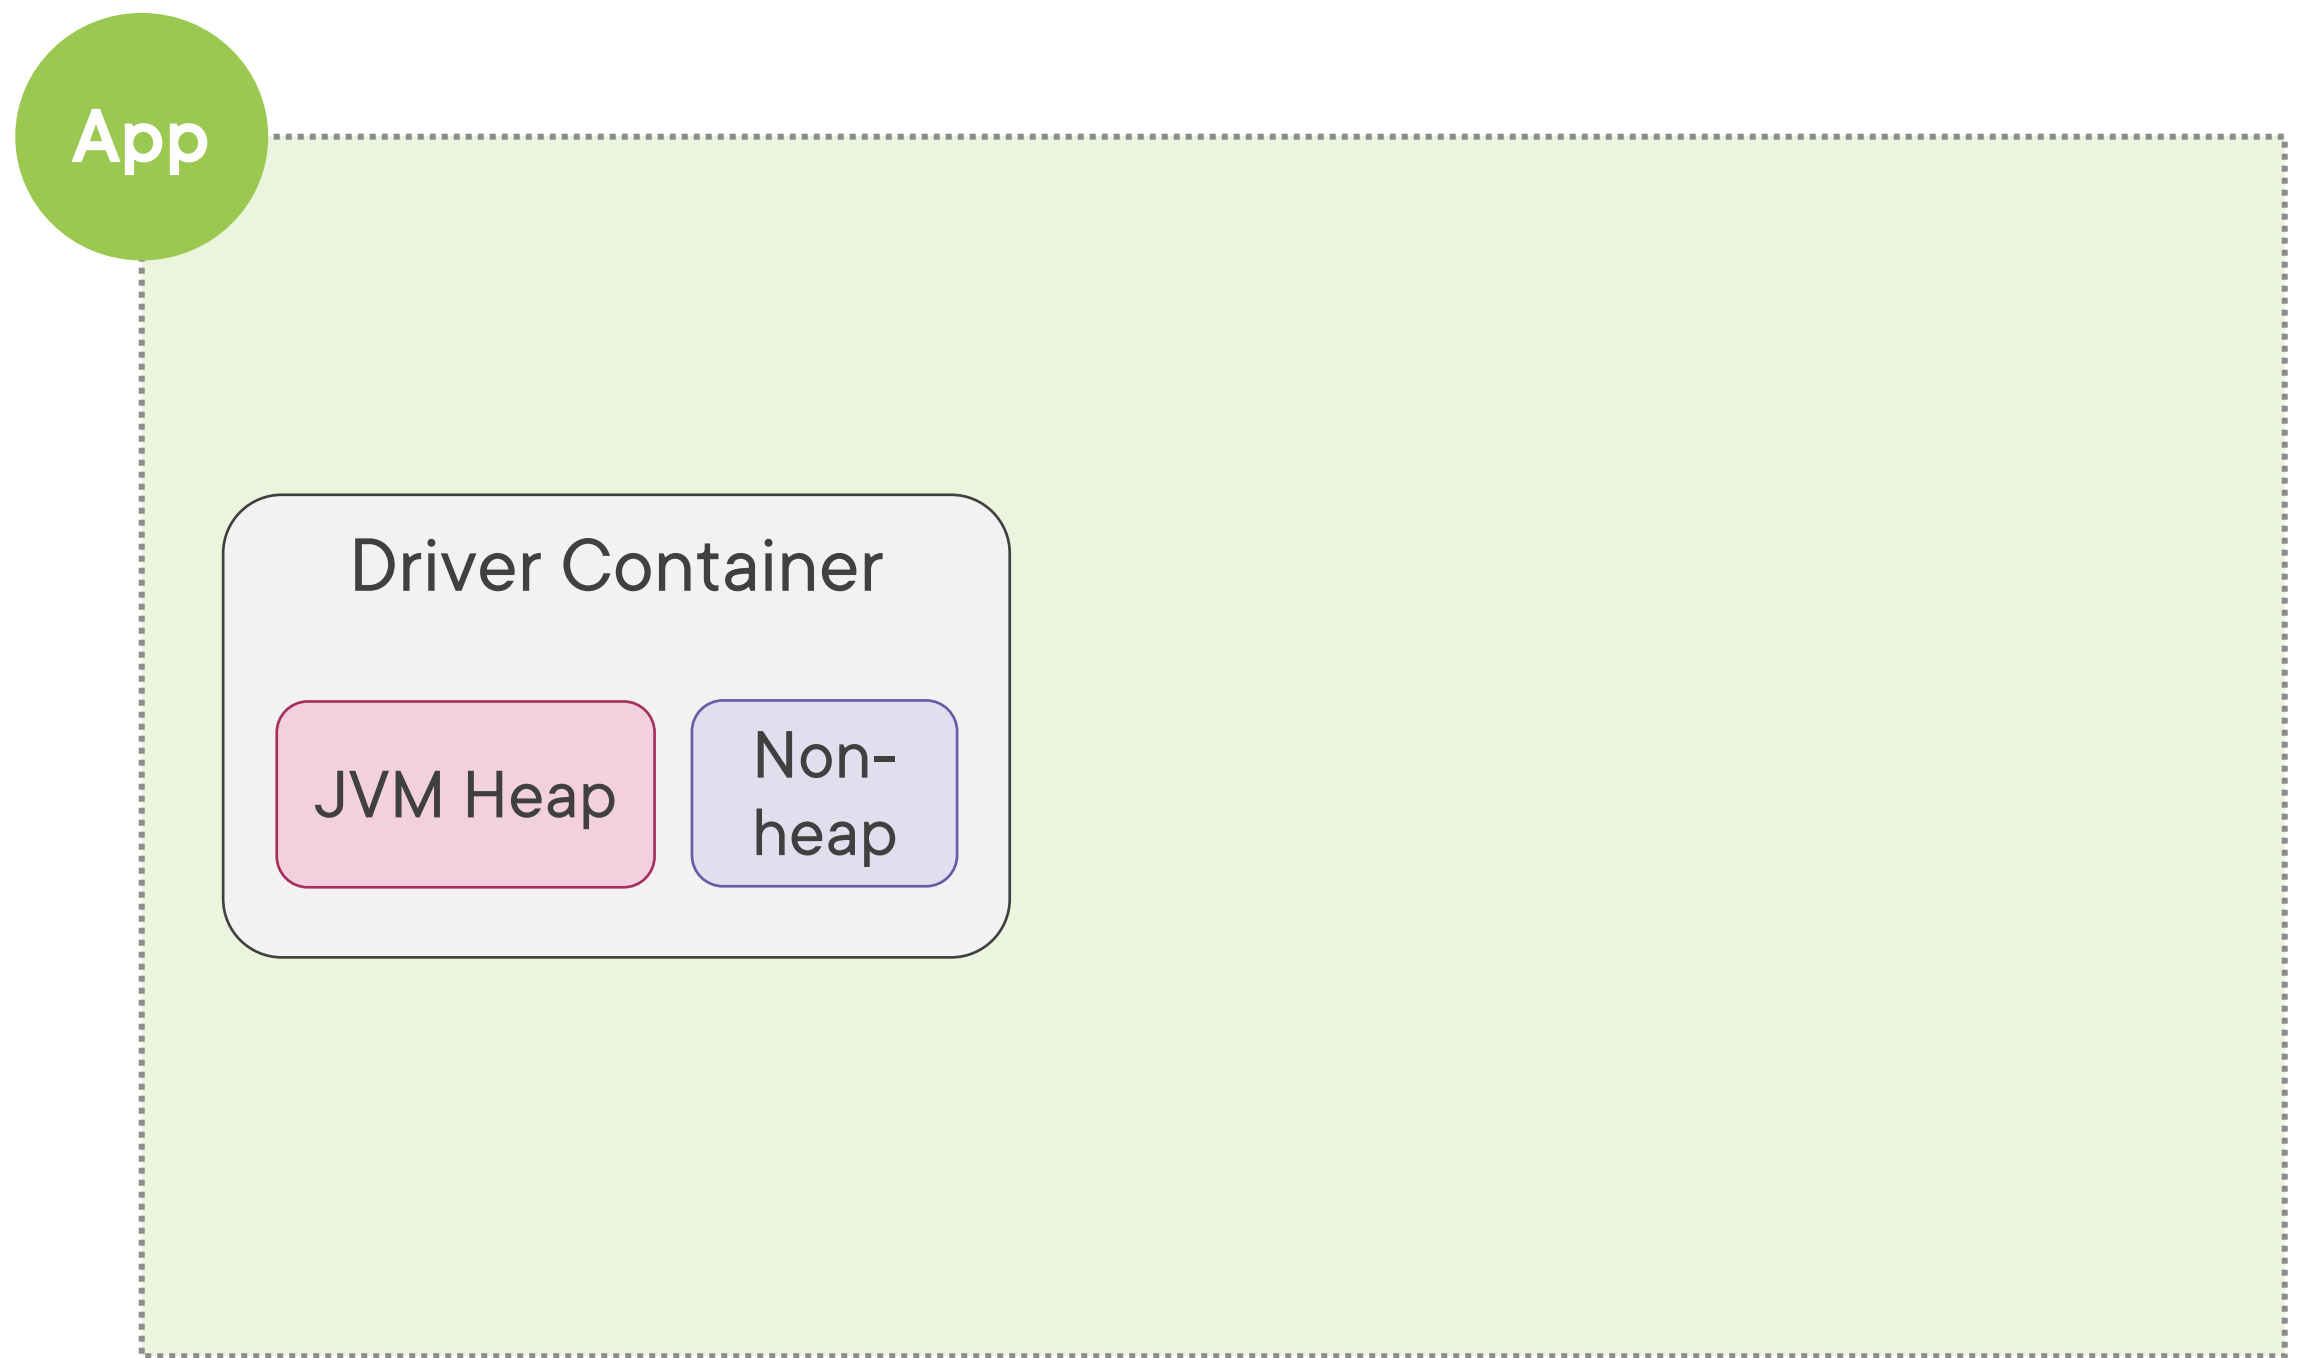
Can result in Out-of-Memory exceptions

- Eg: Existing DF = 100 partitions * 100 MB
- Coalesce(1) → New DF = 10,000 MB

Use Cases

- Partitions are very small
- Output is required in lesser number of files

Memory Management



JVM Heap memory is used for Spark activity

Overhead memory (non-heap memory) is used by non-JVM processes like VM overheads, buffers etc.

Driver Container

spark.driver.cores = 2


spark.driver.memory = 8 GB

App

Driver Container

JVM Heap

Non-
heap



Driver Container

spark.driver.cores = 2

spark.driver.memory = 8 GB

spark.driver.memoryOverhead

= max (10% of memory or 384 MB)

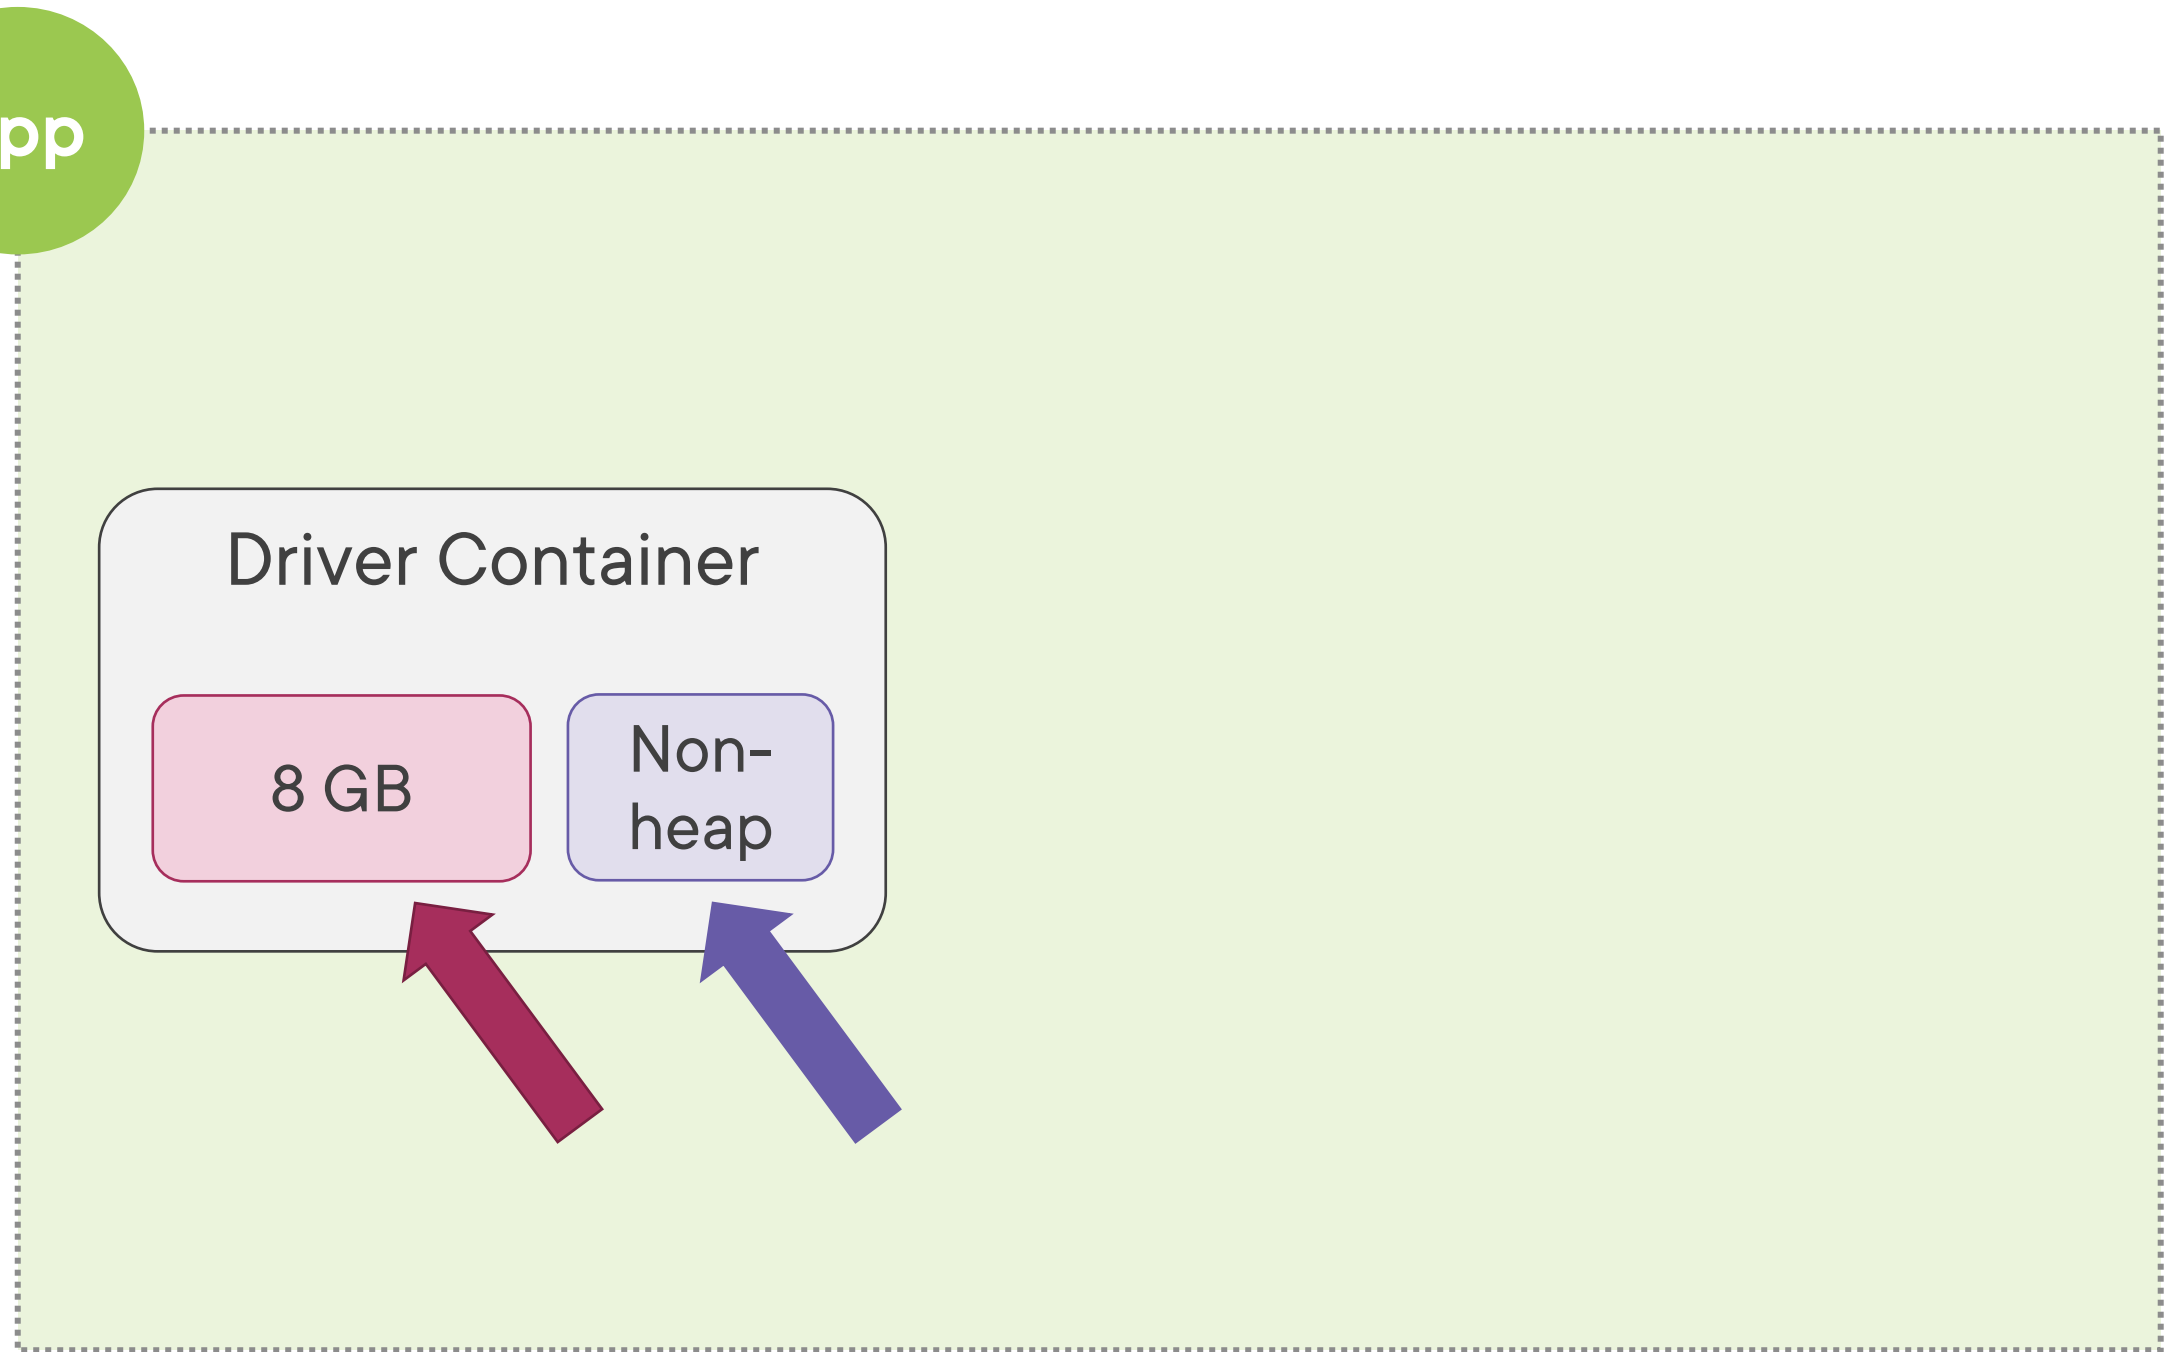
= 800MB

App

Driver Container

8 GB

Non-
heap



Driver Container

spark.driver.cores = 2

spark.driver.memory = 8 GB

spark.driver.memoryOverhead

= max (10% of memory or 384 MB)

= 800MB

App

Driver Container

8 GB

800
MB



Driver Container

spark.driver.cores = 2

spark.driver.memory = 8 GB

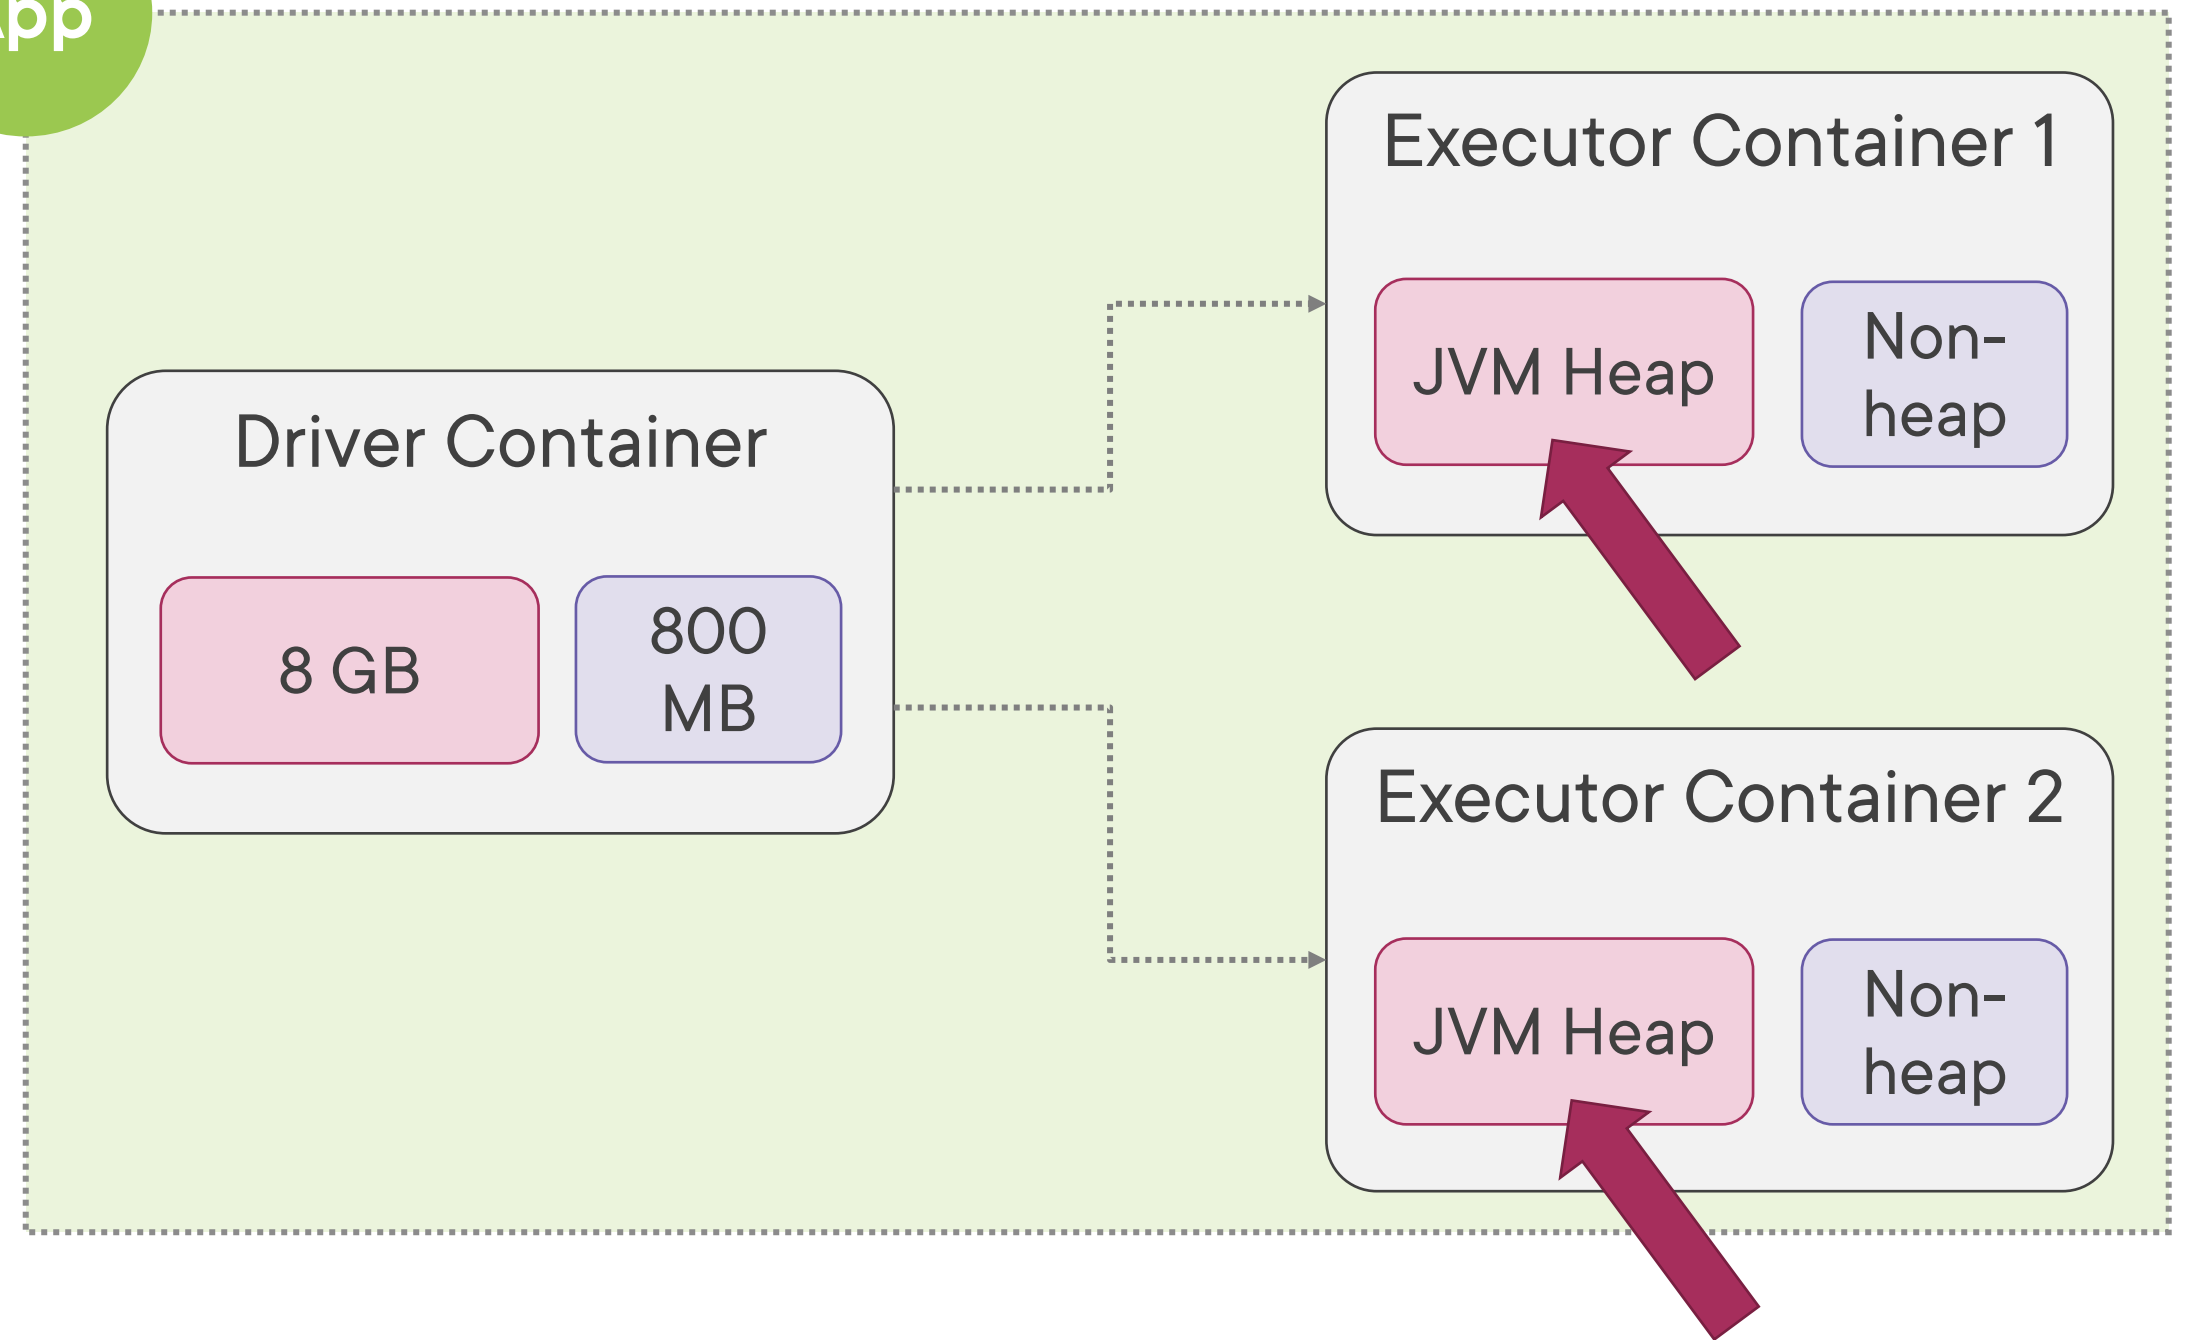
spark.driver.memoryOverhead
= max (10% of memory or 384 MB)
= 800MB

Executor Container

spark.executor.cores = 4

spark.executor.memory = 14 GB

App



Driver Container

`spark.driver.cores` = 2

`spark.driver.memory` = 8 GB

`spark.driver.memoryOverhead`
= max (10% of memory or 384 MB)
= 800MB

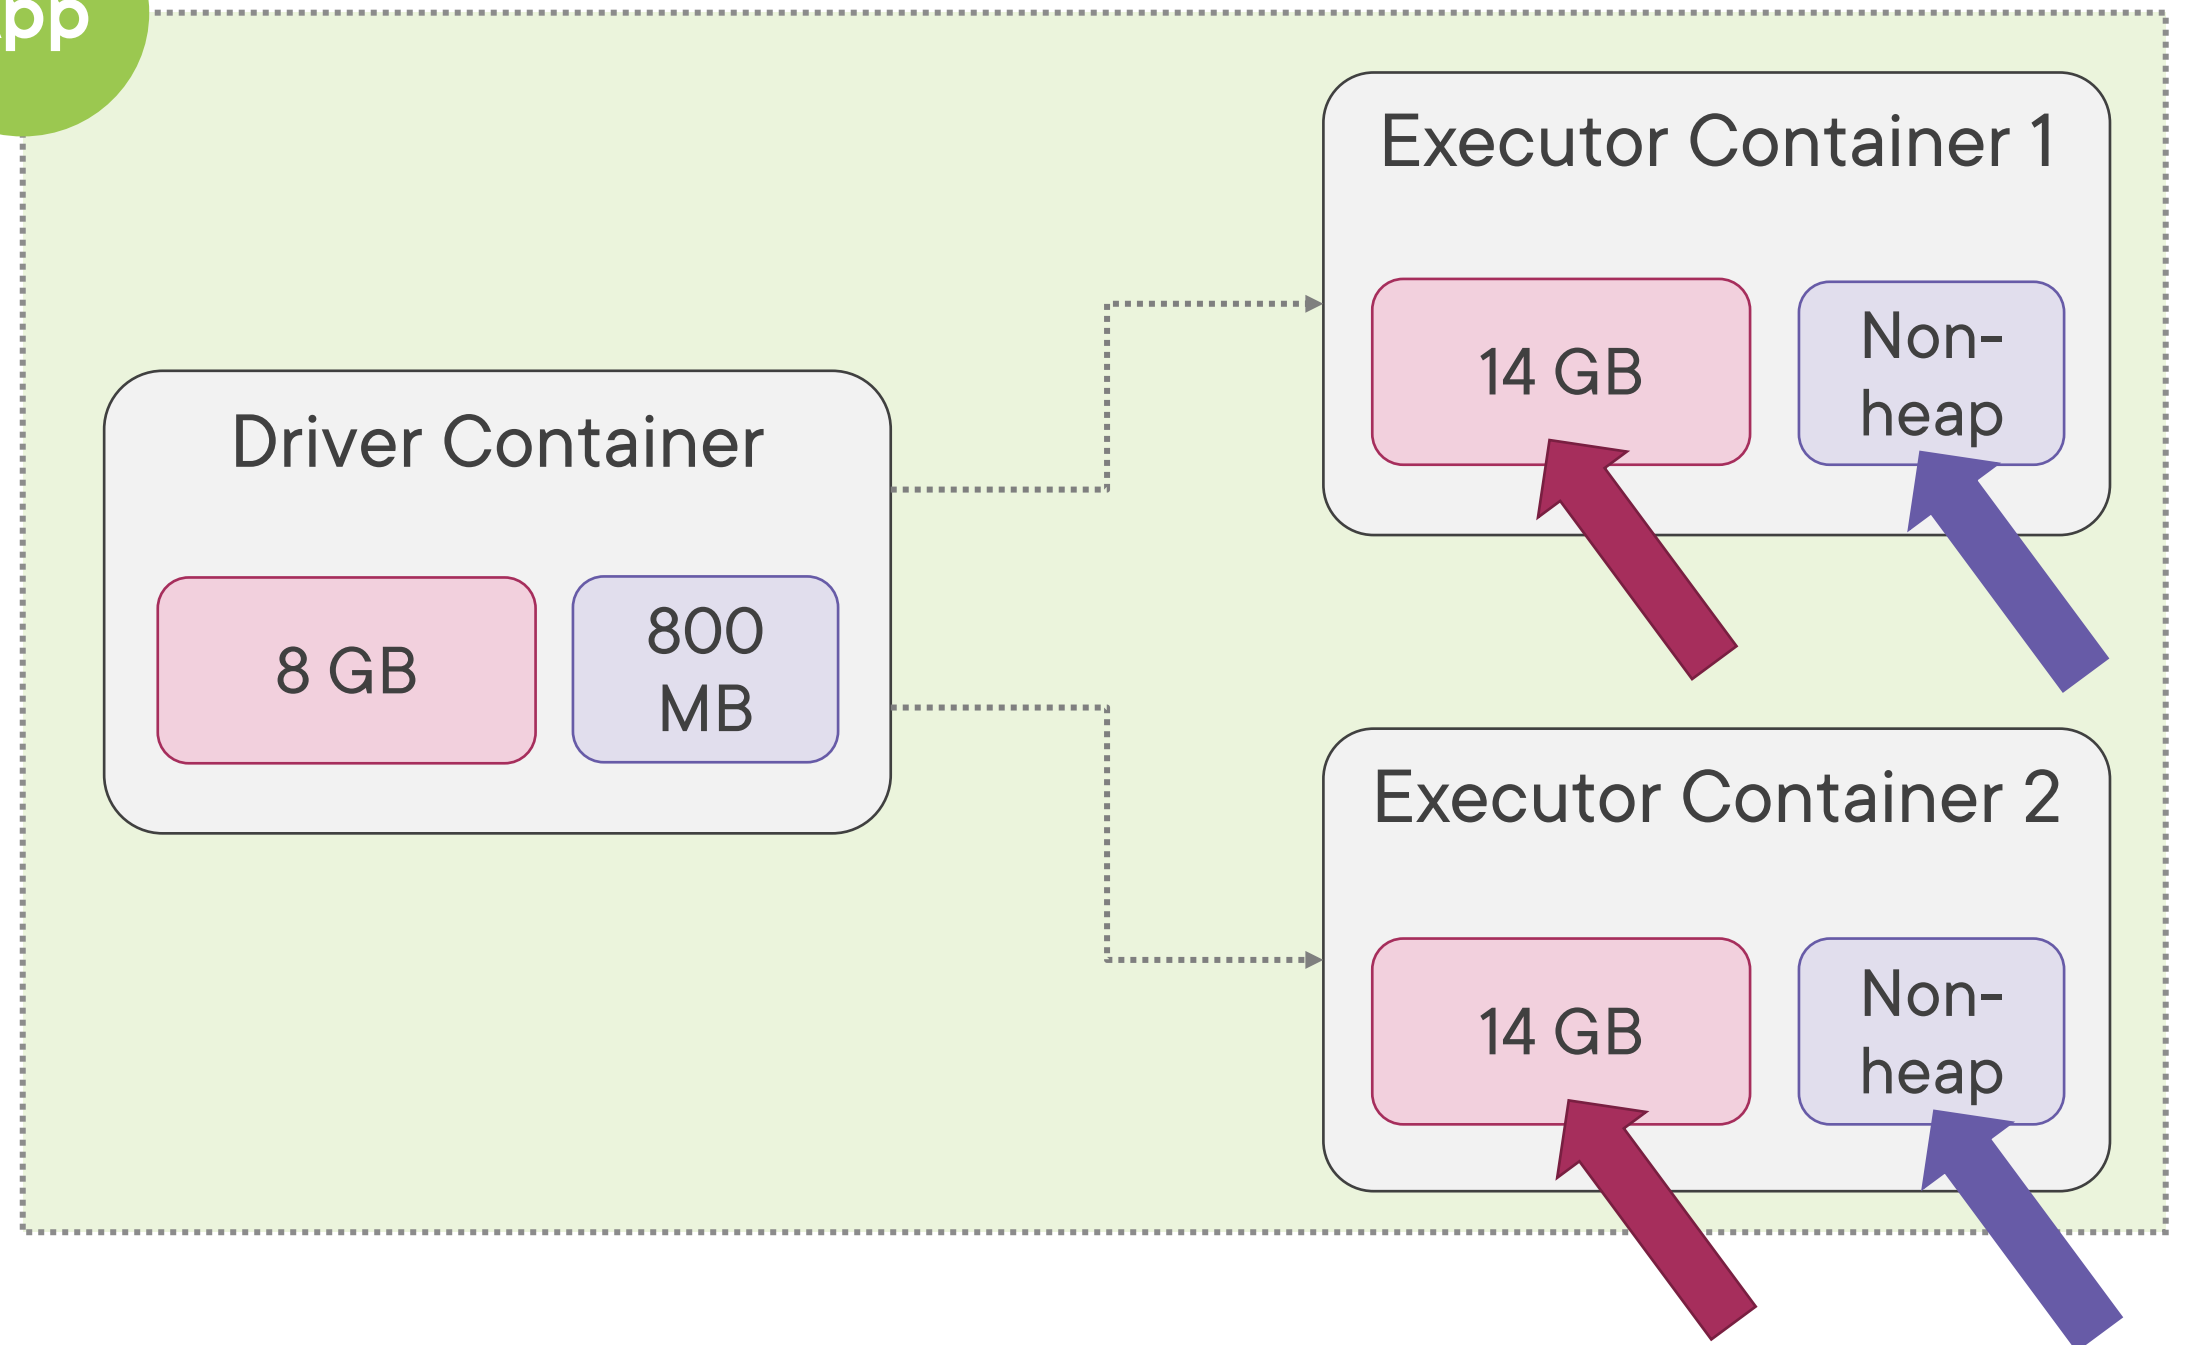
Executor Container

`spark.executor.cores` = 4

`spark.executor.memory` = 14 GB

`spark.executor.memoryOverhead`
= max (10% of memory or 384 MB)
= 1.4 GB

App



Driver Container

`spark.driver.cores` = 2

`spark.driver.memory` = 8 GB

`spark.driver.memoryOverhead`

= max (10% of memory or 384 MB)

= 800MB

Executor Container

`spark.executor.cores` = 4

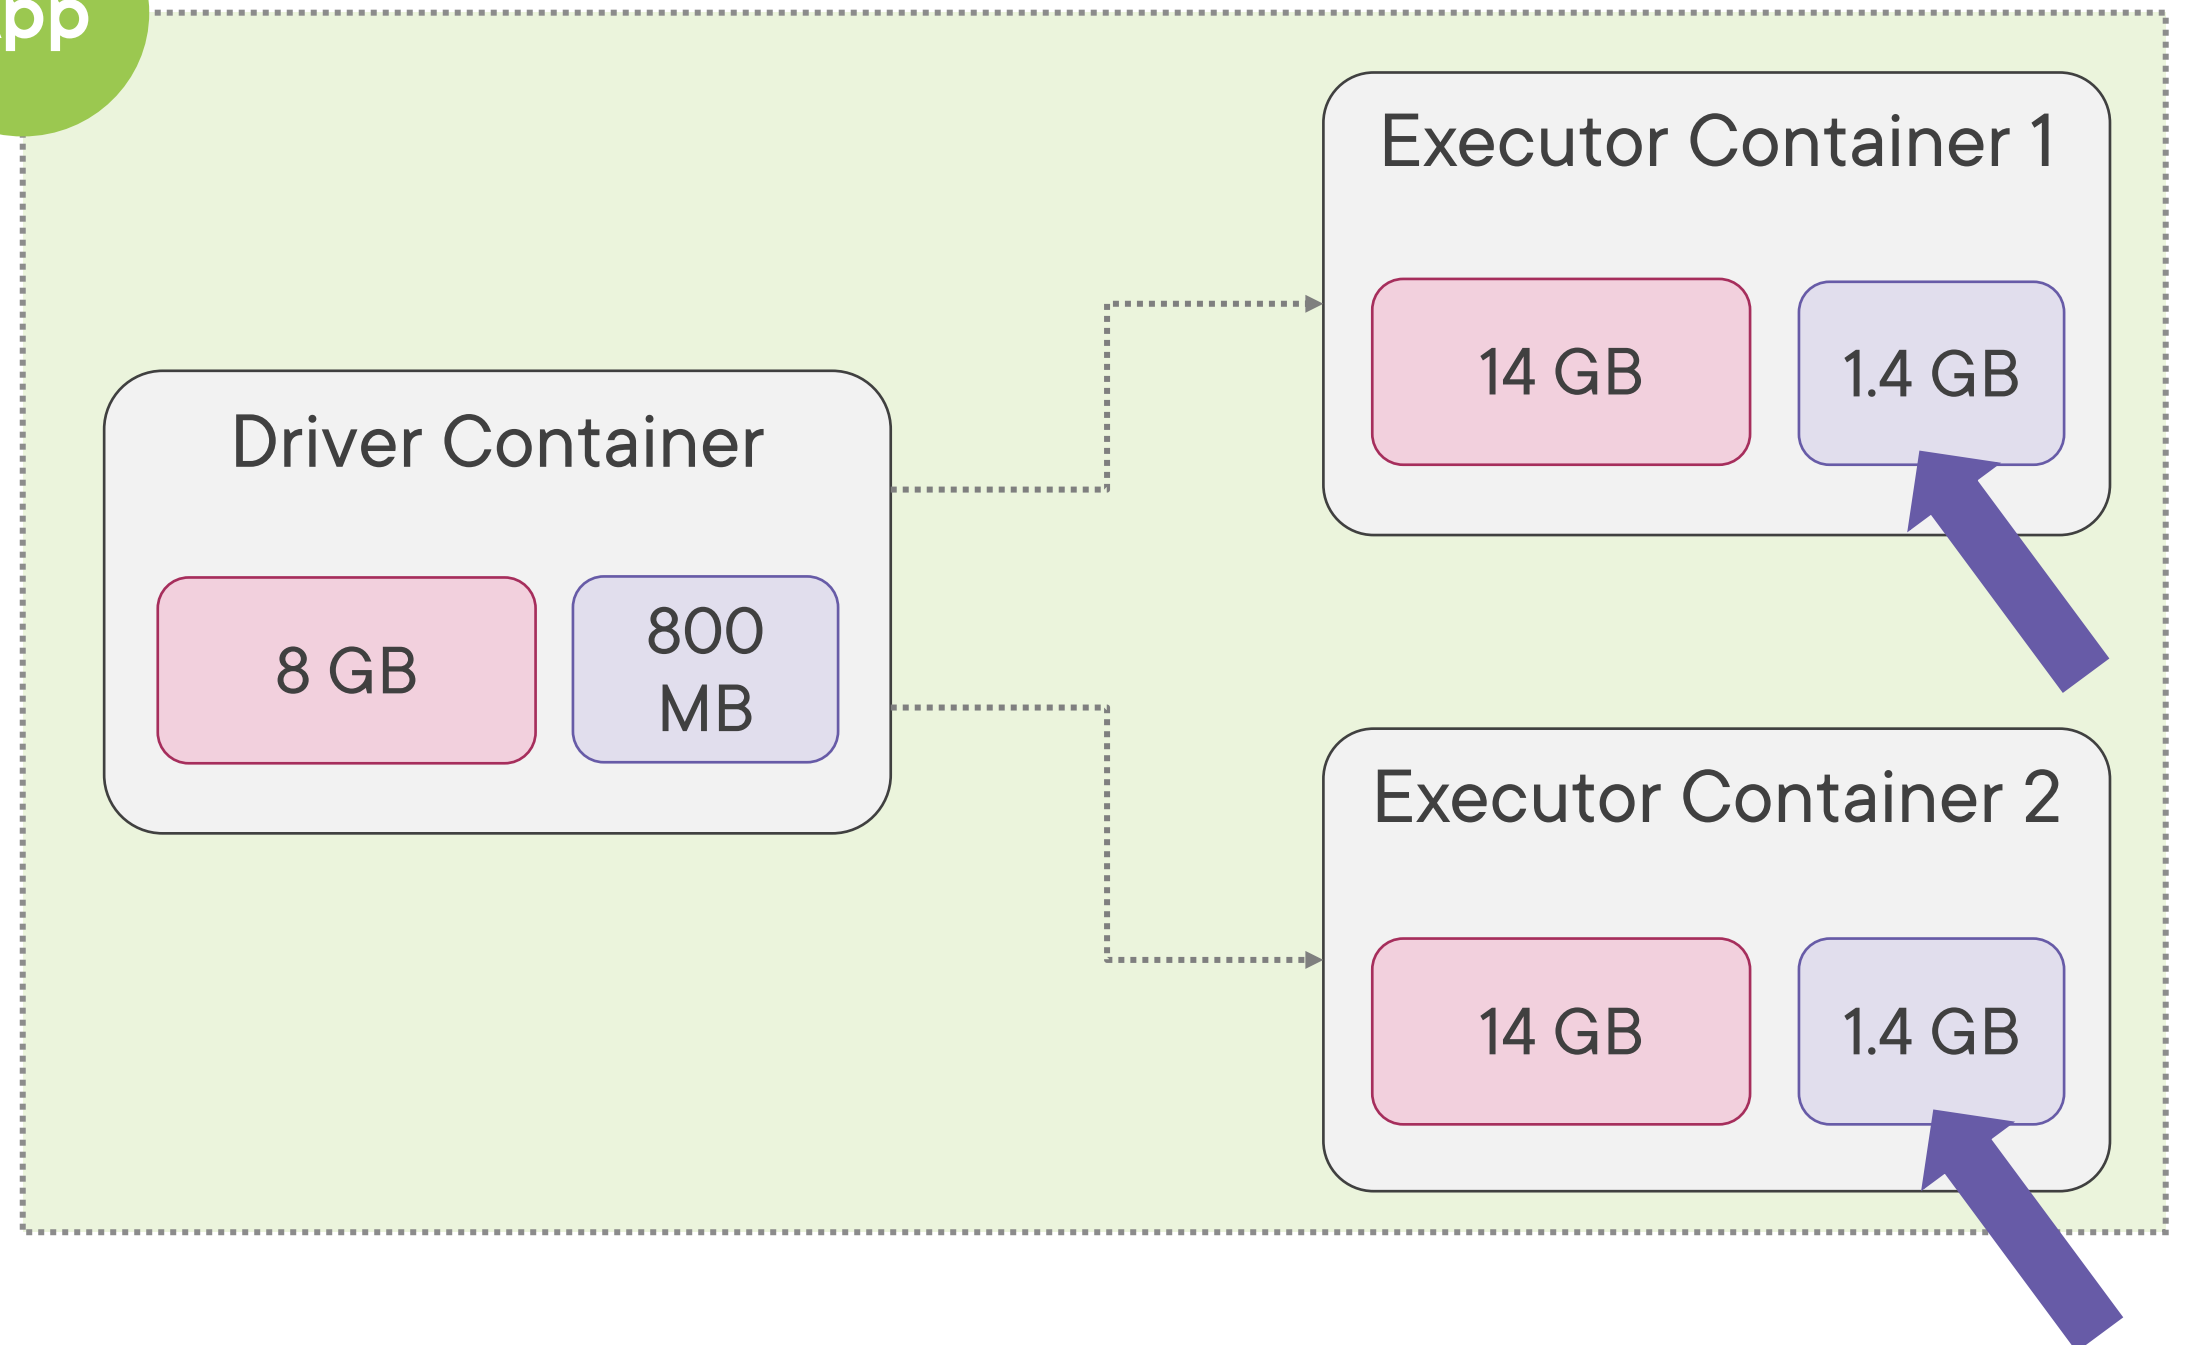
`spark.executor.memory` = 14 GB

`spark.executor.memoryOverhead`

= max (10% of memory or 384 MB)

= 1.4 GB

App



Executor JVM Heap

14,000 MB JVM Heap Memory

spark.executor.cores

= 4

spark.executor.memory

= 14 GB

Executor JVM Heap

Reserved Memory = 300 MB

Reserved by Spark for internal purposes

For task execution, caching, shuffling,
DataFrame operations etc.

Spark Memory
`spark.memory.fraction`

$(14000 \text{ MB} - 300 \text{ MB}) * 60\% = 8220 \text{ MB}$

<code>spark.executor.cores</code>	= 4
<code>spark.executor.memory</code>	= 14 GB
<code>spark.memory.fraction</code>	= 0.6

User Memory

$(14000 \text{ MB} - 300 \text{ MB}) * 40\% = 5480 \text{ MB}$

For storing data structures created by
user, metadata, RDD operations, UDFs etc.

Executor JVM Heap

Reserved Memory = 300 MB

Spark Memory = 8220 MB

Execution Memory (50%) = 4110 MB

Storage Memory (50%) = 4110 MB

User Memory = 5480 MB

For DataFrame operations, task execution, shuffling etc.

For storing cached data

<code>spark.executor.cores</code>	<code>= 4</code>
<code>spark.executor.memory</code>	<code>= 14 GB</code>
<code>spark.memory.fraction</code>	<code>= 0.6</code>
<code>spark.memory.storageFraction</code>	<code>= 0.5</code>



Executor JVM Heap

Reserved Memory = 300 MB

Execution Memory = 4110 MB

Storage Memory = 4110 MB

User Memory = 5480 MB

spark.executor.cores = 4

Task 1

Task 2

Task 3

Executor JVM Heap

Reserved Memory = 300 MB

Execution Memory = 4110 MB

Task 1Task 2Task 3

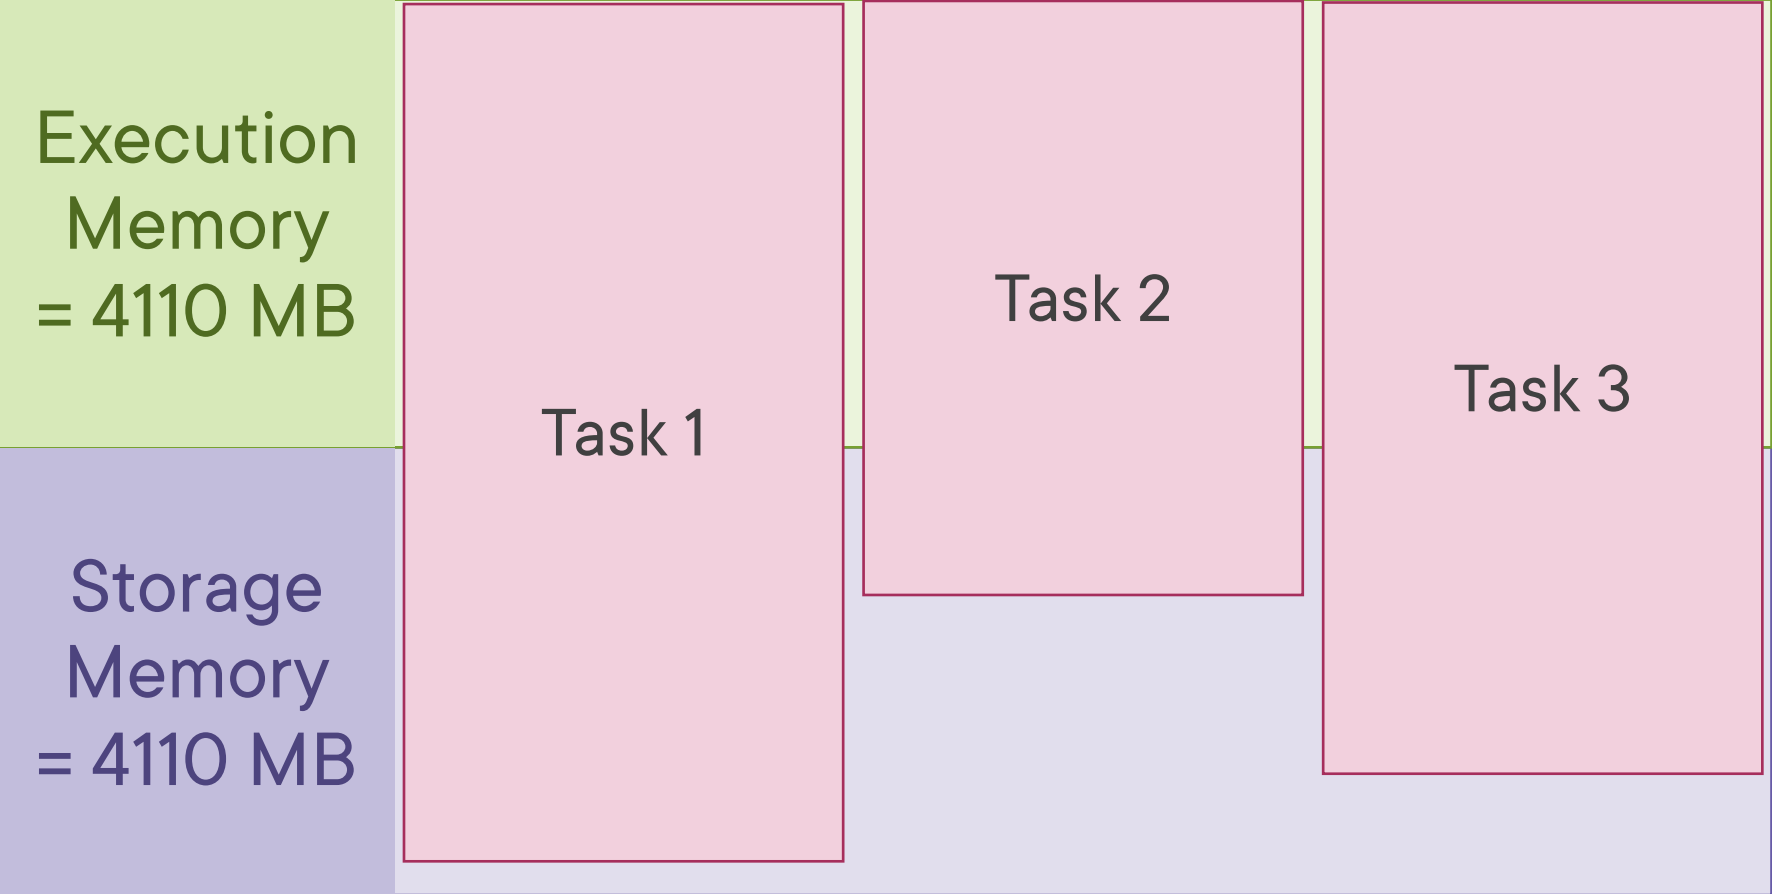
Storage Memory = 4110 MB

User Memory = 5480 MB

spark.executor.cores = 4

Executor JVM Heap

Reserved Memory = 300 MB

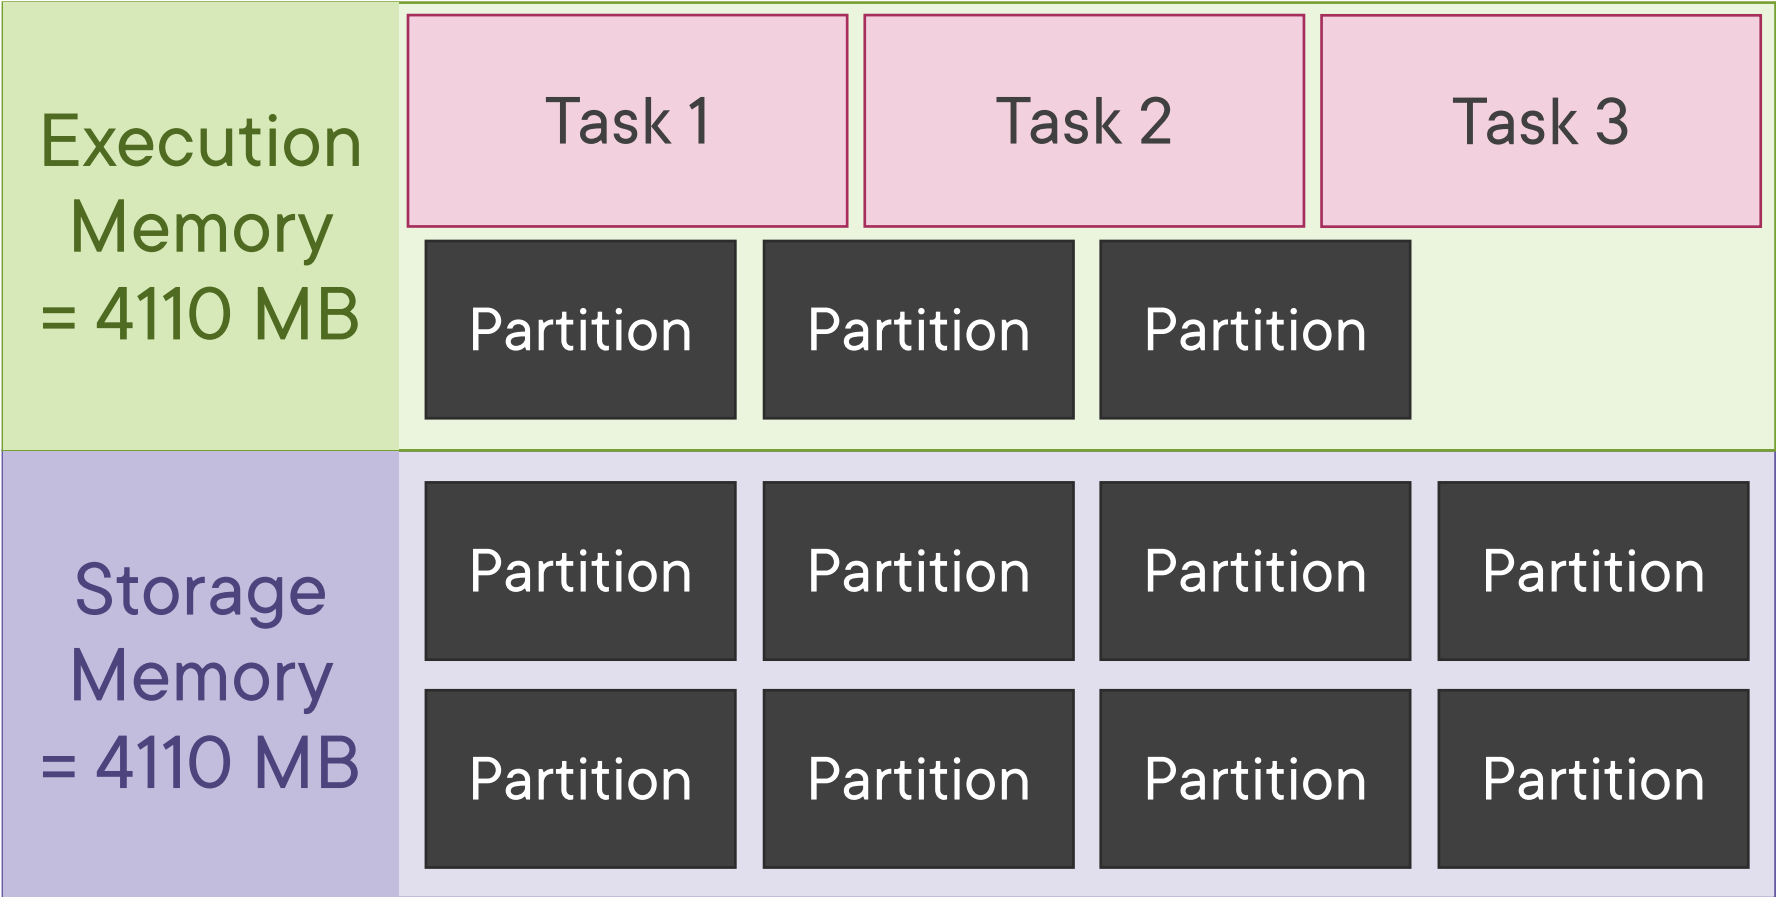


spark.executor.cores

= 4

Executor JVM Heap

Reserved Memory = 300 MB

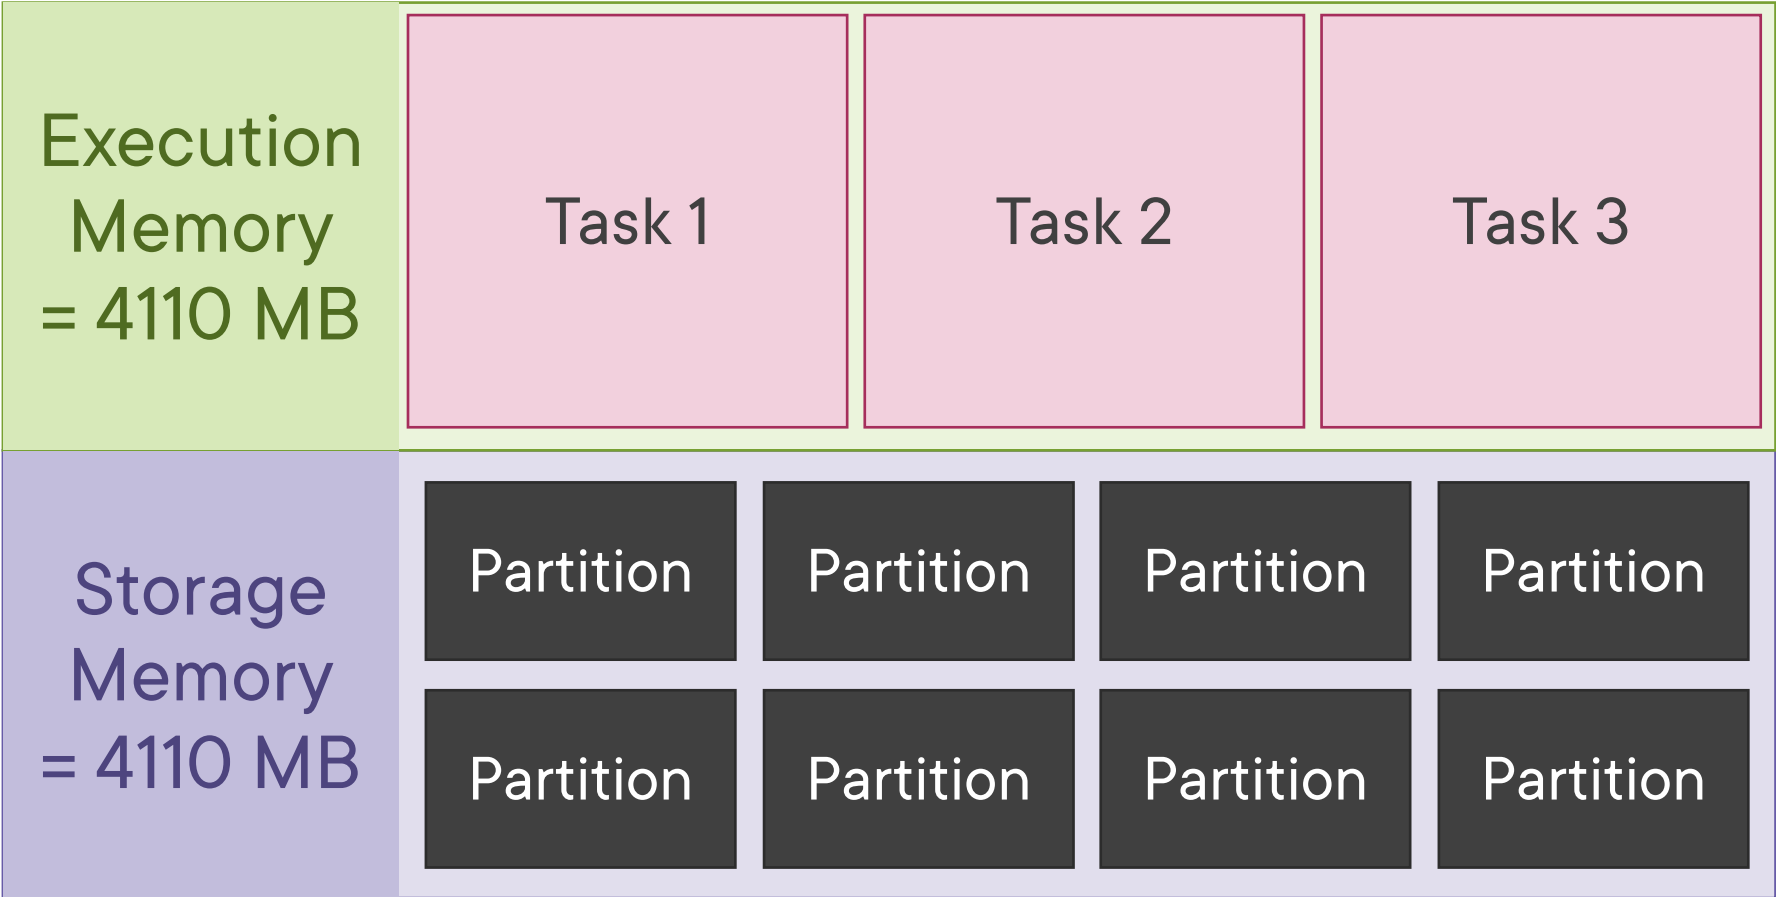


User Memory = 5480 MB

spark.executor.cores = 4

Executor JVM Heap

Reserved Memory = 300 MB

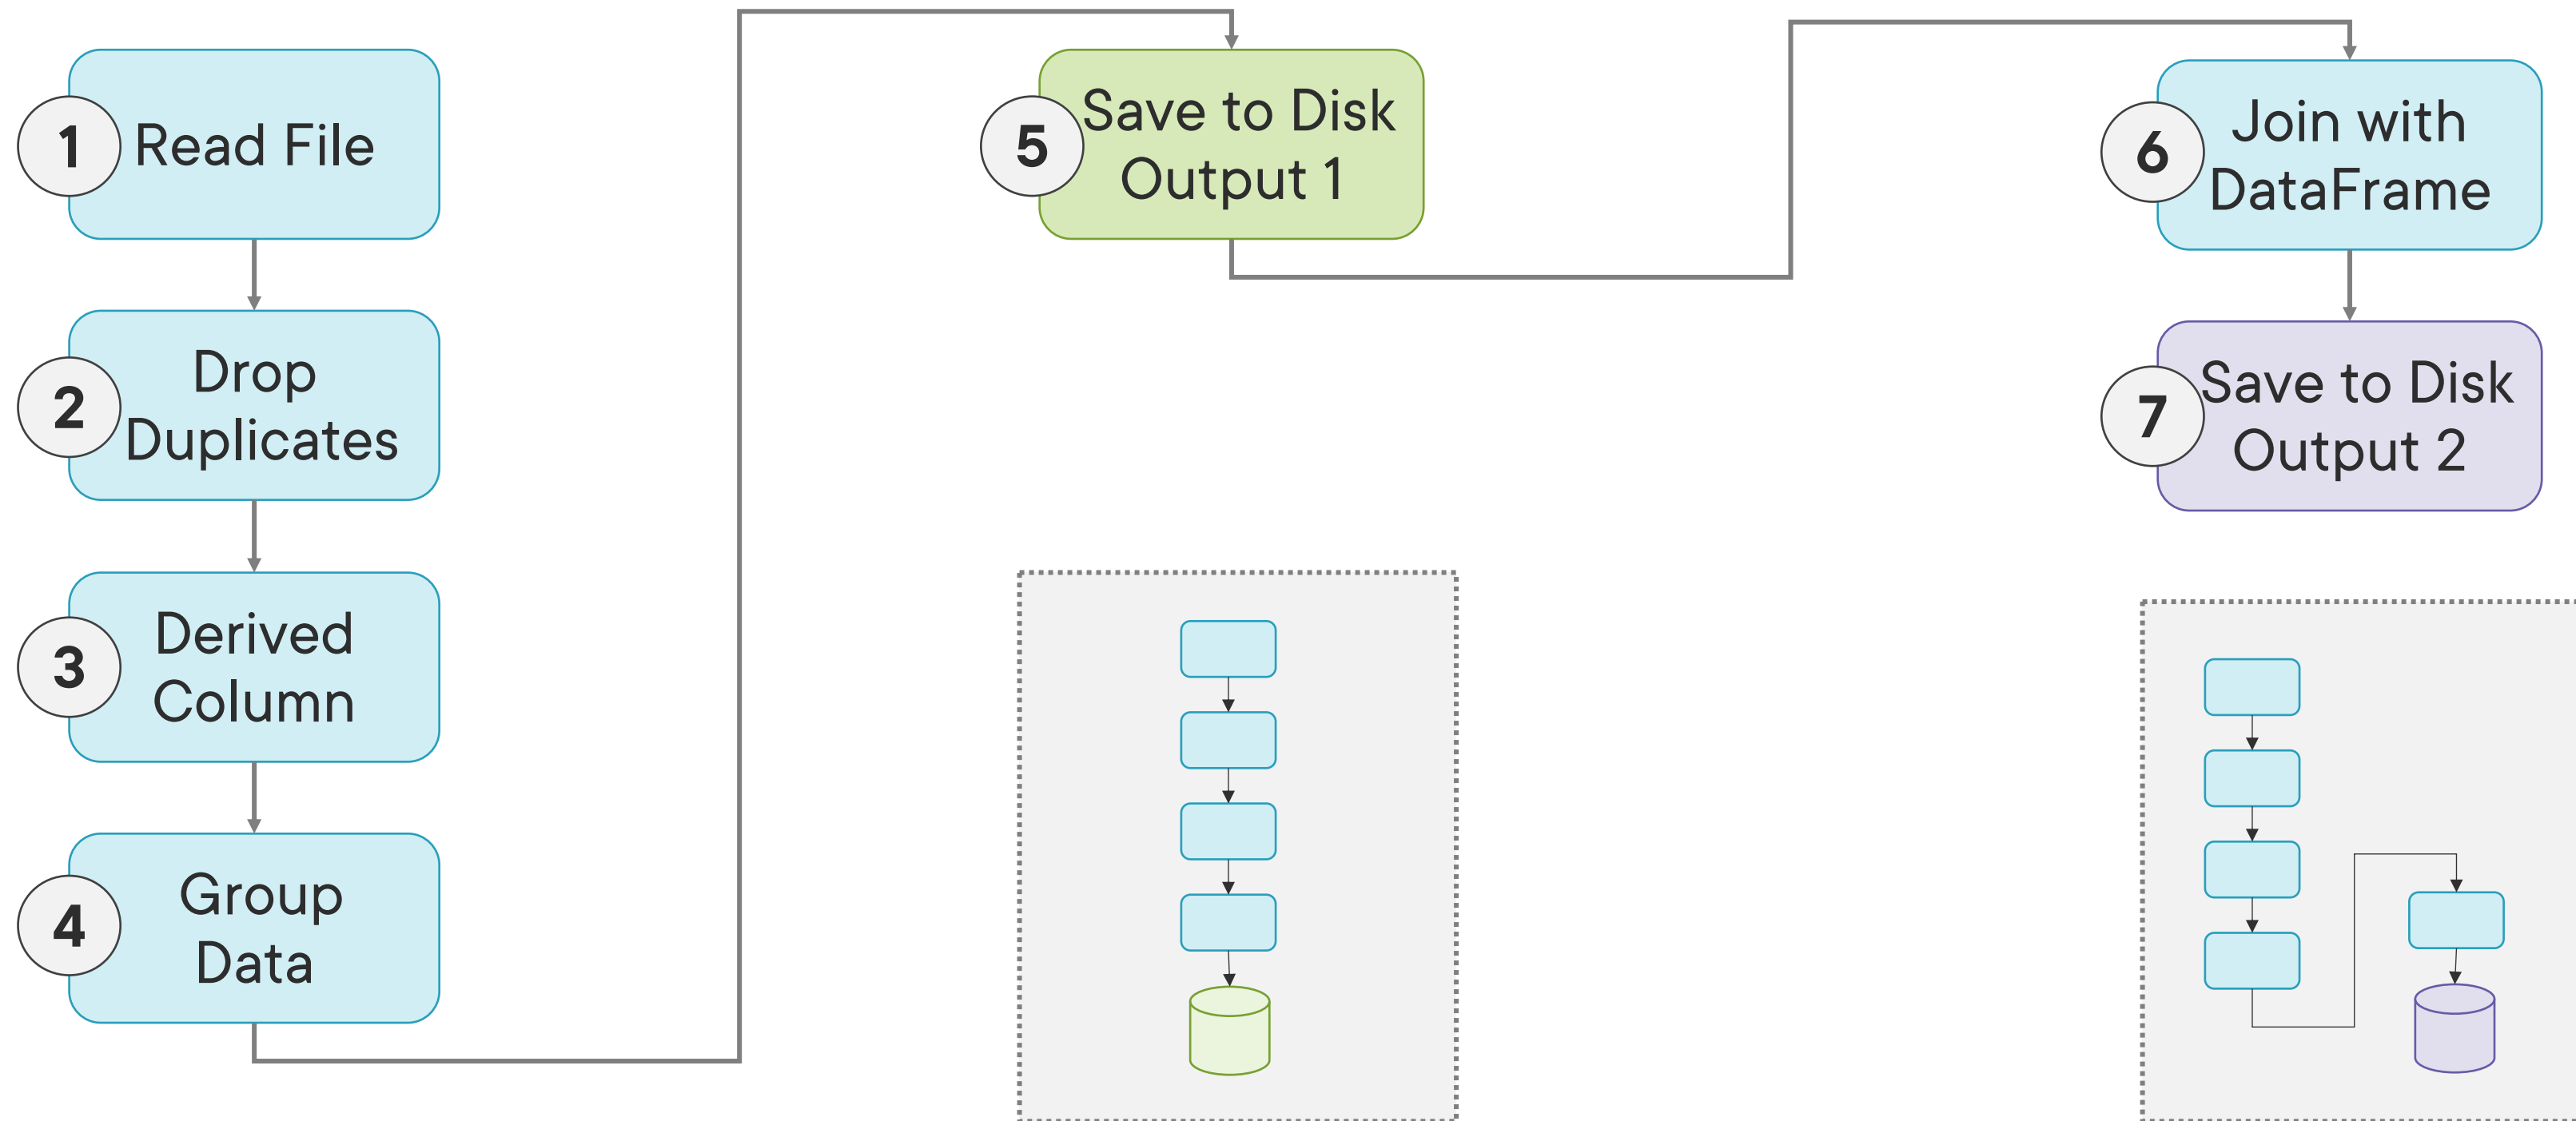


User Memory = 5480 MB

spark.executor.cores = 4

If Tasks need more memory, it may result in Out-of-Memory exceptions

Persisting Data



Transformations Operations

Action 1

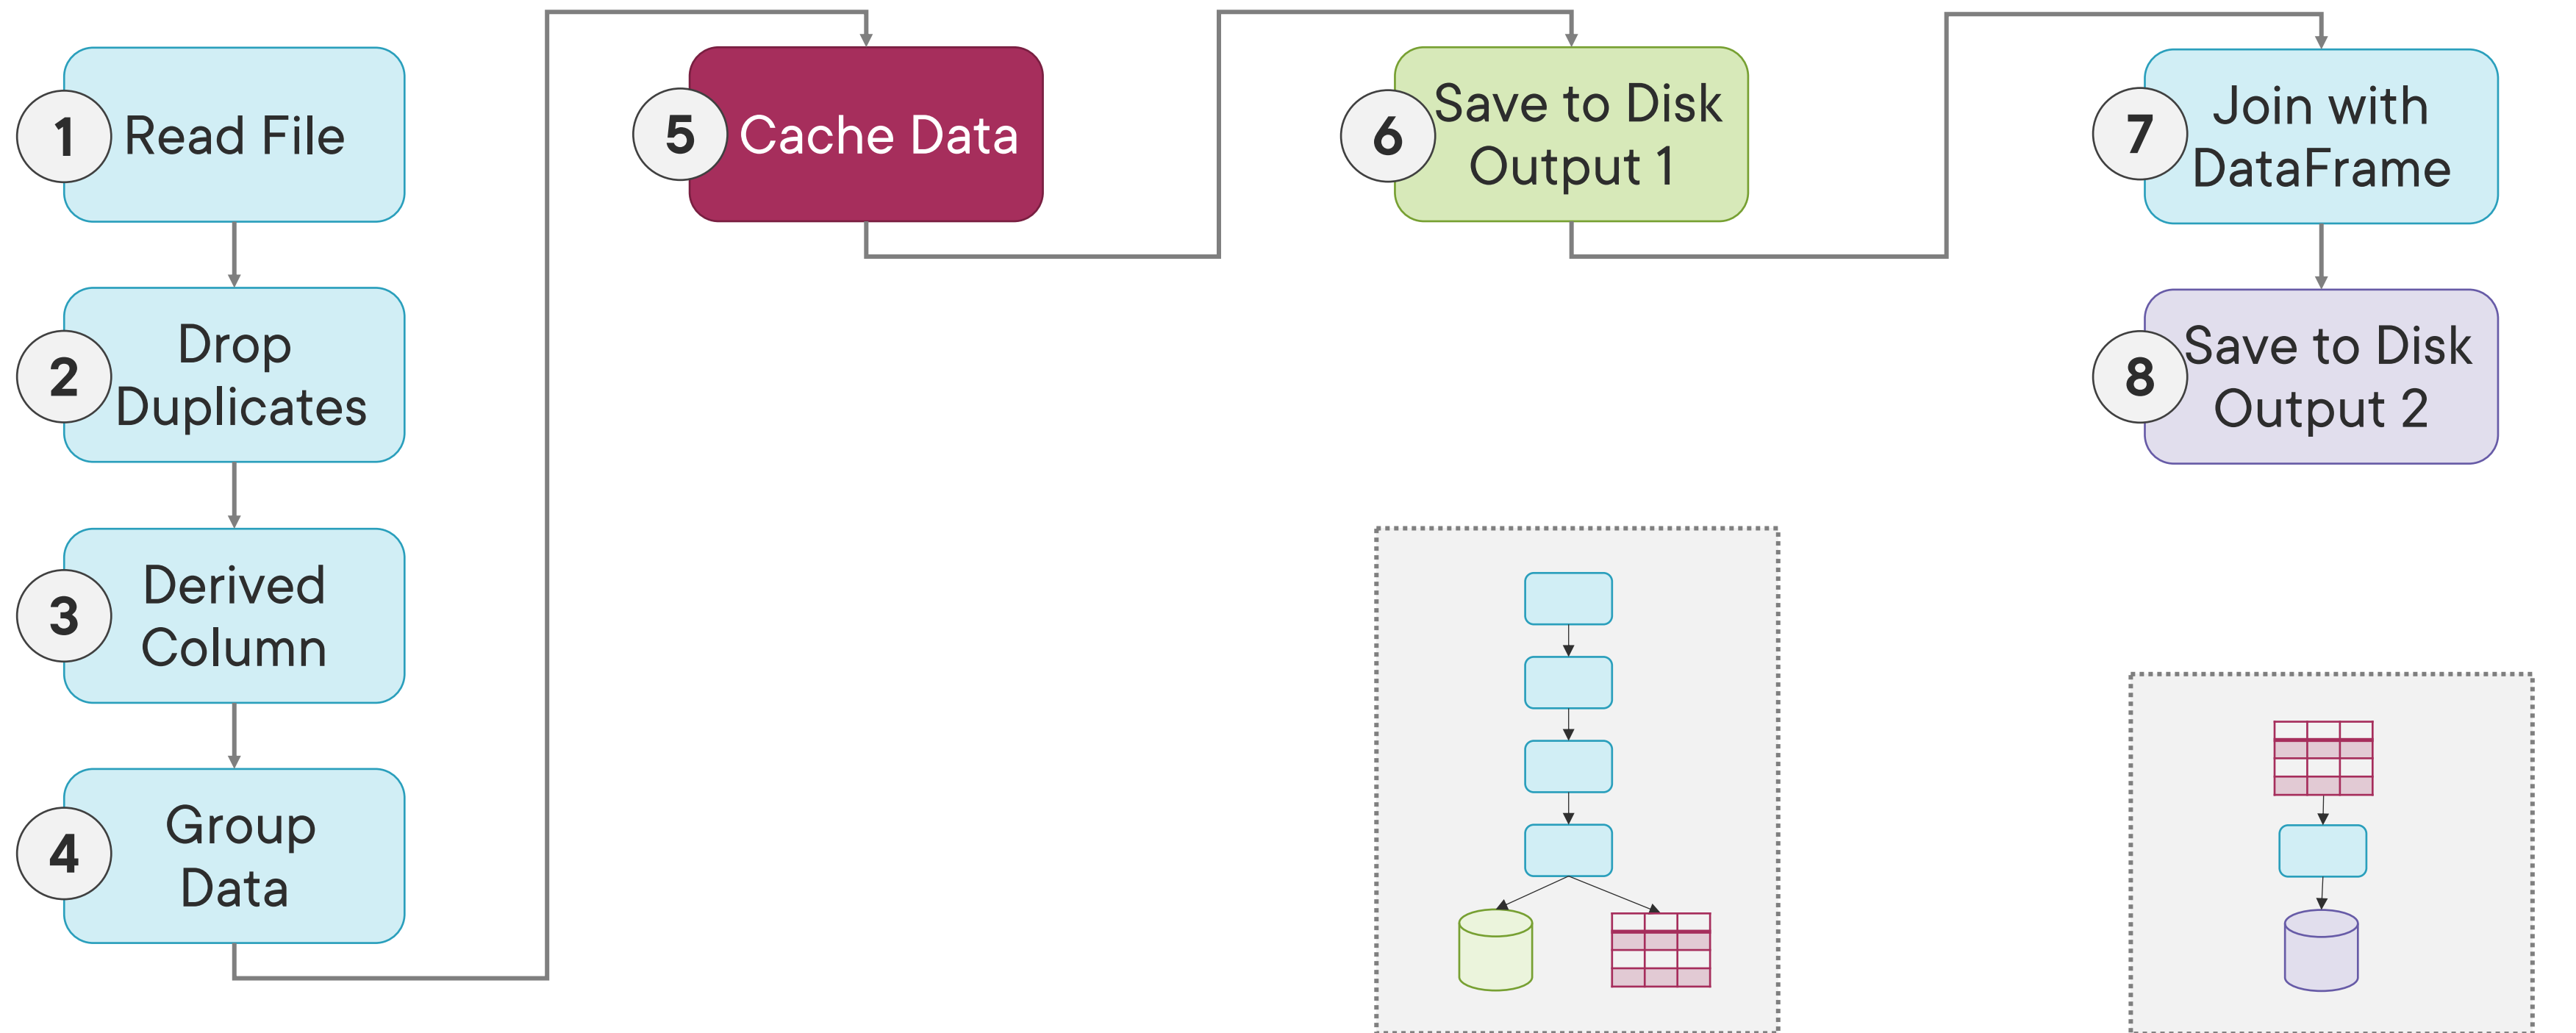
Executes all 4
transformations,
& writes to disk

Action 2

Executes all 5
transformations,
& writes to disk

Persist data to **avoid re-computation**
of complex transformations with
every Action operation

Persisting data is a Lazy operation



**Transformations
Operations**

No data is cached
Lazy operation

Action 1
Executes all 4
transformations,
caches data,
& writes to disk

Action 2
Uses cached data,
applies join,
& writes to disk

Executor JVM Heap

Reserved Memory = 300 MB

Execution Memory = 4110 MB

Task 1

Task 2

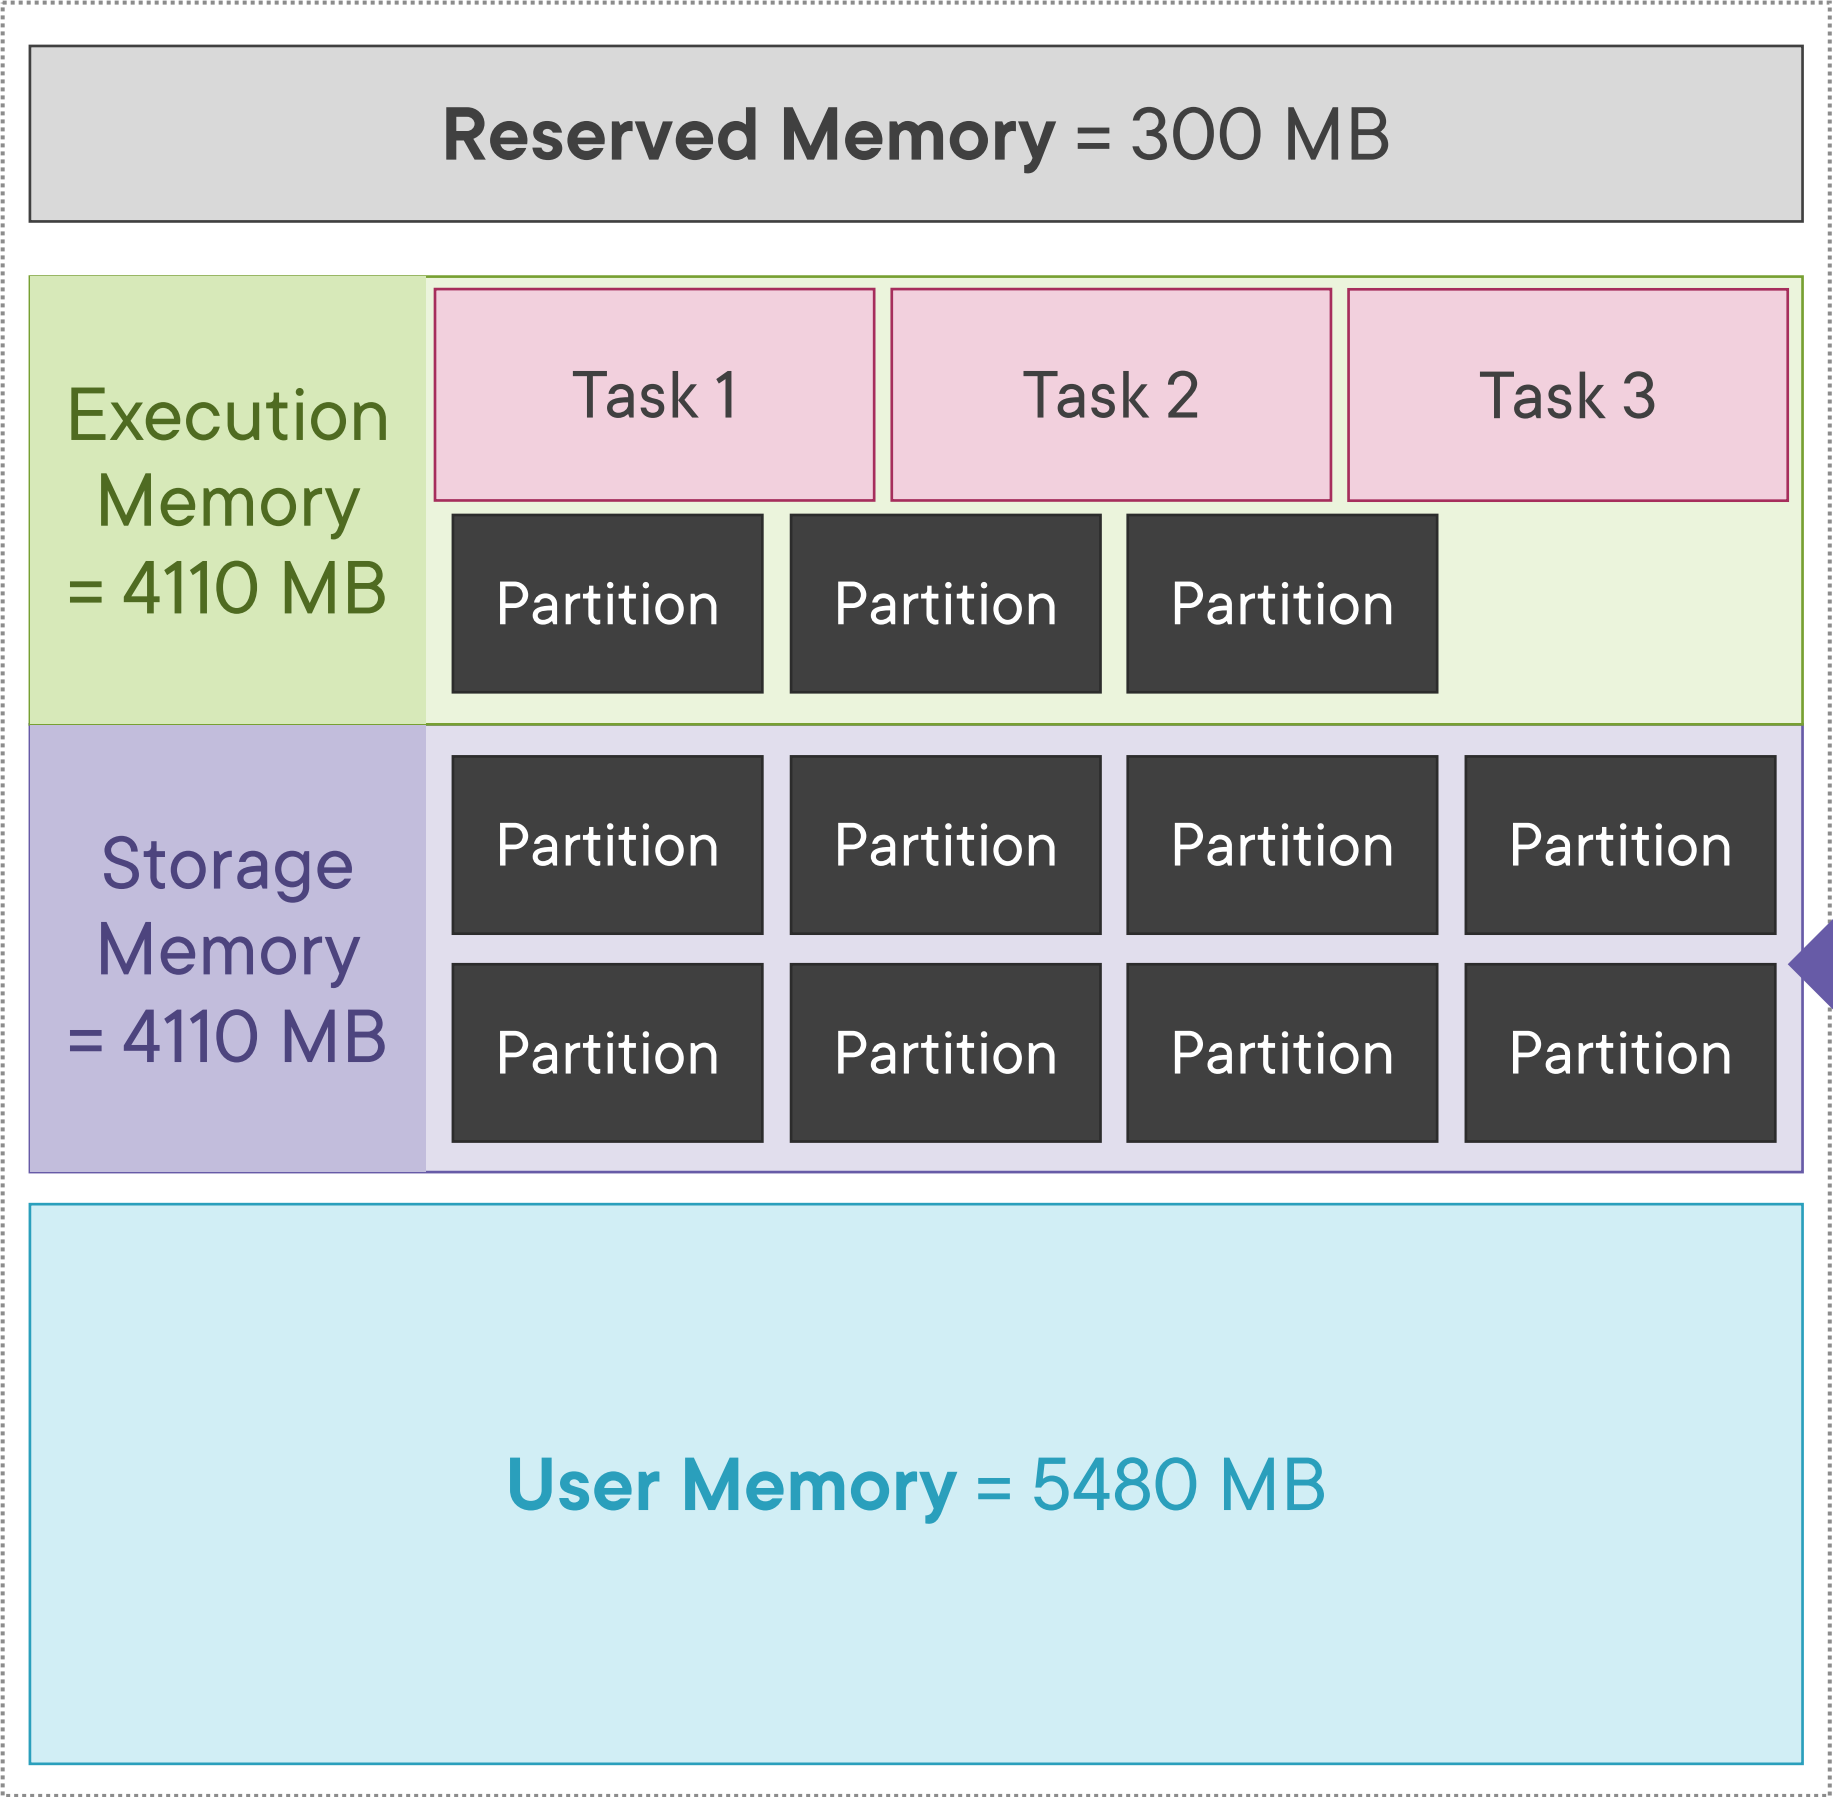
Task 3

Storage Memory = 4110 MB



User Memory = 5480 MB

Executor JVM Heap



Partitions can be cached in Storage Memory

If Execution Memory is free, partitions can be cached there too

Partitions will be evicted from Execution Memory if tasks need them

Partitions can be spilled over to disk if Storage Memory does not have enough space

Cache can be manually evicted

Partitions are evicted from cache in
LRU (Least-Recently-Used) fashion

Caching Methods

Two methods – `cache()` and `persist()`

Applying `cache()` and `persist()` without arguments

- On RDD: `MEMORY_ONLY`
- On DataFrame: `MEMORY_AND_DISK`

`cache()` has no arguments

`persist()` supports Storage Level as argument

- `MEMORY_ONLY`: Partitions that can fit in memory are cached; others are recomputed each time
- `MEMORY_AND_DISK`: Partitions that can fit in memory are cached; others are spilled to local disk of Worker
- `DISK_ONLY`: Only stored on disk; pulled when required
- *more...*

Recommended to unpersist data if not required

Spark Join Strategies and Broadcast Joins

Spark Join Strategy determines how to
move, shuffle, sort, group & merge data
across executors during join operation

Join Strategies

**Broadcast Hash
Join**


**Shuffle Sort Merge
Join**

Shuffle Hash Join

Cartesian Join


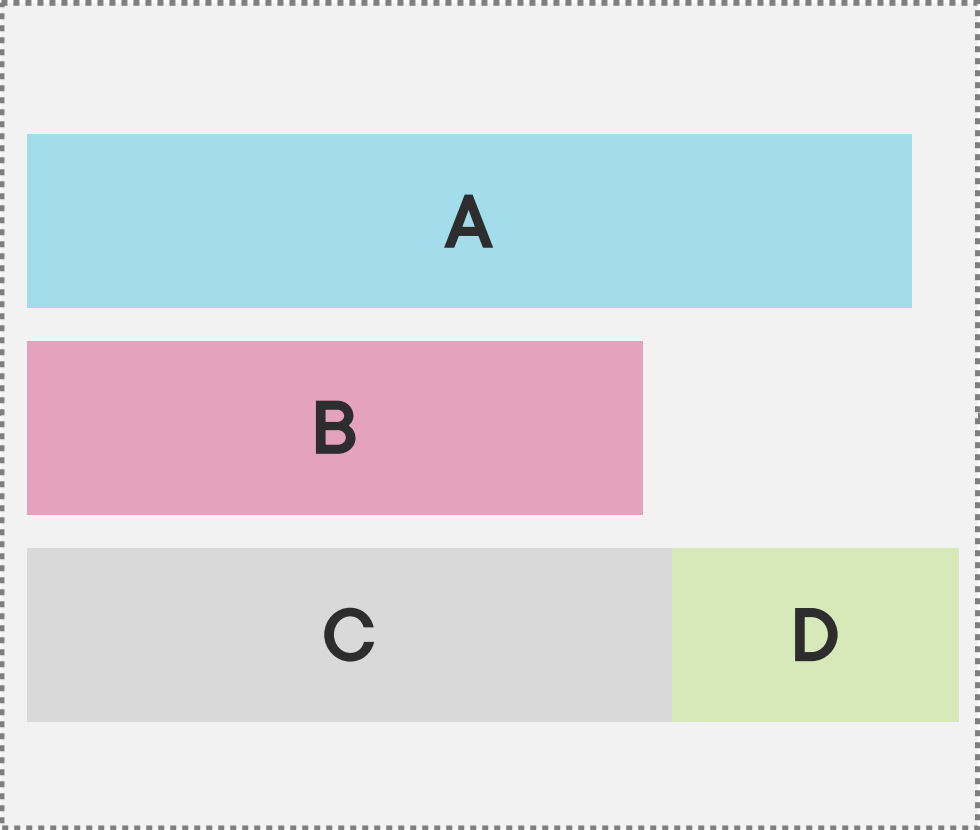
**Broadcast Nested
Loop Join**

Shuffle Sort Merge Join





Sales DF
100 mn records

Shuffle
& Sort




Products DF
10K records

Shuffle
& Sort


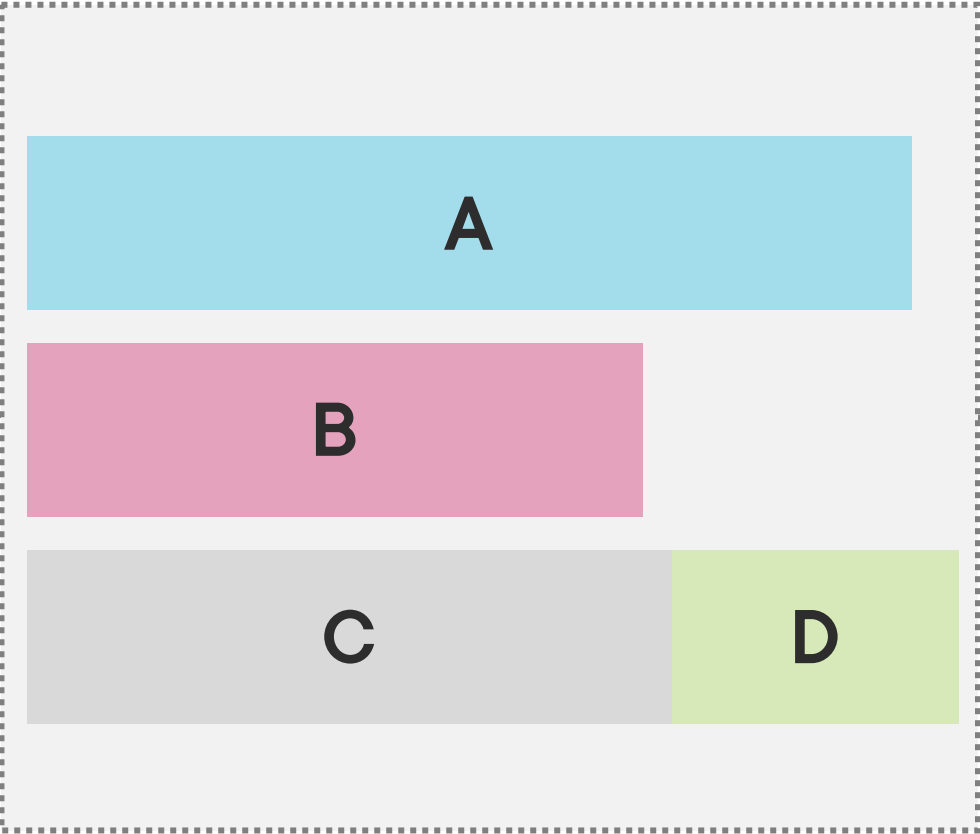


Shuffle Sort Merge Join



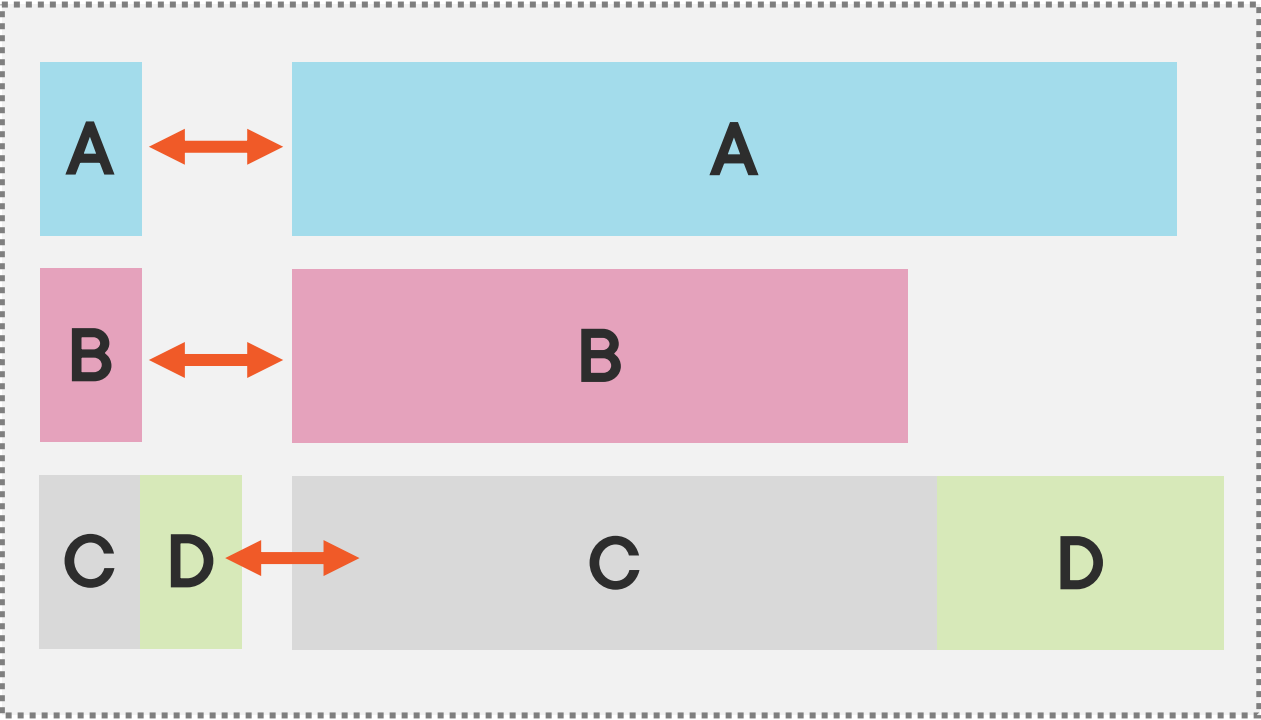

Sales DF
100 mn records

Shuffle
& Sort




Products DF
10K records

Shuffle
& Sort



Sort Merge Join

Shuffle Sort Merge Join




Involves shuffling & sorting of data for both datasets

Great for joining two large datasets


But even if one dataset is small, shuffling still happens

Expensive join strategy

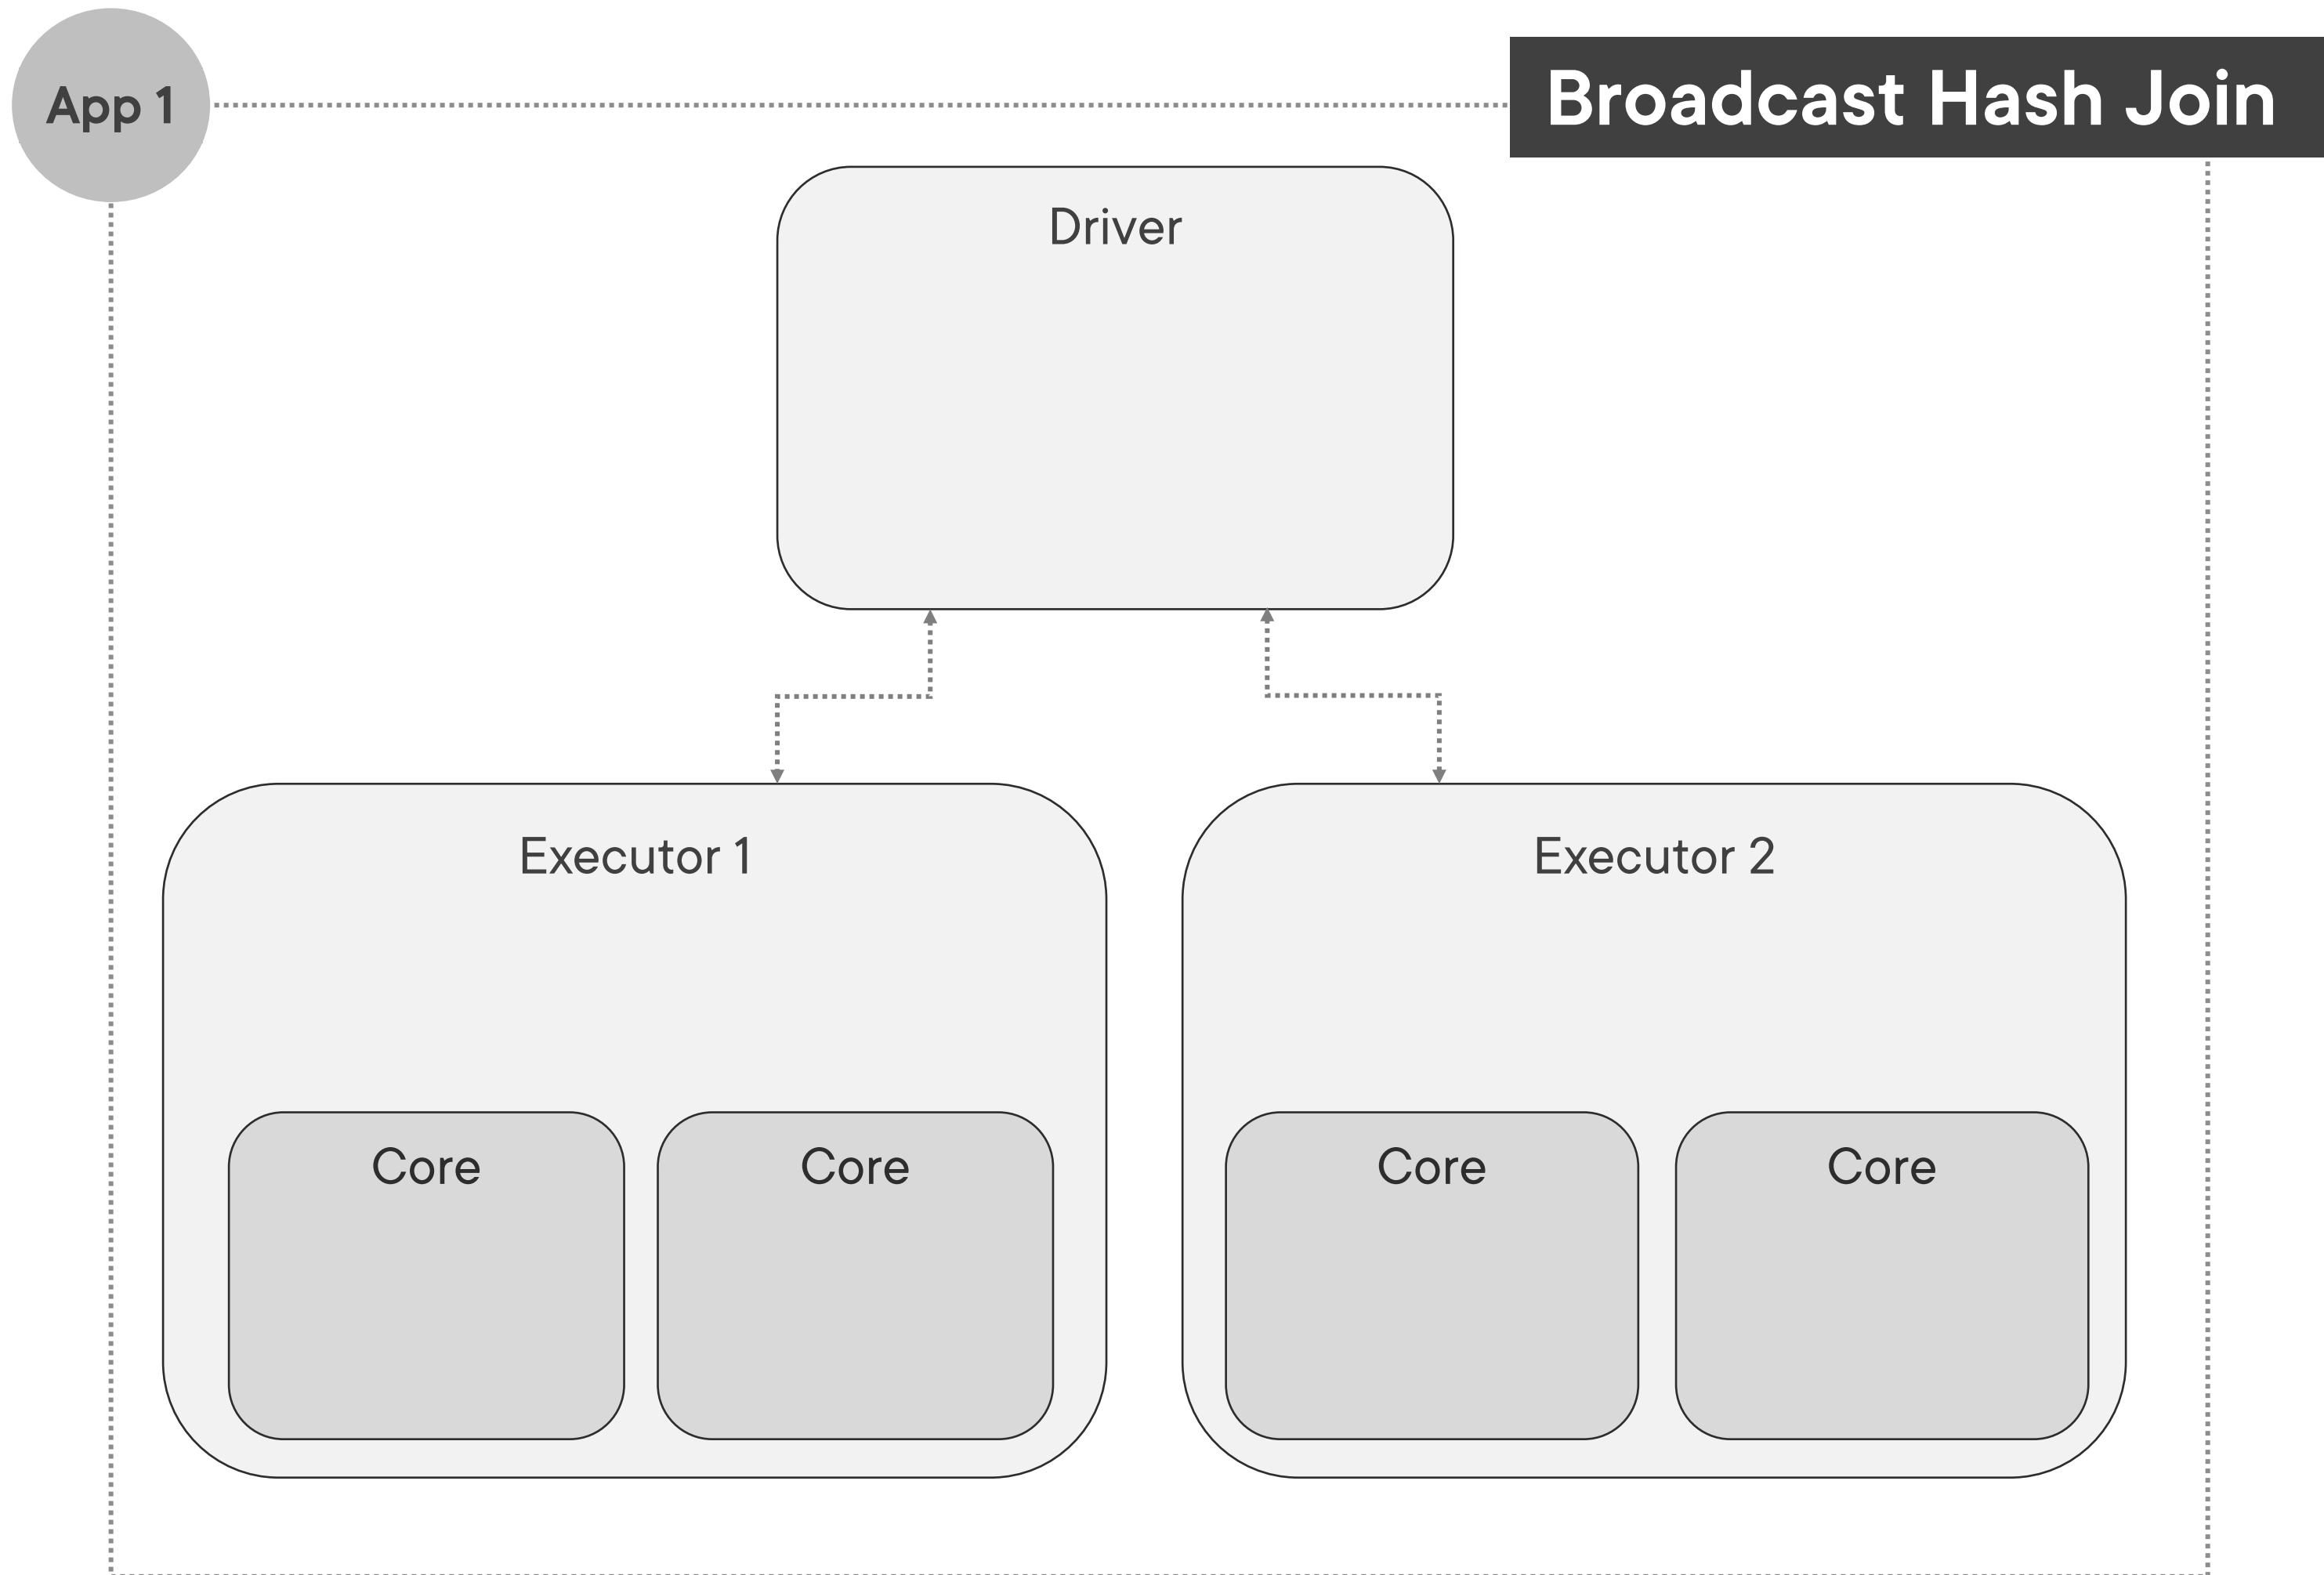
If one dataset in a join operation
is **small**, avoid shuffling using
Broadcast Hash Join



Sales DF
100 mn records



Products DF
10K records



App 1

Broadcast Hash Join

Driver

Executor 1

Executor 2

Core

Core

Core

Core

A

B

C

A

B

C

D

A

C

D

A

B

C

Sales P1

Sales P2

Sales P3

Sales P4

Sales DF

100 mn records

A

C

D

B

Products DF

10K records

App 1

Broadcast Hash Join

Driver

Executor 1

Executor 2

Products

A

C

D

B

Products

A

C

D

B

Core

Core

Core

Core

A

B

C

A

B

C

D

A

C

D

A

B

C

Sales P1

Sales P2

Sales P3

Sales P4

Sales DF

100 mn records

Products DF

10K records

Broadcast Hash Join



Useful when one of datasets is small

Copy of smaller dataset is sent to each executor once


No data shuffling required

Non-expensive & fastest join operation


Spark can perform auto broadcast or can be forced

Optimizing Shuffle Sort Join with Bucketing

Shuffle Sort Merge Join

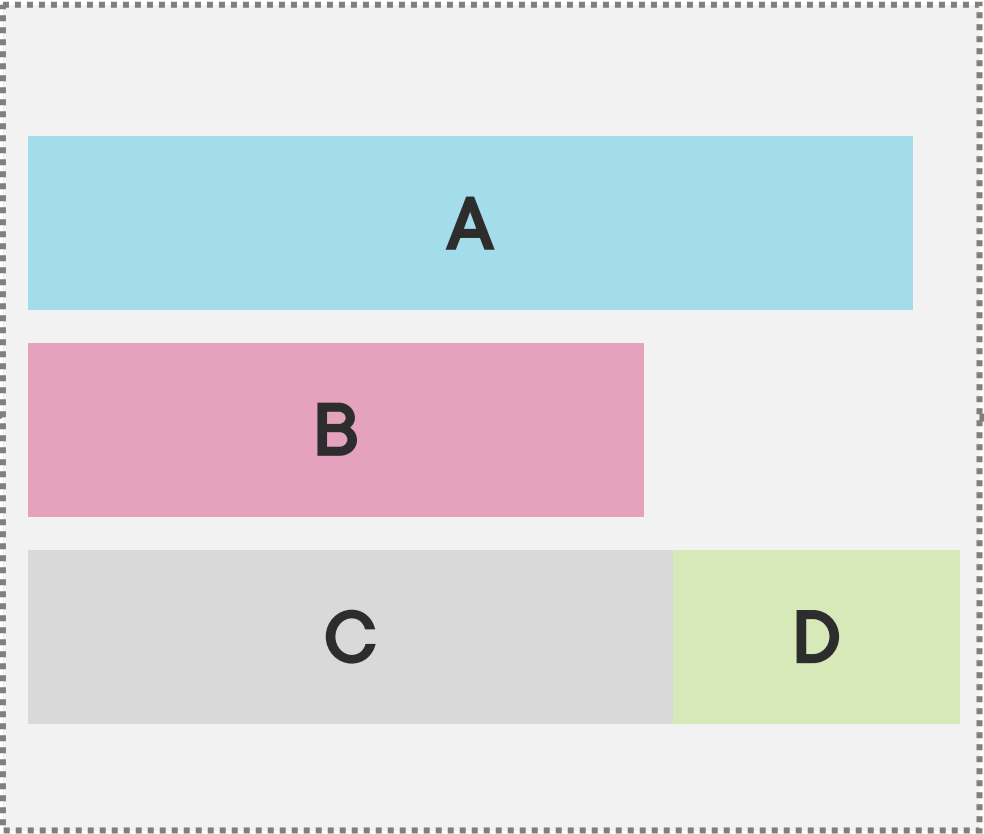


Sales DF
100 mn records

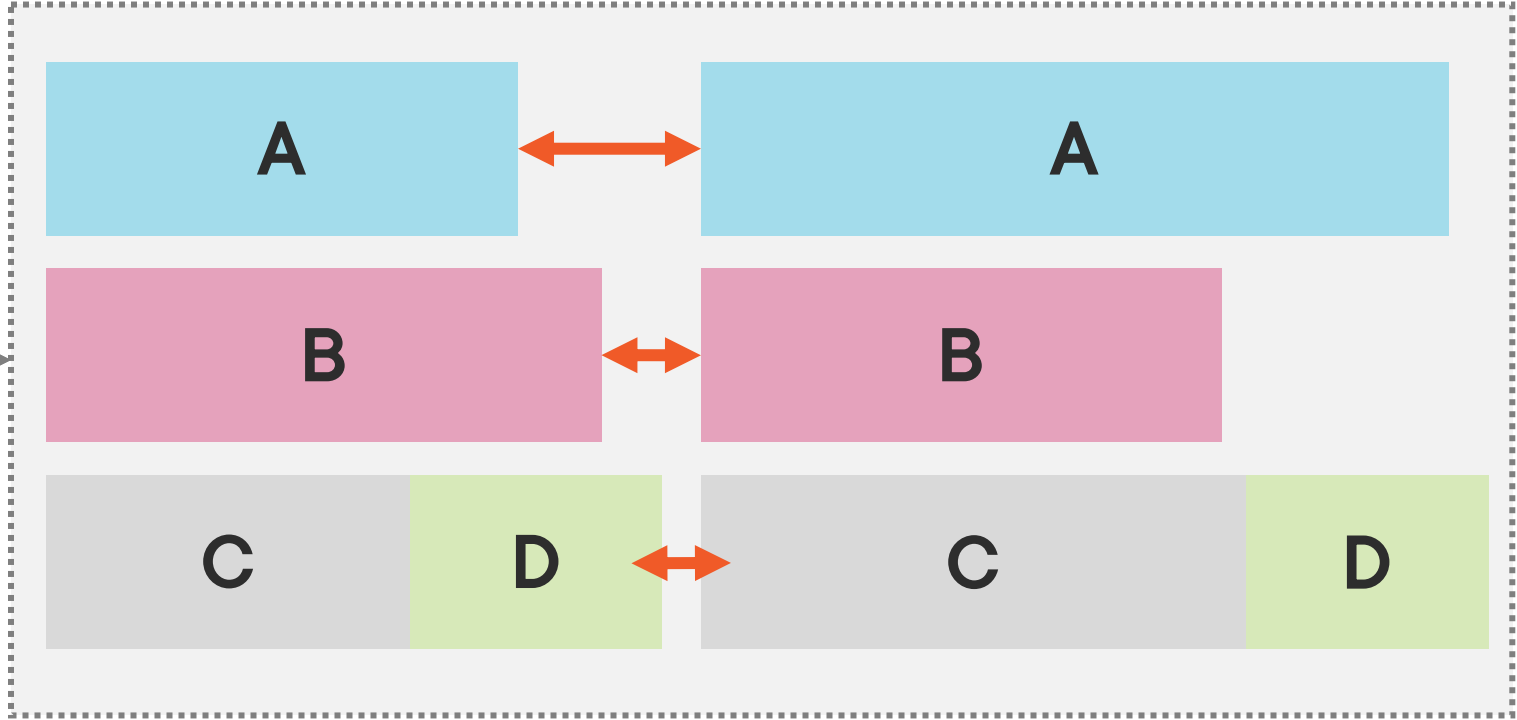



Orders DF
10 mn records

Shuffle
& Sort



Shuffle
& Sort



Sort Merge Join

Bucketing is an optimization technique
to **avoid shuffles** during joins

What is Bucketing?

Id	City	Amount
1	Seattle	600
2	London	300
3	Seattle	700
4	Delhi	400
5	Paris	900
6	Seattle	900
7	Delhi	200

DataFrame

Hash function



What is Bucketing?

Id	City	Amount
1	Seattle	600
2	London	300
3	Seattle	700
4	Delhi	400
5	Paris	900
6	Seattle	900
7	Delhi	200

DataFrame

Hash function



1	Seattle	600
---	---------	-----



What is Bucketing?

Id	City	Amount
1	Seattle	600
2	London	300
3	Seattle	700
4	Delhi	400
5	Paris	900
6	Seattle	900
7	Delhi	200

DataFrame

Hash function
.....→



1	Seattle	600
---	---------	-----



2	London	300
---	--------	-----



What is Bucketing?

Id	City	Amount
1	Seattle	600
2	London	300
3	Seattle	700
4	Delhi	400
5	Paris	900
6	Seattle	900
7	Delhi	200

DataFrame

Hash function



1	Seattle	600
3	Seattle	700



2	London	300
---	--------	-----



What is Bucketing?

Id	City	Amount
1	Seattle	600
2	London	300
3	Seattle	700
4	Delhi	400
5	Paris	900
6	Seattle	900
7	Delhi	200

DataFrame

Hash function



1	Seattle	600
3	Seattle	700
6	Seattle	900



2	London	300
5	Paris	900



4	Delhi	400
7	Delhi	200

**Bucketed data is
written to disk**

Bucketing



Pre-shuffle large datasets & write it to disk in buckets

Use bucketed data in subsequent queries

Use when dataset is frequently used in joins, aggregations or window operations

Conditions for joining bucketed datasets

- Bucketed datasets must be stored as tables
- Both datasets must be bucketed on join columns
- Number of buckets must be the same

Dynamic Resource Allocation

Resource Scheduling for a Job

**Scheduling Across
Applications**

**Scheduling Within an
Application**

Resource Scheduling Across Applications

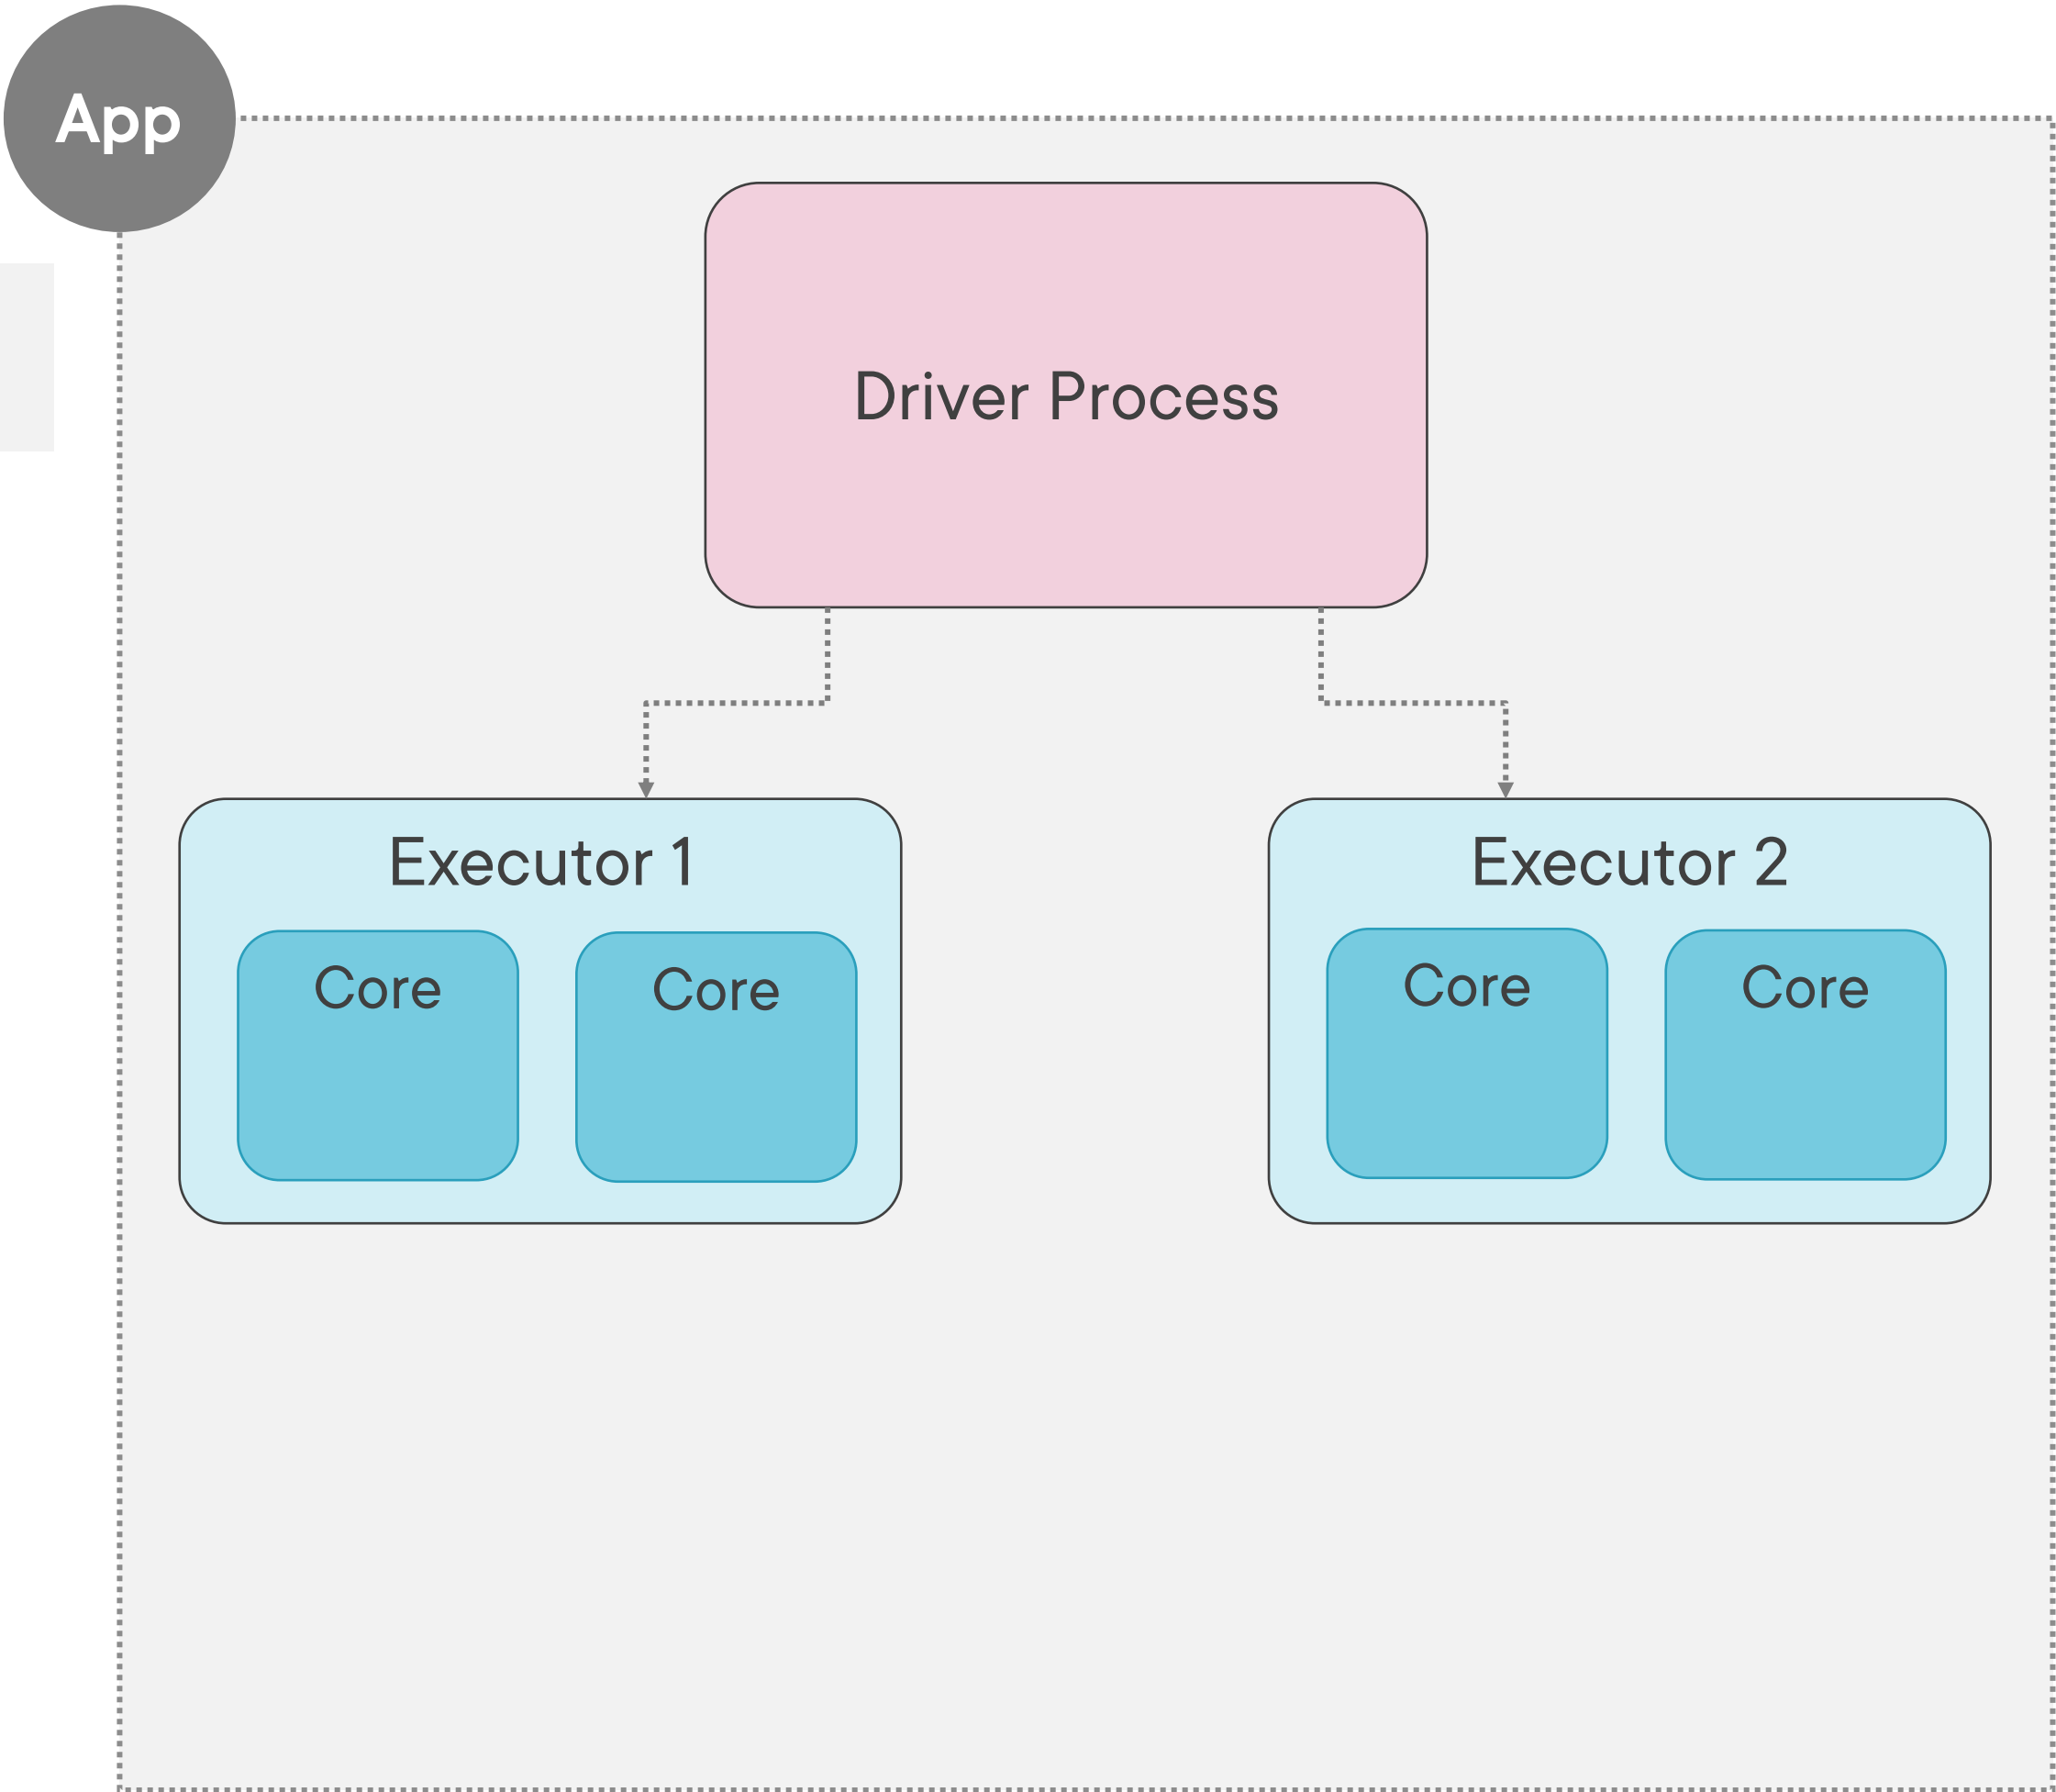
Static Resource Allocation

Dynamic Resource Allocation

Static Resource Allocation

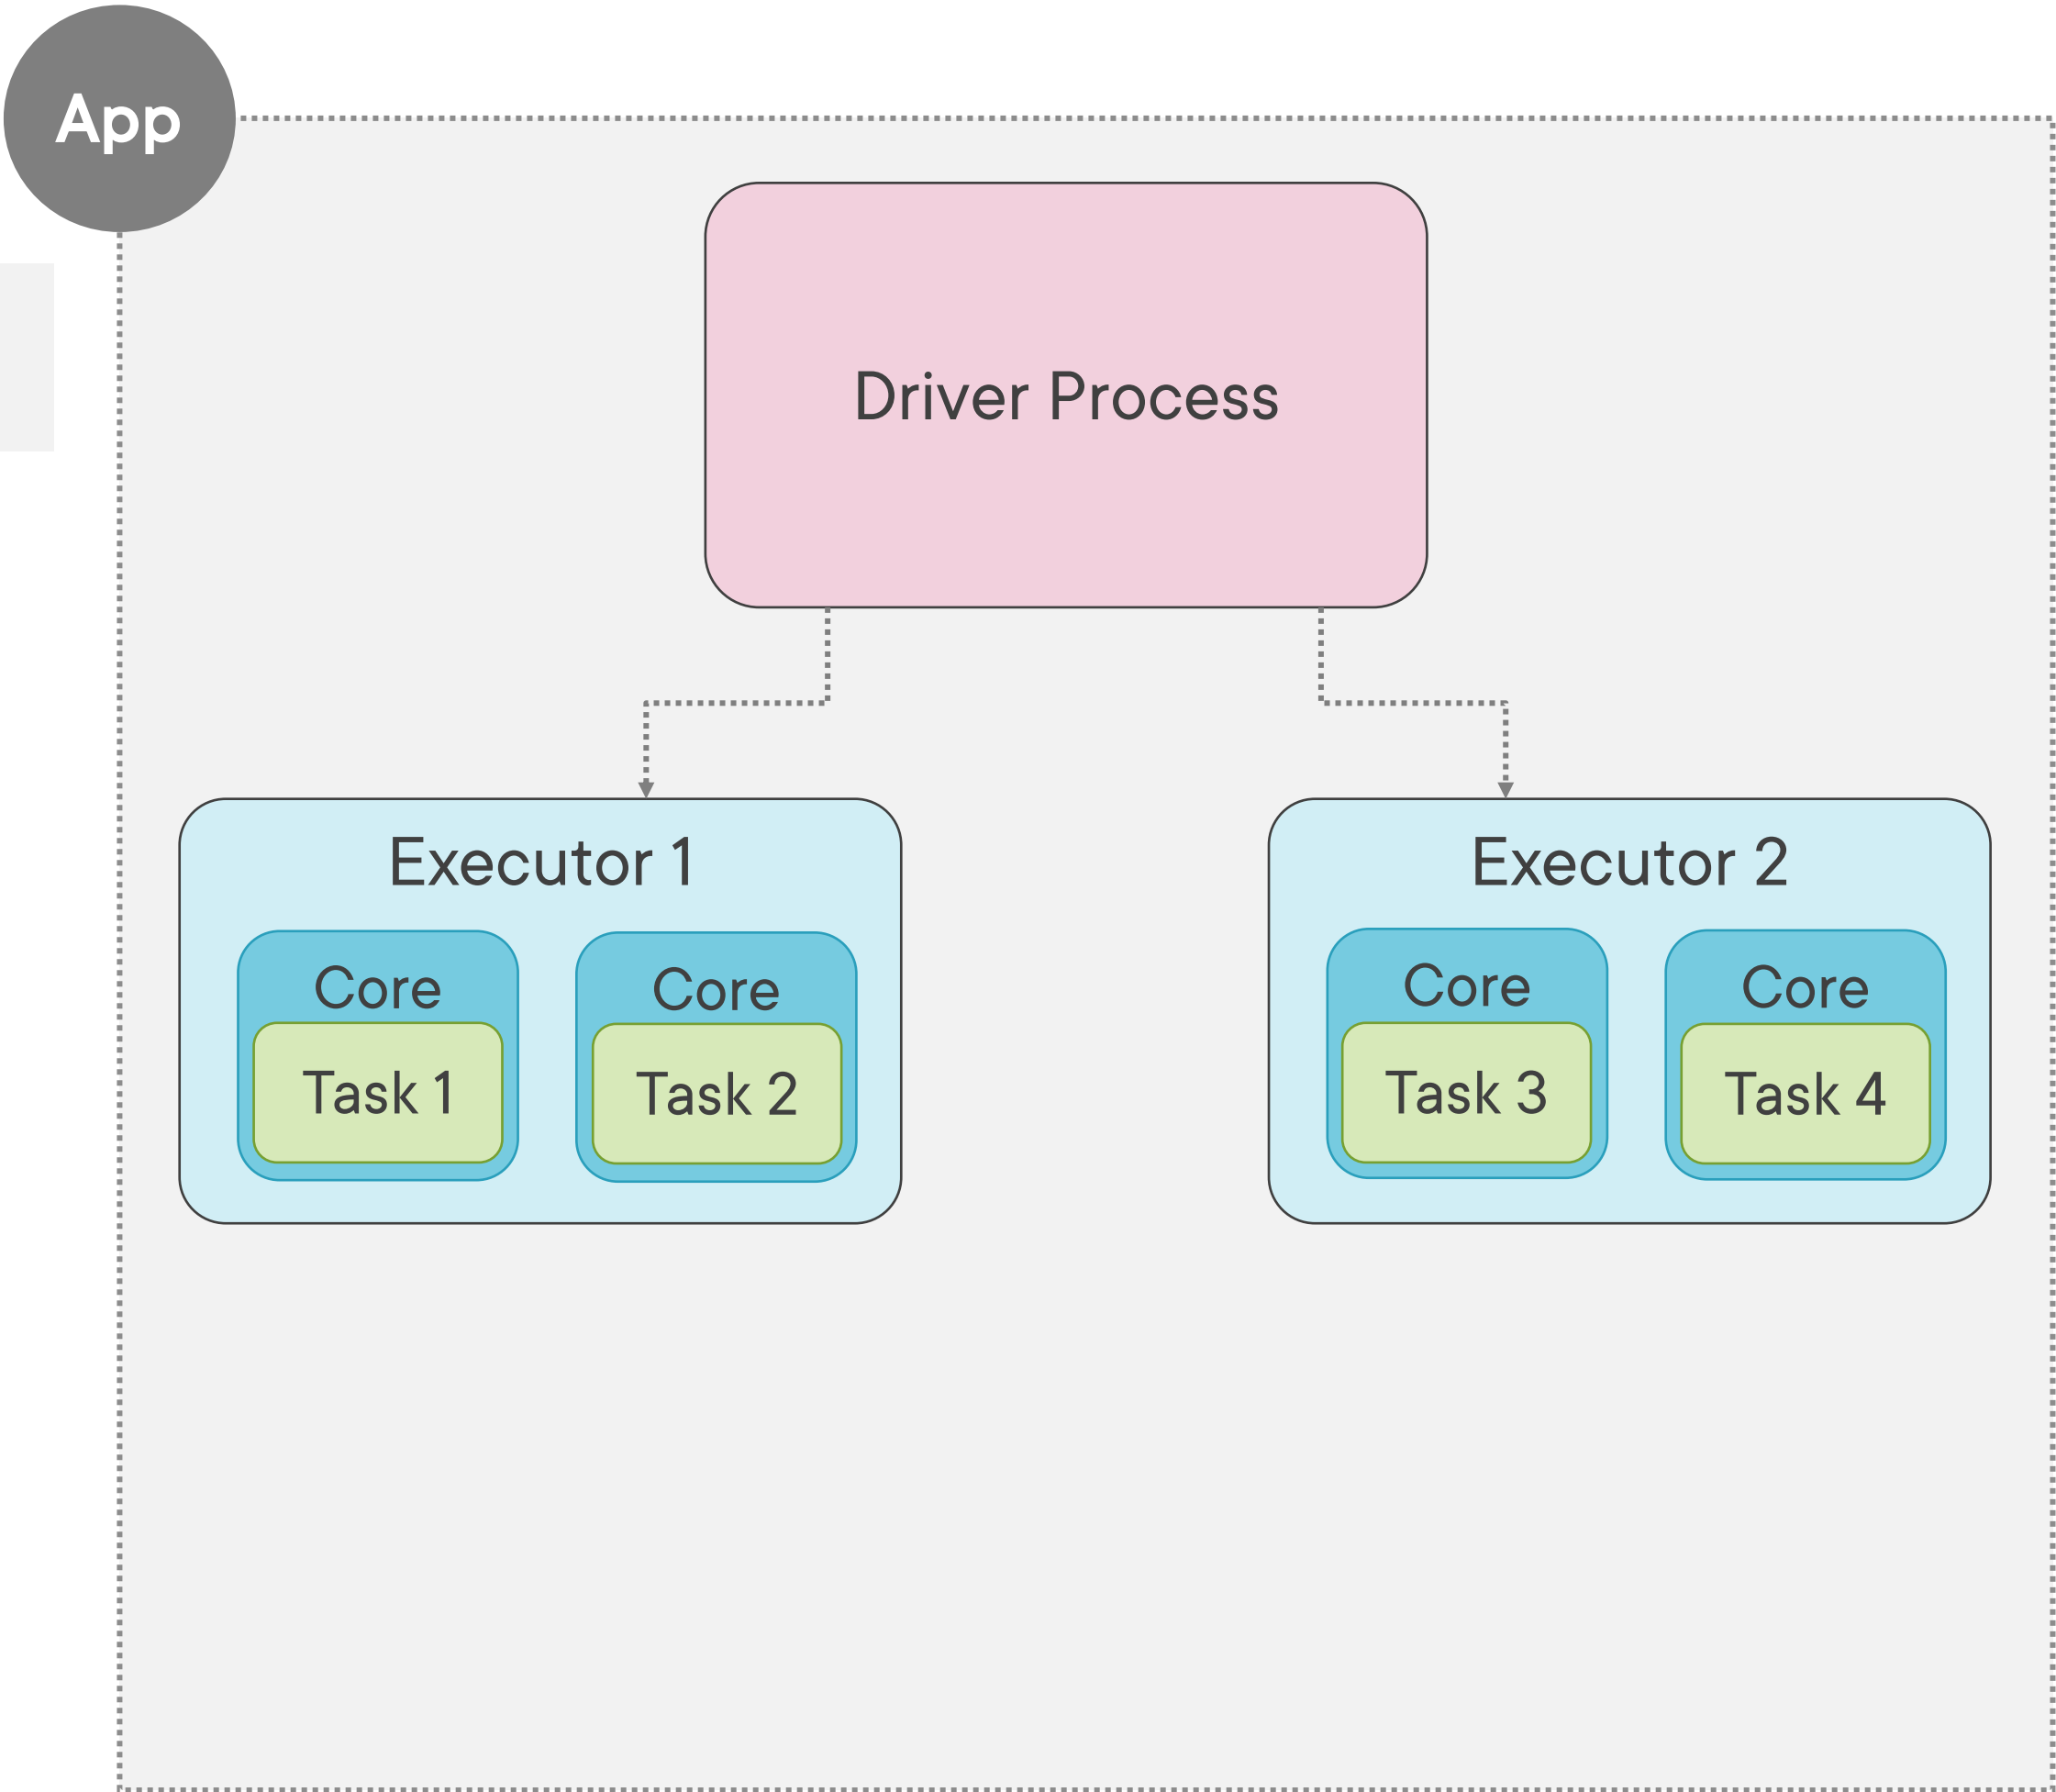

```
spark.dynamicAllocation.enabled = false
```

- | | |
|--------|--------|
| Task 1 | Task 2 |
| Task 3 | Task 4 |
| Task 5 | Task 6 |



Static Resource Allocation

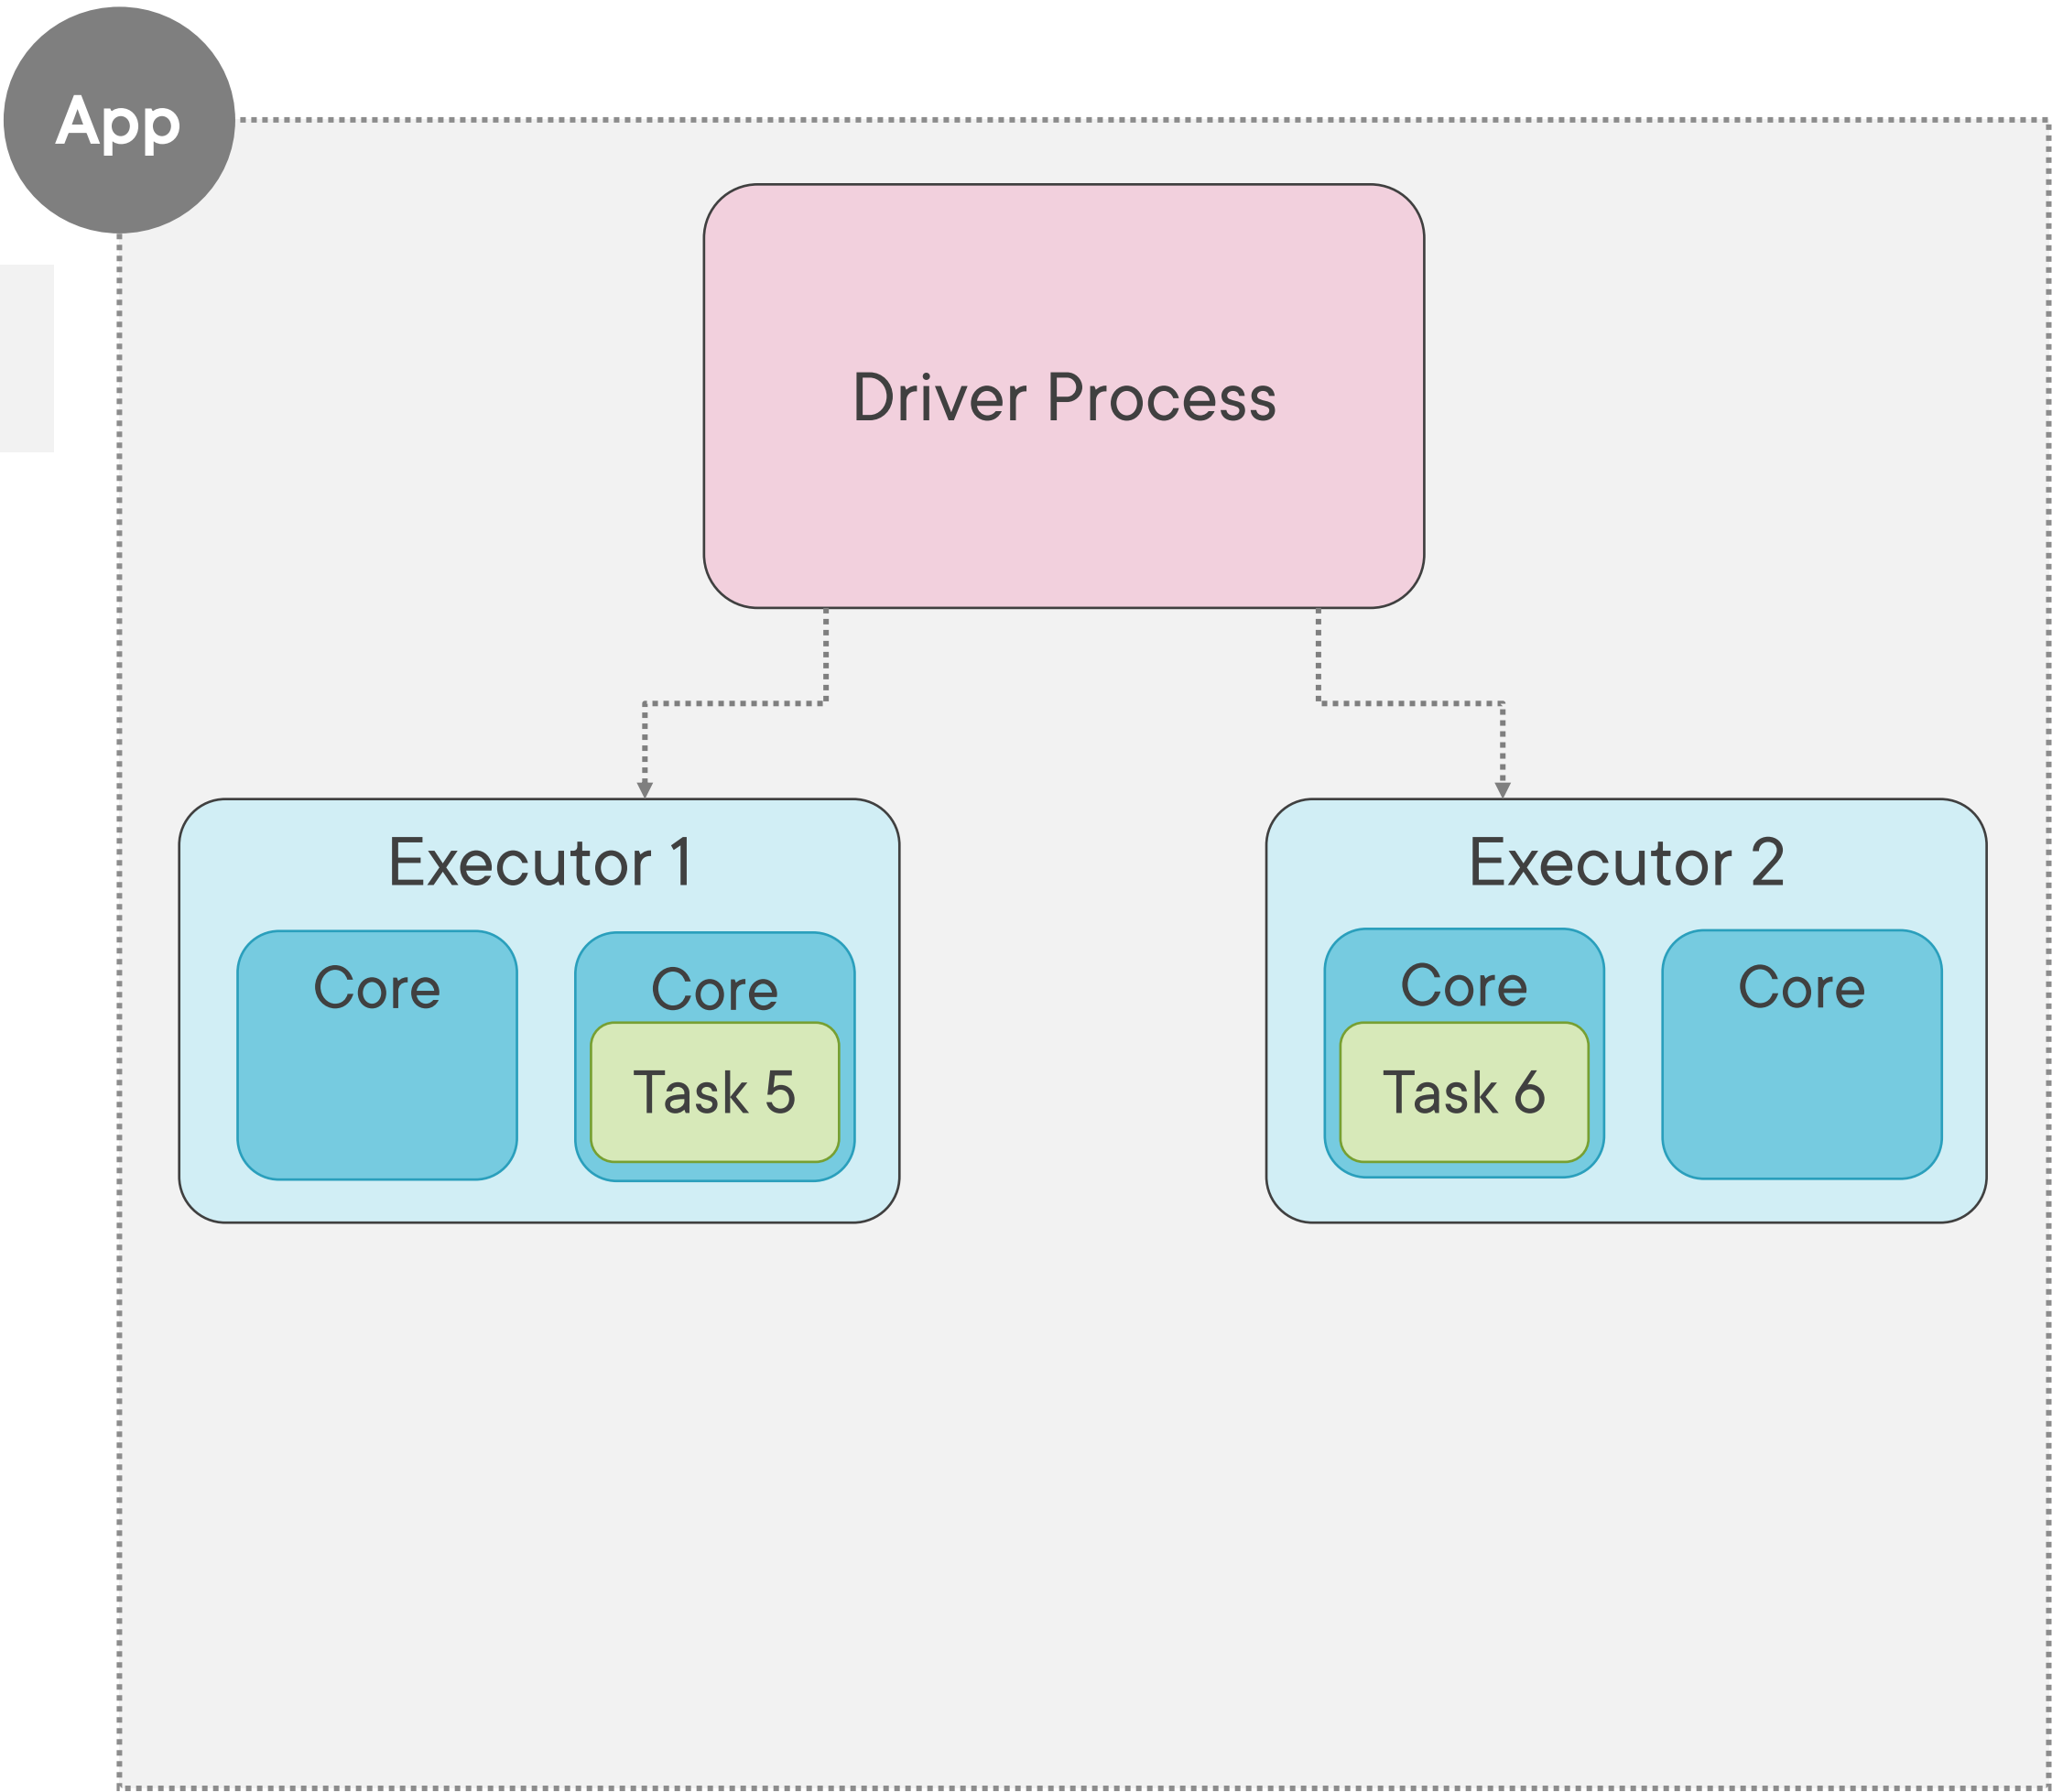
```
spark.dynamicAllocation.enabled = false
```



Static Resource Allocation

`spark.dynamicAllocation.enabled = false`

CPU Cores & Memory are pre-defined for an application



Dynamic Resource Allocation

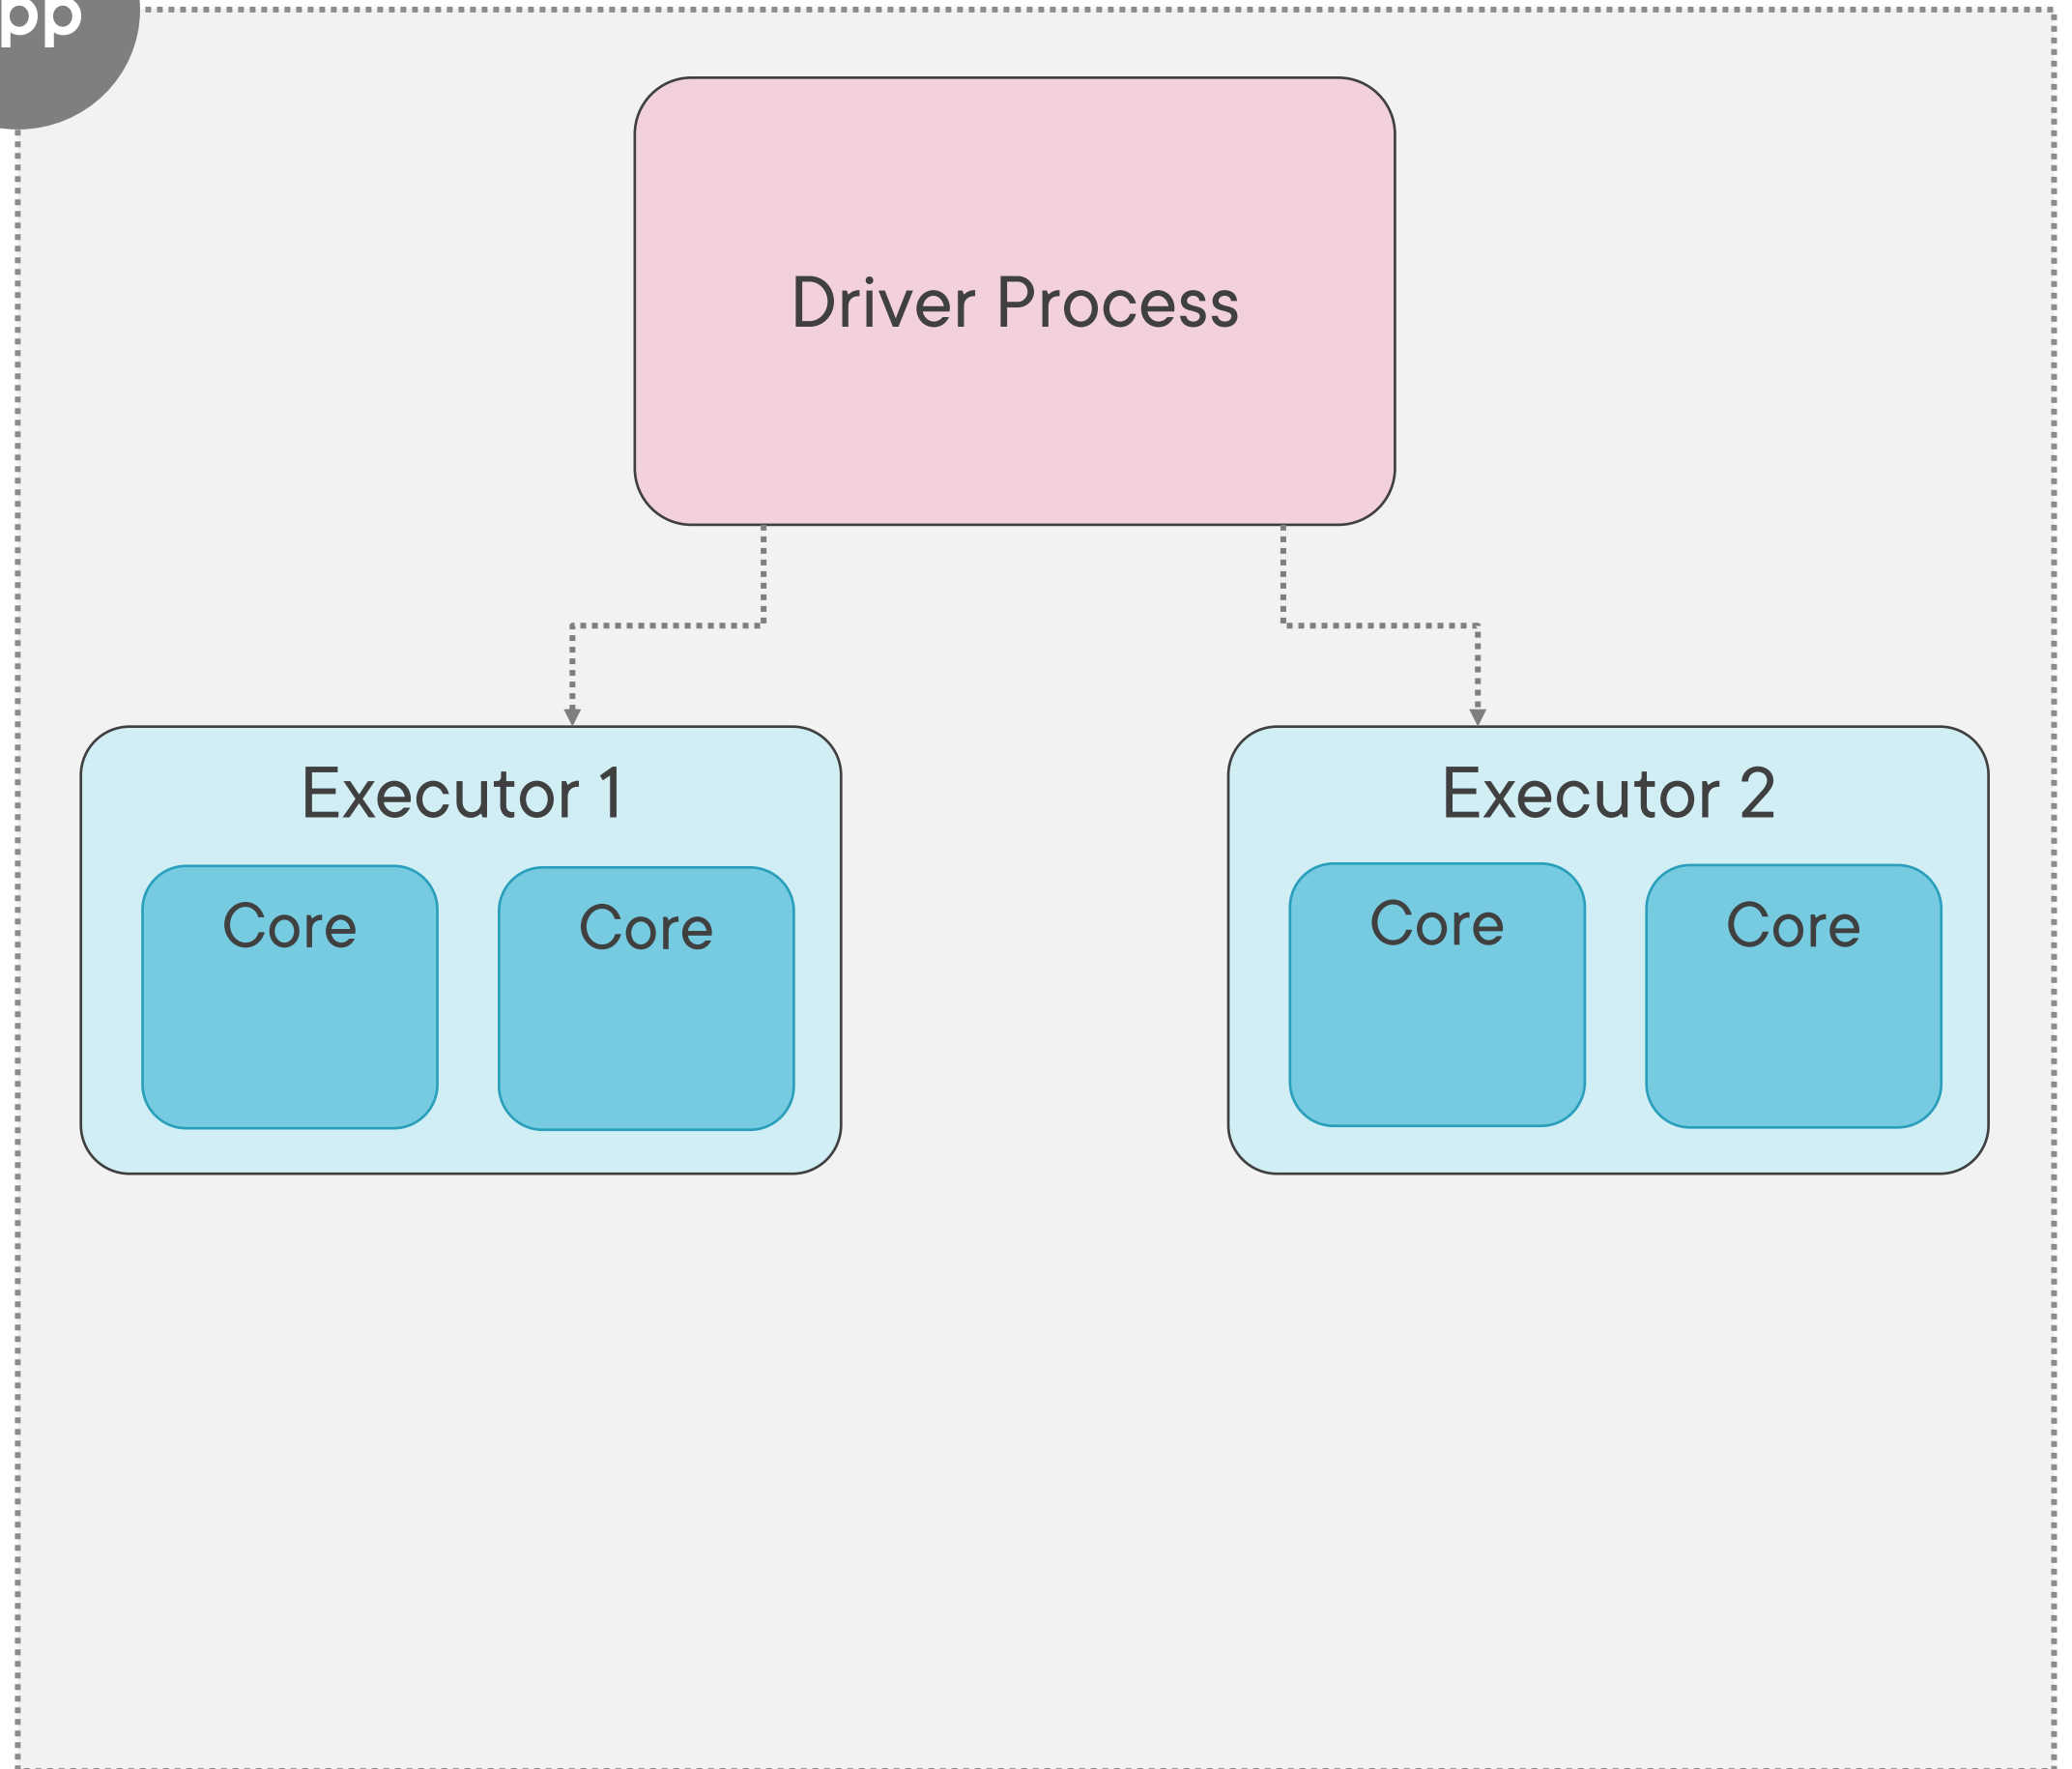


spark.dynamicAllocation :

`.enabled = true`

`.shuffleTracking.enabled = true`

`.minExecutors = 0`

`.maxExecutors = 5`



Dynamic Resource Allocation

spark.dynamicAllocation :

`.enabled = true`

`.shuffleTracking.enabled = true`

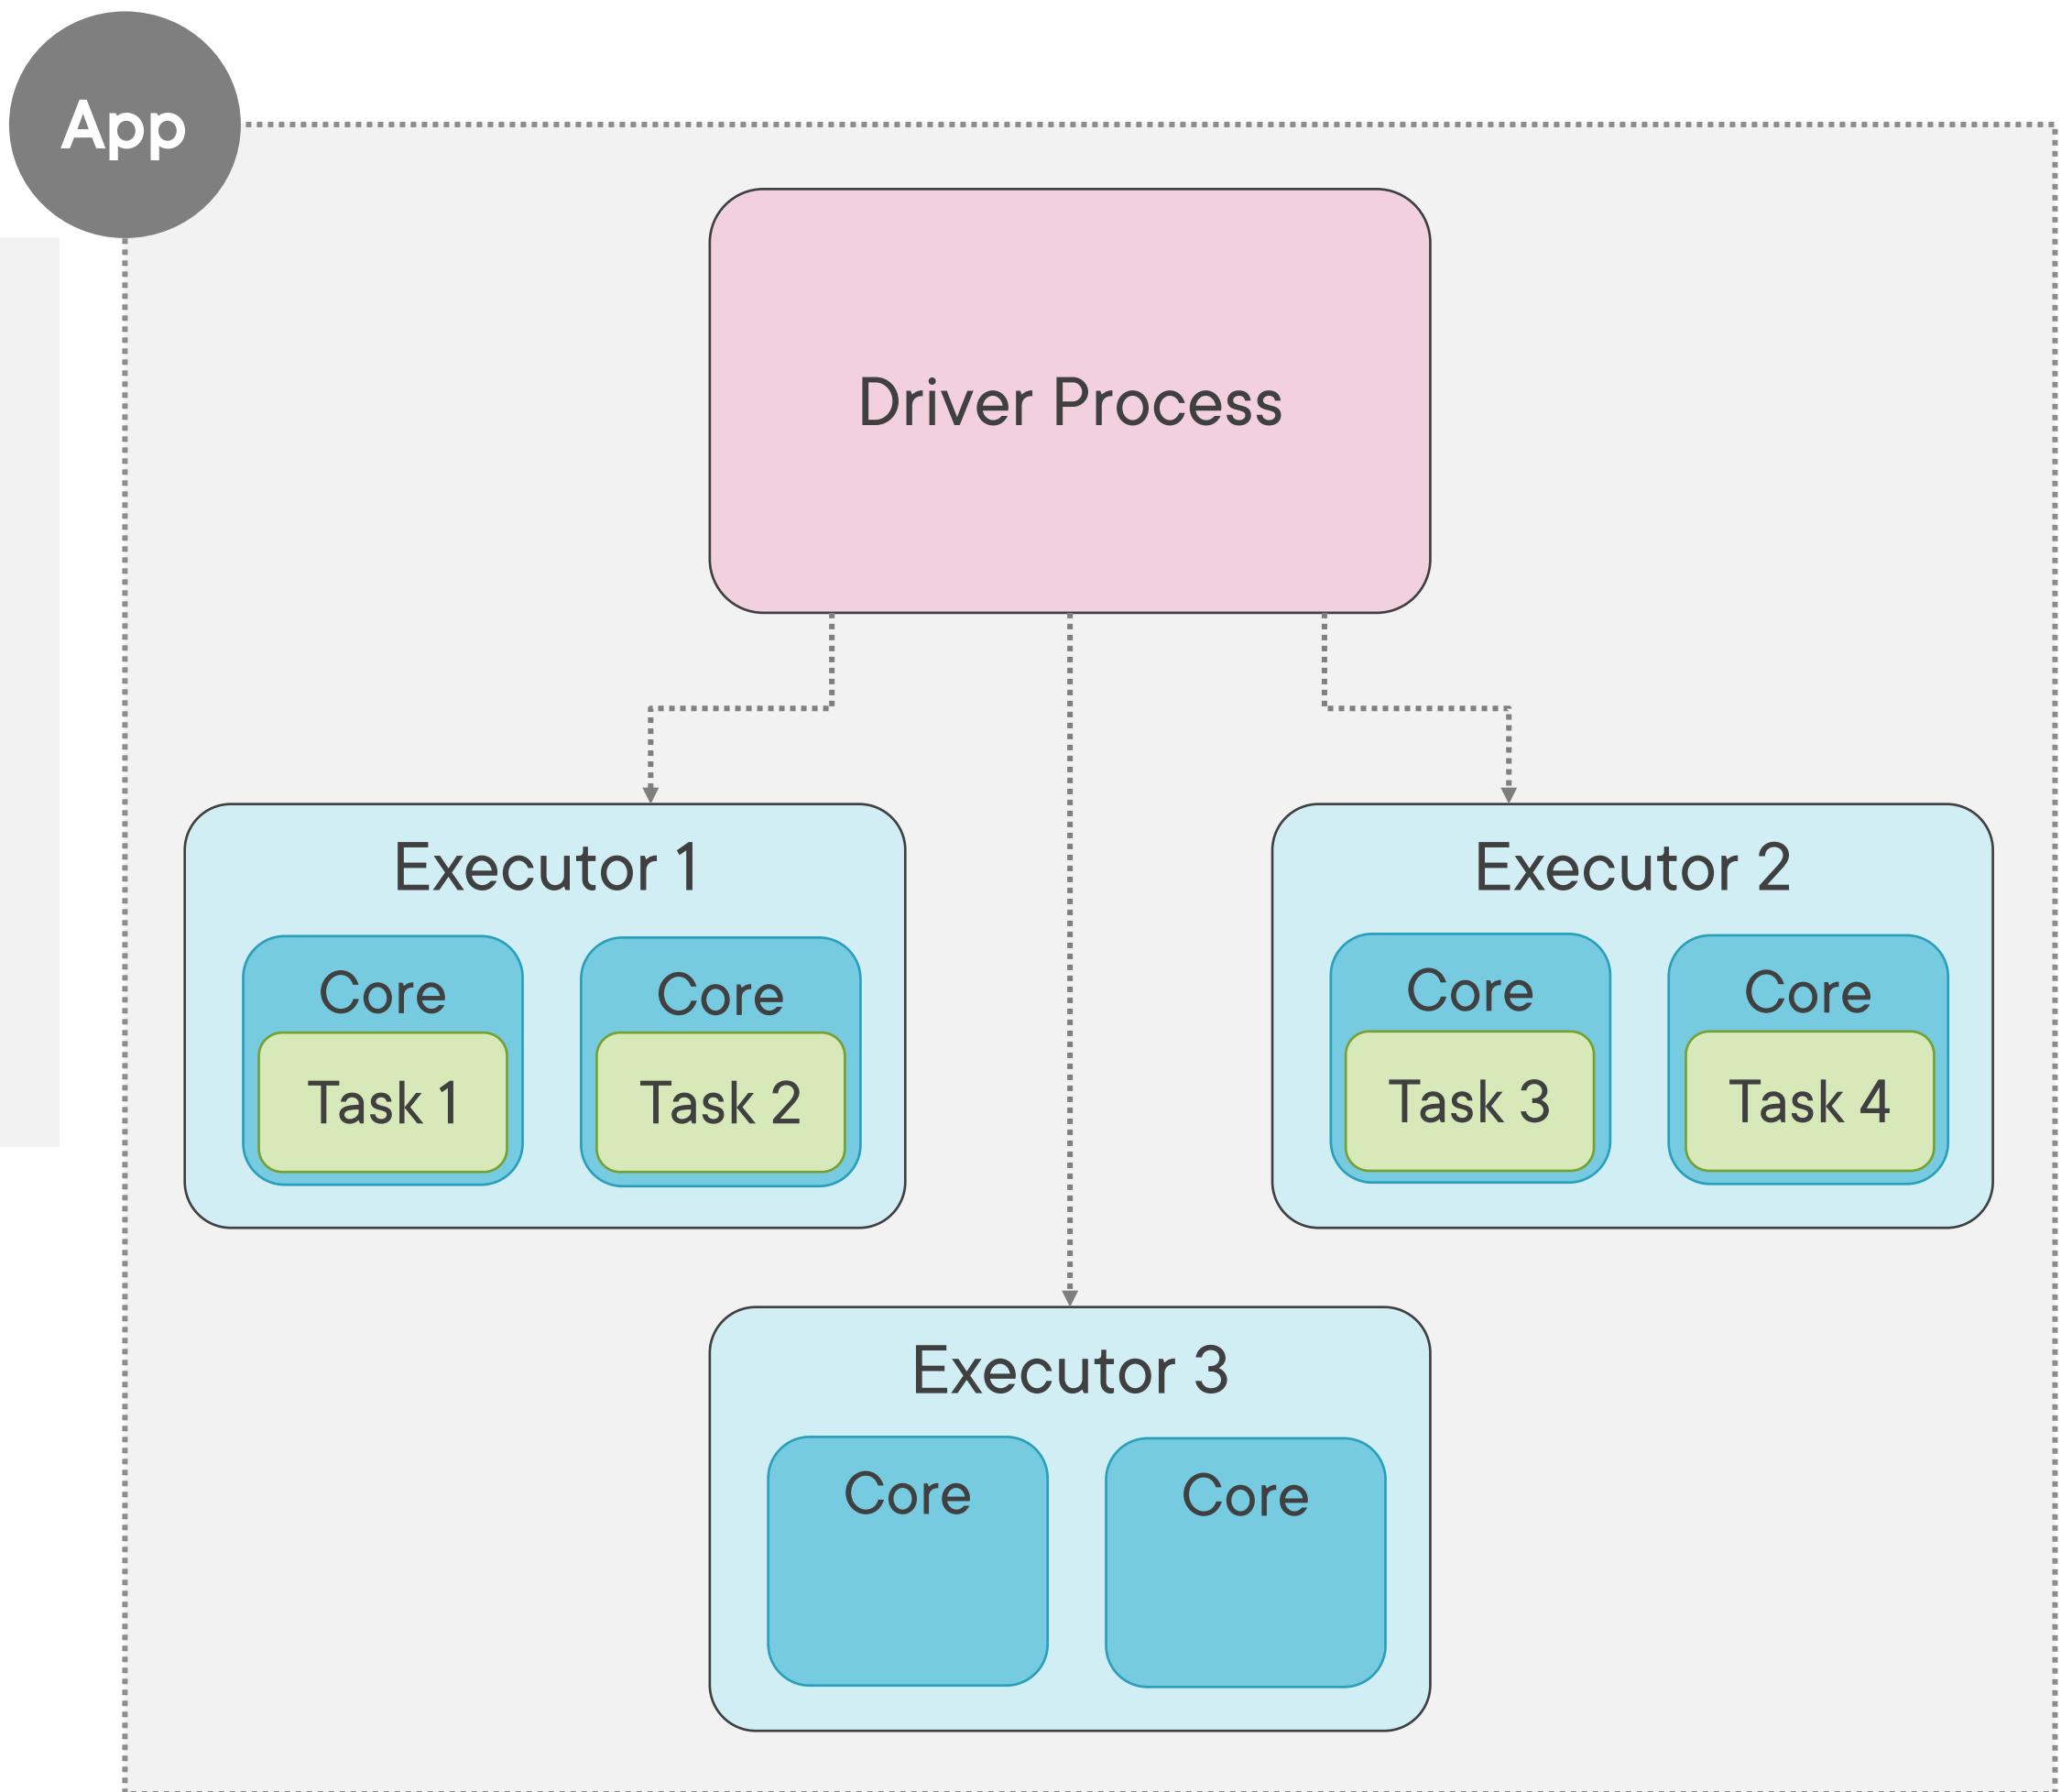
`.minExecutors = 0`

`.maxExecutors = 5`

`.schedulerBacklogTimeout = 1s`

Task 5

Task 6



Dynamic Resource Allocation

spark.dynamicAllocation :

`.enabled = true`

`.shuffleTracking.enabled = true`

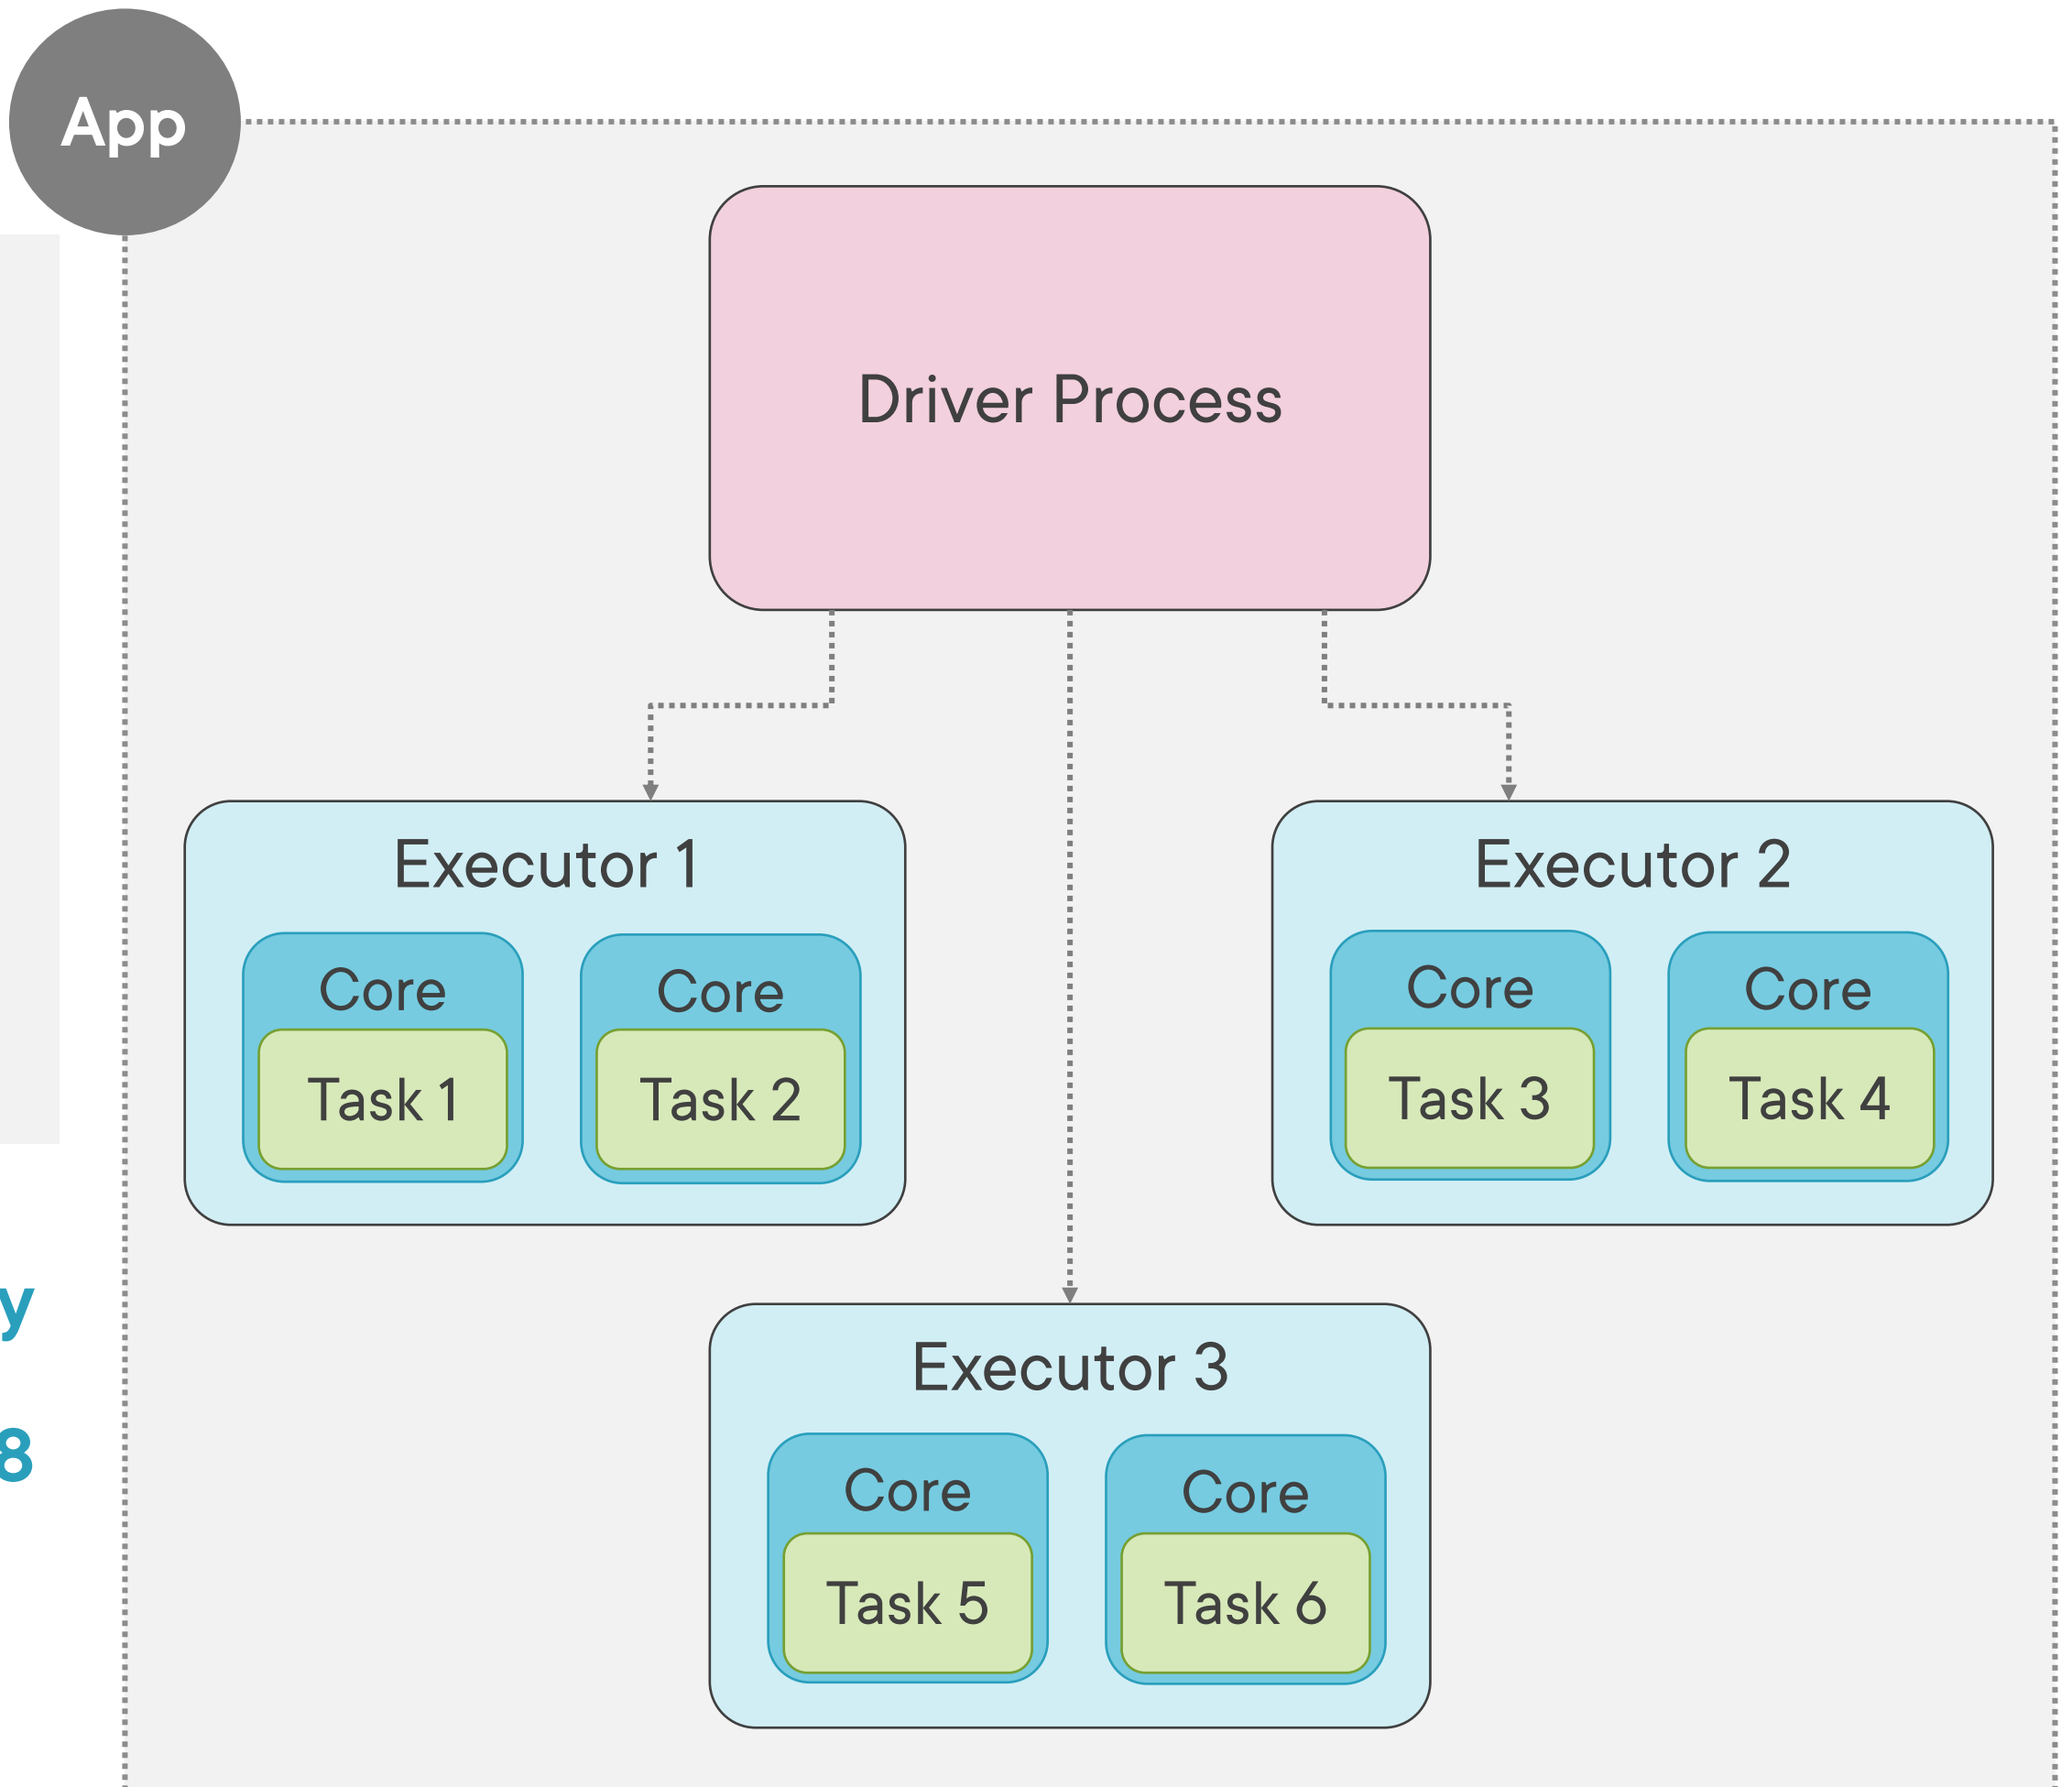
`.minExecutors = 0`

`.maxExecutors = 5`

`.schedulerBacklogTimeout = 1s`

Executors will increase exponentially

Starting with 1 executor – then 2, 4, 8



Dynamic Resource Allocation

spark.dynamicAllocation :

`.enabled = true`

`.shuffleTracking.enabled = true`

`.minExecutors = 0`

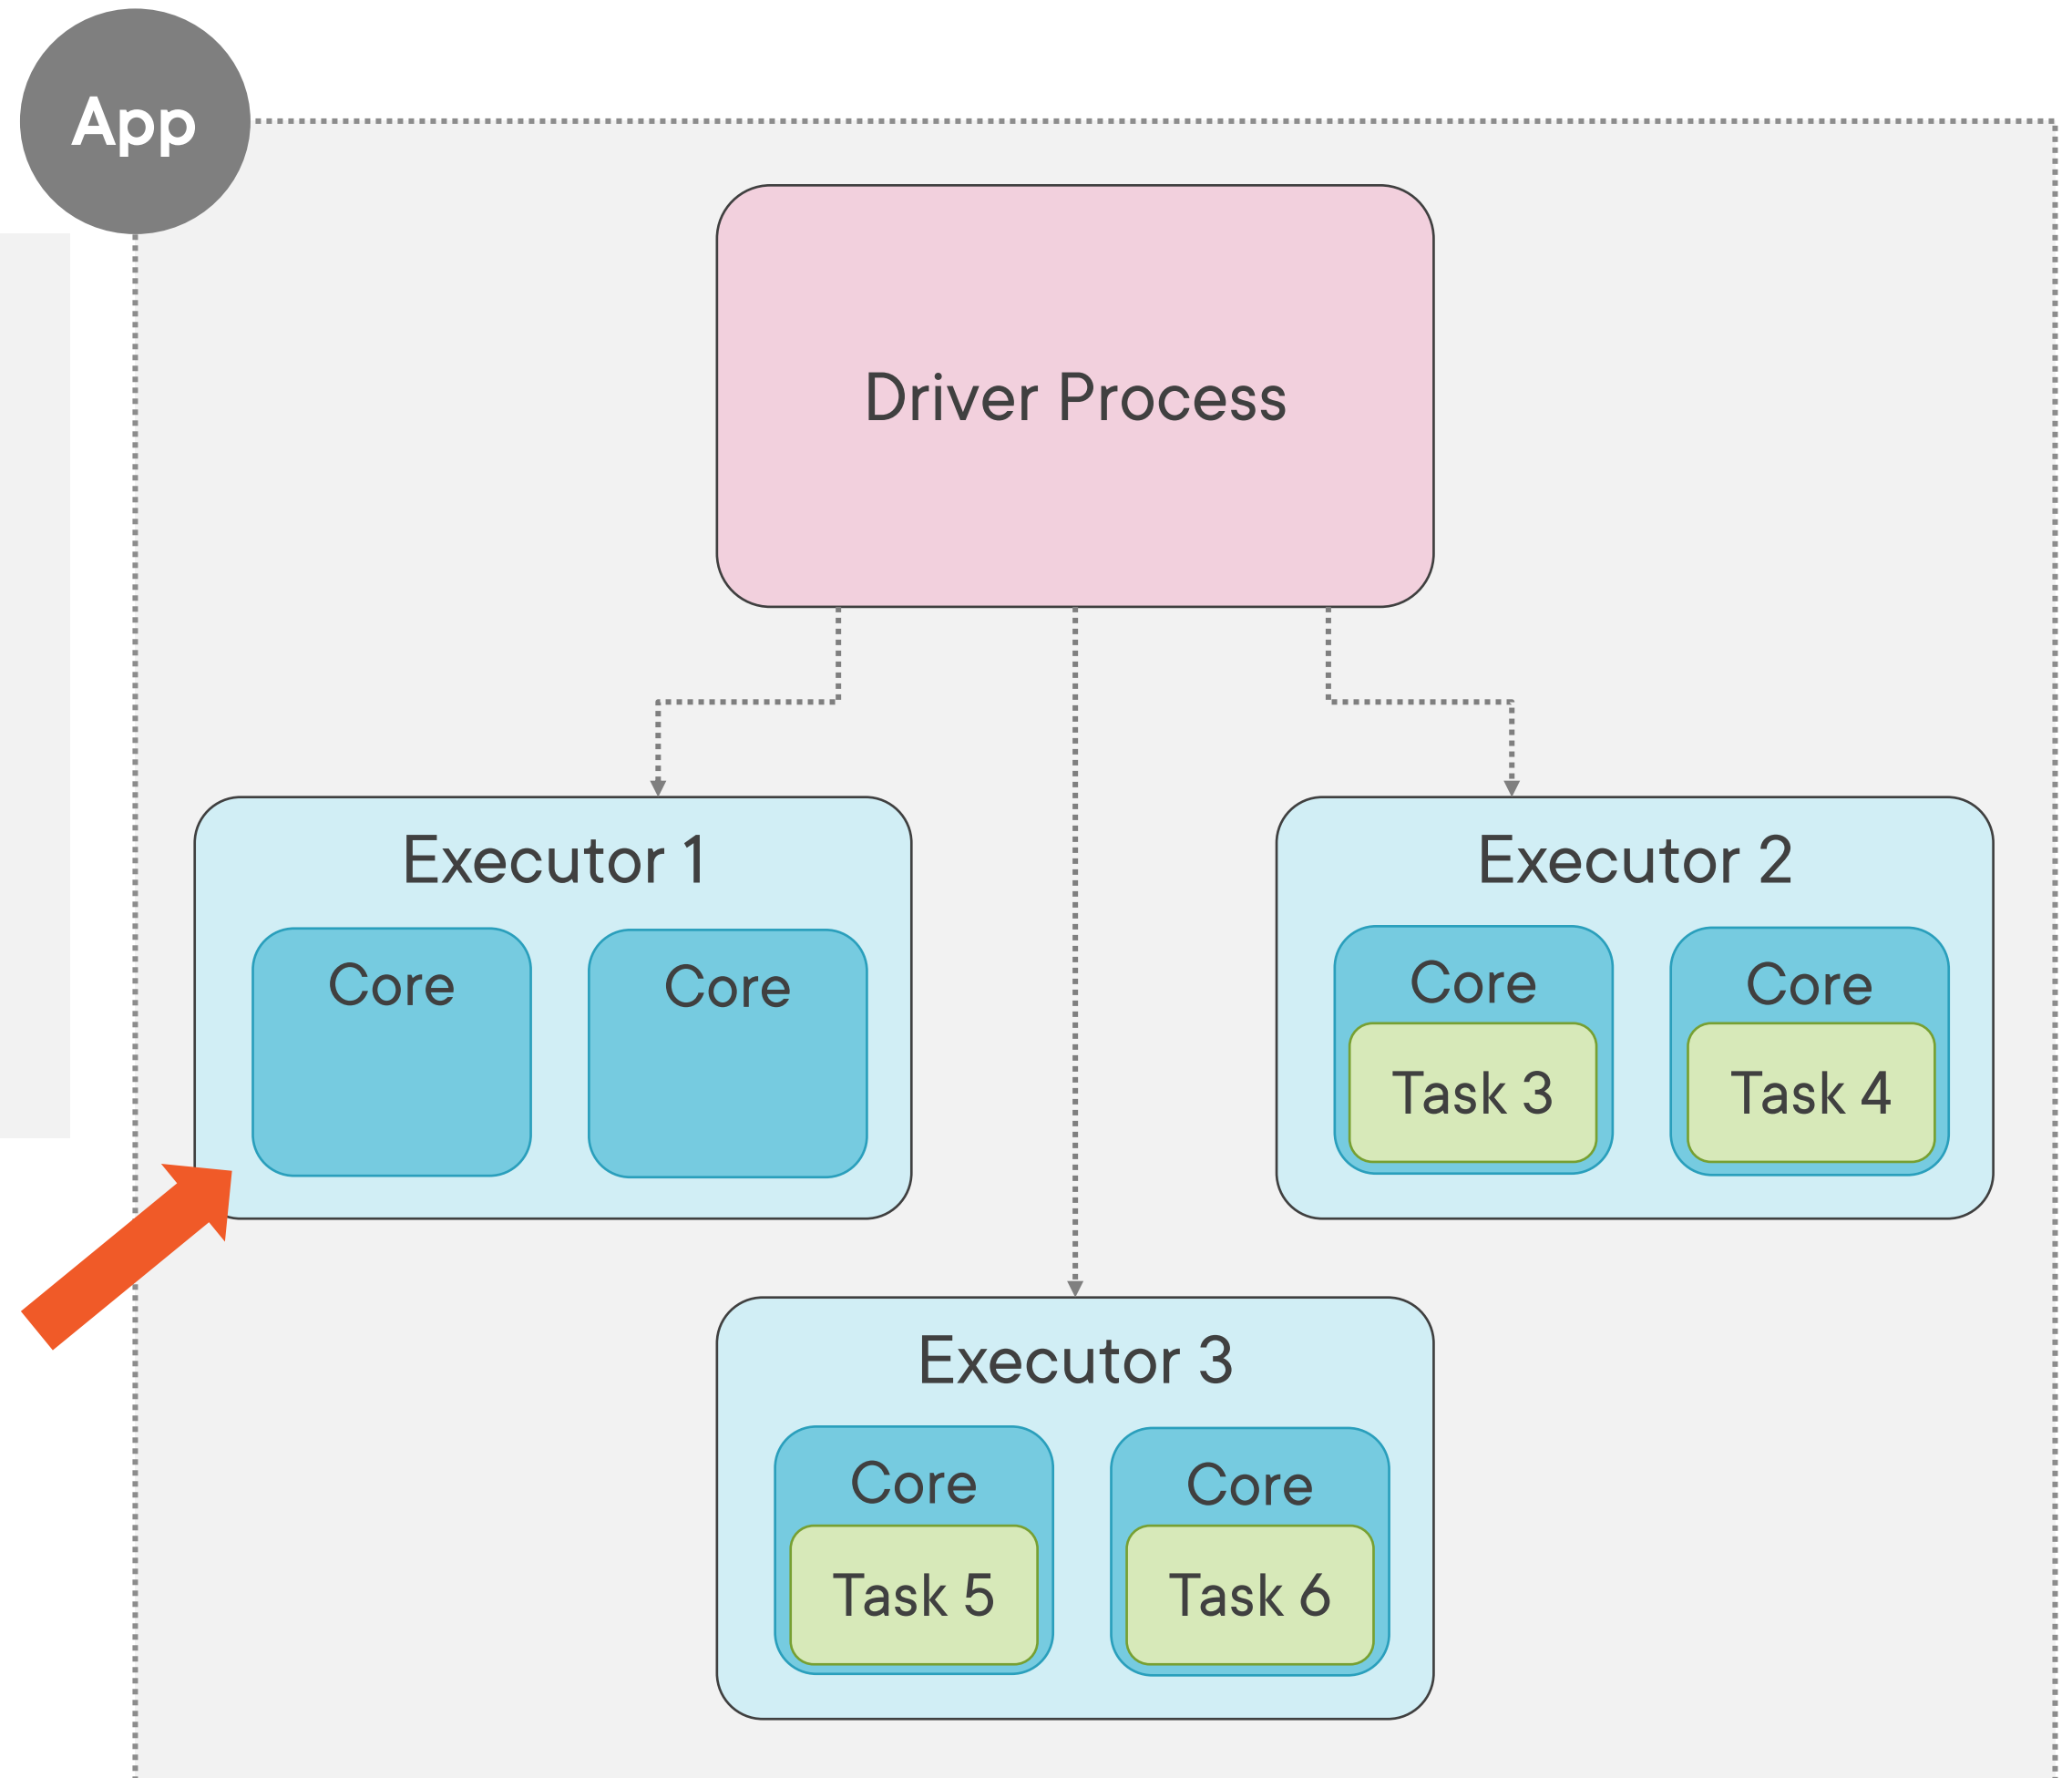
`.maxExecutors = 5`

`.schedulerBacklogTimeout = 1s`

`.executorIdleTimeout = 60s`

Any cached data on executor will
also be removed

No shuffle data is lost



**Request more executors
to complete pending tasks!**

**Remove executors
when they are idle!**

Use Cases



Running ad-hoc interactive applications

Long running ETL jobs to process multiple entities

Jobs with large shuffle operations

Resource Allocation Using Fair Scheduling

Resource Scheduling for a Job

**Scheduling Across
Applications**

**Scheduling Within an
Application**

Resource Scheduling Within Application

FIFO Scheduling

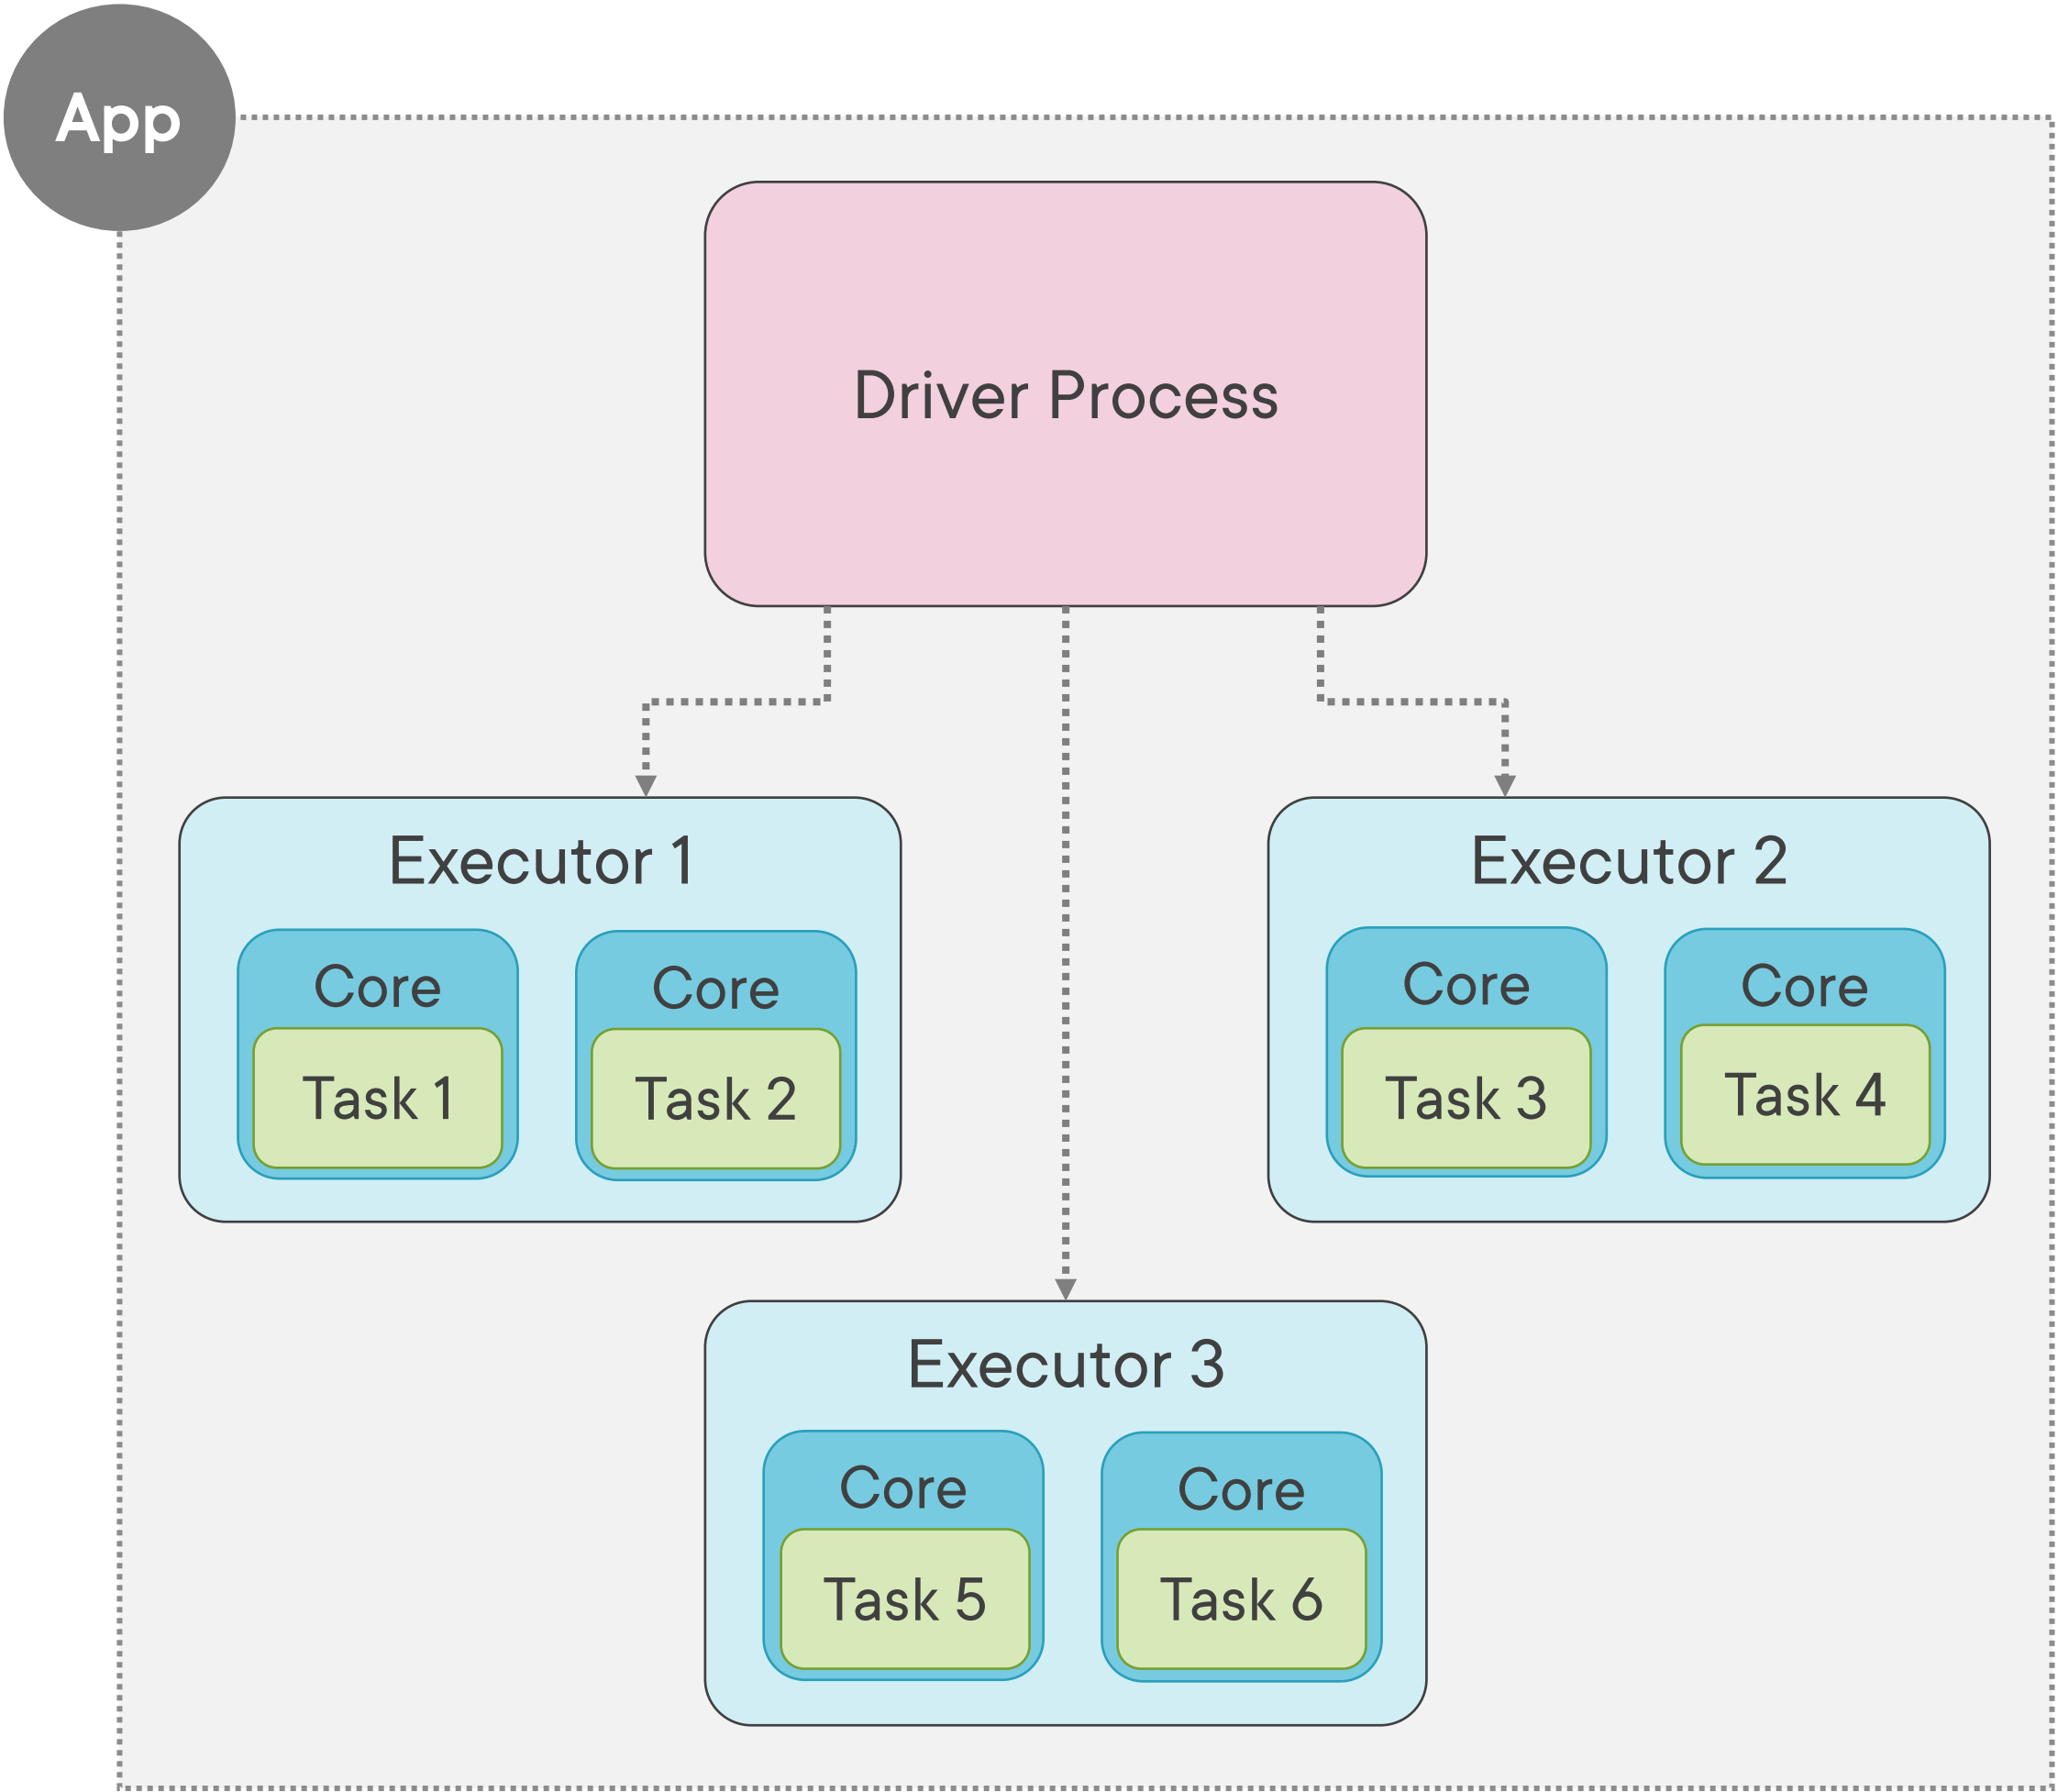

FAIR Scheduling

FIFO Scheduling

Job 1



Job 2




FIFO Scheduling

Job 1



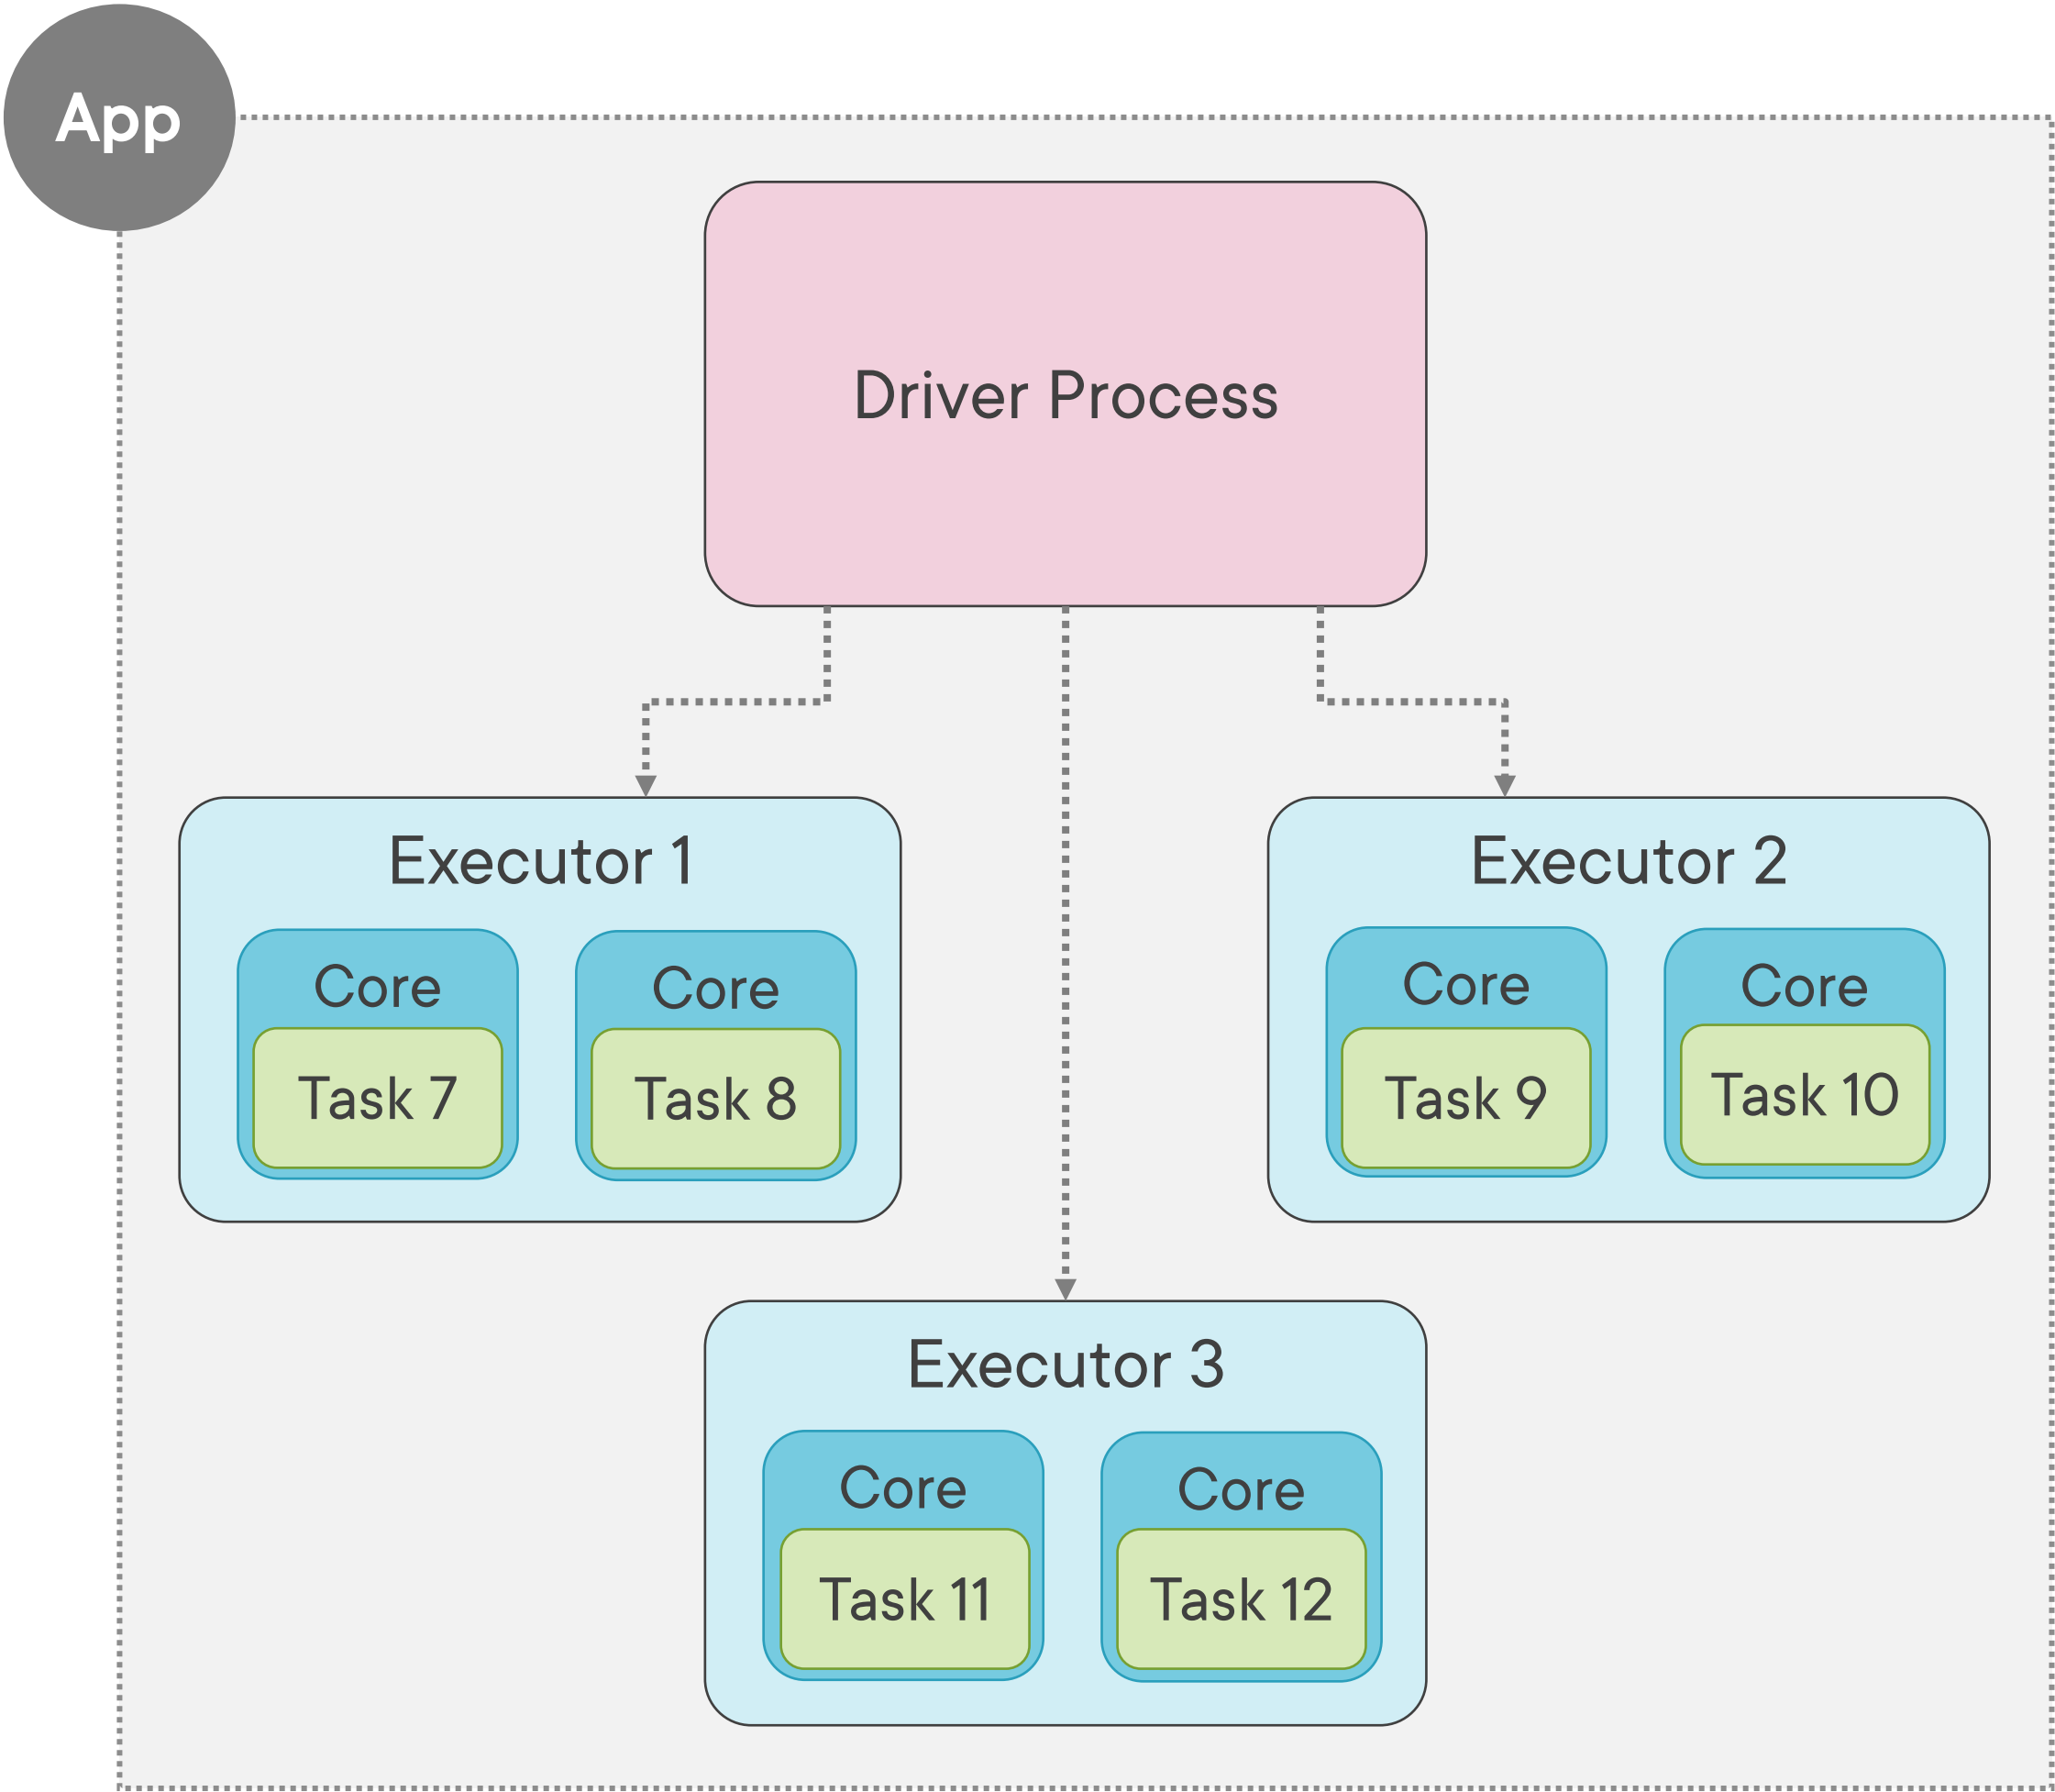
Job 2



Default scheduling

First job gets priority on resources

Other jobs in queue are delayed





FAIR Scheduling

Job 1



Job 2

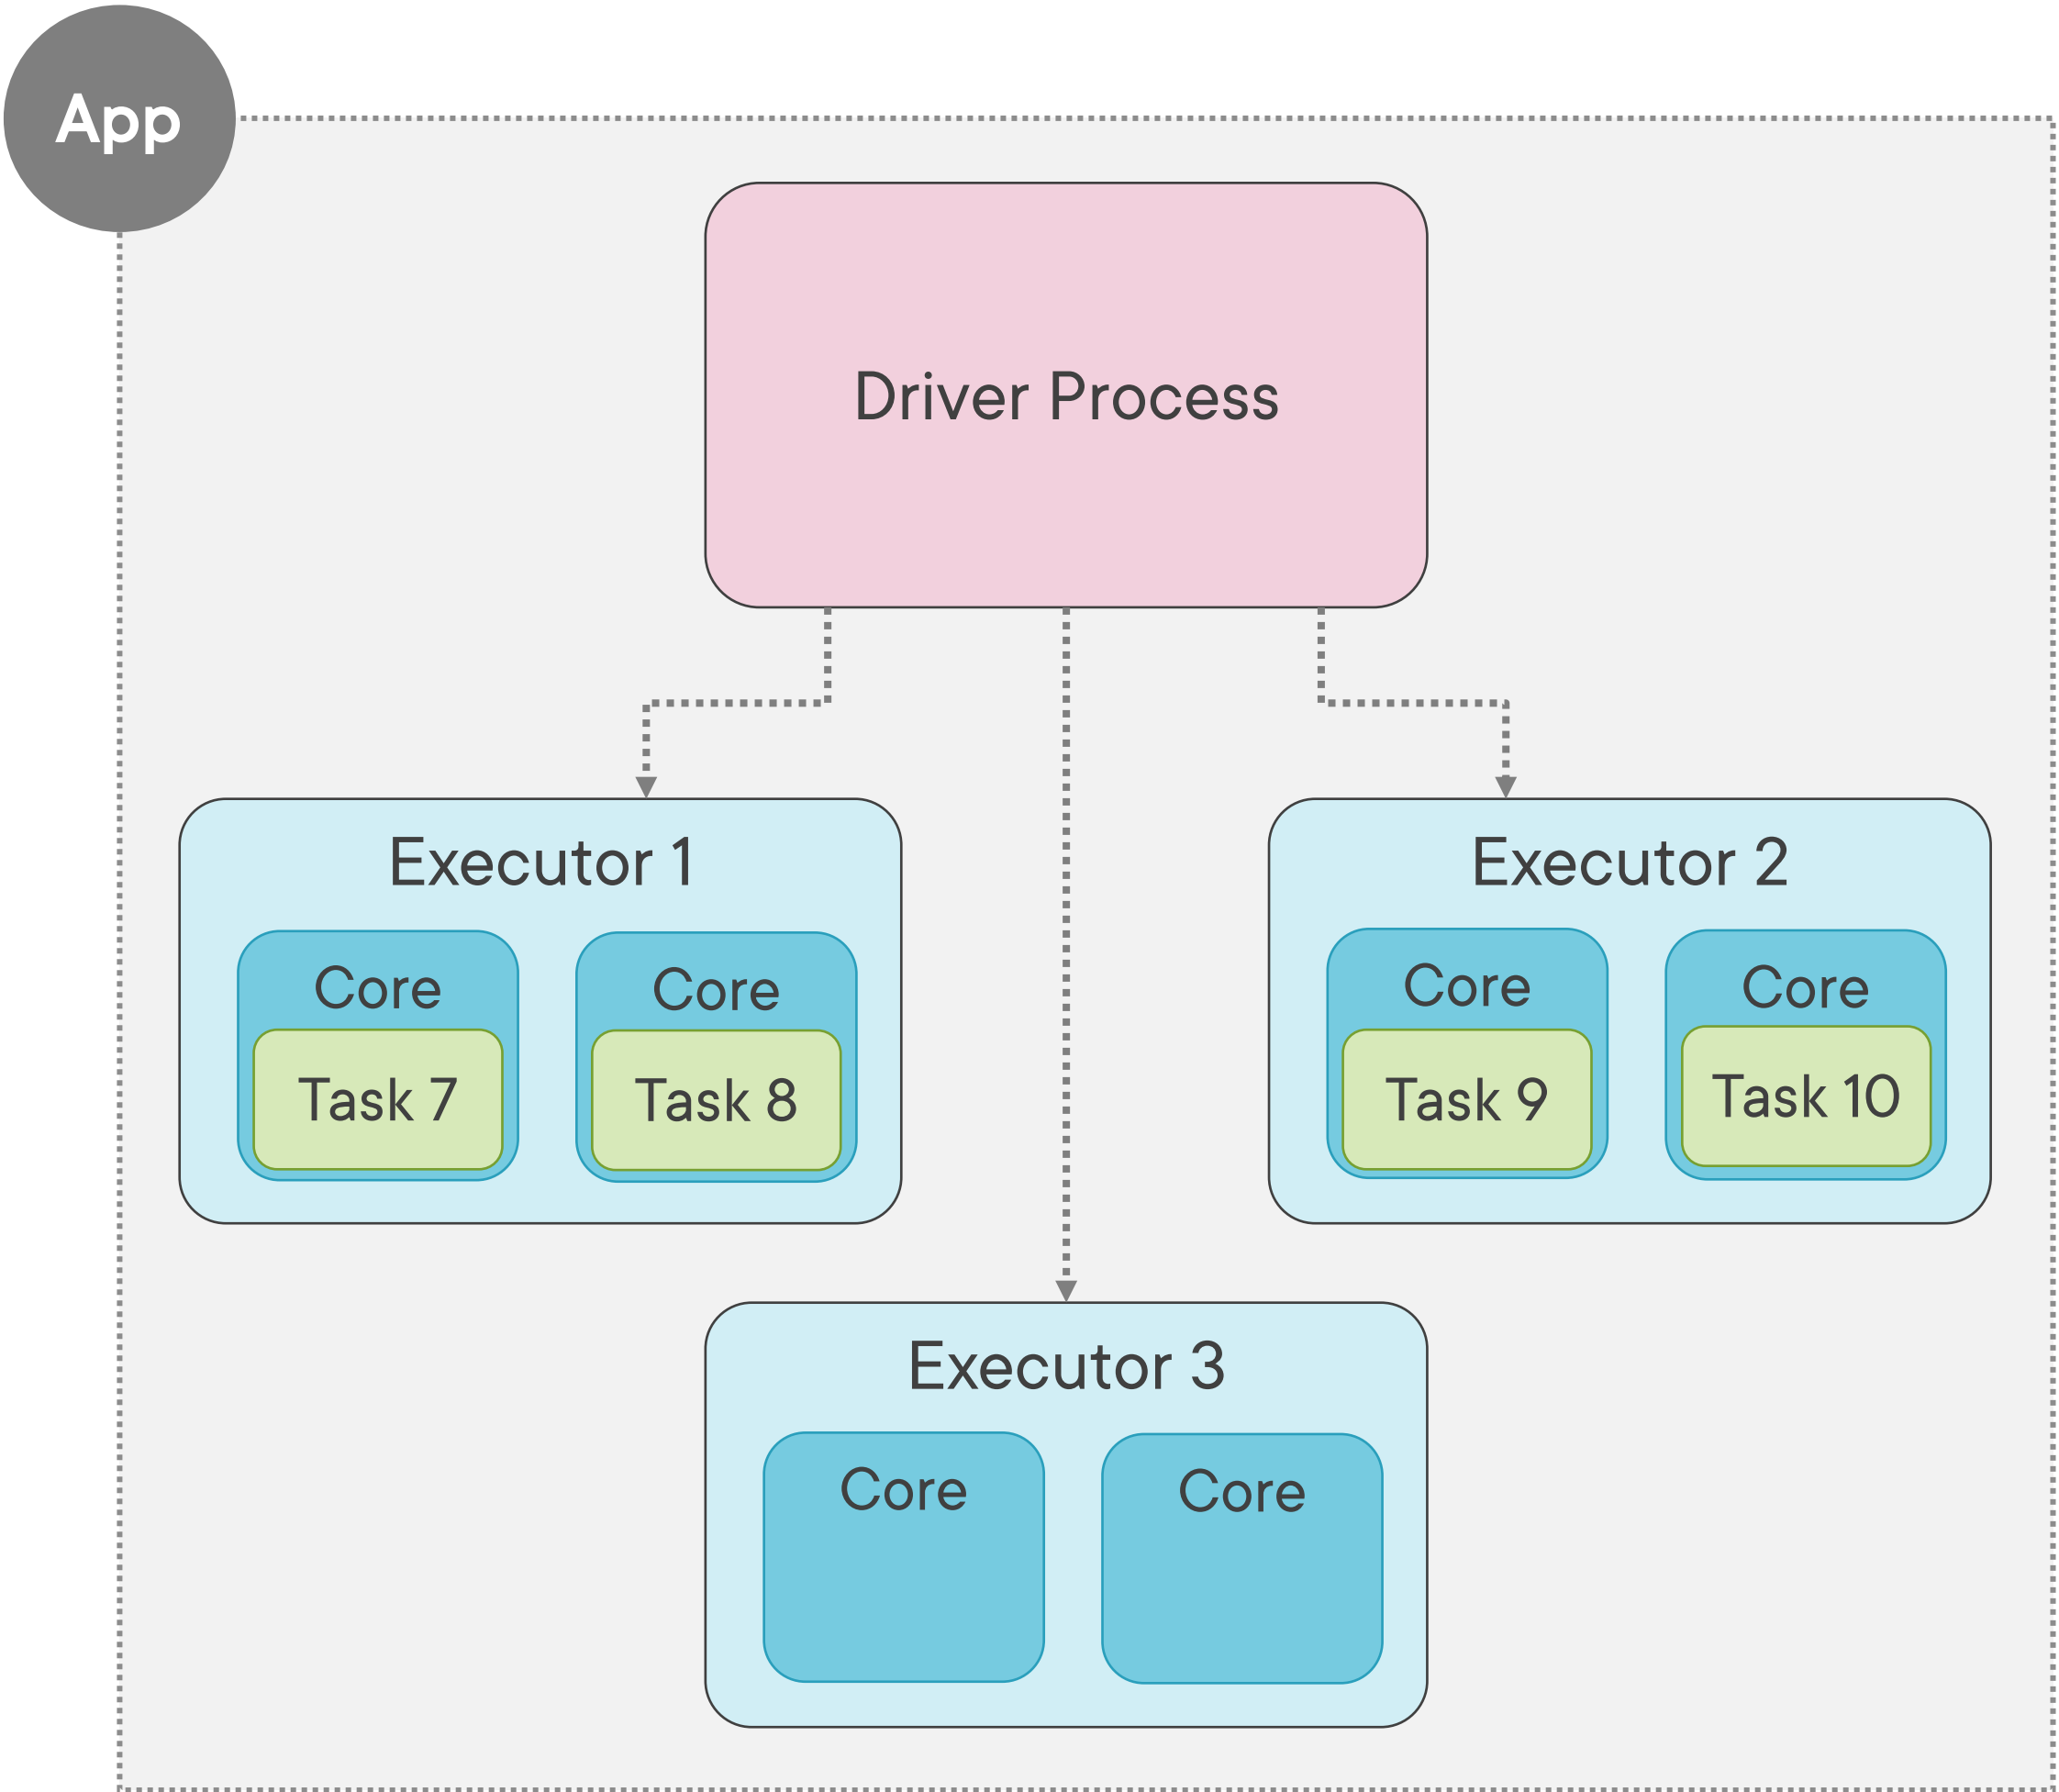



FAIR Scheduling

Job 1



Job 2

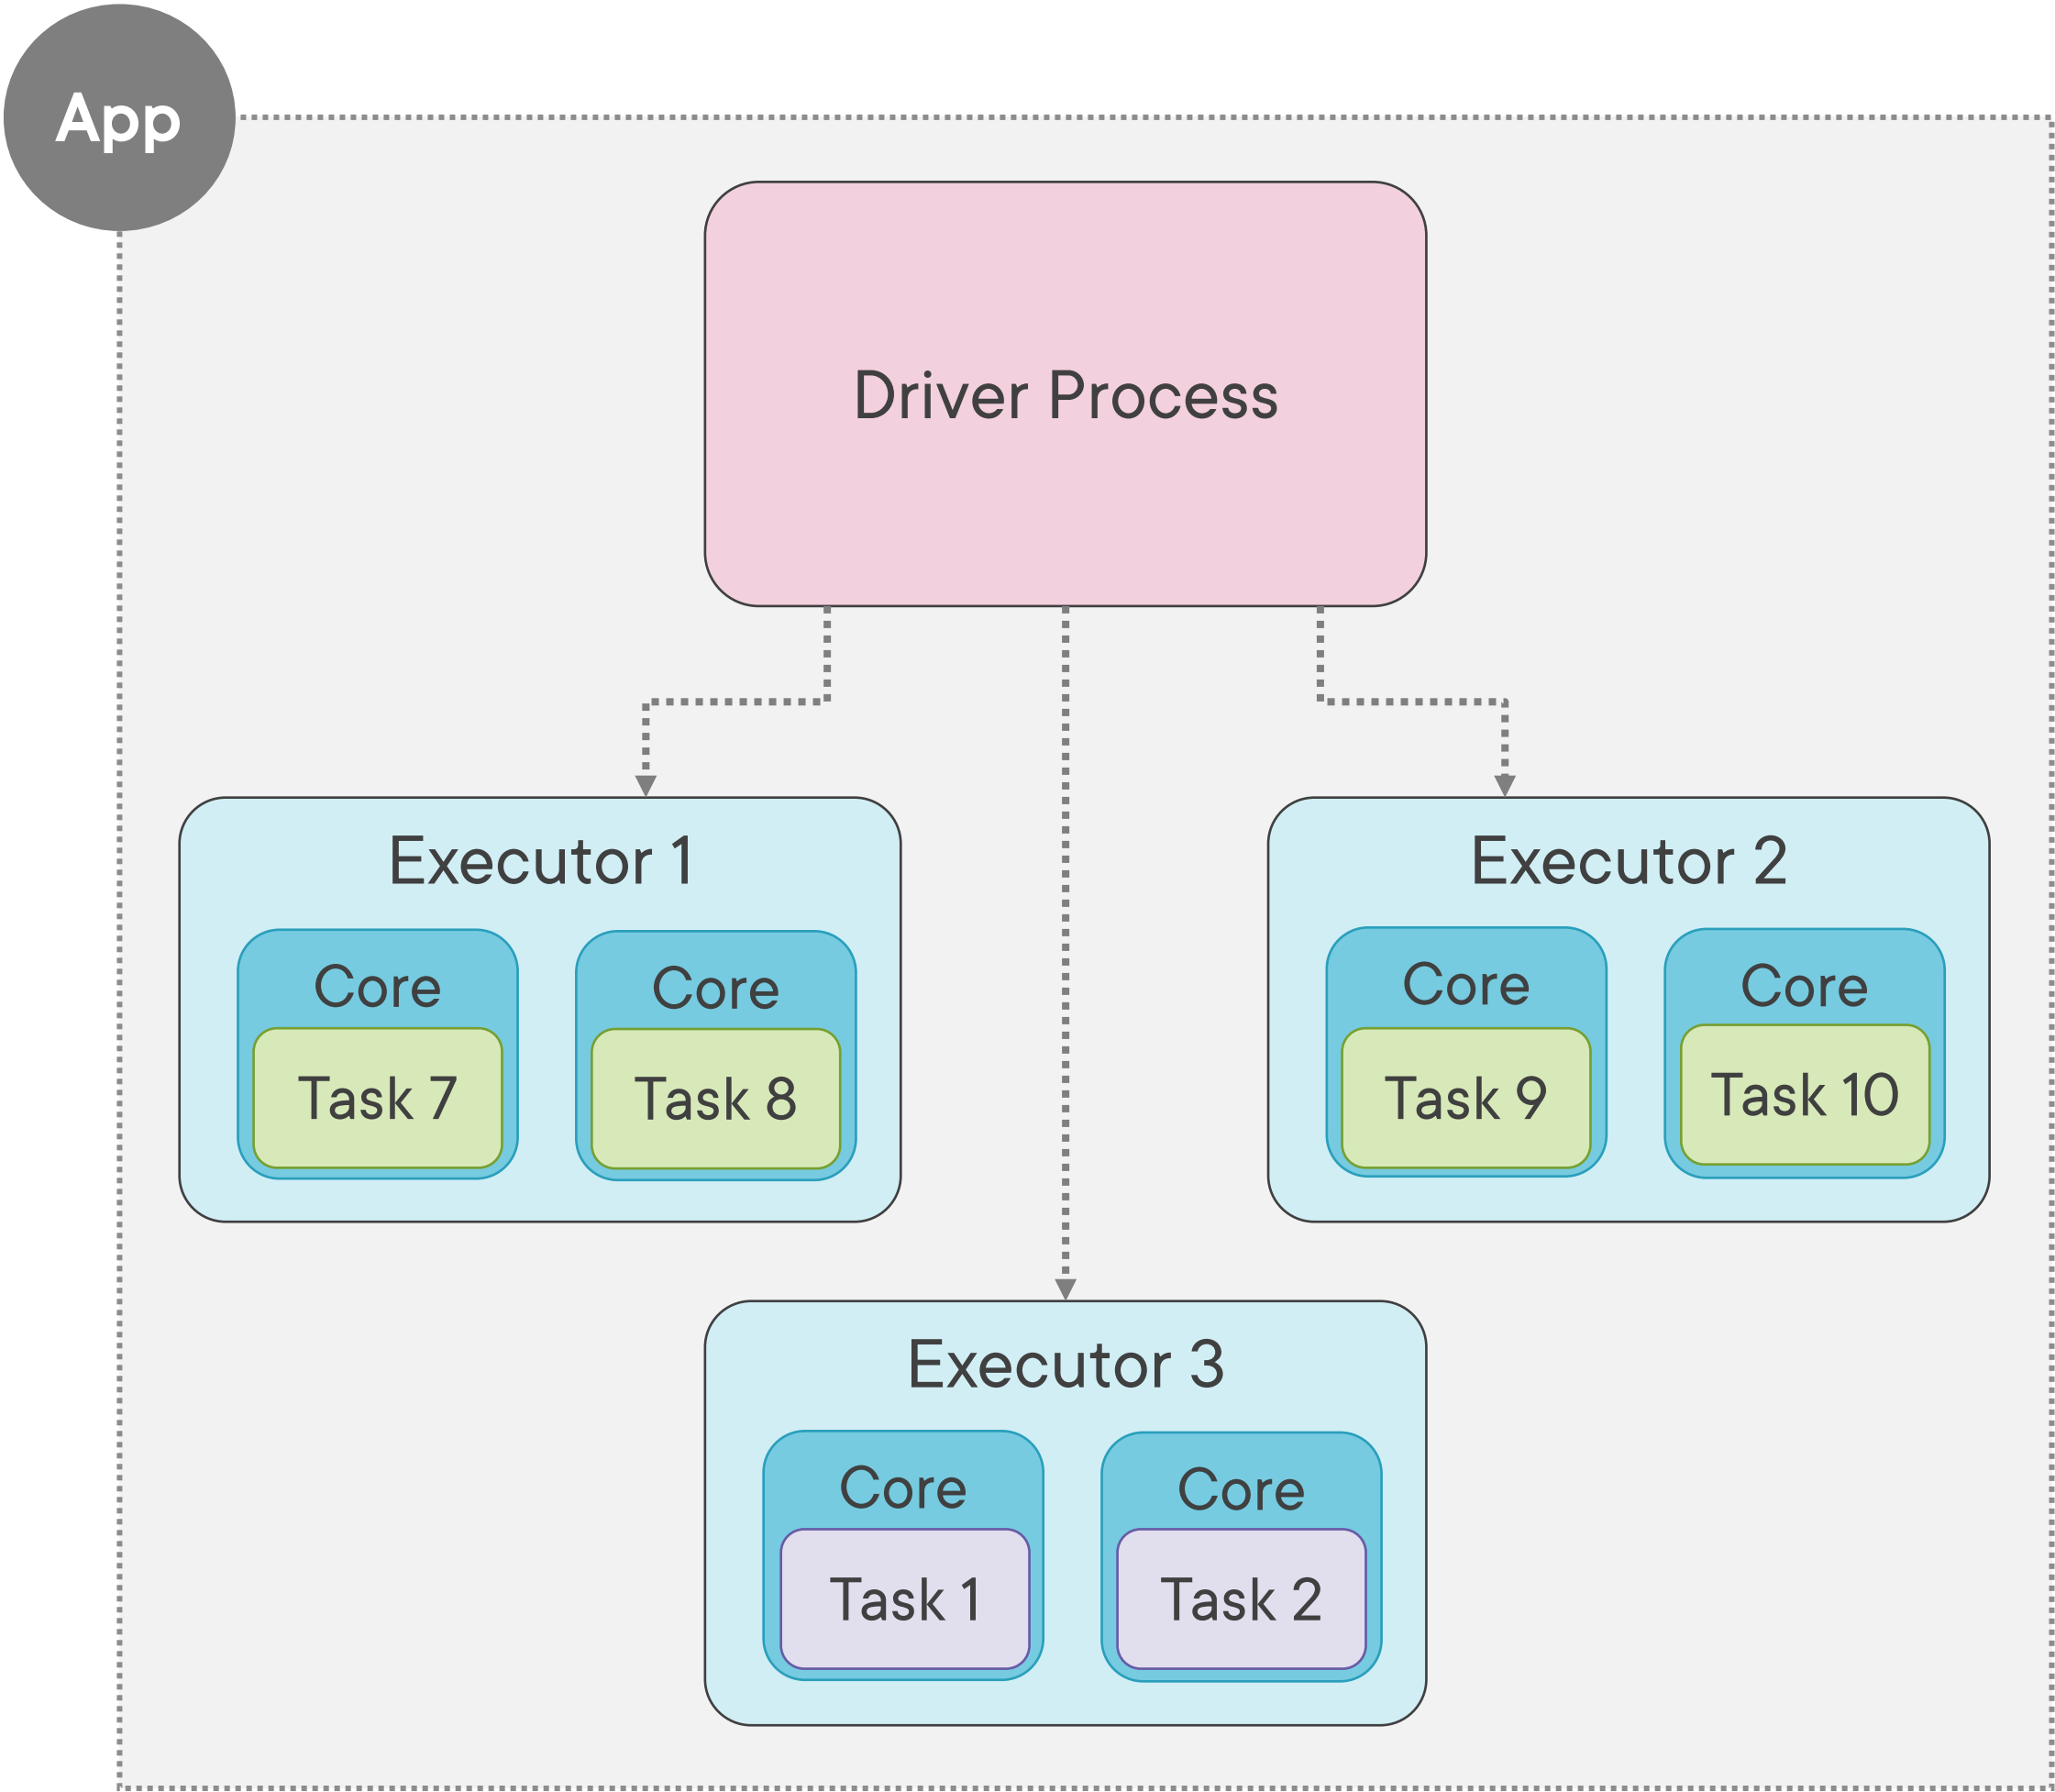



FAIR Scheduling

Job 1




Job 2

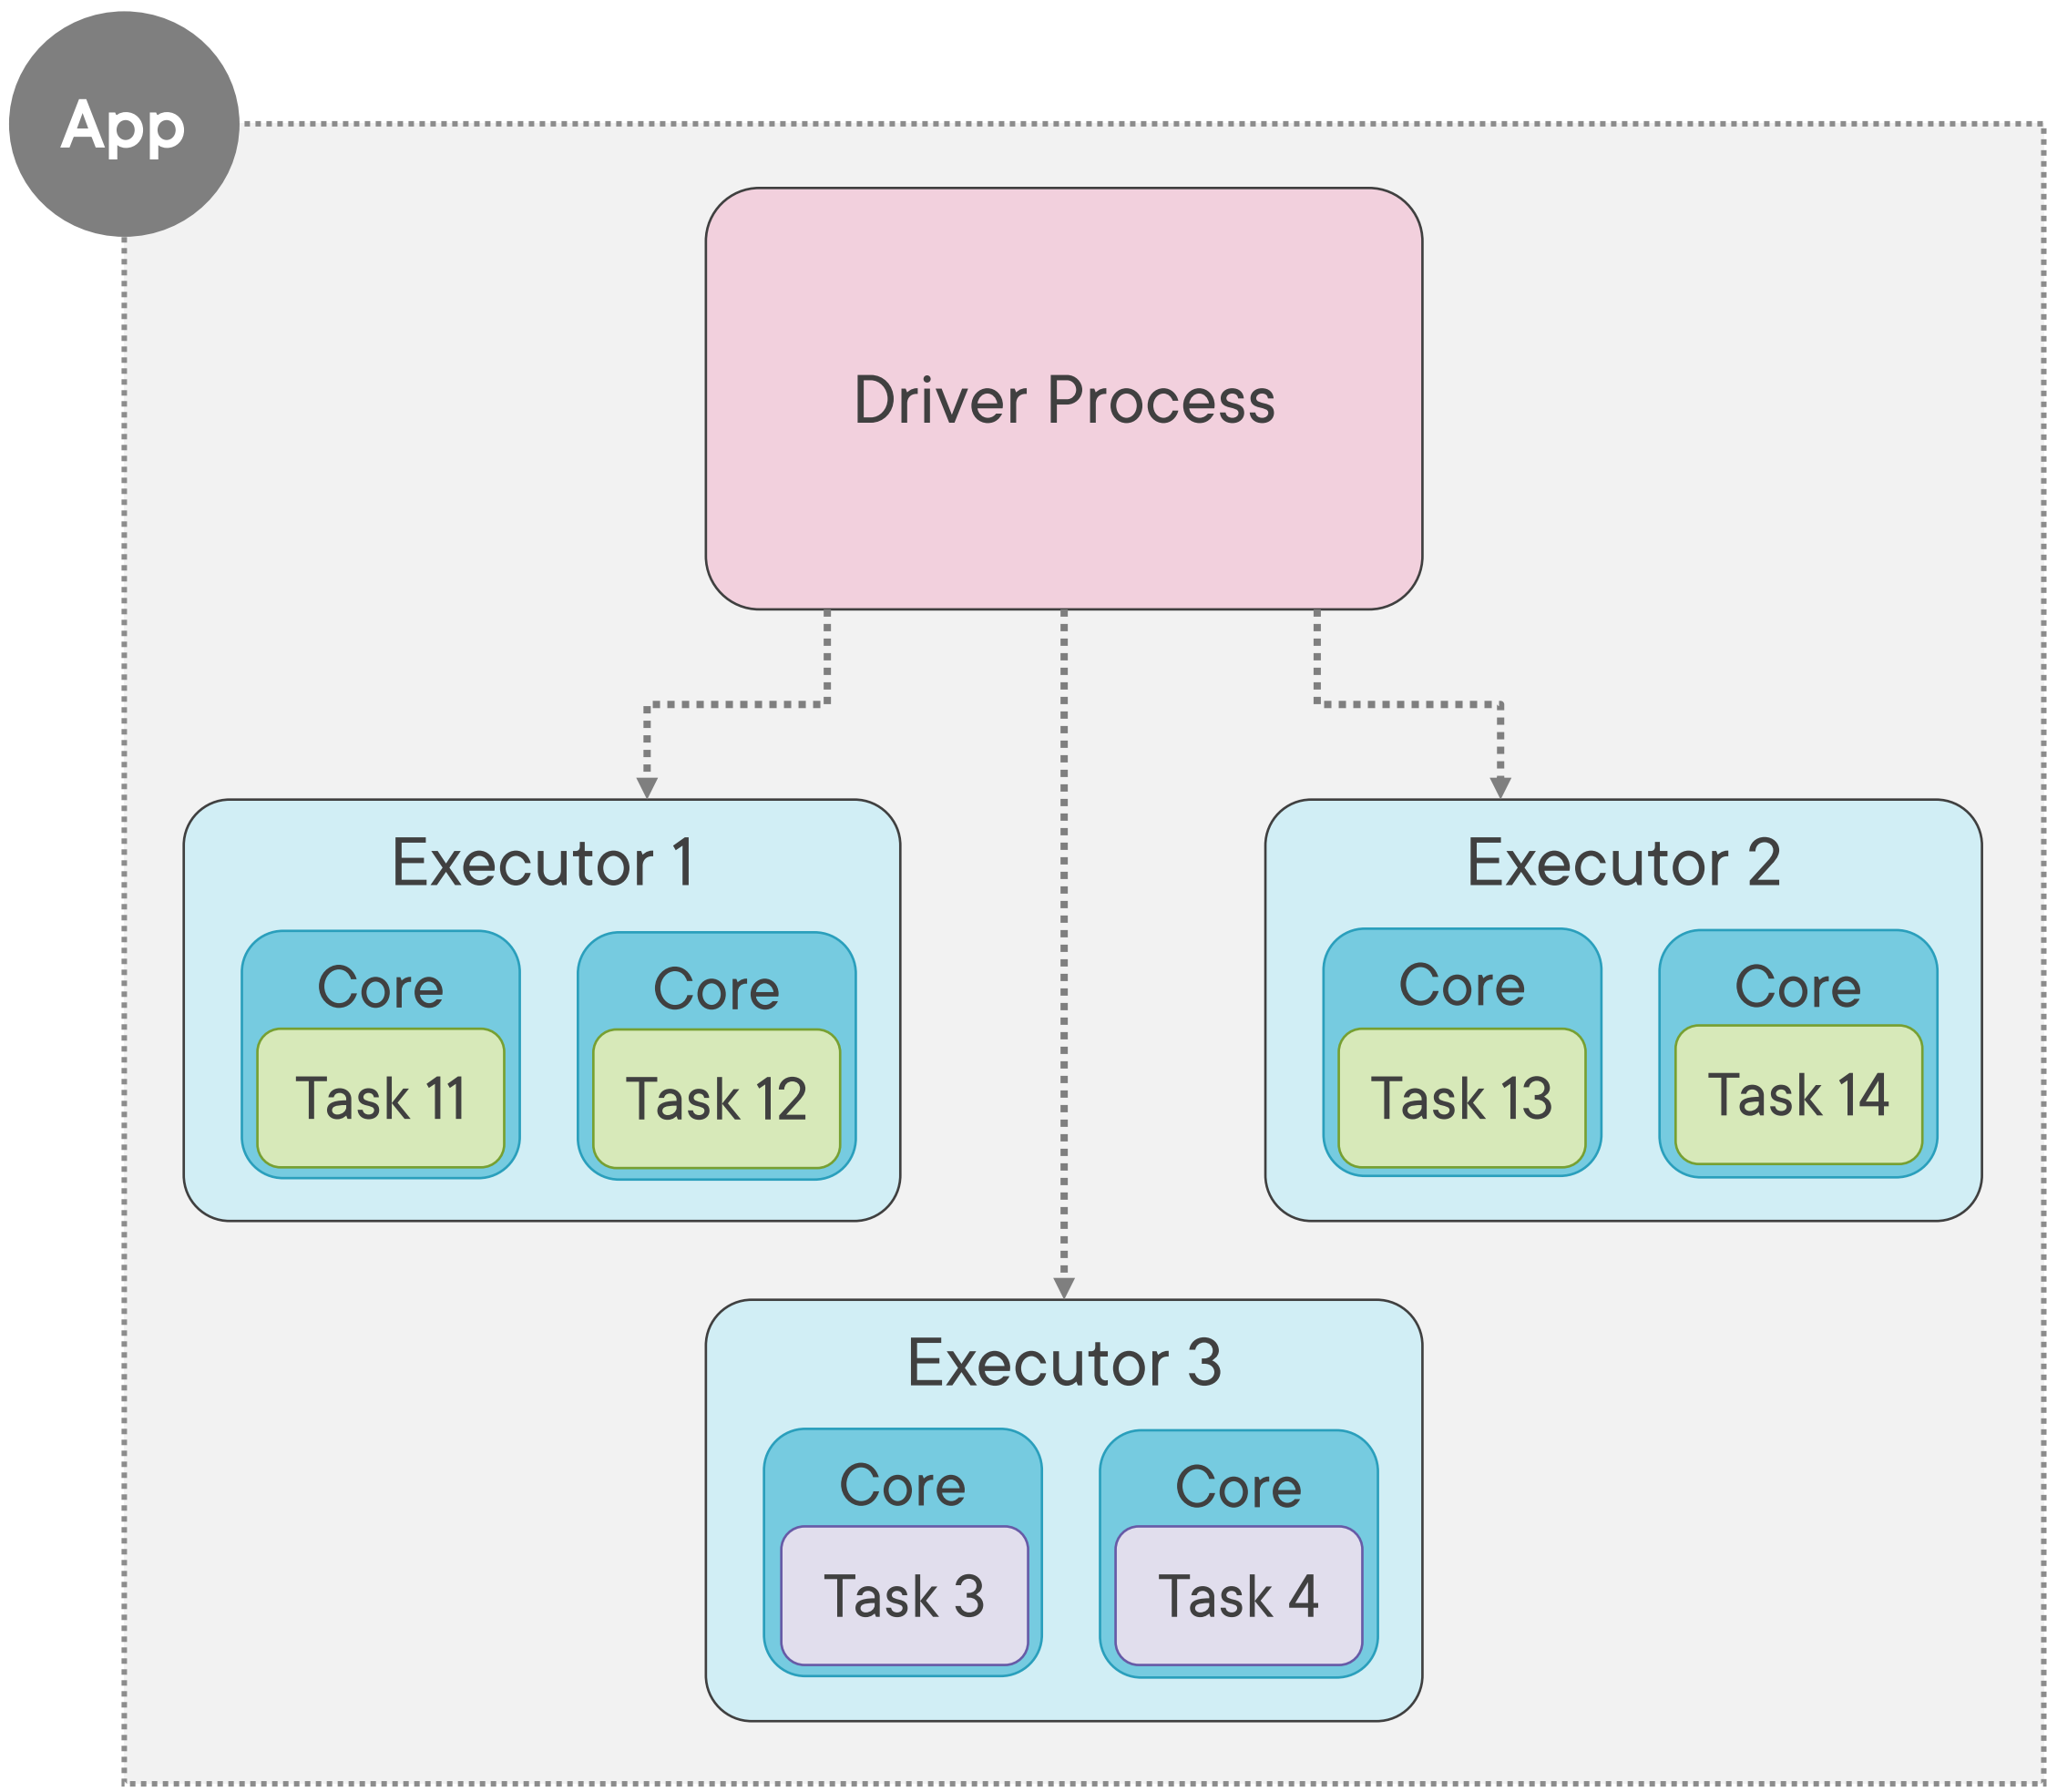



FAIR Scheduling

Job 1



Job 2



FAIR Scheduling

Job 1



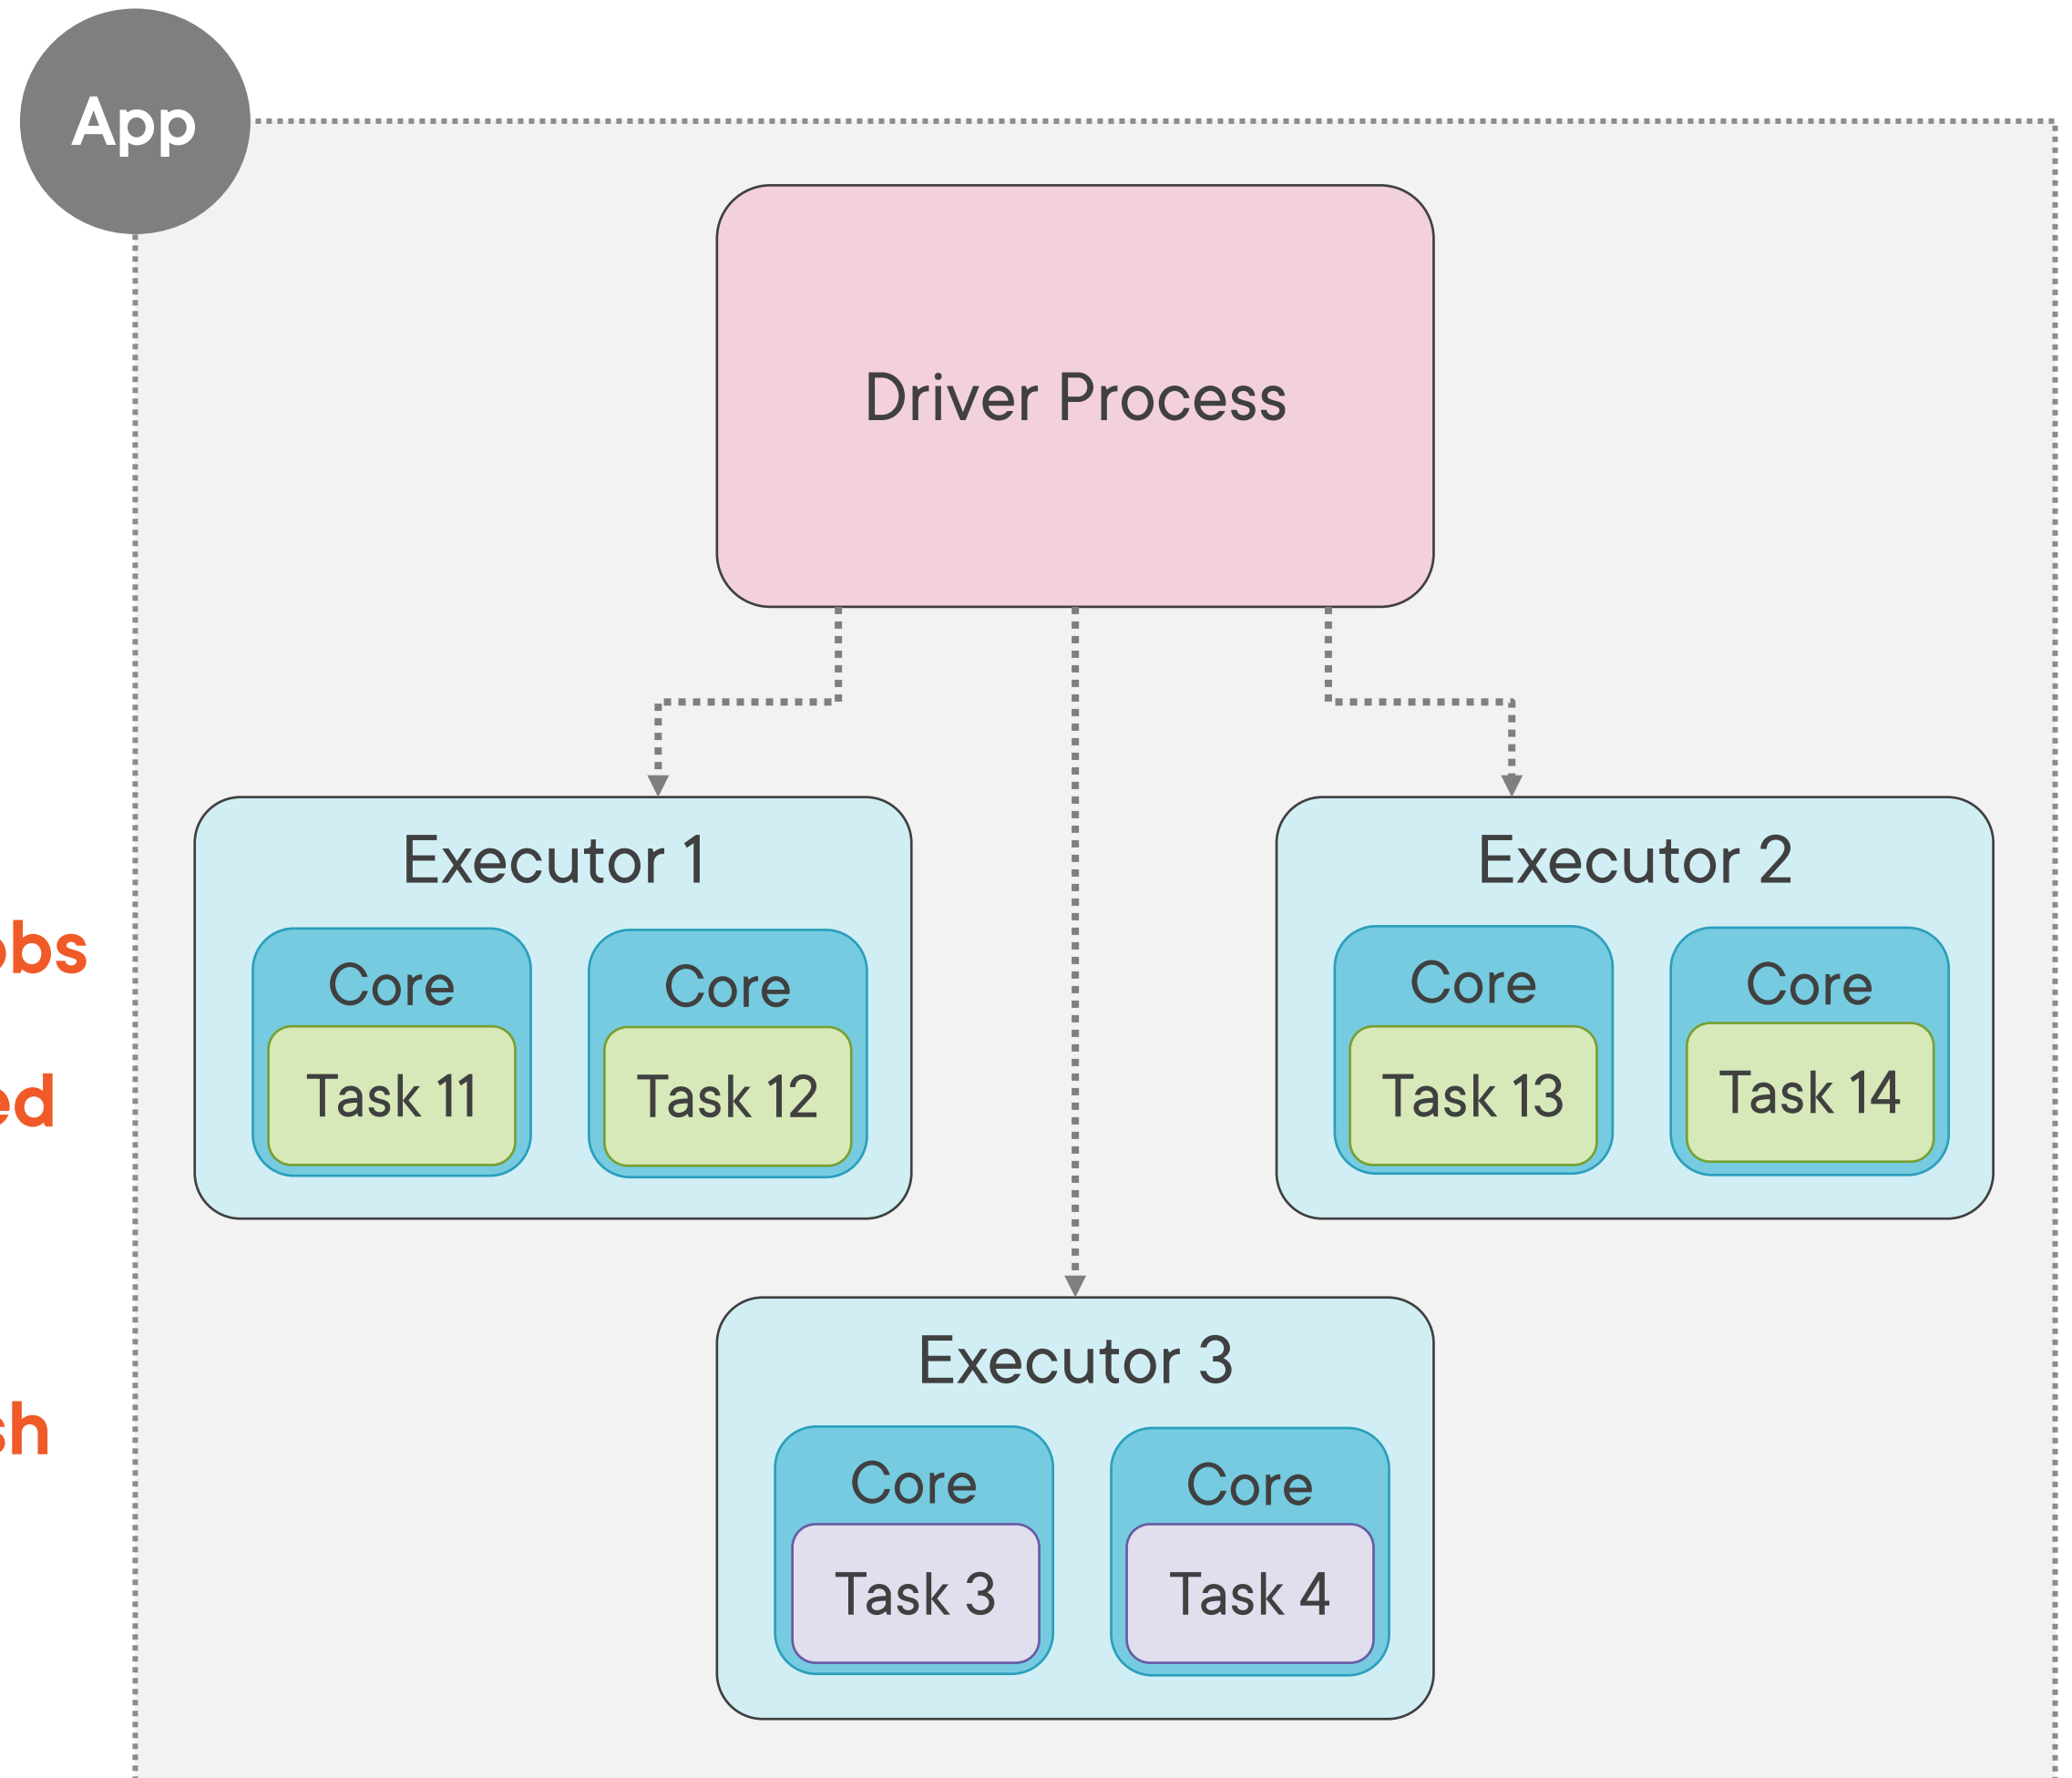
Fair sharing of resources between jobs

Tasks from different jobs are assigned in a round-robin fashion

Small jobs can finish faster, without waiting for long-running jobs to finish

To configure, set:

- `spark.scheduler.mode = FAIR`



Summary



In-memory partitions

- Settings to control partitions while reading & shuffling
- Change DF partitions using coalesce & repartition

Each driver & executor container has memory allocated

- Allocated JVM heap & non-heap memory
- Flexible memory usage b/w execution & caching

Persist data using `cache()` and `persist()`

- Avoids re-running transformations with every Action

Spark supports multiple join strategies

- Broadcast hash join: Large-small dataset join
- Shuffle sort merge join: Large-large dataset join
- Shuffle sort merge join can be improved by Bucketing

Spark supports dynamic allocation of resources

- Resource scheduling across applications
- Resource scheduling within application

Up Next:
Features in Apache Spark 3
